

CSC 252: Computer Organization

Spring 2019: Lecture 2

Instructor: Yuhao Zhu

Department of Computer Science
University of Rochester

Action Items:

- **Programming Assignment 1 is out**
- **Trivia 1 is due on Friday, midnight**

Announcement

- Programming Assignment 1 is out
 - Details: <http://cs.rochester.edu/courses/252/spring2019/labs/assignment1.html>
 - Due on Feb 1, 11:59 PM
 - Trivia due Friday, 1/25, 11:59 PM
 - You have 3 slip days (not for trivia)

20	21	22	23	24	25	26
					Trivia	
27	28	29	30	31	Feb 1	2
					Due	

Announcement

- TA office hours are all posted. Start from this week.
- TA review sessions schedule to be posted soon...
- Programming assignment 1 is in C language. Seek help from TAs.
- Programming assignments do NOT repeat the lecture materials. They ask you to synthesize what you have learned from the lectures and work out something new.

Previously in 252...

Problem

Algorithm

Program

Instruction Set
Architecture (ISA)

Microarchitecture

Circuit

Previously in 252...

Problem

Algorithm

Program

Instruction Set
Architecture (ISA)

ISA is the contract
between software and
hardware.

Microarchitecture

Circuit

Previously in 252...

Problem

Algorithm

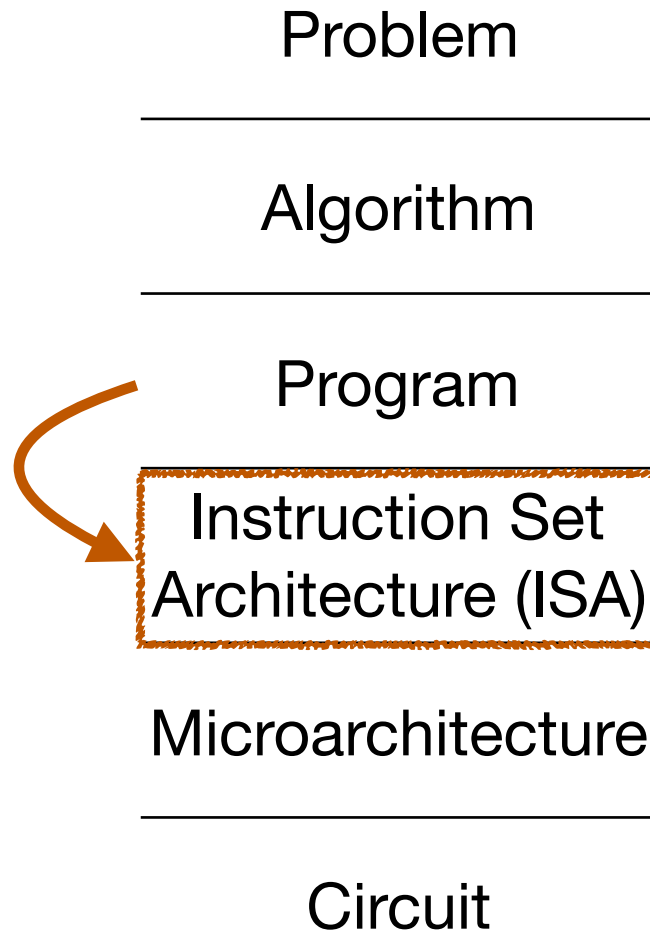
	Renting	Computing
Service provider	Landlord	Hardware
Service receiver	YOU	Software
Contract	Lease	Assembly Program
Contract's language	Natural language (e.g., English)	ISA

Circuit

ct
e and

Previously in 252...

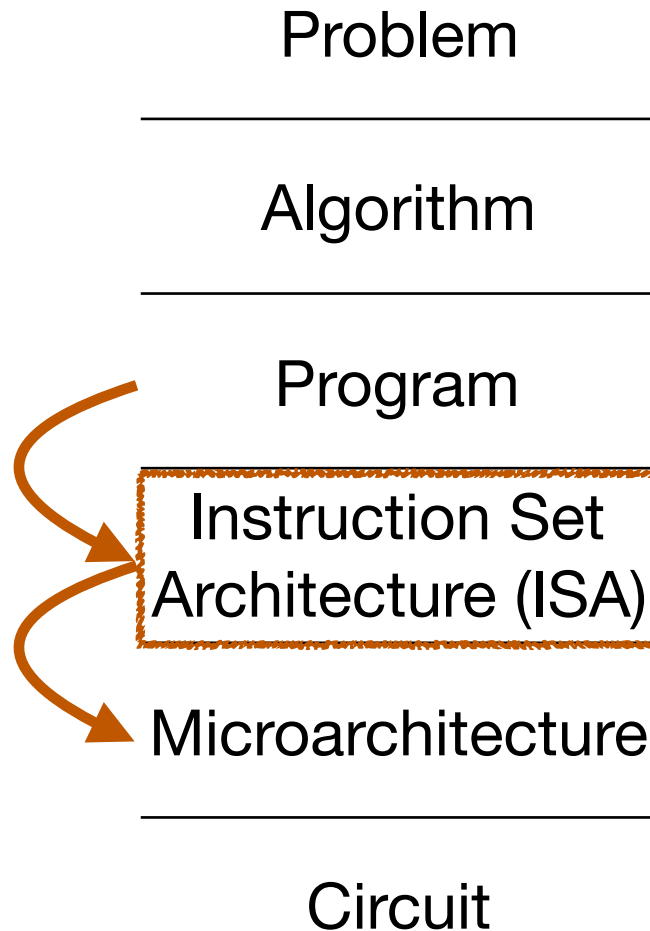
- How is a human-readable program translated to a representation that computers can understand?



ISA is the contract between software and hardware.

Previously in 252...

- How is a human-readable program translated to a representation that computers can understand?
- How does a modern computer execute that program?

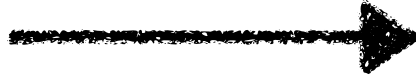


ISA is the contract between software and hardware.

Previously in 252...

C Program

```
void add() {  
    int a = 1;  
    int b = 2;  
    int c = a + b;  
}
```



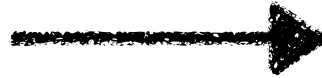
Assembly program

```
movl    $1, -4(%rbp)  
movl    $2, -8(%rbp)  
movl    -4(%rbp), %eax  
addl    -8(%rbp), %eax
```

Previously in 252...

Assembly program

```
movl    $1, -4(%rbp)
movl    $2, -8(%rbp)
movl    -4(%rbp), %eax
addl    -8(%rbp), %eax
```



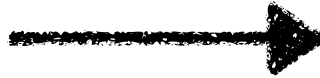
Executable Binary

```
00011001 ...
01101010 ...
11010101 ...
01110001 ...
```

Previously in 252...

Assembly program

```
movl    $1, -4(%rbp)
movl    $2, -8(%rbp)
movl    -4(%rbp), %eax
addl    -8(%rbp), %eax
```



Executable Binary

```
00011001 ...
01101010 ...
11010101 ...
01110001 ...
```

- What's the difference between an assembly program and an executable binary?
 - They refer to the same thing — a list of instructions that the software asks the hardware to perform
 - They are just different representations
- Instruction = Operator + Operand(s)

Previously in 252...

Assembly program

```
movl $1, -4(%rbp)
movl $2, -8(%rbp)
movl -4(%rbp), %eax
addl -8(%rbp), %eax
```



Executable Binary

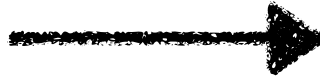
```
00011001 ...
01101010 ...
11010101 ...
01110001 ...
```

- What's the difference between an assembly program and an executable binary?
 - They refer to the same thing — a list of instructions that the software asks the hardware to perform
 - They are just different representations
- Instruction = Operator + Operand(s)

Previously in 252...

Assembly program

```
movl  $1, -4(%rbp)
movl  $2, -8(%rbp)
movl  -4(%rbp), %eax
addl  -8(%rbp), %eax
```



Executable Binary

```
00011001 ...
01101010 ...
11010101 ...
01110001 ...
```

- What's the difference between an assembly program and an executable binary?
 - They refer to the same thing — a list of instructions that the software asks the hardware to perform
 - They are just different representations
- Instruction = Operator + Operand(s)

Today: Representing Information in Binary

- Why Binary (bits)?
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary
- Representations in memory, pointers, strings

Everything is bits

Everything is bits

- Each bit is 0 or 1. Bits are how programs talk to the hardware

Everything is bits

- Each bit is 0 or 1. Bits are how programs talk to the hardware
- Programs encode instructions in bits

Everything is bits

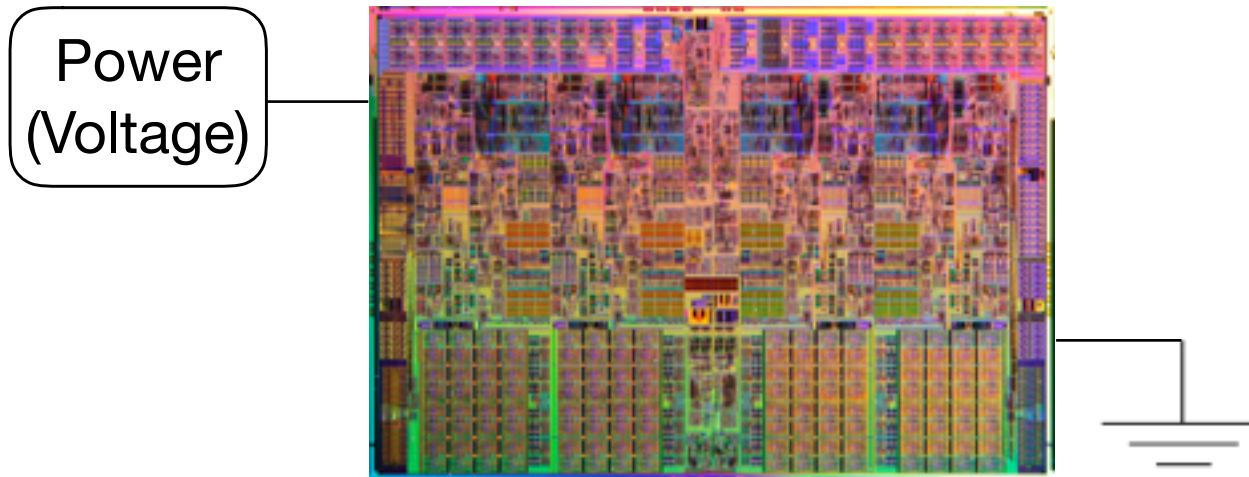
- Each bit is 0 or 1. Bits are how programs talk to the hardware
- Programs encode instructions in bits
- Hardware then interprets the bits

Everything is bits

- Each bit is 0 or 1. Bits are how programs talk to the hardware
- Programs encode instructions in bits
- Hardware then interprets the bits
- Why bits? Electronic Implementation

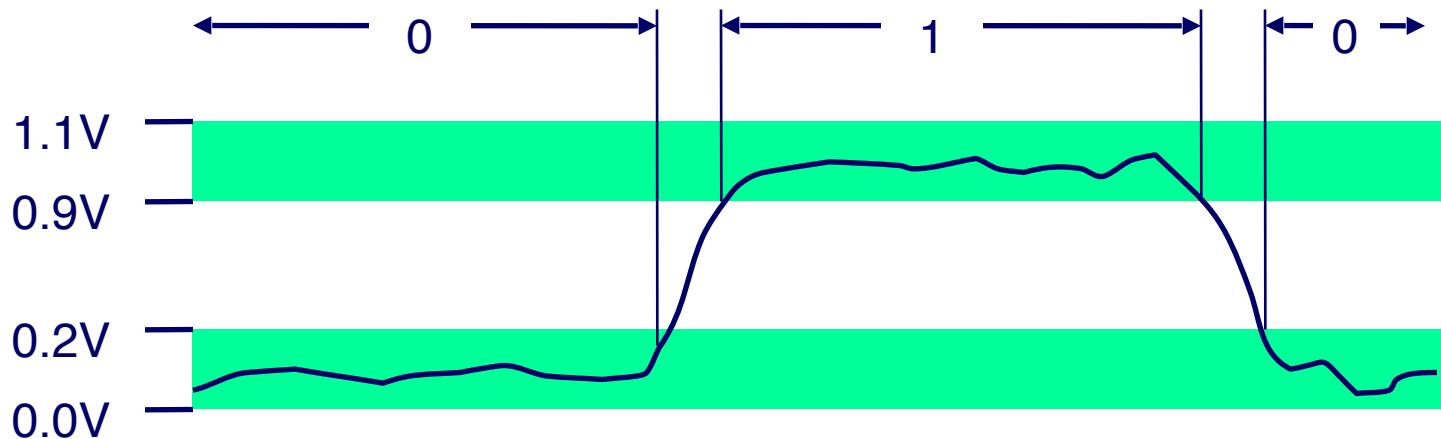
Everything is bits

- Each bit is 0 or 1. Bits are how programs talk to the hardware
- Programs encode instructions in bits
- Hardware then interprets the bits
- Why bits? Electronic Implementation



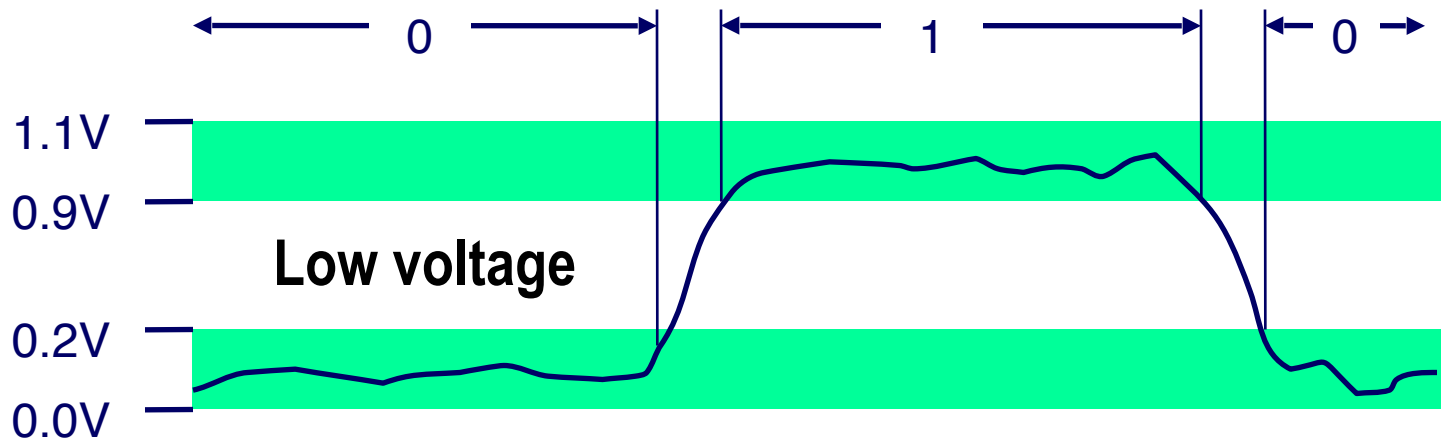
Why Bits?

- Each bit is 0 or 1. Bits are how programs talk to the hardware
- Programs encode instructions in bits
- Hardware then interprets the bits
- Why bits? Electronic Implementation
 - Use high voltage to represent 1
 - Use low voltage to represent 0



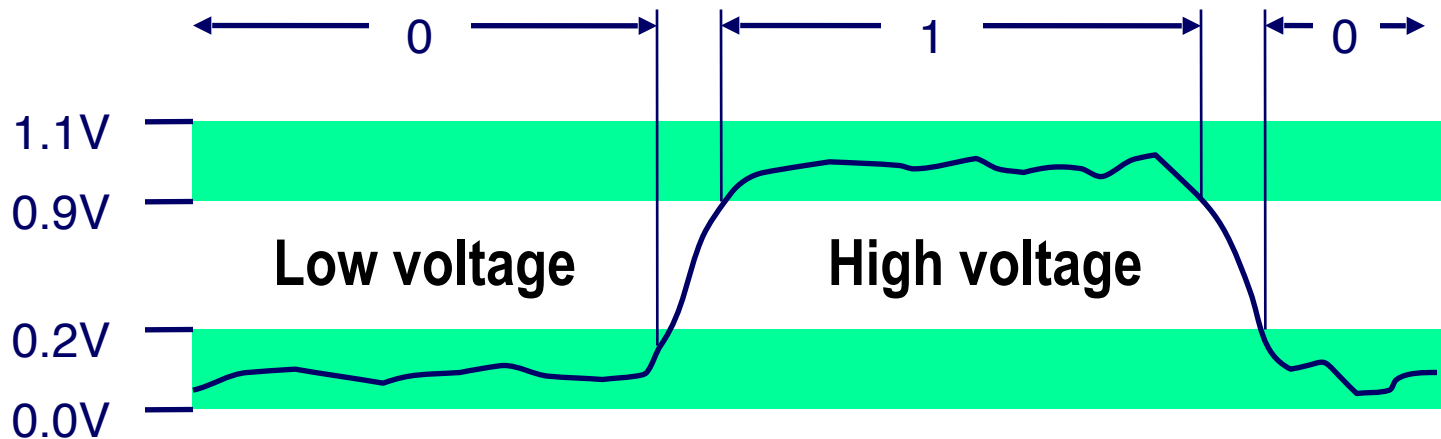
Why Bits?

- Each bit is 0 or 1. Bits are how programs talk to the hardware
- Programs encode instructions in bits
- Hardware then interprets the bits
- Why bits? Electronic Implementation
 - Use high voltage to represent 1
 - Use low voltage to represent 0



Why Bits?

- Each bit is 0 or 1. Bits are how programs talk to the hardware
- Programs encode instructions in bits
- Hardware then interprets the bits
- Why bits? Electronic Implementation
 - Use high voltage to represent 1
 - Use low voltage to represent 0



Why Bits?

Why Bits?

Processors are made of transistors, which are Metal Oxide Semiconductor (MOS)

Why Bits?

Processors are made of transistors, which are Metal Oxide Semiconductor (MOS)

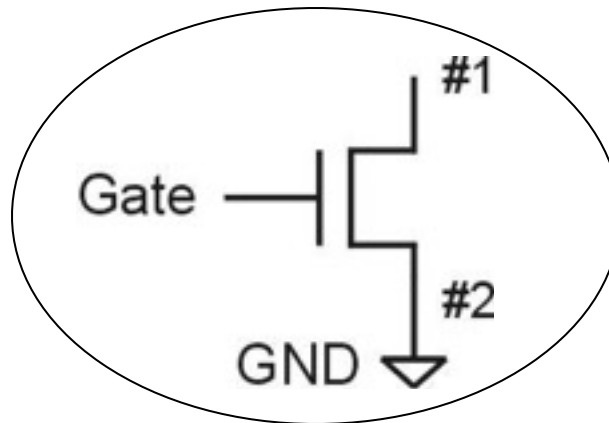
- two types: n-type and p-type

Why Bits?

Processors are made of transistors, which are Metal Oxide Semiconductor (MOS)

- two types: n-type and p-type

n-type (NMOS)



Terminal #2 must be connected to GND (0V).

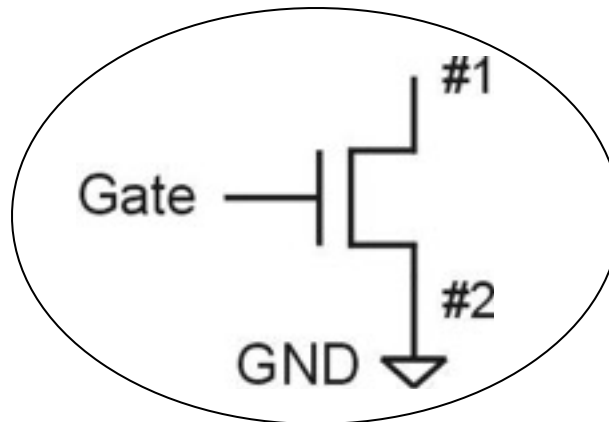
Why Bits?

Processors are made of transistors, which are Metal Oxide Semiconductor (MOS)

- two types: n-type and p-type

n-type (NMOS)

- when Gate has high voltage,
short circuit between #1 and #2
(switch closed)



Terminal #2 must be connected to GND (0V).

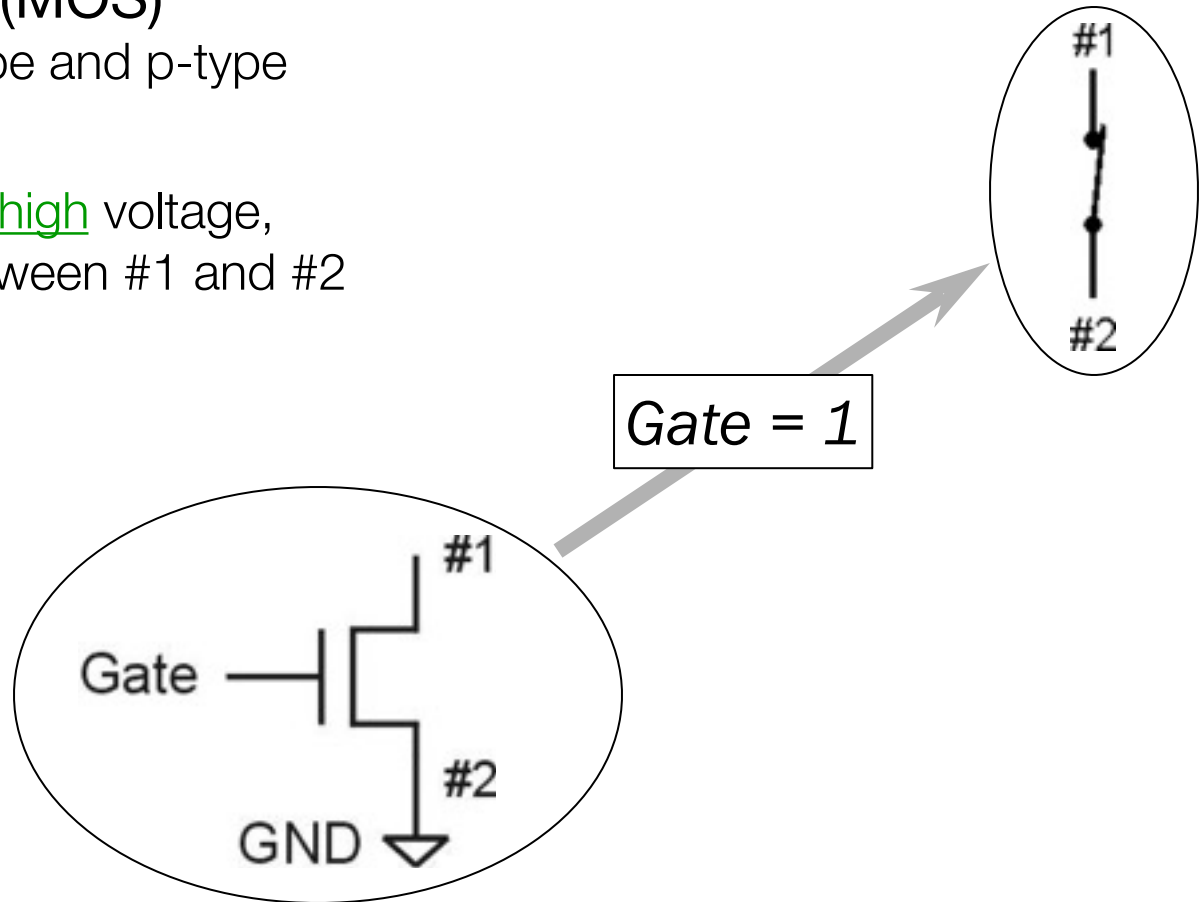
Why Bits?

Processors are made of transistors, which are Metal Oxide Semiconductor (MOS)

- two types: n-type and p-type

n-type (NMOS)

- when Gate has high voltage, short circuit between #1 and #2 (switch closed)



Terminal #2 must be connected to GND (0V).

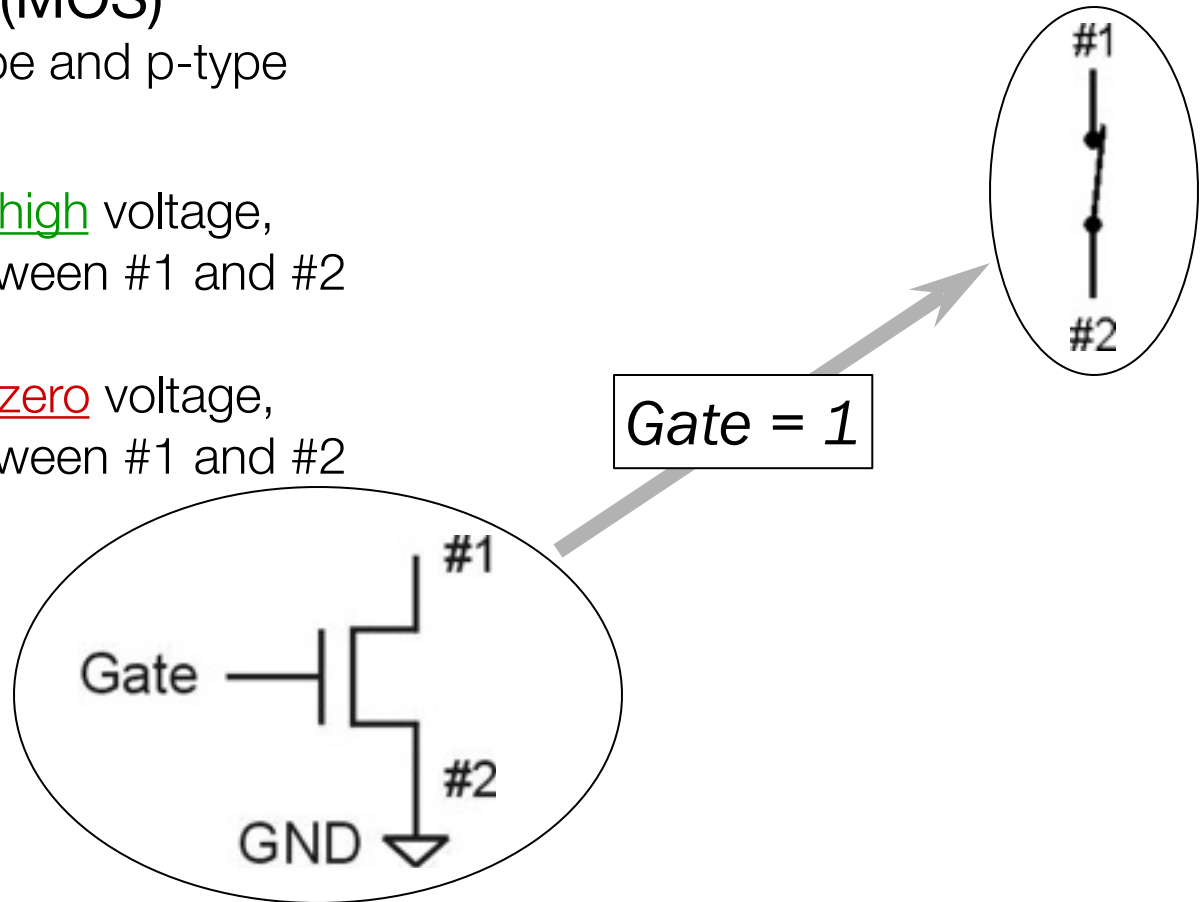
Why Bits?

Processors are made of transistors, which are Metal Oxide Semiconductor (MOS)

- two types: n-type and p-type

n-type (NMOS)

- when Gate has high voltage, short circuit between #1 and #2 (switch closed)
- when Gate has zero voltage, open circuit between #1 and #2 (switch open)



Terminal #2 must be connected to GND (0V).

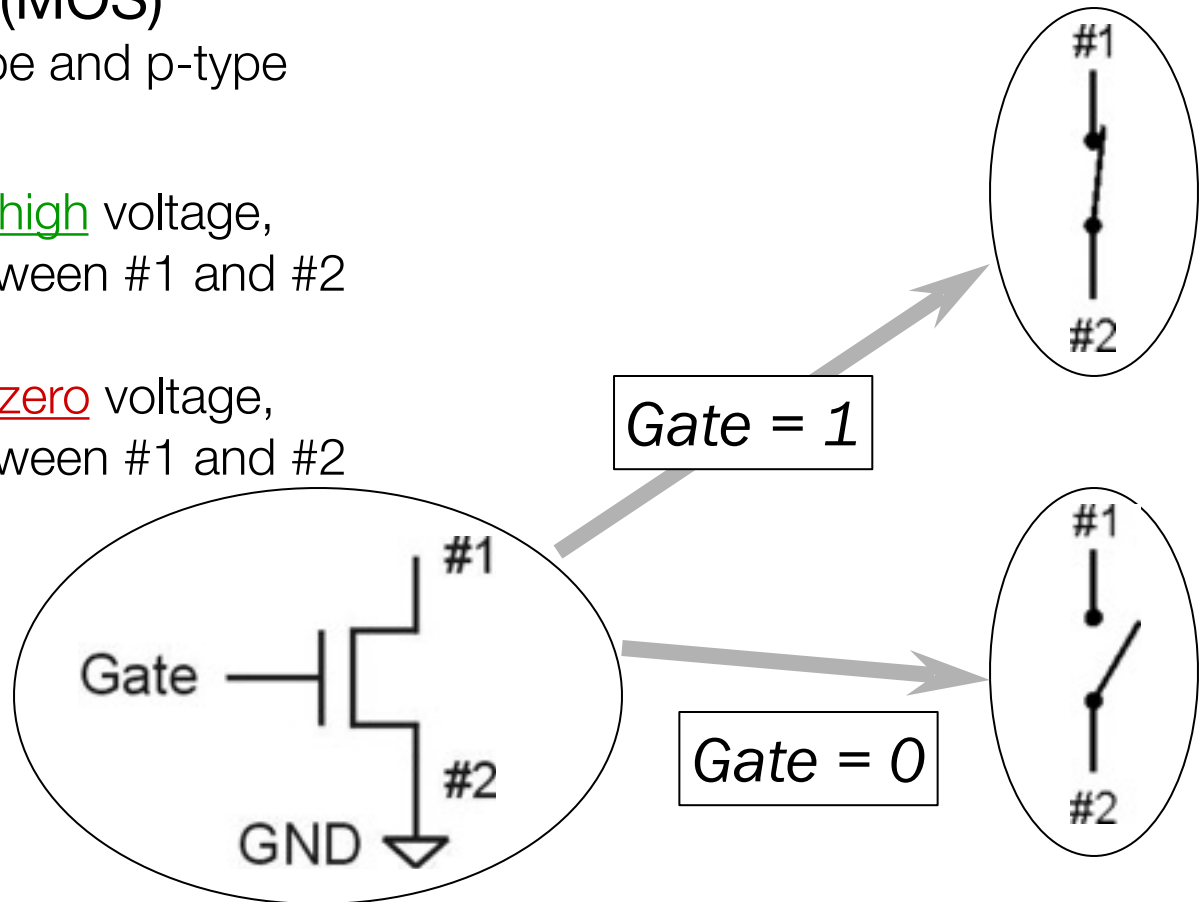
Why Bits?

Processors are made of transistors, which are Metal Oxide Semiconductor (MOS)

- two types: n-type and p-type

n-type (NMOS)

- when Gate has high voltage, short circuit between #1 and #2 (switch closed)
- when Gate has zero voltage, open circuit between #1 and #2 (switch open)

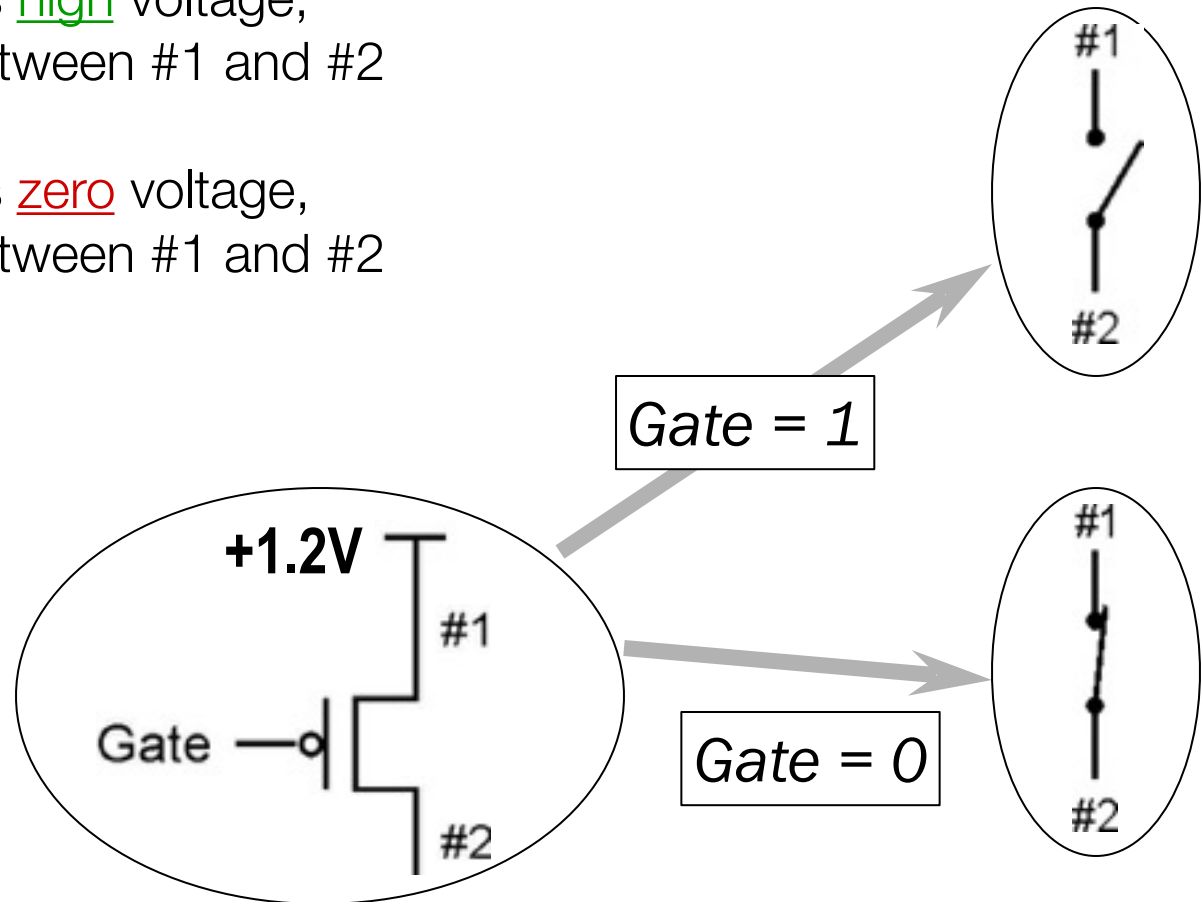


Terminal #2 must be connected to GND (0V).

Why Bits?

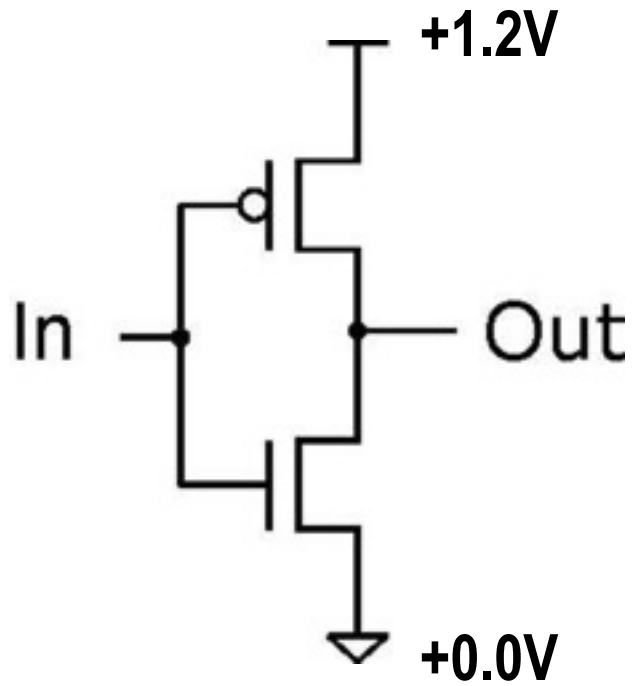
p-type is *complementary* to n-type (**PMOS**)

- when Gate has high voltage, open circuit between #1 and #2 (switch open)
- when Gate has zero voltage, short circuit between #1 and #2 (switch closed)

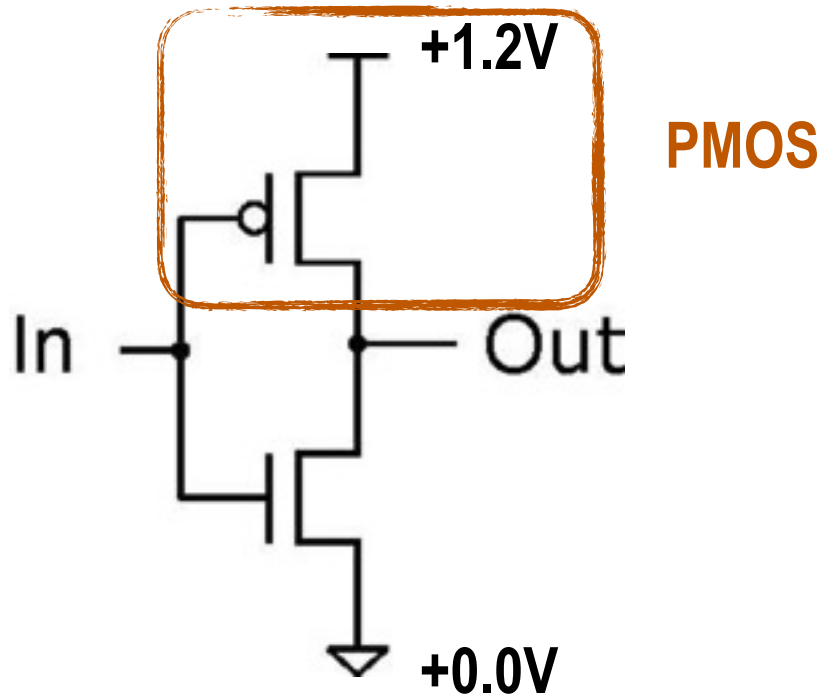


Terminal #1 must be connected to +1.2V

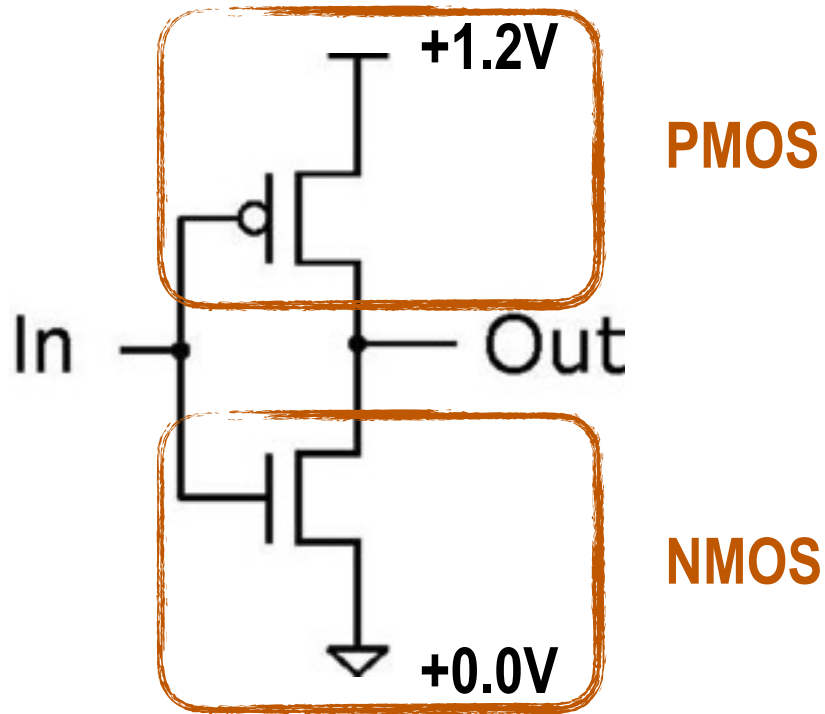
Inverter



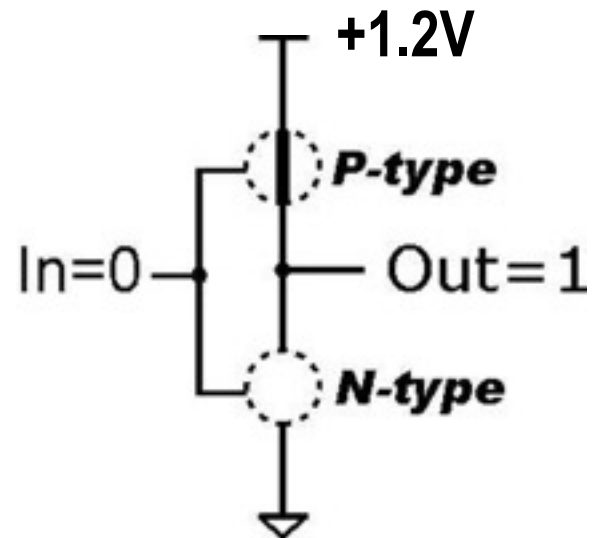
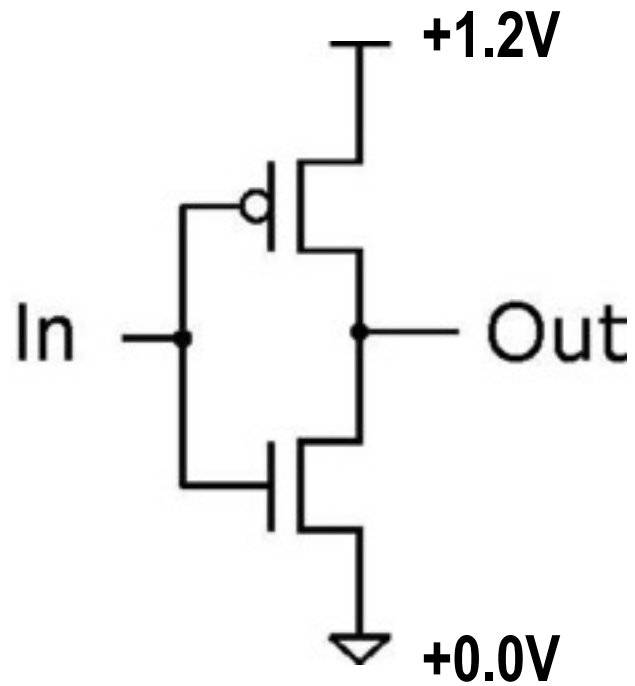
Inverter



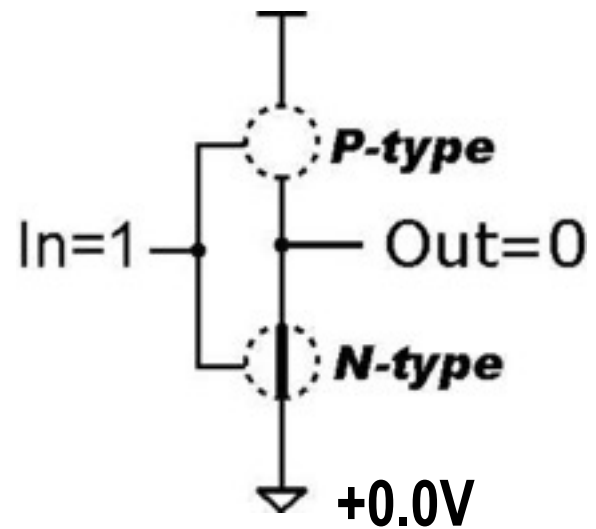
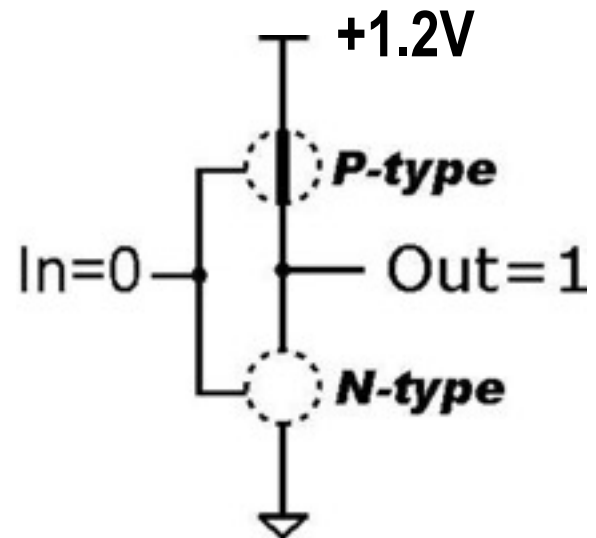
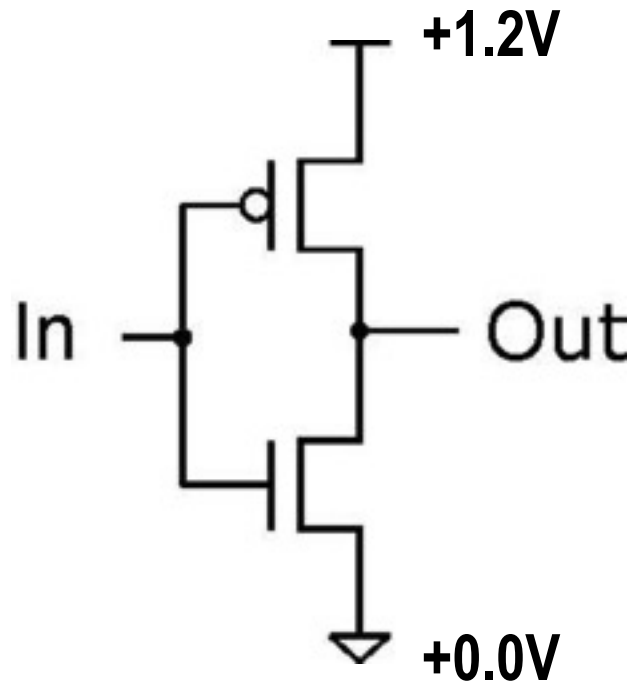
Inverter



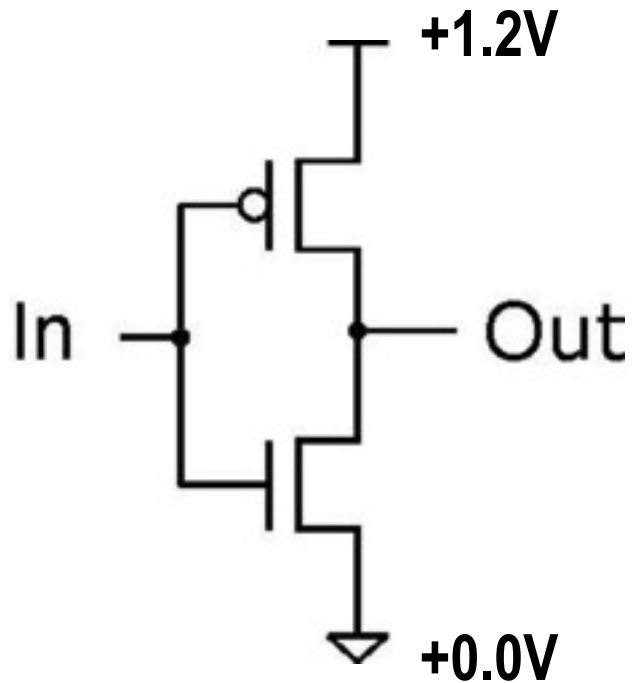
Inverter



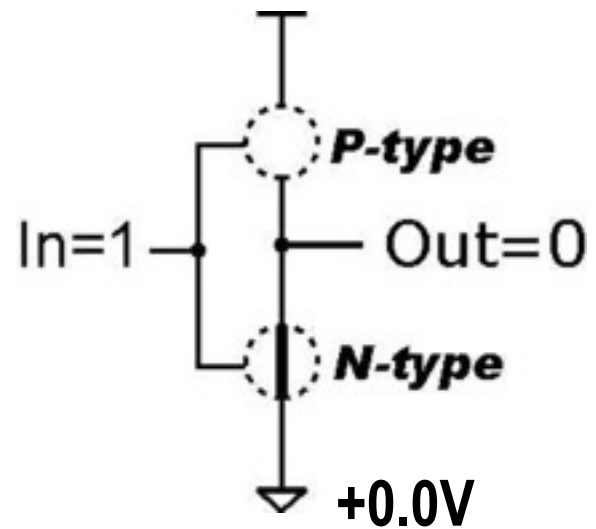
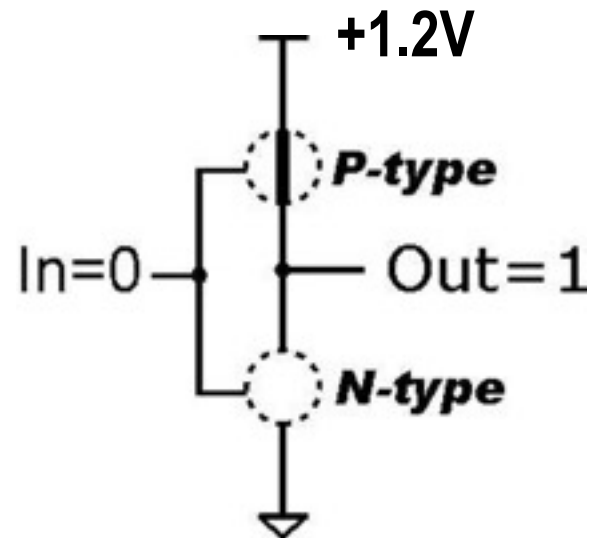
Inverter



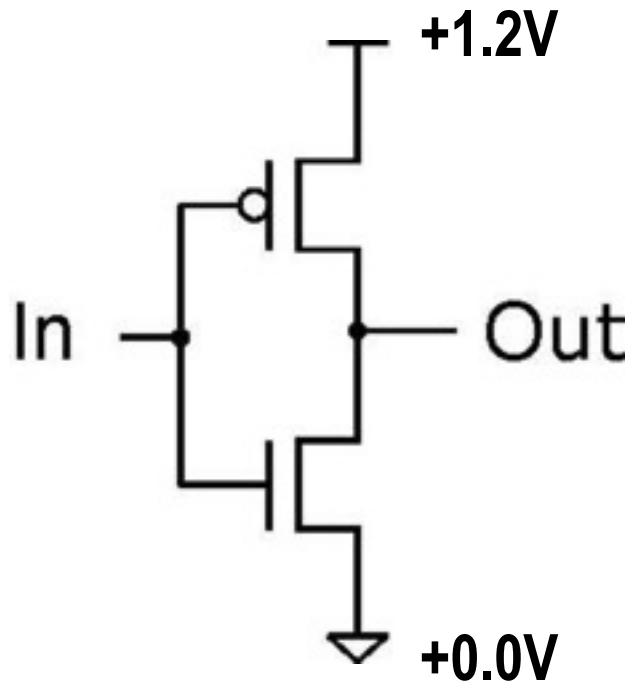
Inverter



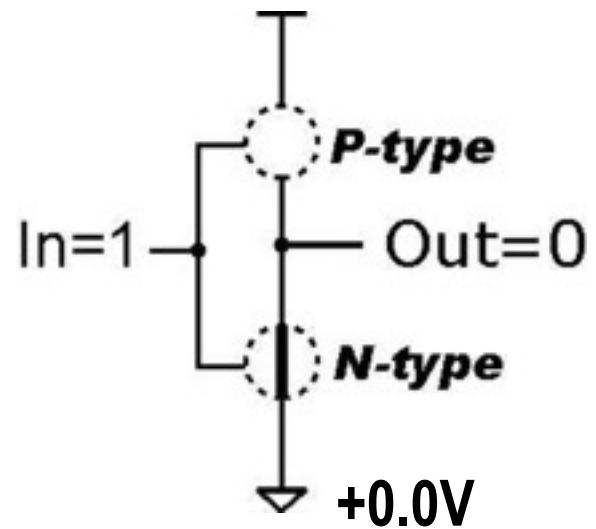
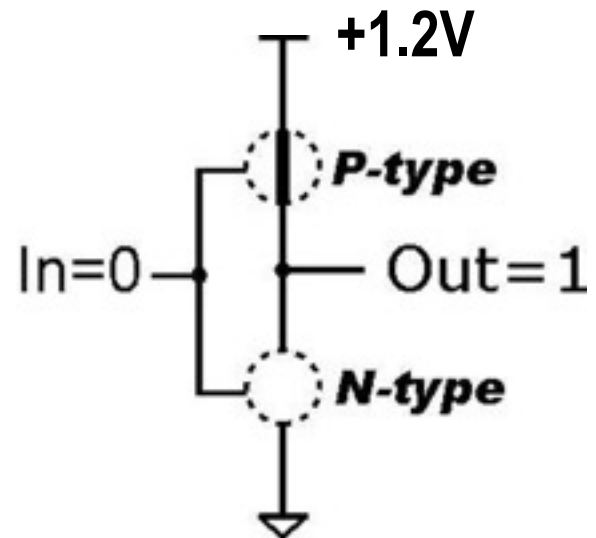
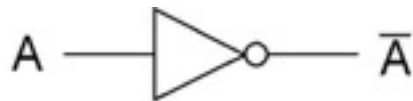
In	Out
0	1
1	0



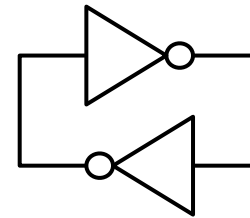
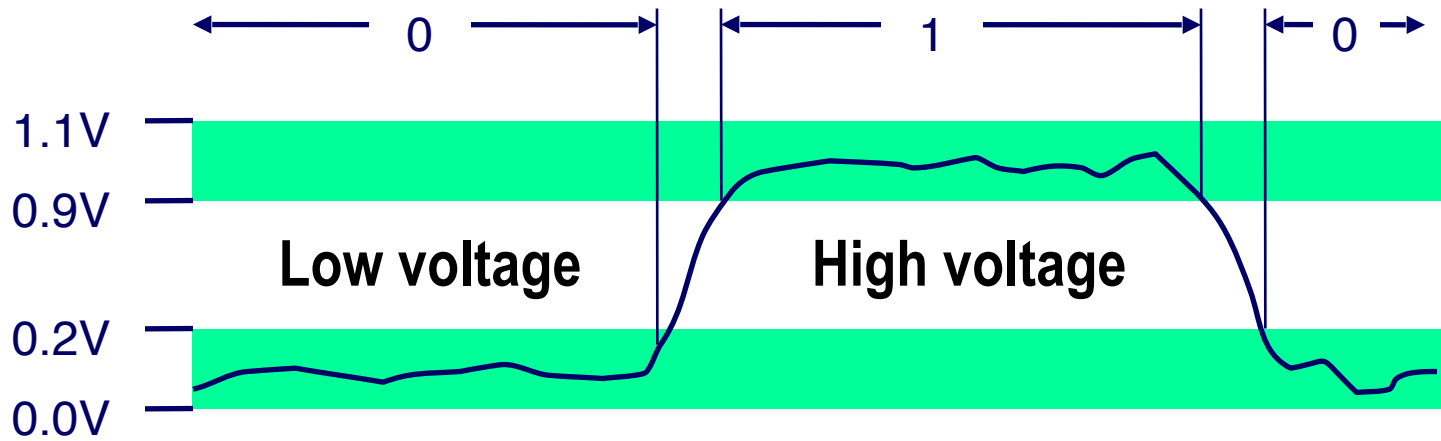
Inverter



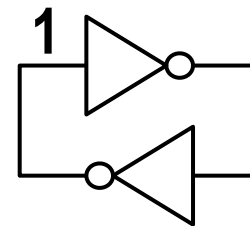
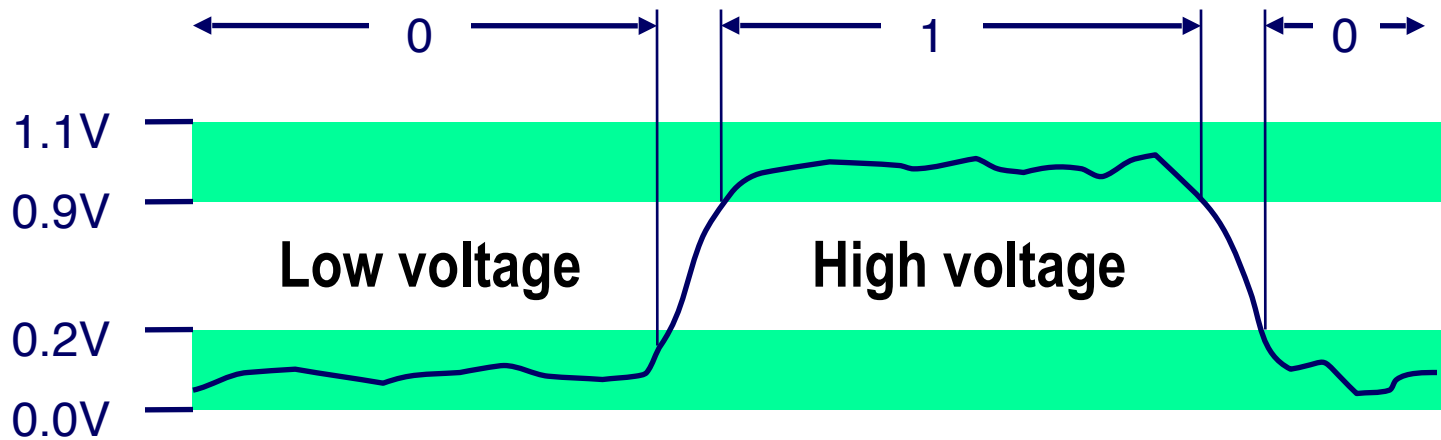
In	Out
0	1
1	0



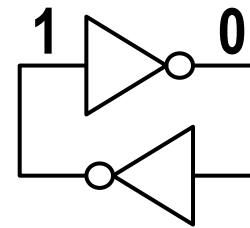
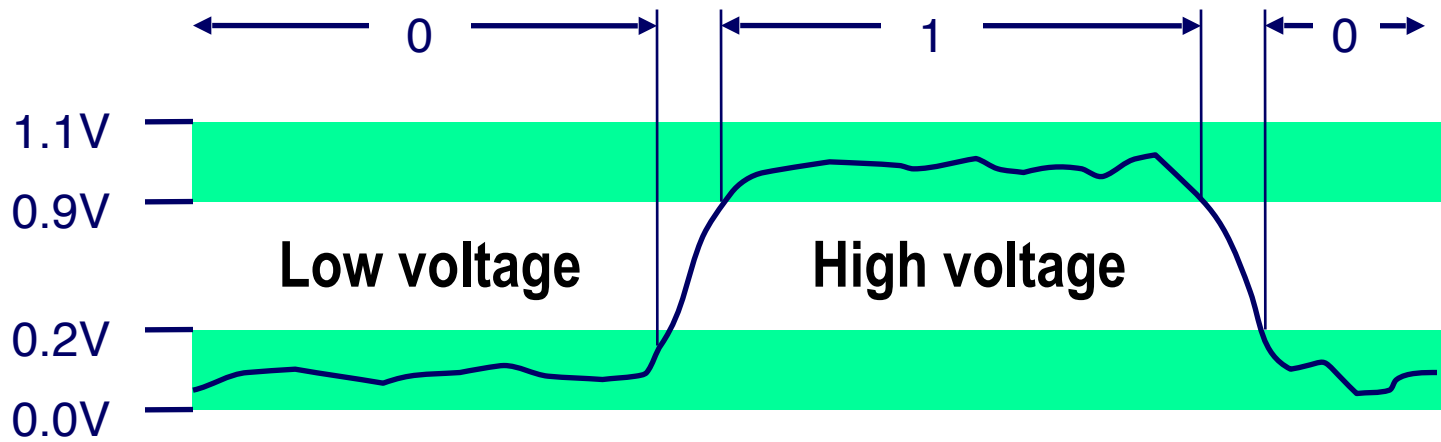
Store/Access Data



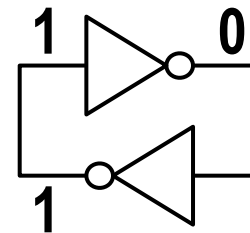
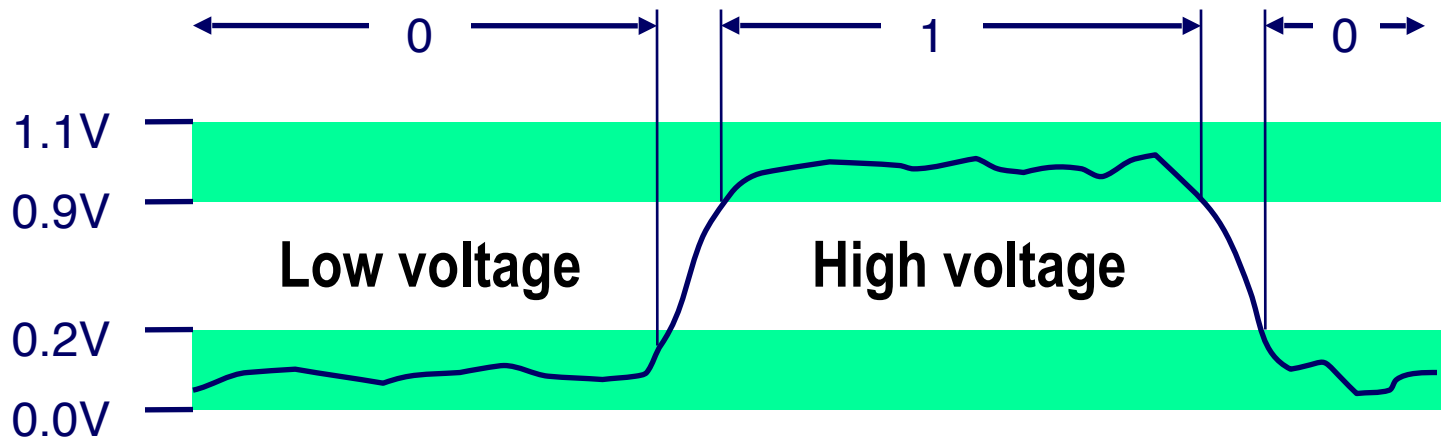
Store/Access Data



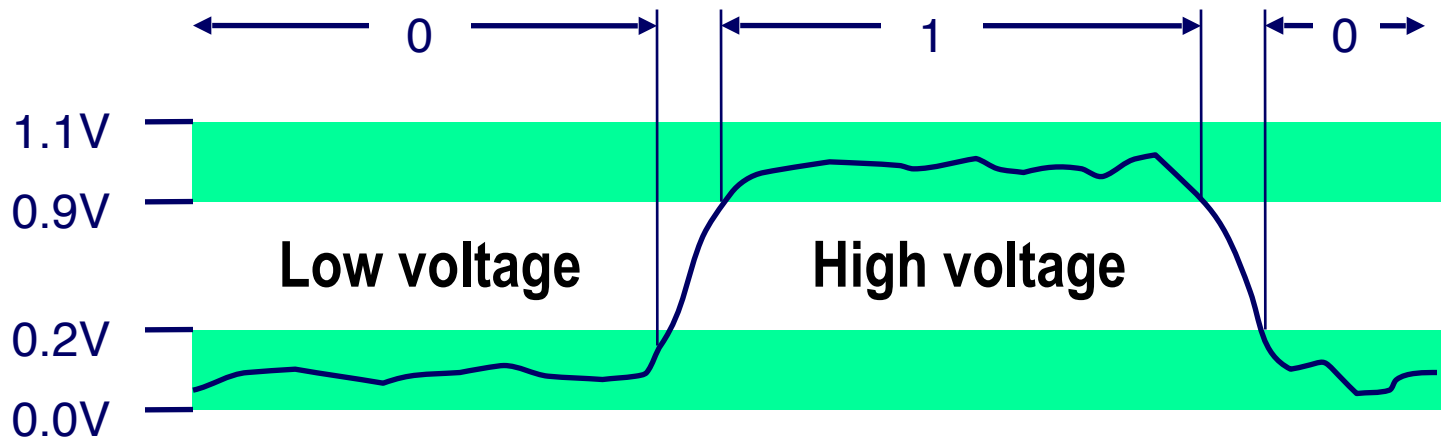
Store/Access Data



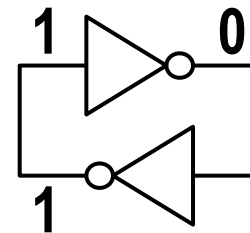
Store/Access Data



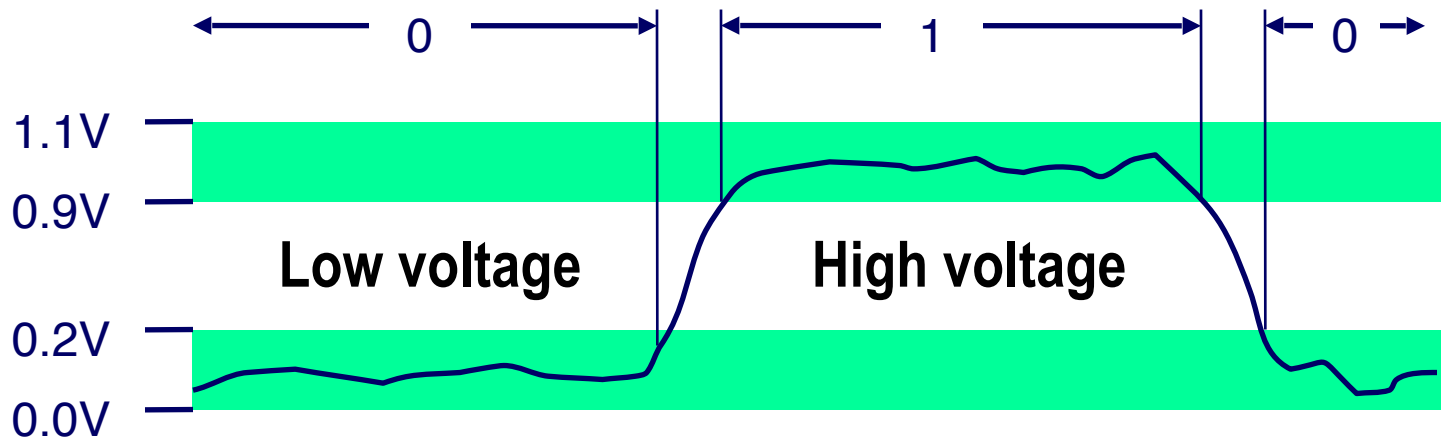
Store/Access Data



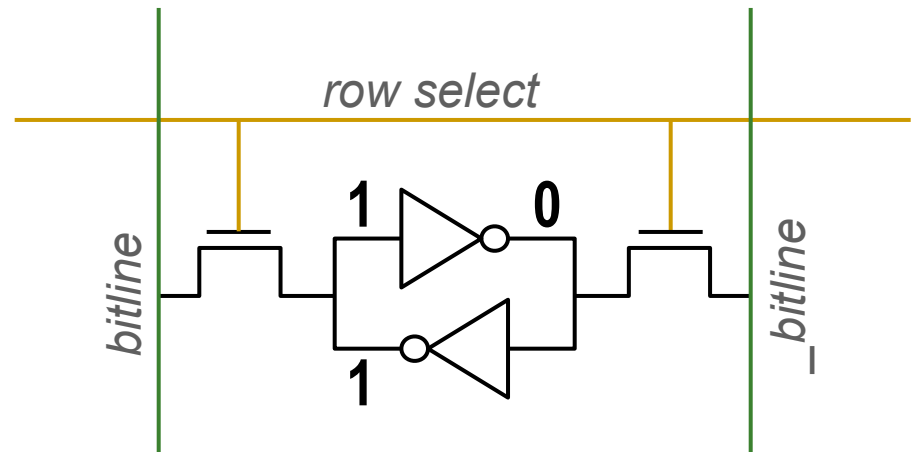
- Two cross coupled inverters store a single bit
 - Feedback path persists the value in the “cell”



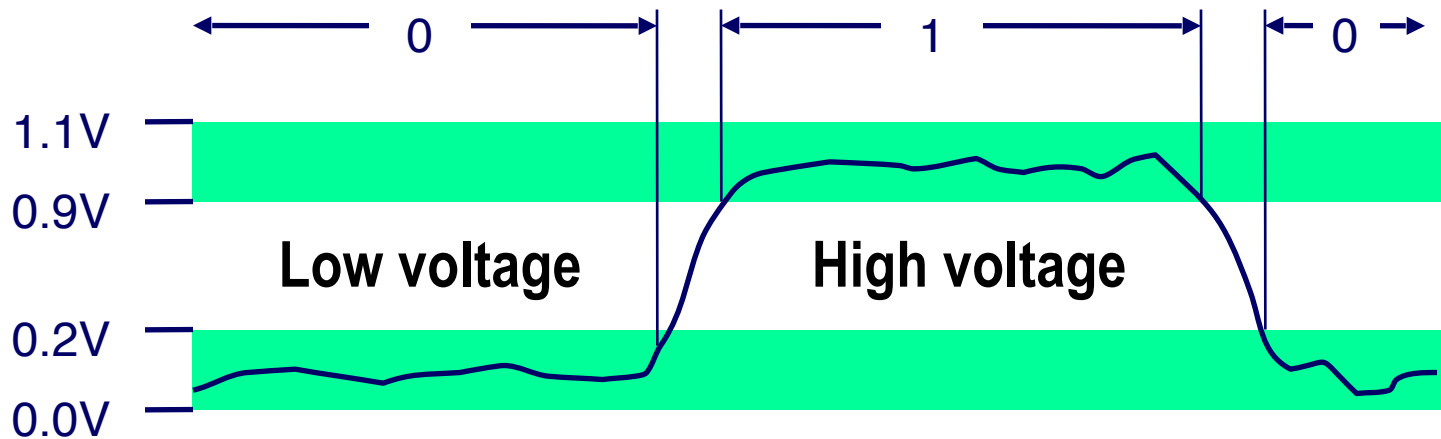
Store/Access Data



- Two cross coupled inverters store a single bit
 - Feedback path persists the value in the “cell”

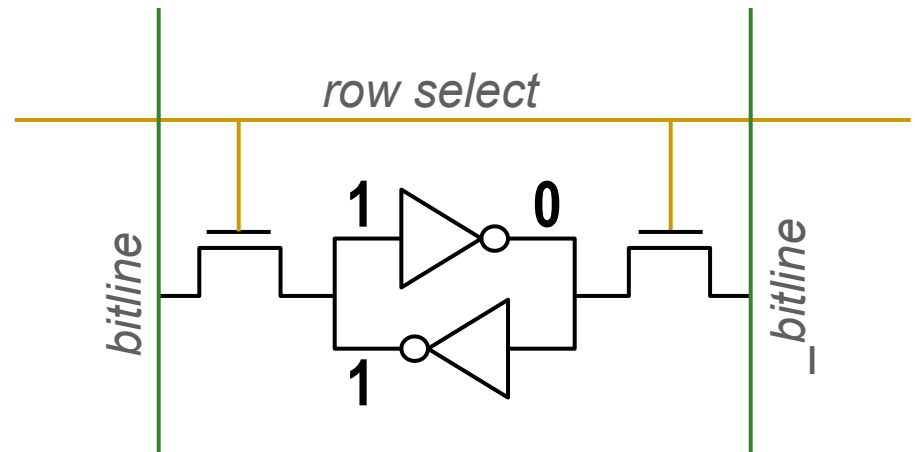


Store/Access Data

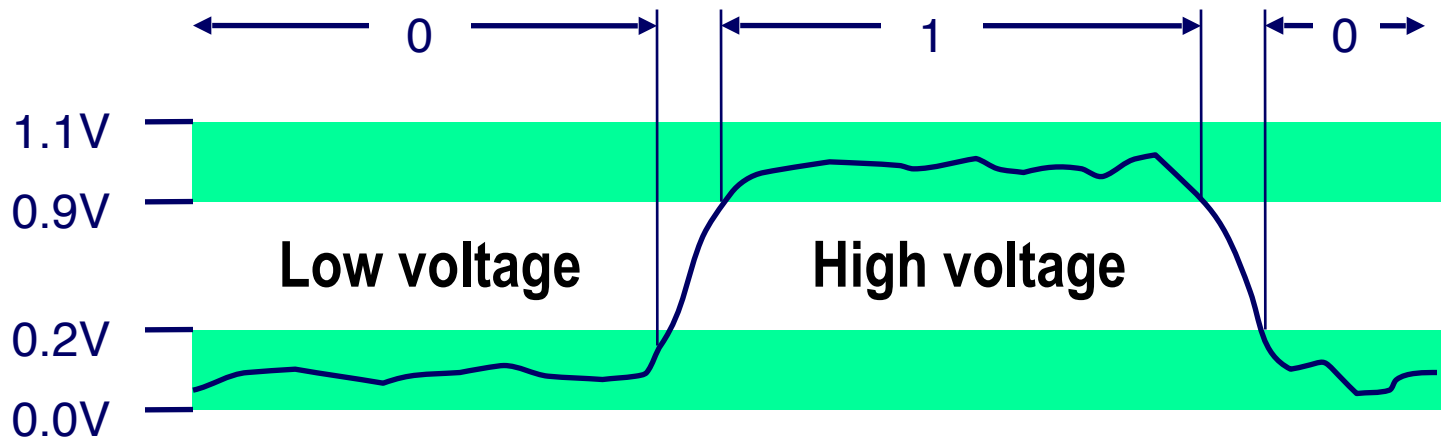


- Two cross coupled inverters store a single bit

- Feedback path persists the value in the “cell”
- 4 transistors for storage

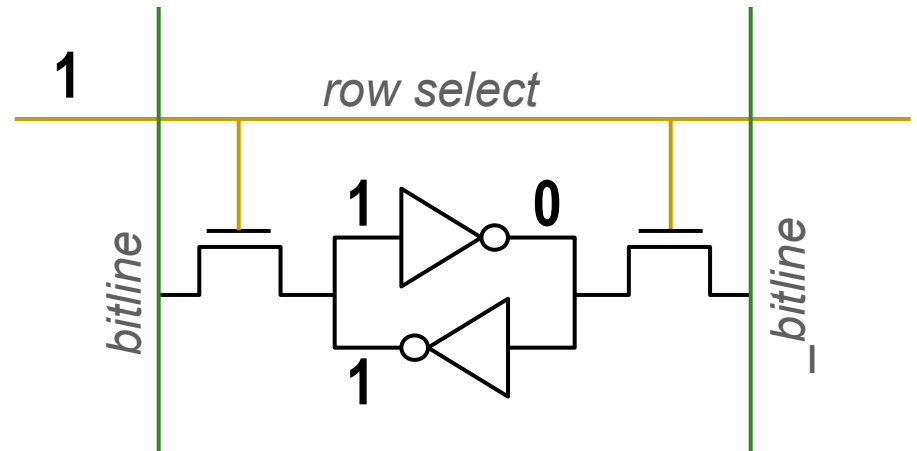


Store/Access Data

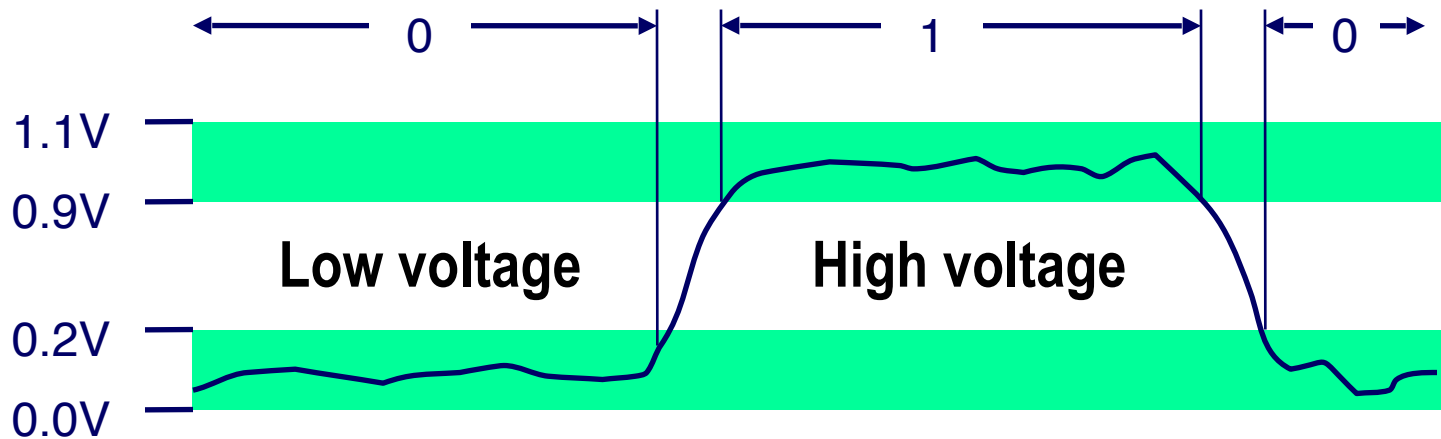


- Two cross coupled inverters store a single bit

- Feedback path persists the value in the “cell”
- 4 transistors for storage

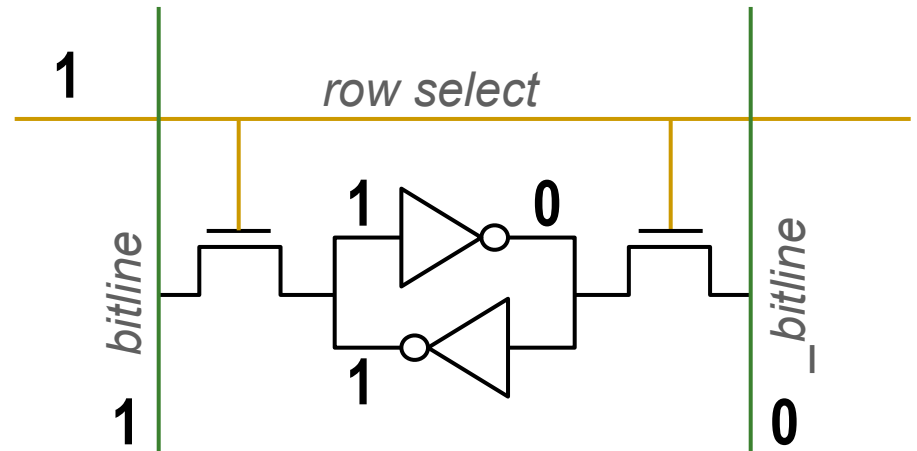


Store/Access Data

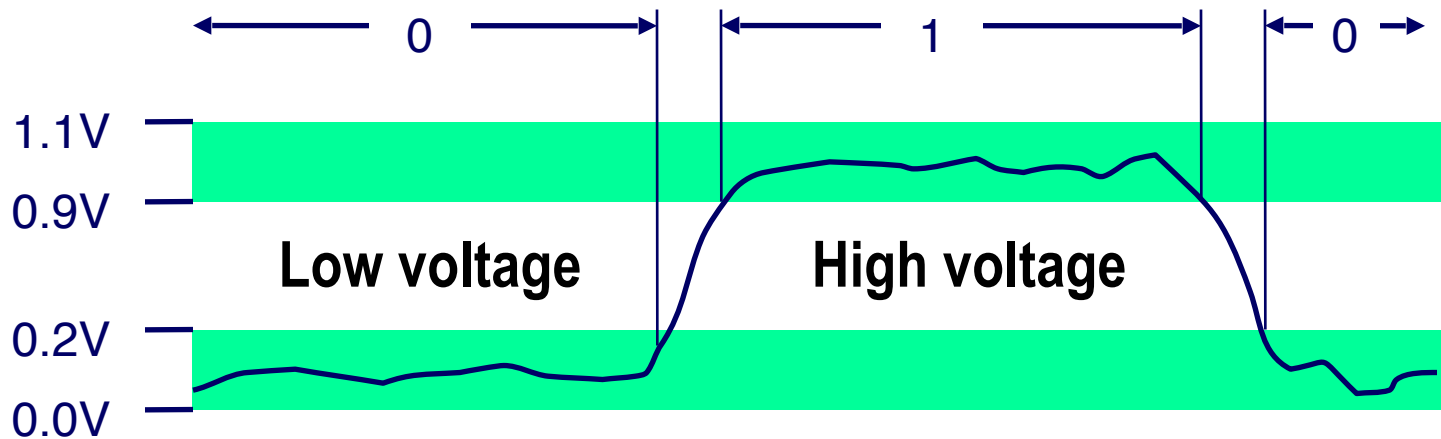


- Two cross coupled inverters store a single bit

- Feedback path persists the value in the “cell”
- 4 transistors for storage

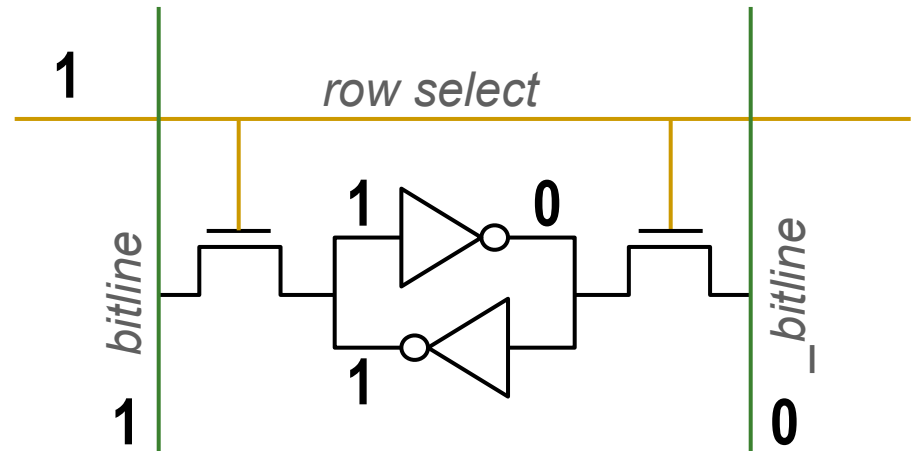


Store/Access Data



- Two cross coupled inverters store a single bit

- Feedback path persists the value in the “cell”
- 4 transistors for storage
- 2 transistors for access



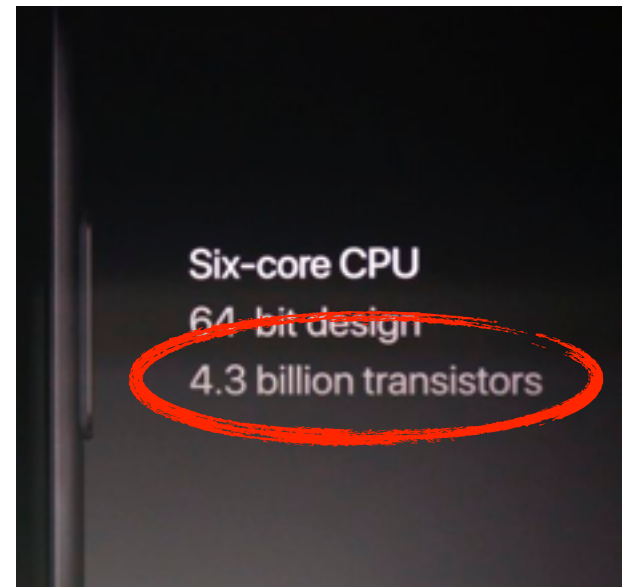
Transistors

- Computers are made of transistors
- Transistors have become smaller over the years
 - Not so much anymore...



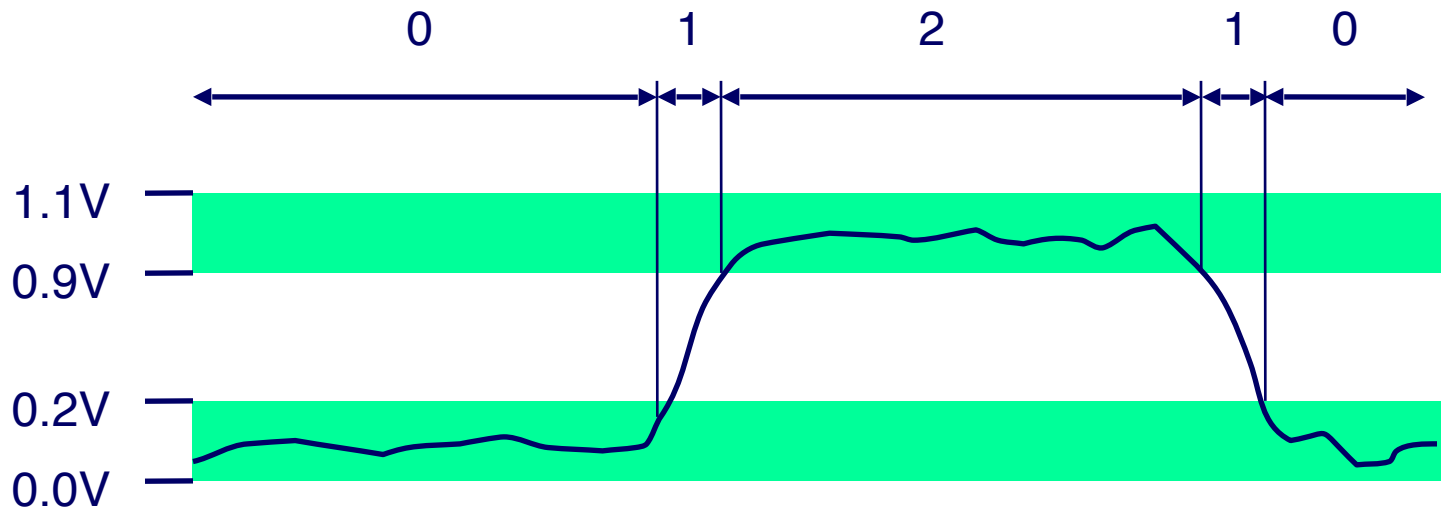
Transistors

- Computers are made of transistors
- Transistors have become smaller over the years
 - Not so much anymore...



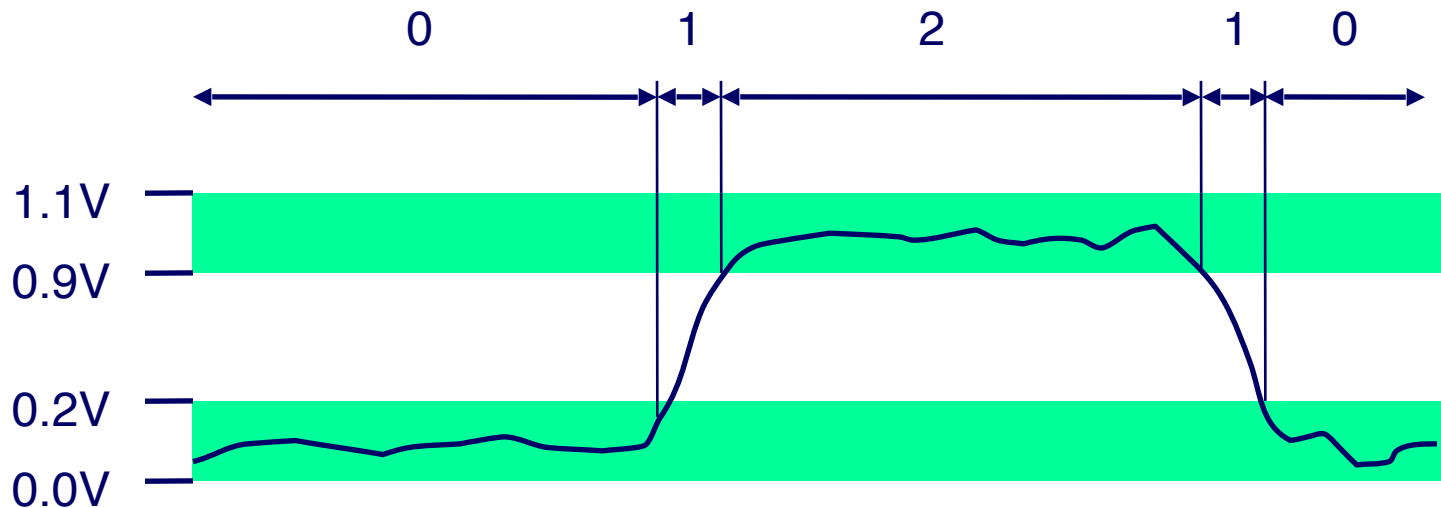
Aside: Why Limit Ourselves Only to Binary?

- Voltage is continuous. Why interpret it only as 0s and 1s?



Aside: Why Limit Ourselves Only to Binary?

- Voltage is continuous. Why interpret it only as 0s and 1s?
- Answer: Noise

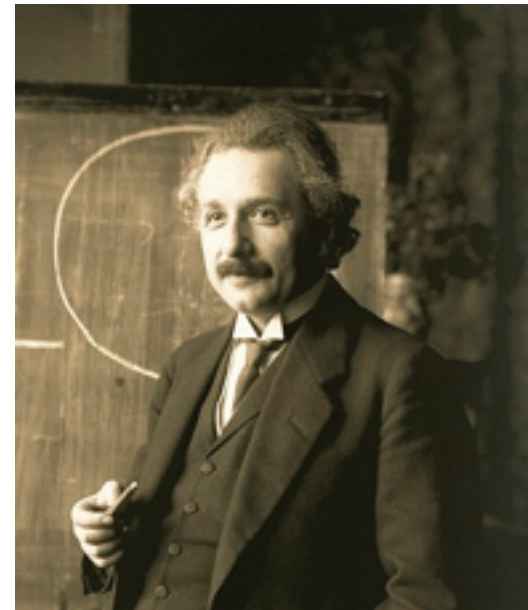
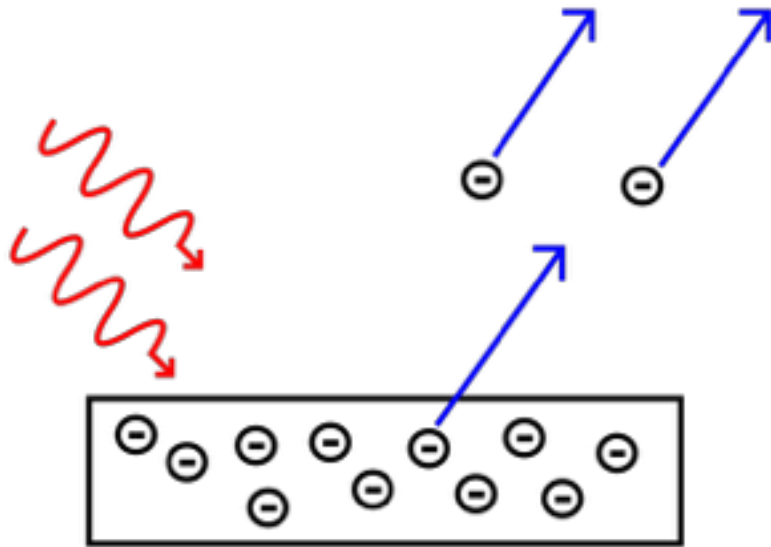


Aside: Why Limit Ourselves Only to Binary?

- But, there are applications that can tolerate noise

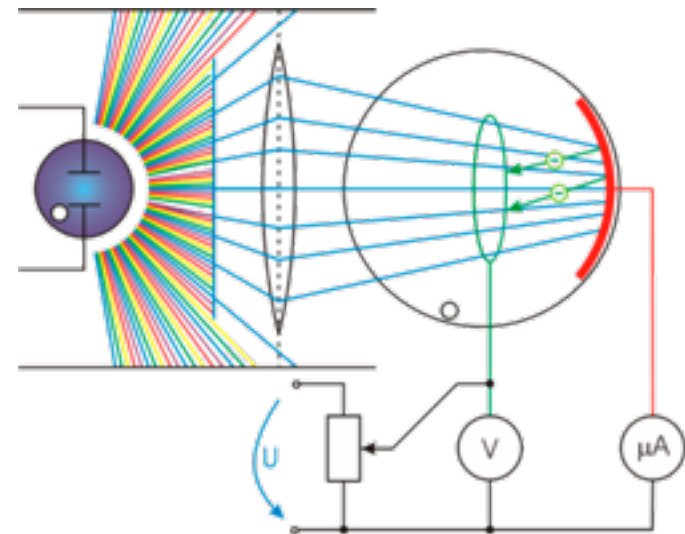
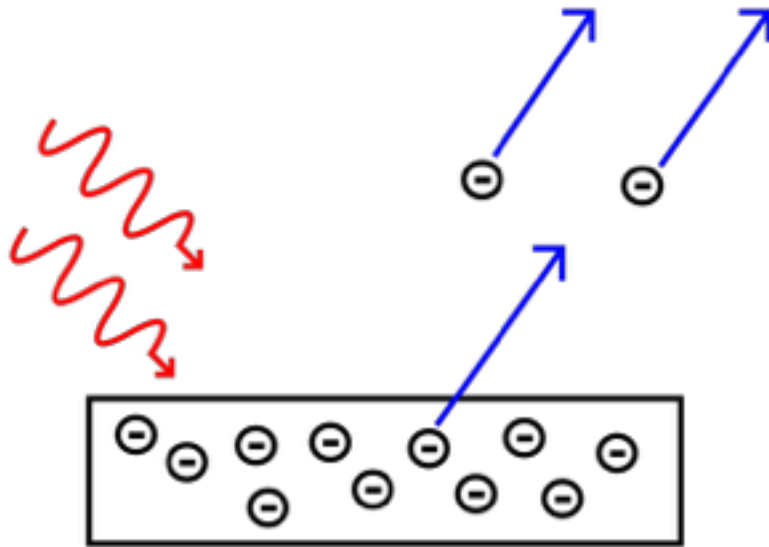
Aside: Why Limit Ourselves Only to Binary?

- But, there are applications that can tolerate noise
- Classic Example: Camera Sensor
 - Photoelectric Effect



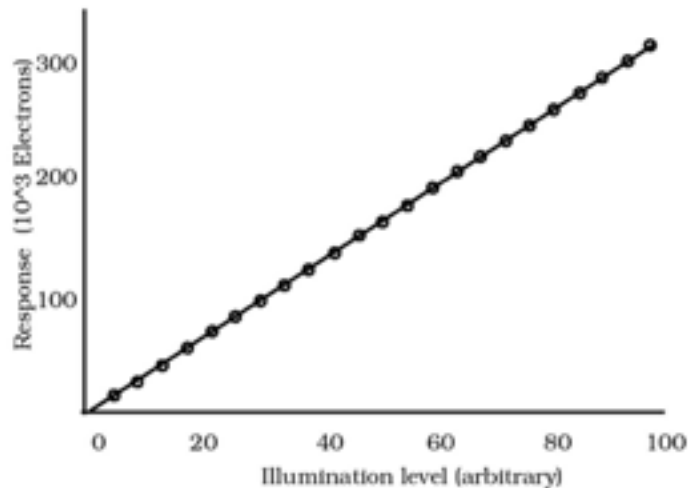
Aside: Why Limit Ourselves Only to Binary?

- But, there are applications that can tolerate noise
- Classic Example: Camera Sensor
 - Photoelectric Effect

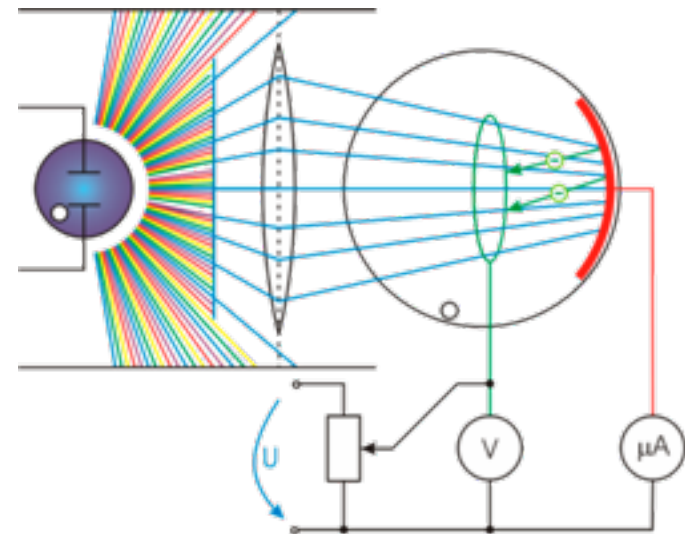


Aside: Why Limit Ourselves Only to Binary?

- But, there are applications that can tolerate noise
- Classic Example: Camera Sensor
 - Photoelectric Effect

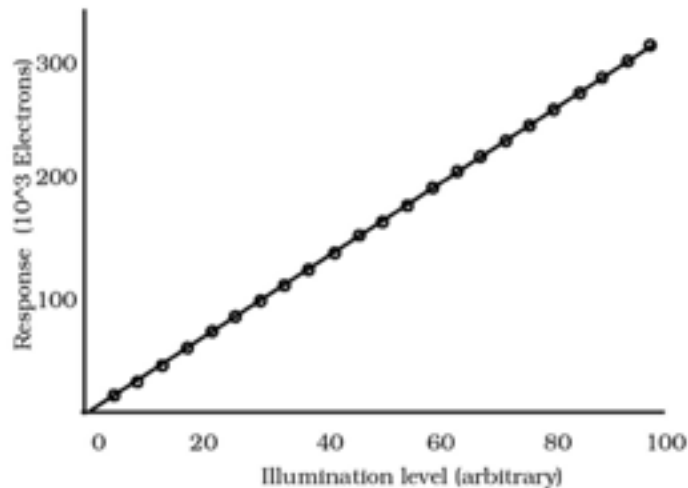


(Epperson, P.M. et al. Electro-optical characterization of the Tektronix TK5 ..., *Opt Eng.*, 25, 1987)



Aside: Why Limit Ourselves Only to Binary?

- But, there are applications that can tolerate noise
- Classic Example: Camera Sensor
 - Photoelectric Effect



(Epperson, P.M. et al. Electro-optical characterization of the Tektronix TK5 ..., Opt Eng., 25, 1987)



Binary Notation

Binary Notation

- Base 2 Number Representation (Binary)

Binary Notation

- Base 2 Number Representation (Binary)
- C.f., Base 10 number representation (Decimal)

Binary Notation

- Base 2 Number Representation (Binary)
- C.f., Base 10 number representation (Decimal)
- $21_{10} = 1*10^0 + 2*10^1 = 21$

Binary Notation

- **Base 2** Number Representation (Binary)
- C.f., Base 10 number representation (Decimal)
- $21_{10} = 1*10^0 + 2*10^1 = 21$
- Weighted Positional Notation
 - Each bit has a weight depending on its position

Binary Notation

- **Base 2** Number Representation (Binary)
- C.f., Base 10 number representation (Decimal)
- $21_{10} = 1*10^0 + 2*10^1 = 21$
- Weighted Positional Notation
 - Each bit has a weight depending on its position
- $1011_2 = 1*2^0 + 1*2^1 + 0*2^2 + 1*2^3 = 11_{10}$

Binary Notation

- Base 2 Number Representation (Binary)
- C.f., Base 10 number representation (Decimal)
- $21_{10} = 1*10^0 + 2*10^1 = 21$
- Weighted Positional Notation
 - Each bit has a weight depending on its position
- $1011_2 = 1*2^0 + 1*2^1 + 0*2^2 + 1*2^3 = 11_{10}$
- $b_3b_2b_1b_0 = b^0*2^0 + b^1*2^1 + b^2*2^2 + b^3*2^3$

Binary Notation

- Base 2 Number Representation (Binary)
- C.f., Base 10 number representation (Decimal)
- $21_{10} = 1*10^0 + 2*10^1 = 21$
- Weighted Positional Notation
 - Each bit has a weight depending on its position
- $1011_2 = 1*2^0 + 1*2^1 + 0*2^2 + 1*2^3 = 11_{10}$
- $b_3b_2b_1b_0 = b^0*2^0 + b^1*2^1 + b^2*2^2 + b^3*2^3$
- Binary Arithmetic

Binary Notation

- **Base 2** Number Representation (Binary)
- C.f., Base 10 number representation (Decimal)
- $21_{10} = 1*10^0 + 2*10^1 = 21$
- Weighted Positional Notation
 - Each bit has a weight depending on its position
- $1011_2 = 1*2^0 + 1*2^1 + 0*2^2 + 1*2^3 = 11_{10}$
- $b_3b_2b_1b_0 = b^0*2^0 + b^1*2^1 + b^2*2^2 + b^3*2^3$
- Binary Arithmetic

Decimal	Binary
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111
8	1000
9	1001
10	1010
11	1011
12	1100
13	1101
14	1110
15	1111

Binary Notation

- **Base 2** Number Representation (Binary)
- C.f., Base 10 number representation (Decimal)
- $21_{10} = 1*10^0 + 2*10^1 = 21$
- Weighted Positional Notation
 - Each bit has a weight depending on its position
- $1011_2 = 1*2^0 + 1*2^1 + 0*2^2 + 1*2^3 = 11_{10}$
- $b_3b_2b_1b_0 = b^0*2^0 + b^1*2^1 + b^2*2^2 + b^3*2^3$
- Binary Arithmetic

$$\begin{array}{r} 0110 \\ + 0101 \\ \hline 1011 \end{array}$$

Decimal	Binary
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111
8	1000
9	1001
10	1010
11	1011
12	1100
13	1101
14	1110
15	1111

Binary Notation

- **Base 2** Number Representation (Binary)
- C.f., Base 10 number representation (Decimal)
- $21_{10} = 1*10^0 + 2*10^1 = 21$
- Weighted Positional Notation
 - Each bit has a weight depending on its position
- $1011_2 = 1*2^0 + 1*2^1 + 0*2^2 + 1*2^3 = 11_{10}$
- $b_3b_2b_1b_0 = b^0*2^0 + b^1*2^1 + b^2*2^2 + b^3*2^3$
- Binary Arithmetic

$$\begin{array}{r} 0110 \\ + 0101 \\ \hline 1011 \end{array}$$

$$\begin{array}{r} 6 \\ + 5 \\ \hline 11 \end{array}$$

Decimal	Binary
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111
8	1000
9	1001
10	1010
11	1011
12	1100
13	1101
14	1110
15	1111

Hexadecimal (Hex) Notation

- **Base 16** Number Representation
 - Use characters '0' to '9' and 'A' to 'F'
 - Four bits per Hex digit
 - $11111110_2 = FE_{16}$
- Write $FA1D37B_{16}$ in C as
 - `0xFA1D37B`
 - `0xfa1d37b`

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

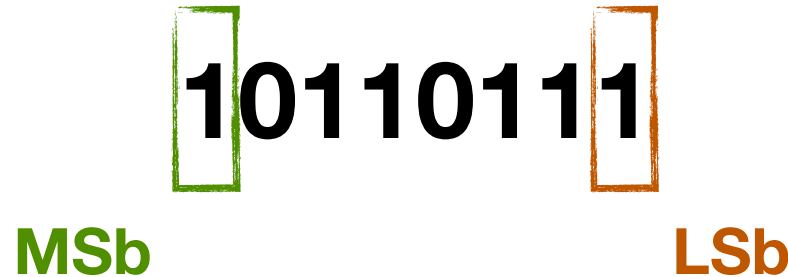
Bit, Byte, Word

- Byte = 8 bits

- Binary 00000000_2 to 11111111_2 ; Decimal: 0_{10} to 255_{10} ; Hex: 00_{16} to FF_{16}
- Least Significant Bit (LSb) vs. Most Significant Bit (MSb)

10110111

MSb **LSb**

The diagram shows the binary sequence '10110111'. The first bit '1' is enclosed in a green rectangular box, and the last bit '1' is enclosed in an orange rectangular box. Below the green box is the label 'MSb' in green text, and below the orange box is the label 'LSb' in orange text.

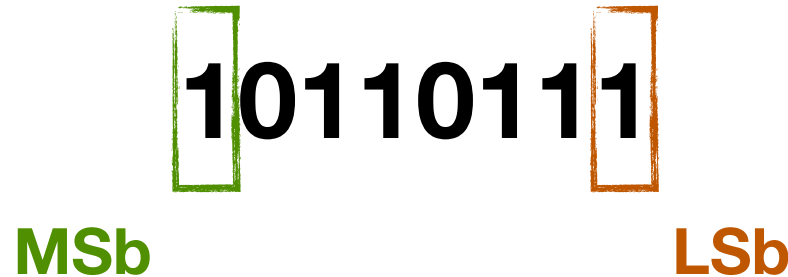
Bit, Byte, Word

- Byte = 8 bits

- Binary 00000000_2 to 11111111_2 ; Decimal: 0_{10} to 255_{10} ; Hex: 00_{16} to FF_{16}
- Least Significant Bit (LSb) vs. Most Significant Bit (MSb)

10110111

MSb **LSb**

A diagram showing the binary sequence '10110111'. The first bit '1' is enclosed in a green rectangular box, and the last bit '1' is enclosed in an orange rectangular box. Below the green box is the label 'MSb' in green text, and below the orange box is the label 'LSb' in orange text.

- Word = 4 Bytes (32-bit machine) / 8 Bytes (64-bit machine)
- Least Significant Byte (LSB) vs. Most Significant Byte (MSB)

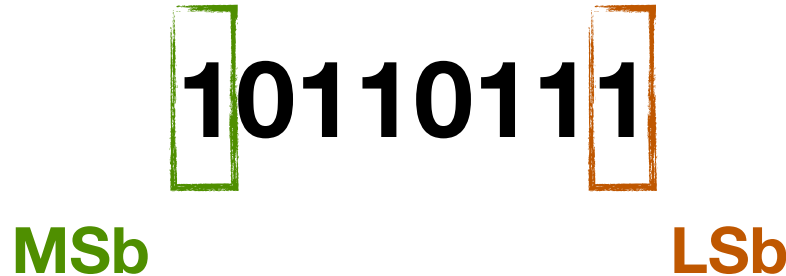
Bit, Byte, Word

- Byte = 8 bits

- Binary 00000000_2 to 11111111_2 ; Decimal: 0_{10} to 255_{10} ; Hex: 00_{16} to FF_{16}
- Least Significant Bit (LSb) vs. Most Significant Bit (MSb)

10110111

MSb **LSb**

A diagram showing the binary sequence '10110111'. The first bit '1' is enclosed in a green rectangular box, and the last bit '1' is enclosed in an orange rectangular box. Below the green box is the label 'MSb' in green text, and below the orange box is the label 'LSb' in orange text.

- Word = 4 Bytes (32-bit machine) / 8 Bytes (64-bit machine)
- Least Significant Byte (LSB) vs. Most Significant Byte (MSB)

Questions?

Today: Representing Information in Binary

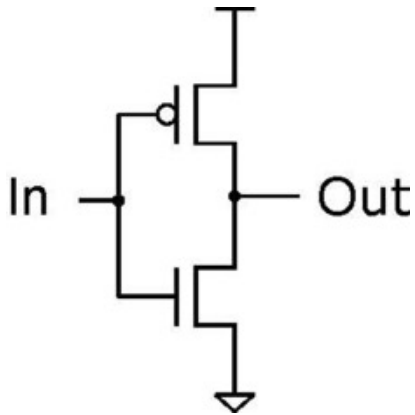
- Why Binary (bits)?
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary
- Representations in memory, pointers, strings

Bit-level manipulations

Not

- $\sim A = 1$ when $A=0$

\sim	
0	1
1	0



Bit-level manipulations

Not

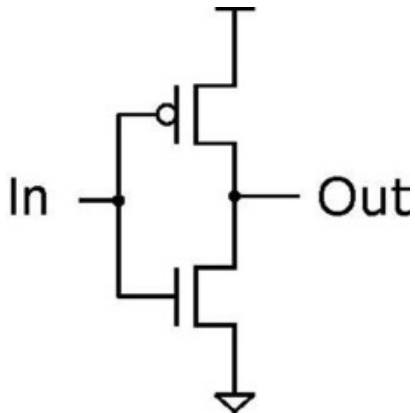
- $\sim A = 1$ when $A=0$

\sim	
0	1
1	0

Or

- $A|B = 1$ when either $A=1$ or $B=1$

	0	1
0	0	1
1	1	1



Bit-level manipulations

Not

- $\sim A = 1$ when $A=0$

\sim	
0	1
1	0

Or

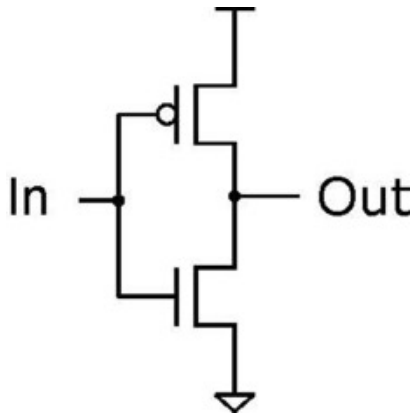
- $A|B = 1$ when either $A=1$ or $B=1$

And

- $A\&B = 1$ when both $A=1$ and $B=1$

$ $	0	1
0	0	1
1	1	1

$\&$	0	1
0	0	0
1	0	1

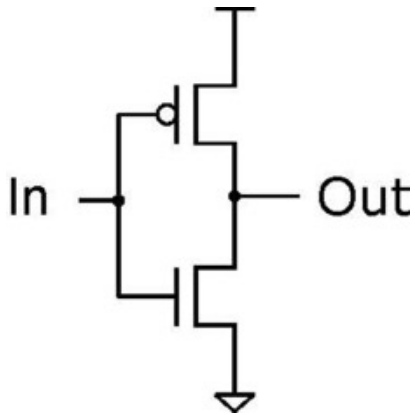


Bit-level manipulations

Not

- $\sim A = 1$ when $A=0$

\sim	
0	1
1	0



Or

- $A|B = 1$ when either $A=1$ or $B=1$

And

- $A\&B = 1$ when both $A=1$ and $B=1$

Exclusive-Or (Xor)

- $A\wedge B = 1$ when either $A=1$ or $B=1$, but not both

$ $	0	1
0	0	1
1	1	1

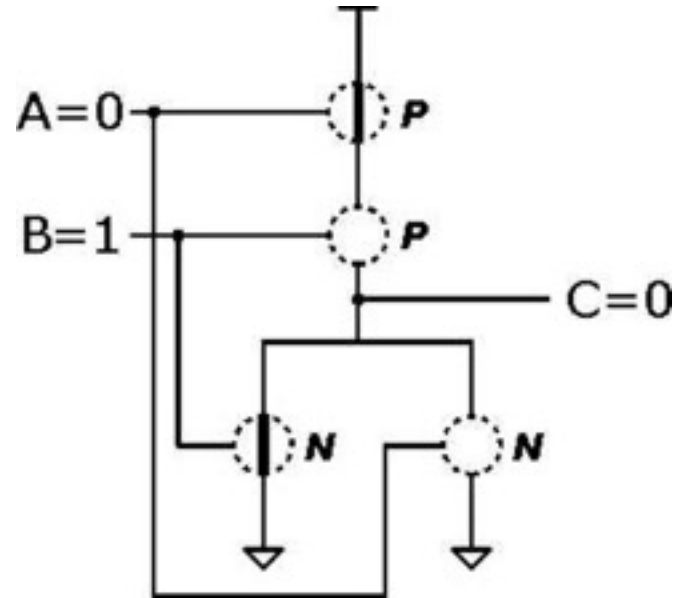
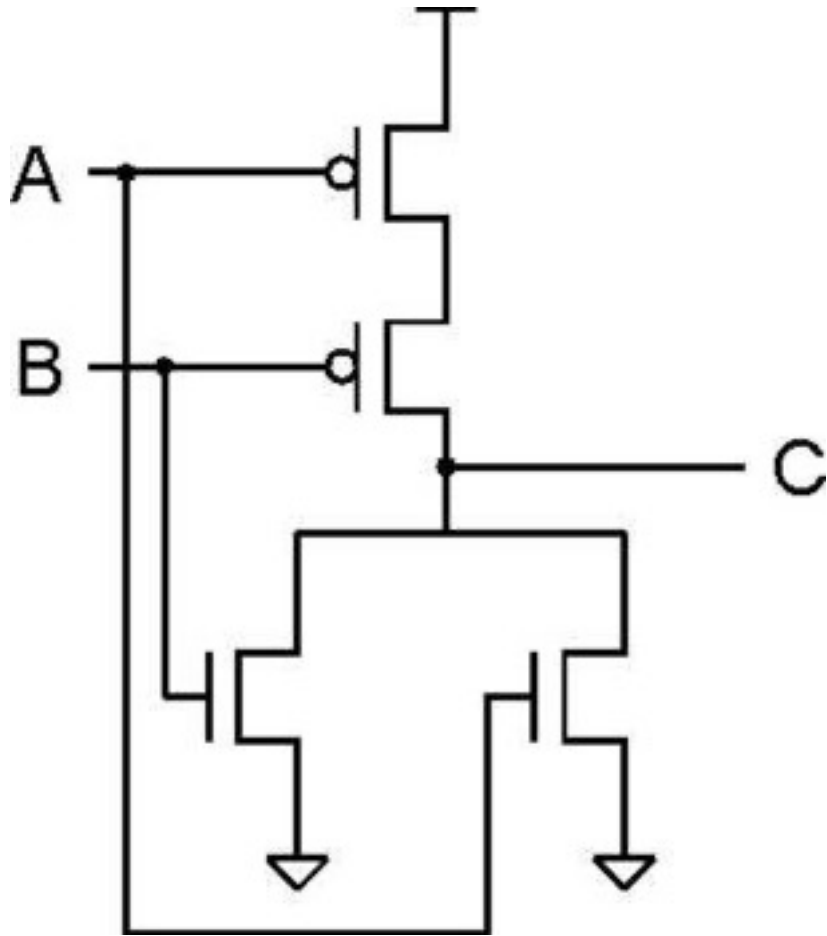
$\&$	0	1
0	0	0
1	0	1

\wedge	0	1
0	0	1
1	1	0

NOR (OR + NOT)

A	B	C
0	0	1
0	1	0
1	0	0
1	1	0

NOR (OR + NOT)



A	B	C
0	0	1
0	1	0
1	0	0
1	1	0

Bit Vector Operations

- Operate on Bit Vectors
 - Operations applied bitwise

01101001	01101001	01101001	
& 01010101	01010101	^ 01010101	~ 01010101
<u> </u>	<u> </u>	<u> </u>	<u> </u>

Bit Vector Operations

- Operate on Bit Vectors
 - Operations applied bitwise

01101001	01101001	01101001	
& 01010101	01010101	^ 01010101	~ 01010101
<u> </u>	<u> </u>	<u> </u>	<u> </u>
01000001			

Bit Vector Operations

- Operate on Bit Vectors
 - Operations applied bitwise

$\begin{array}{r} 01101001 \\ \& 01010101 \\ \hline 01000001 \end{array}$	$\begin{array}{r} 01101001 \\ 01010101 \\ \hline 01111101 \end{array}$	$\begin{array}{r} 01101001 \\ \wedge 01010101 \\ \hline \end{array}$	$\begin{array}{r} \sim 01010101 \\ \hline \end{array}$
---	--	--	--

Bit Vector Operations

- Operate on Bit Vectors
 - Operations applied bitwise

$\begin{array}{r} 01101001 \\ \& 01010101 \\ \hline 01000001 \end{array}$	$\begin{array}{r} 01101001 \\ 01010101 \\ \hline 01111101 \end{array}$	$\begin{array}{r} 01101001 \\ \wedge 01010101 \\ \hline 00111100 \end{array}$	$\begin{array}{r} \sim 01010101 \\ \hline \end{array}$
---	--	---	--

Bit Vector Operations

- Operate on Bit Vectors
 - Operations applied bitwise

$\begin{array}{r} 01101001 \\ \& 01010101 \\ \hline 01000001 \end{array}$	$\begin{array}{r} 01101001 \\ 01010101 \\ \hline 01111101 \end{array}$	$\begin{array}{r} 01101001 \\ \wedge 01010101 \\ \hline 00111100 \end{array}$	$\begin{array}{r} \sim 01010101 \\ \hline 10101010 \end{array}$
---	--	---	---

Bit-Level Operations in C

- Operations $\&$, $|$, \sim , \wedge Available in C
 - Apply to any “integral” data type
 - long, int, short, char, unsigned
 - View arguments as bit vectors
 - Arguments applied bit-wise
- Examples (Char data type)
 - $\sim 0x41 \rightarrow 0xBE$
 - $\sim 01000001_2 \rightarrow 10111110_2$
 - $\sim 0x00 \rightarrow 0xFF$
 - $\sim 00000000_2 \rightarrow 11111111_2$
 - $0x69 \& 0x55 \rightarrow 0x41$
 - $01101001_2 \& 01010101_2 \rightarrow 01000001_2$
 - $0x69 | 0x55 \rightarrow 0x7D$
 - $01101001_2 | 01010101_2 \rightarrow 01111101_2$

Contrast: Logic Operations in C

- Contrast to Logical Operators
 - `&&`, `||`, `!`
 - View 0 as “False”
 - Anything nonzero as “True”
 - Always return 0 or 1
 - Early termination (e.g., `0 && 1 && 1`)
- Examples (char data type)
 - `!0x41` → `0x00`
 - `!0x00` → `0x01`
 - `!!0x41` → `0x01`
 - `0x69 && 0x55` → `0x01`
 - `0x69 || 0x55` → `0x01`
 - `p && *p` (avoids null pointer access)

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
<< 3	
Log. >> 2	
Arith. >> 2	

Argument x	10100010
<< 3	
Log. >> 2	
Arith. >> 2	

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
$\ll 3$	00010
Log. $\gg 2$	
Arith. $\gg 2$	

Argument x	10100010
$\ll 3$	
Log. $\gg 2$	
Arith. $\gg 2$	

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
<< 3	00010000
Log. >> 2	
Arith. >> 2	

Argument x	10100010
<< 3	
Log. >> 2	
Arith. >> 2	

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
<< 3	00010000
Log. >> 2	011000
Arith. >> 2	

Argument x	10100010
<< 3	
Log. >> 2	
Arith. >> 2	

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
<< 3	00010000
Log. >> 2	00011000
Arith. >> 2	

Argument x	10100010
<< 3	
Log. >> 2	
Arith. >> 2	

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
<< 3	00010000
Log. >> 2	00011000
Arith. >> 2	011000

Argument x	10100010
<< 3	
Log. >> 2	
Arith. >> 2	

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
<< 3	00010000
Log. >> 2	00011000
Arith. >> 2	00011000

Argument x	10100010
<< 3	
Log. >> 2	
Arith. >> 2	

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
<< 3	00010000
Log. >> 2	00011000
Arith. >> 2	00011000

Argument x	10100010
<< 3	00010
Log. >> 2	
Arith. >> 2	

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
<< 3	00010000
Log. >> 2	00011000
Arith. >> 2	00011000

Argument x	10100010
<< 3	00010000
Log. >> 2	
Arith. >> 2	

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
<< 3	00010000
Log. >> 2	00011000
Arith. >> 2	00011000

Argument x	10100010
<< 3	00010000
Log. >> 2	101000
Arith. >> 2	

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
<< 3	00010000
Log. >> 2	00011000
Arith. >> 2	00011000

Argument x	10100010
<< 3	00010000
Log. >> 2	00101000
Arith. >> 2	

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
<< 3	00010000
Log. >> 2	00011000
Arith. >> 2	00011000

Argument x	10100010
<< 3	00010000
Log. >> 2	00101000
Arith. >> 2	101000

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
<< 3	00010000
Log. >> 2	00011000
Arith. >> 2	00011000

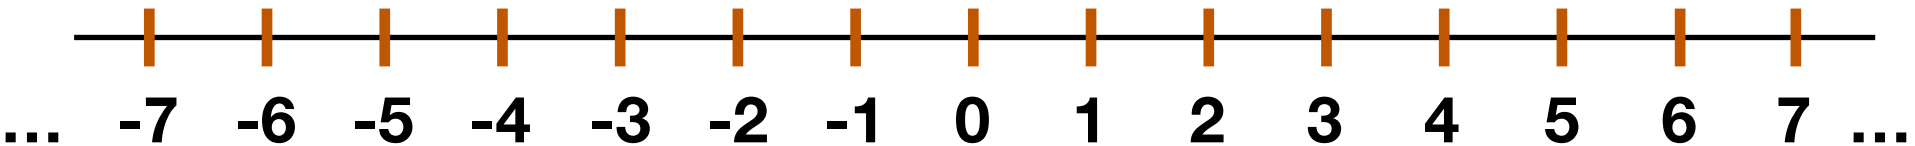
Argument x	10100010
<< 3	00010000
Log. >> 2	00101000
Arith. >> 2	11101000

Today: Representing Information in Binary

- Why Binary (bits)?
- Bit-level manipulations
- **Integers**
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary
- Representations in memory, pointers, strings

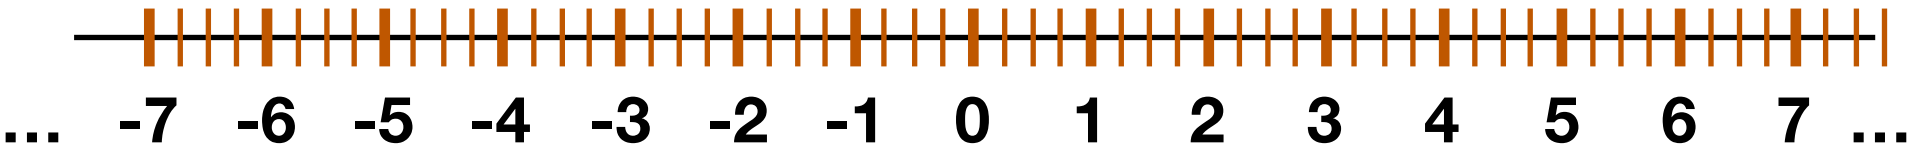
Representing Numbers in Binary

- What numbers do we need to represent in bits?
 - Integer (Negative and Non-negative)
 - Fractions
 - Irrationals



Representing Numbers in Binary

- What numbers do we need to represent in bits?
 - Integer (Negative and Non-negative)
 - Fractions
 - Irrationals



Encoding Negative Numbers

Encoding Negative Numbers

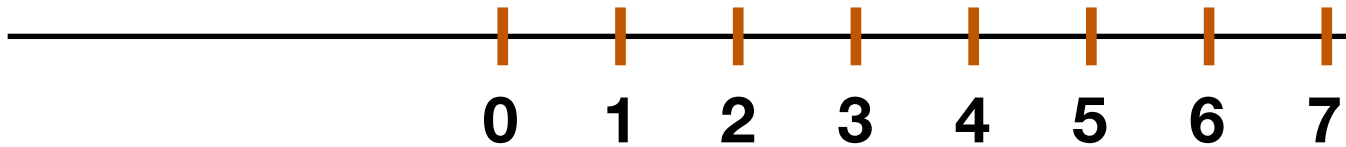
- So far we have been discussing non-negative numbers: so called **unsigned**. How about negative numbers?

Encoding Negative Numbers

- So far we have been discussing non-negative numbers: so called **unsigned**. How about negative numbers?
- Solution 1: Sign-magnitude
 - First bit represents sign; 0 for positive; 1 for negative
 - The rest represents magnitude

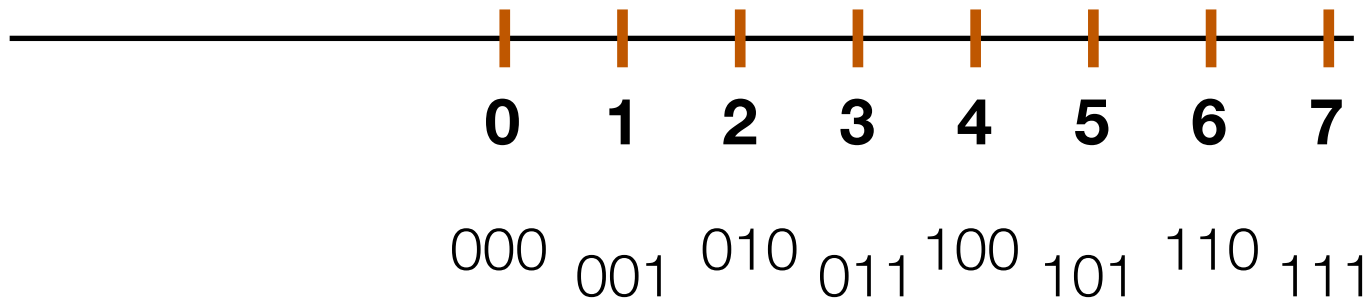
Encoding Negative Numbers

- So far we have been discussing non-negative numbers: so called **unsigned**. How about negative numbers?
- Solution 1: Sign-magnitude
 - First bit represents sign; 0 for positive; 1 for negative
 - The rest represents magnitude



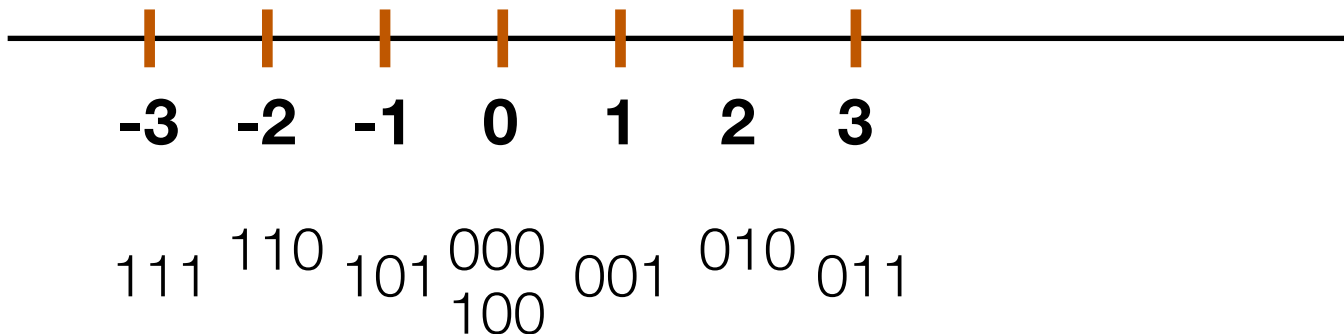
Encoding Negative Numbers

- So far we have been discussing non-negative numbers: so called **unsigned**. How about negative numbers?
- Solution 1: Sign-magnitude
 - First bit represents sign; 0 for positive; 1 for negative
 - The rest represents magnitude



Encoding Negative Numbers

- So far we have been discussing non-negative numbers: so called **unsigned**. How about negative numbers?
- Solution 1: Sign-magnitude
 - First bit represents sign; 0 for positive; 1 for negative
 - The rest represents magnitude



Sign-Magnitude Implications

- Bits have different semantics
 - Two zeros...
 - Normal arithmetic doesn't work
 - Make hardware design harder

Signed Value	Binary
0	000
1	001
2	010
3	011
-0	100
-1	101
-2	110
-3	111

Sign-Magnitude Implications

- Bits have different semantics
 - Two zeros...
 - Normal arithmetic doesn't work
 - Make hardware design harder

$$\begin{array}{r} 010 \\ +) 101 \\ \hline 111 \end{array}$$

Signed Value	Binary
0	000
1	001
2	010
3	011
-0	100
-1	101
-2	110
-3	111

Sign-Magnitude Implications

- Bits have different semantics
 - Two zeros...
 - Normal arithmetic doesn't work
 - Make hardware design harder

$$\begin{array}{r} 010 \\ +) 101 \\ \hline 111 \end{array}$$

$$\begin{array}{r} 2 \\ +) -1 \\ \hline -3 \end{array}$$

Signed Value	Binary
0	000
1	001
2	010
3	011
-0	100
-1	101
-2	110
-3	111

Sign-Magnitude Implications

- Bits have different semantics
 - Two zeros...
 - Normal arithmetic doesn't work
 - Make hardware design harder

The diagram shows two arithmetic operations, each crossed out with a large red 'X'.
Left operation:
$$\begin{array}{r} 010 \\ +) 101 \\ \hline 111 \end{array}$$

Right operation:
$$\begin{array}{r} 2 \\ +) -1 \\ \hline -3 \end{array}$$

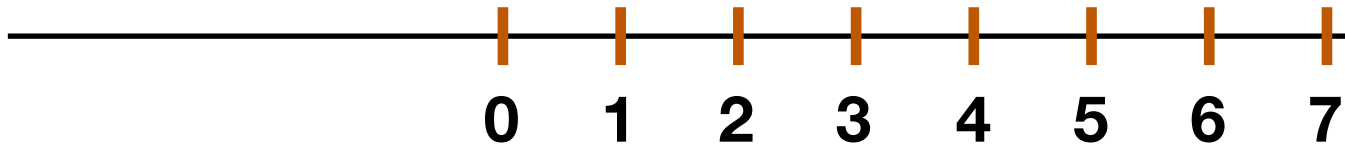
Signed Value	Binary
0	000
1	001
2	010
3	011
-0	100
-1	101
-2	110
-3	111

Encoding Negative Numbers

- Solution 2: Two's Complement

Encoding Negative Numbers

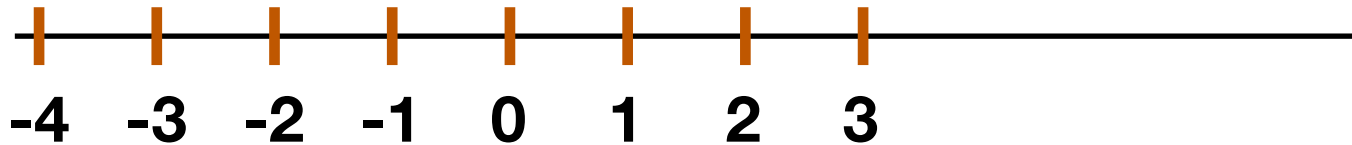
- Solution 2: Two's Complement



Unsigned	Binary
0	000
1	001
2	010
3	011
4	100
5	101
6	110
7	111

Encoding Negative Numbers

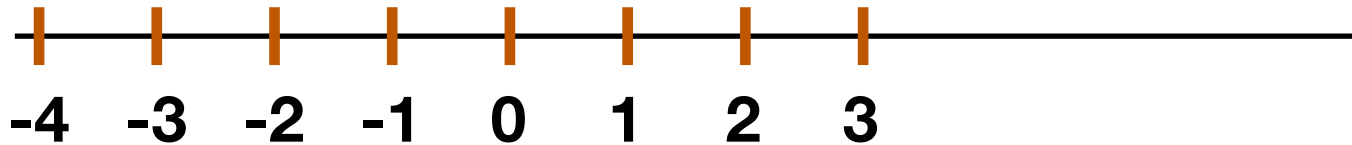
- Solution 2: Two's Complement



Unsigned	Binary
0	000
1	001
2	010
3	011
4	100
5	101
6	110
7	111

Encoding Negative Numbers

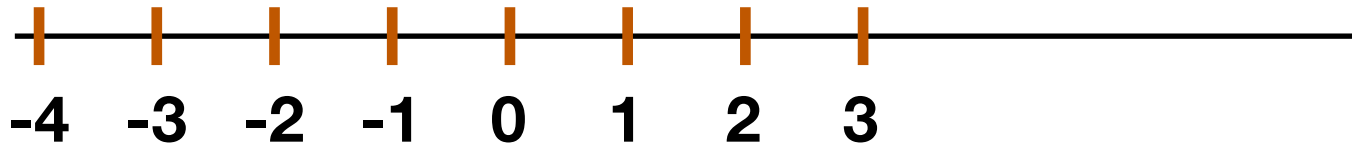
- Solution 2: Two's Complement



Signed	Unsigned	Binary
0	0	000
1	1	001
2	2	010
3	3	011
-4	4	100
-3	5	101
-2	6	110
-1	7	111

Encoding Negative Numbers

- Solution 2: Two's Complement

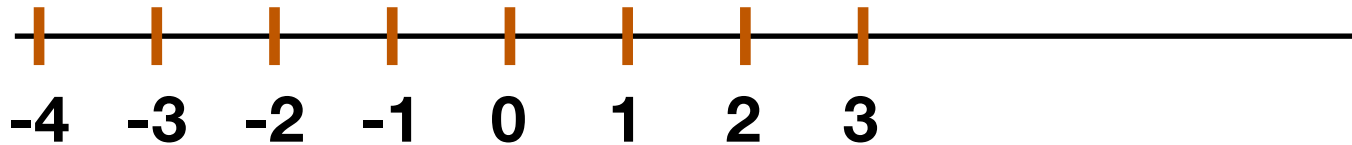


Signed Weight	Unsigned Weight	Bit Position
2^0	2^0	0
2^1	2^1	1
-2^2	2^2	2

Signed	Unsigned	Binary
0	0	000
1	1	001
2	2	010
3	3	011
-4	4	100
-3	5	101
-2	6	110
-1	7	111

Encoding Negative Numbers

- Solution 2: Two's Complement

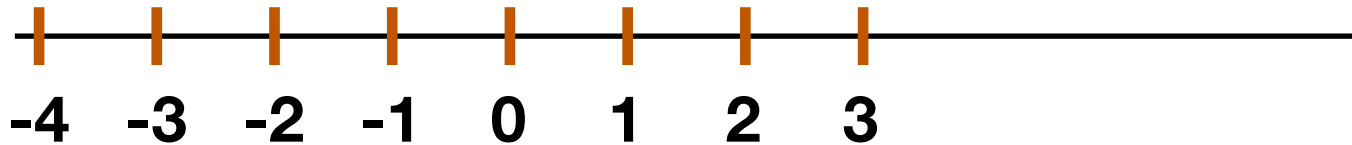


Signed Weight	Unsigned Weight	Bit Position
2^0	2^0	0
2^1	2^1	1
-2^2	2^2	2

Signed	Unsigned	Binary
0	0	000
1	1	001
2	2	010
3	3	011
-4	4	100
-3	5	101
-2	6	110
-1	7	111

Encoding Negative Numbers

- Solution 2: Two's Complement



Signed Weight	Unsigned Weight	Bit Position
2^0	2^0	0
2^1	2^1	1
-2^2	2^2	2

Signed	Unsigned	Binary
0	0	000
1	1	001
2	2	010
3	3	011
-4	4	100
-3	5	101
-2	6	110
-1	7	111

$$101_2 = 1 \cdot 2^0 + 0 \cdot 2^1 - 1 \cdot 2^2 = -3_{10}$$

Two-Complement Encoding Example

x = 15213: 00111011 01101101
y = -15213: 11000100 10010011

Weight	15213		-15213	
1	1	1	1	1
2	0	0	1	2
4	1	4	0	0
8	1	8	0	0
16	0	0	1	16
32	1	32	0	0
64	1	64	0	0
128	0	0	1	128
256	1	256	0	0
512	1	512	0	0
1024	0	0	1	1024
2048	1	2048	0	0
4096	1	4096	0	0
8192	1	8192	0	0
16384	0	0	1	16384
-32768	0	0	1	-32768
Sum	15213		-15213	

Two-Complement Implications

- Only 1 zero
- Usual arithmetic still works
- There is a bit that represents sign!
- Most widely used in today's machines

Signed	Binary
0	000
1	001
2	010
3	011
-4	100
-3	101
-2	110
-1	111

Two-Complement Implications

- Only 1 zero
- Usual arithmetic still works
- There is a bit that represents sign!
- Most widely used in today's machines

$$\begin{array}{r} 010 \\ +) 101 \\ \hline 111 \end{array}$$

Signed	Binary
0	000
1	001
2	010
3	011
-4	100
-3	101
-2	110
-1	111

Two-Complement Implications

- Only 1 zero
- Usual arithmetic still works
- There is a bit that represents sign!
- Most widely used in today's machines

$$\begin{array}{r} 010 \\ +) 101 \\ \hline 111 \end{array}$$

$$\begin{array}{r} 2 \\ +) -3 \\ \hline -1 \end{array}$$

Signed	Binary
0	000
1	001
2	010
3	011
-4	100
-3	101
-2	110
-1	111

Numeric Ranges

Numeric Ranges

- Unsigned Values

- $UMin = 0$
000...0

- $UMax = 2^w - 1$
111...1

Numeric Ranges

- Unsigned Values

- $UMin = 0$
000...0

- $UMax = 2^w - 1$
111...1

- Two's Complement Values

- $TMin = -2^{w-1}$
100...0

- $TMax = 2^{w-1} - 1$
011...1

Numeric Ranges

- Unsigned Values

- $UMin = 0$
000...0

- $UMax = 2^w - 1$
111...1

- Two's Complement Values

- $TMin = -2^{w-1}$
100...0

- $TMax = 2^{w-1} - 1$
011...1

Values for $W = 16$

	Decimal	Hex	Binary
UMax	65535	FF FF	11111111 11111111
TMax	32767	7F FF	01111111 11111111
TMin	-32768	80 00	10000000 00000000
-1	-1	FF FF	11111111 11111111
0	0	00 00	00000000 00000000

Numeric Ranges

- Unsigned Values

- $UMin = 0$
000...0

- $UMax = 2^w - 1$
111...1

- Two's Complement Values

- $TMin = -2^{w-1}$
100...0

- $TMax = 2^{w-1} - 1$
011...1

- Other Values

- Minus 1
111...1

Values for $W = 16$

	Decimal	Hex	Binary
UMax	65535	FF FF	11111111 11111111
TMax	32767	7F FF	01111111 11111111
TMin	-32768	80 00	10000000 00000000
-1	-1	FF FF	11111111 11111111
0	0	00 00	00000000 00000000

Data Representations in C (in Bytes)

- By default variables are signed
- Unless explicitly declared as unsigned (e.g., `unsigned int`)
- Signed variables use two-complement encoding

C Data Type	32-bit	64-bit
<code>char</code>	1	1
<code>short</code>	2	2
<code>int</code>	4	4
<code>long</code>	4	8

Data Representations in C (in Bytes)

	W			
	8	16	32	64
UMax	255	65,535	4,294,967,295	18,446,744,073,709,551,615
TMax	127	32,767	2,147,483,647	9,223,372,036,854,775,807
TMin	-128	-32,768	-2,147,483,648	-9,223,372,036,854,775,808

C Data Type	32-bit	64-bit
char	1	1
short	2	2
int	4	4
long	4	8

Data Representations in C (in Bytes)

	W			
	8	16	32	64
UMax	255	65,535	4,294,967,295	18,446,744,073,709,551,615
TMax	127	32,767	2,147,483,647	9,223,372,036,854,775,807
TMin	-128	-32,768	-2,147,483,648	-9,223,372,036,854,775,808

C Data Type	32-bit	64-bit
char	1	1
short	2	2
int	4	4
long	4	8

- C Language

- `#include <limits.h>`
- Declares constants, e.g.,
 - `ULONG_MAX`
 - `LONG_MAX`
 - `LONG_MIN`
- Values platform specific

Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
- C.f., Decimal
 - $12.45 = 1*10^1 + 2*10^0 + 4*10^{-1} + 5*10^{-2}$
- $10.01_2 = 1*2^1 + 0*2^0 + 0*2^{-1} + 1*2^{-2} = 2.25_{10}$

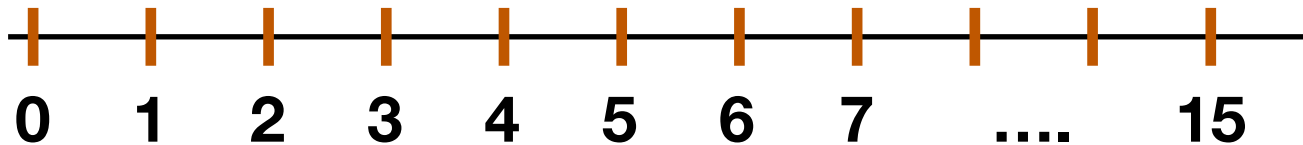
Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
- C.f., Decimal
 - $12.45 = 1 \cdot 10^1 + 2 \cdot 10^0 + 4 \cdot 10^{-1} + 5 \cdot 10^{-2}$
- $10.01_2 = 1 \cdot 2^1 + 0 \cdot 2^0 + 0 \cdot 2^{-1} + 1 \cdot 2^{-2} = 2.25_{10}$

Decimal	Binary
0	00.00
0.25	00.01
0.5	00.10
0.75	00.11
1	01.00
1.25	01.01
1.5	01.10
1.75	01.11
2	10.00
2.25	10.01
2.5	10.10
2.75	10.11
3	11.00
3.25	11.01
3.5	11.10
3.75	11.11

Can We Represent Fractions in Binary?

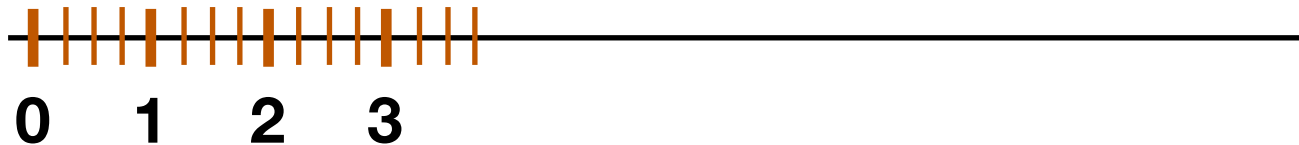
- What does 10.01_2 mean?
- C.f., Decimal
 - $12.45 = 1 \cdot 10^1 + 2 \cdot 10^0 + 4 \cdot 10^{-1} + 5 \cdot 10^{-2}$
- $10.01_2 = 1 \cdot 2^1 + 0 \cdot 2^0 + 0 \cdot 2^{-1} + 1 \cdot 2^{-2} = 2.25_{10}$



Decimal	Binary
0	00.00
0.25	00.01
0.5	00.10
0.75	00.11
1	01.00
1.25	01.01
1.5	01.10
1.75	01.11
2	10.00
2.25	10.01
2.5	10.10
2.75	10.11
3	11.00
3.25	11.01
3.5	11.10
3.75	11.11

Can We Represent Fractions in Binary?

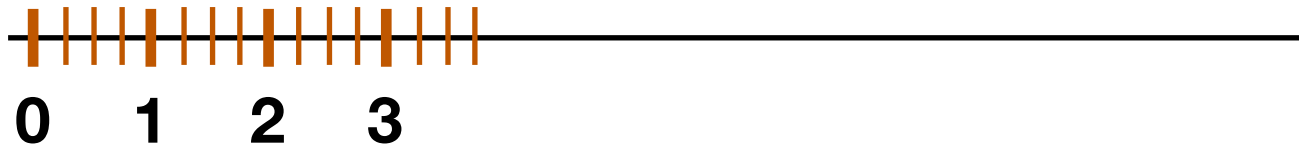
- What does 10.01_2 mean?
- C.f., Decimal
 - $12.45 = 1 \cdot 10^1 + 2 \cdot 10^0 + 4 \cdot 10^{-1} + 5 \cdot 10^{-2}$
- $10.01_2 = 1 \cdot 2^1 + 0 \cdot 2^0 + 0 \cdot 2^{-1} + 1 \cdot 2^{-2} = 2.25_{10}$



Decimal	Binary
0	00.00
0.25	00.01
0.5	00.10
0.75	00.11
1	01.00
1.25	01.01
1.5	01.10
1.75	01.11
2	10.00
2.25	10.01
2.5	10.10
2.75	10.11
3	11.00
3.25	11.01
3.5	11.10
3.75	11.11

Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
- C.f., Decimal
 - $12.45 = 1 \cdot 10^1 + 2 \cdot 10^0 + 4 \cdot 10^{-1} + 5 \cdot 10^{-2}$
- $10.01_2 = 1 \cdot 2^1 + 0 \cdot 2^0 + 0 \cdot 2^{-1} + 1 \cdot 2^{-2} = 2.25_{10}$



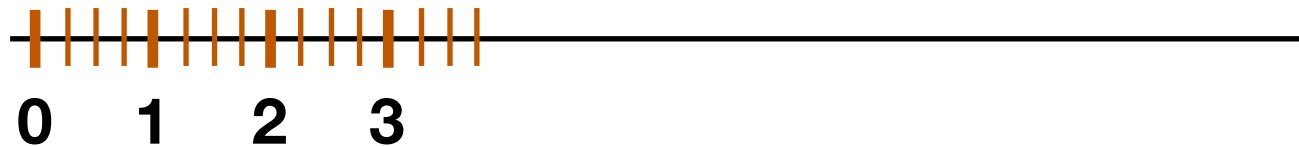
$$\begin{array}{r}
 01.10 \\
 + 01.01 \\
 \hline
 10.11
 \end{array}$$

$$\begin{array}{r}
 1.50 \\
 + 1.25 \\
 \hline
 2.75
 \end{array}$$

Decimal	Binary
0	00.00
0.25	00.01
0.5	00.10
0.75	00.11
1	01.00
1.25	01.01
1.5	01.10
1.75	01.11
2	10.00
2.25	10.01
2.5	10.10
2.75	10.11
3	11.00
3.25	11.01
3.5	11.10
3.75	11.11

Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
- C.f., Decimal
 - $12.45 = 1 \cdot 10^1 + 2 \cdot 10^0 + 4 \cdot 10^{-1} + 5 \cdot 10^{-2}$
- $10.01_2 = 1 \cdot 2^1 + 0 \cdot 2^0 + 0 \cdot 2^{-1} + 1 \cdot 2^{-2} = 2.25_{10}$



Integer Arithmetic Still Works!

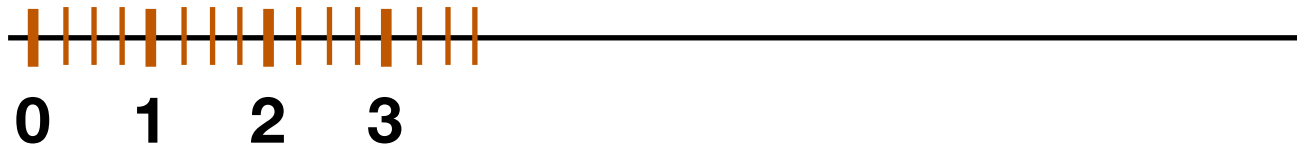
$$\begin{array}{r}
 01.10 \\
 + 01.01 \\
 \hline
 10.11
 \end{array}$$

$$\begin{array}{r}
 1.50 \\
 + 1.25 \\
 \hline
 2.75
 \end{array}$$

Decimal	Binary
0	00.00
0.25	00.01
0.5	00.10
0.75	00.11
1	01.00
1.25	01.01
1.5	01.10
1.75	01.11
2	10.00
2.25	10.01
2.5	10.10
2.75	10.11
3	11.00
3.25	11.01
3.5	11.10
3.75	11.11

Fixed-Point Representation

- Fixed interval between two representable numbers as long as the **binary point stays fixed**
 - Each bit represents 0.25_{10}
- **Fixed-point** representation of numbers
 - Integer is one special case of fixed-point

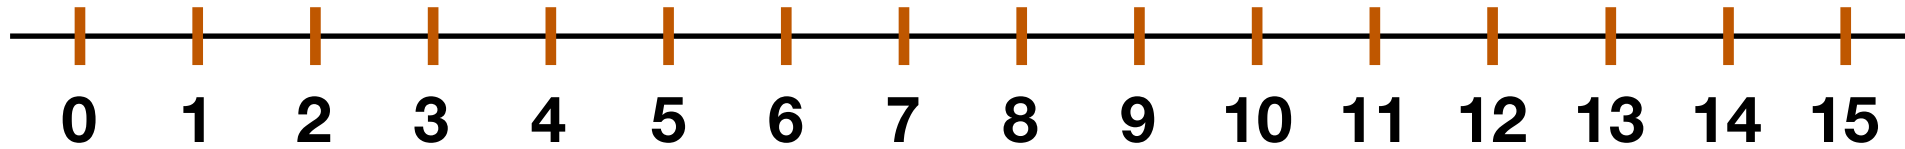


$$\begin{array}{r}
 01.10 \\
 + 01.01 \\
 \hline
 10.11
 \end{array}$$

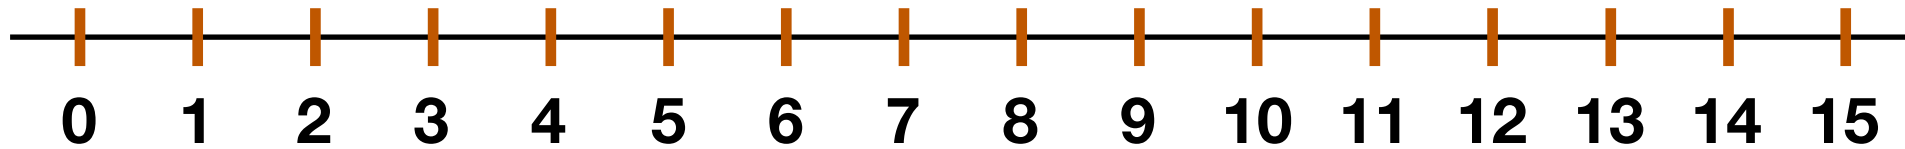
$$\begin{array}{r}
 1.50 \\
 + 1.25 \\
 \hline
 2.75
 \end{array}$$

Decimal	Binary
0	00.00
0.25	00.01
0.5	00.10
0.75	00.11
1	01.00
1.25	01.01
1.5	01.10
1.75	01.11
2	10.00
2.25	10.01
2.5	10.10
2.75	10.11
3	11.00
3.25	11.01
3.5	11.10
3.75	11.11

Aside: Quantization

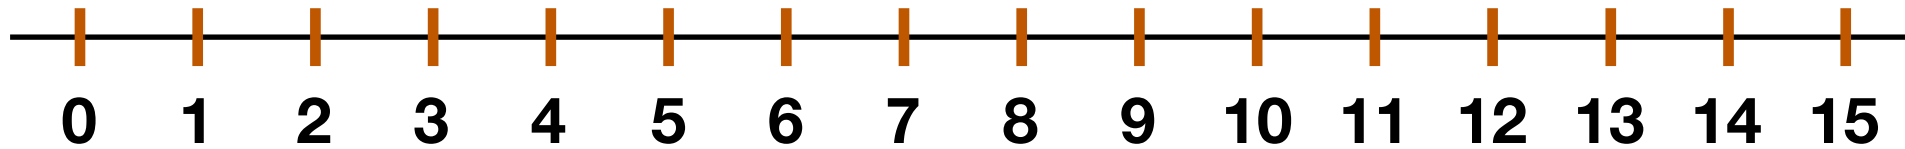


Aside: Quantization



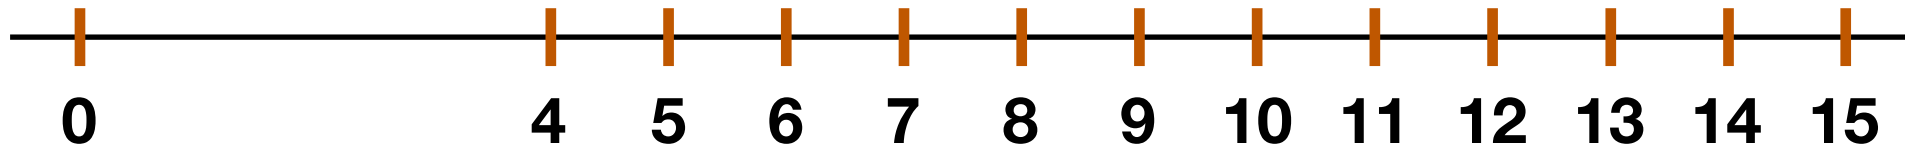
- Representing all integers **precisely** requires 4 bits

Aside: Quantization



- Representing all integers **precisely** requires 4 bits
- What if we can tolerate some imprecisions
 - 1, 2, 3 are approximated by 0
 - 5, 6, 7 are approximated by 4...
 - We would only need 2 bits

Aside: Quantization



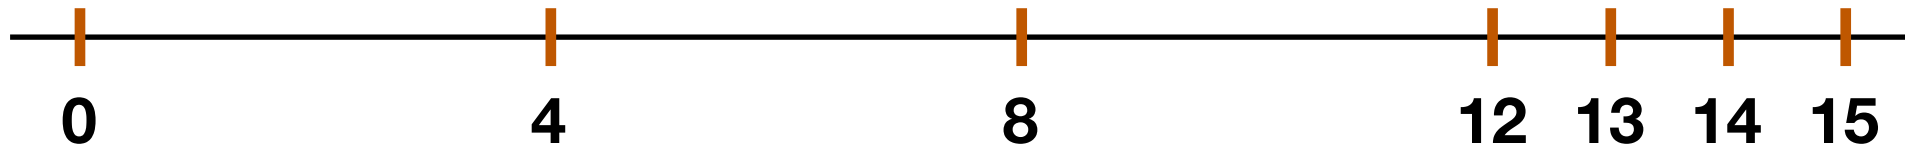
- Representing all integers **precisely** requires 4 bits
- What if we can tolerate some imprecisions
 - 1, 2, 3 are approximated by 0
 - 5, 6, 7 are approximated by 4...
 - We would only need 2 bits

Aside: Quantization



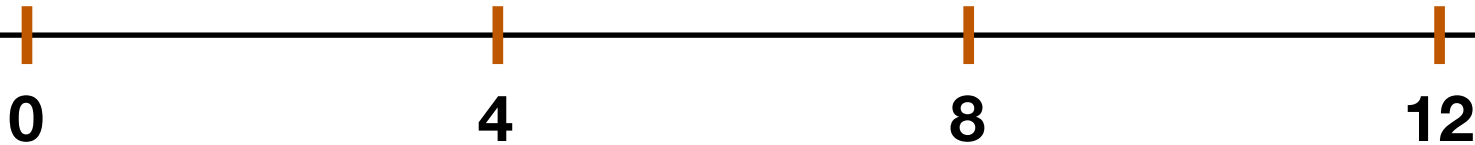
- Representing all integers **precisely** requires 4 bits
- What if we can tolerate some imprecisions
 - 1, 2, 3 are approximated by 0
 - 5, 6, 7 are approximated by 4...
 - We would only need 2 bits

Aside: Quantization



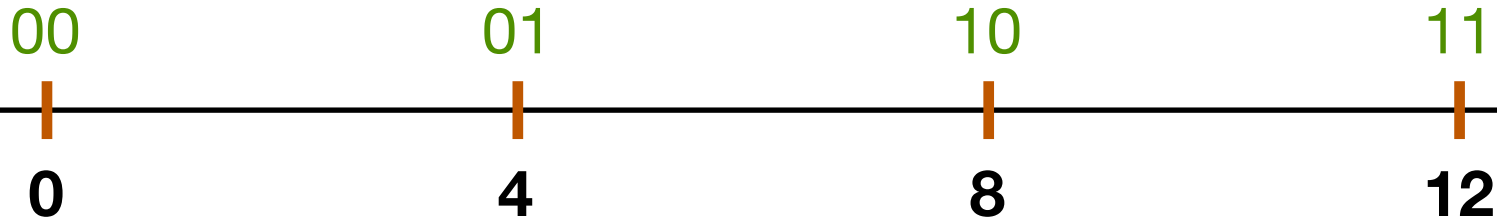
- Representing all integers **precisely** requires 4 bits
- What if we can tolerate some imprecisions
 - 1, 2, 3 are approximated by 0
 - 5, 6, 7 are approximated by 4...
 - We would only need 2 bits

Aside: Quantization



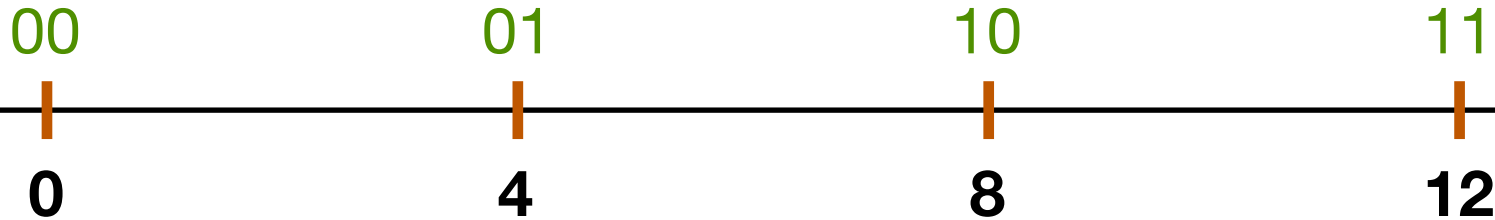
- Representing all integers **precisely** requires 4 bits
- What if we can tolerate some imprecisions
 - 1, 2, 3 are approximated by 0
 - 5, 6, 7 are approximated by 4...
 - We would only need 2 bits

Aside: Quantization



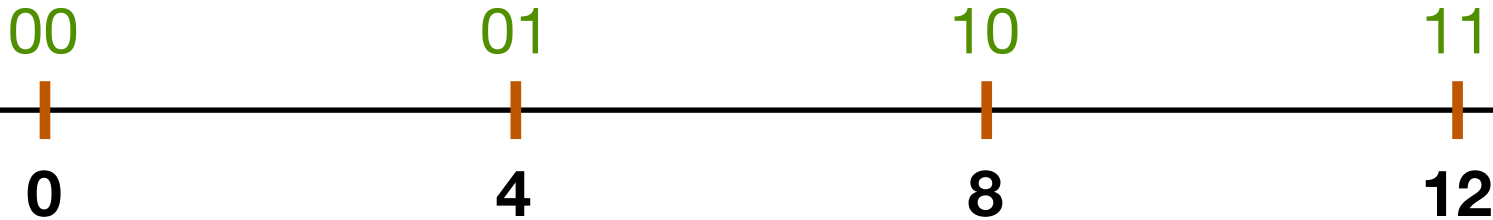
- Representing all integers **precisely** requires 4 bits
- What if we can tolerate some imprecisions
 - 1, 2, 3 are approximated by 0
 - 5, 6, 7 are approximated by 4...
 - We would only need 2 bits

Aside: Quantization



- Representing all integers **precisely** requires 4 bits
- What if we can tolerate some imprecisions
 - 1, 2, 3 are approximated by 0
 - 5, 6, 7 are approximated by 4...
 - We would only need 2 bits
- That is, 1 bit represents 4_{10}
 - 10_2 becomes $4 * (1 * 2^1) = 8$
 - Every time we increment a bit, the value is incremented by 4
 - 1, 2, 3 are represented **approximately** by 00_2

Aside: Quantization



- Representing all integers **precisely** requires 4 bits

- What if

- 1, 2
- 5, 6
- We

Note that this is different from “base 4”

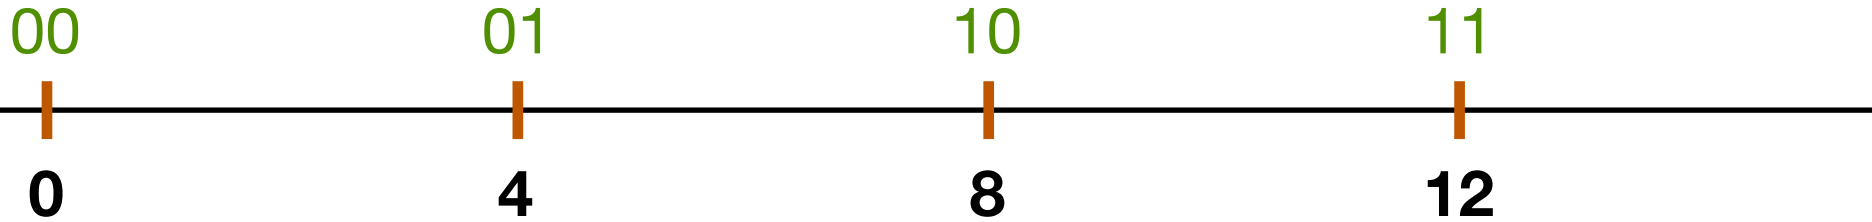
- $10_4 = 1 * 4^1 + 0 * 4^0 = 4$

- Every increment still only increments 1

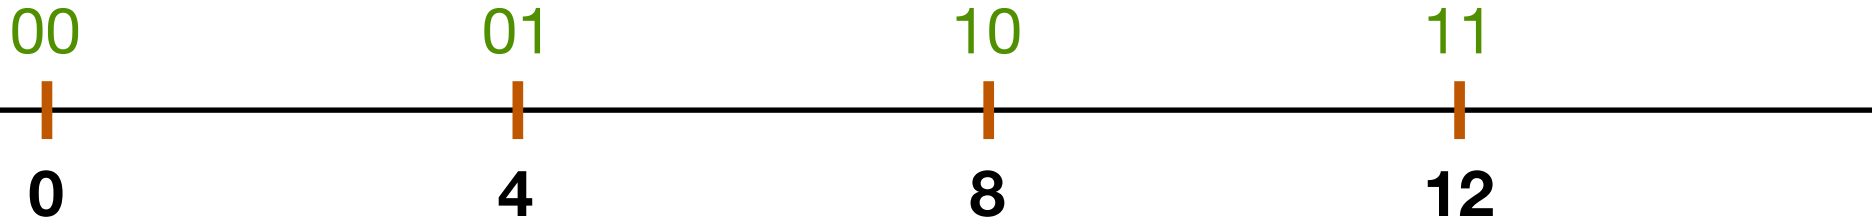
- That is, 1 bit represents 4_{10}

- 10_2 becomes $4 * (1 * 2^1) = 8$
- Every time we increment a bit, the value is incremented by 4
- 1, 2, 3 are represented **approximately** by 00_2

Aside: Quantization

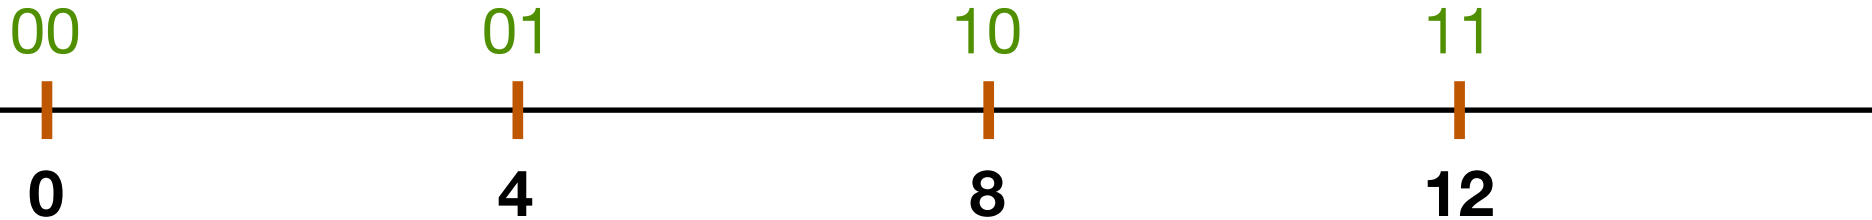


Aside: Quantization



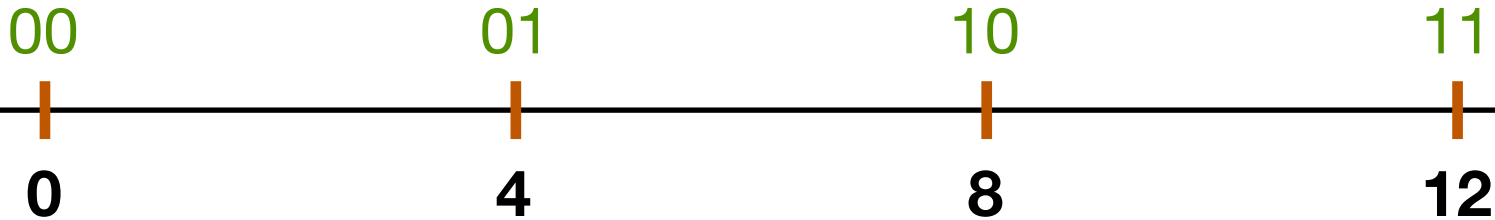
- Saves storage space and improves computation speed
 - 50% space saving
 - 4-bit arithmetic becomes 2-bit arithmetic

Aside: Quantization



- Saves storage space and improves computation speed
 - 50% space saving
 - 4-bit arithmetic becomes 2-bit arithmetic
- Many real-world applications can tolerate imprecisions
 - Image processing
 - Computer vision
 - Real-time graphics
 - Machine learning (Neural networks)

Aside: Quantization



- Saves storage space and improves computation speed
 - 50% space saving



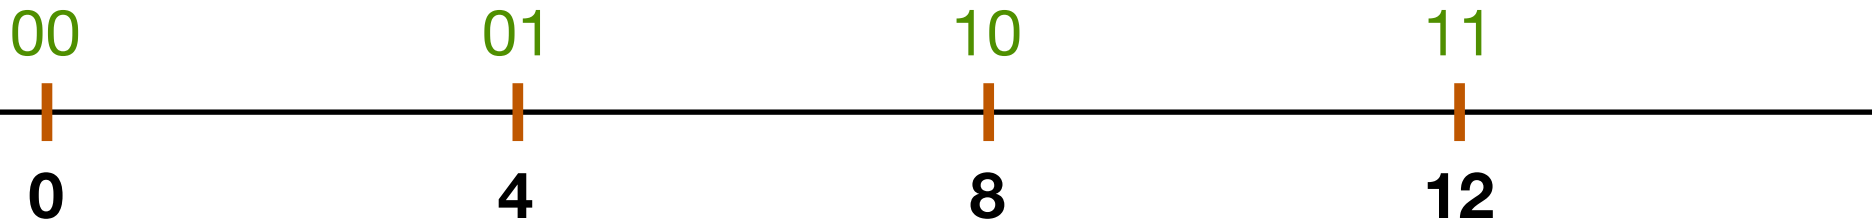
es 2
ons

al ne



Aside: Quantization

Questions?



- Saves storage space and improves computation speed
 - 50% space saving

