

CSC 252: Computer Organization

Spring 2026: Lecture 2

Instructor: Yuhao Zhu

Department of Computer Science
University of Rochester

Announcement

- Make sure you can access CSUG machines!!!
- Programming assignment 1 will be posted this week.
 - I will send an announcement when it's out.
 - It is in C language. Seek help from TAs.
 - TAs are best positioned to answer your questions about programming assignments!!!
- Programming assignments do NOT repeat the lecture materials. They ask you to synthesize what you have learned from the lectures and work out something new.

Previously in 252...

- How is a human-readable program translated to a representation?
- How is a contract modified and executed?

Problem

Algorithm

Renting	
Service provider	Landlord
Service receiver	YOU
Contract	Lease
Contract's language	Natural language (e.g., English)

Circuit

Previously in 252...

C Program

```
void add() {  
    int a = 1;  
    int b = 2;  
    int c = a + b;  
}
```



Assembly program

```
movl $1, -4(%rbp)  
movl $2, -8(%rbp)  
movl -4(%rbp), %eax  
addl -8(%rbp), %eax
```

Previously in 252...

Assembly program

```
movl $1, -4(%rbp)
movl $2, -8(%rbp)
movl -4(%rbp), %eax
addl -8(%rbp), %eax
```



Executable Binary

00011001	...
01101010	...
11010101	...
01110001	...

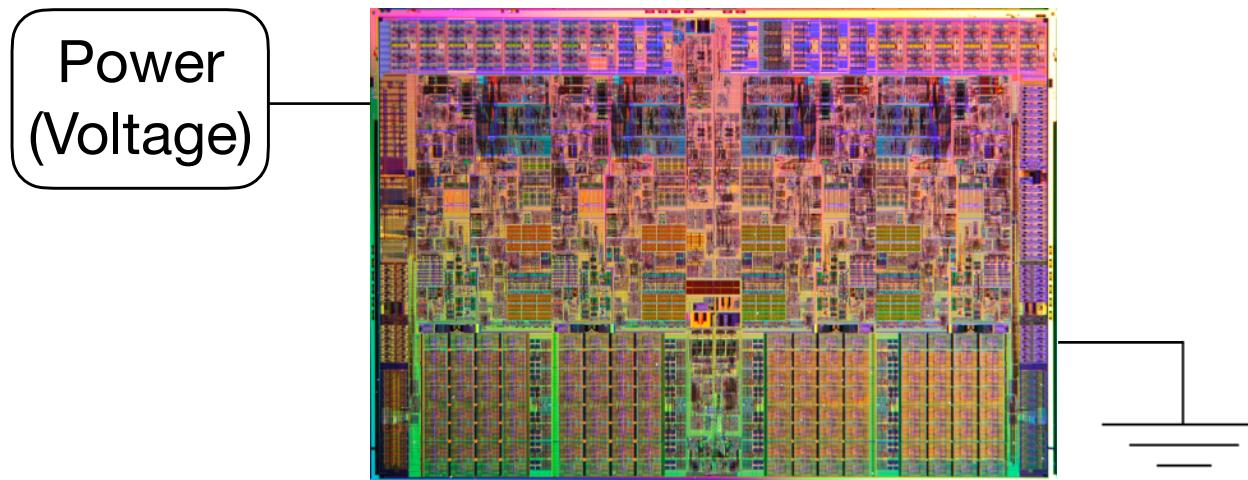
- What's the difference between an assembly program and an executable binary?
 - They refer to the same thing — a list of instructions that the software asks the hardware to perform
 - They are just different representations
- Instruction = Operator + Operand(s)

Today: Representing Information in Binary

- Why Binary (bits)?
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting

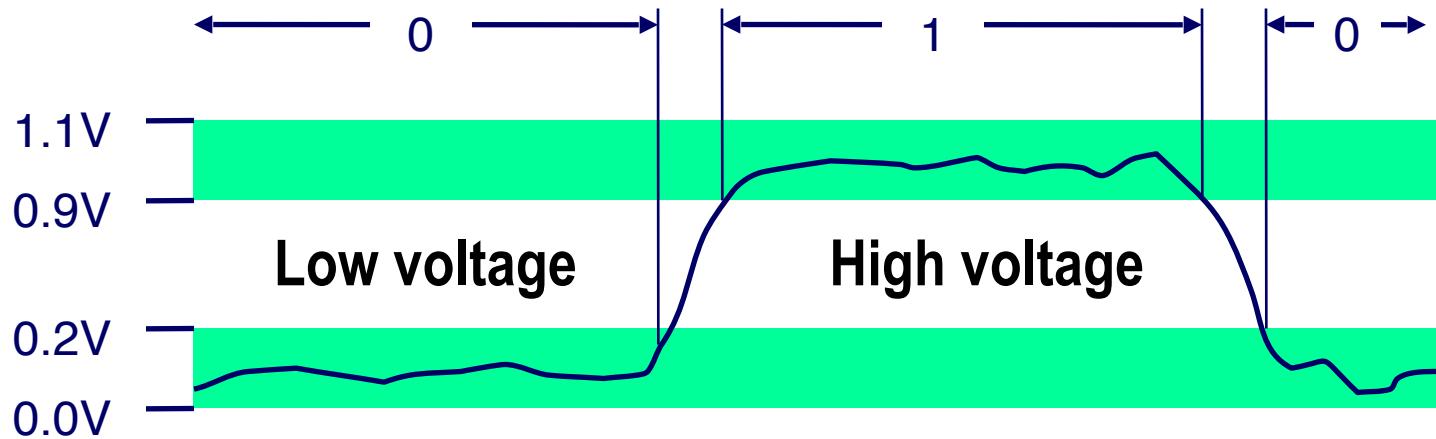
Everything in Computers is bits

- Each bit is 0 or 1. Bits are how programs talk to the hardware
- Programs encode instructions in bits
- Hardware then interprets the bits
- Why bits? Electronic Implementation



Everything in Computers is bits

- Each bit is 0 or 1. Bits are how programs talk to the hardware
- Programs encode instructions in bits
- Hardware then interprets the bits
- Why bits? Electronic Implementation
 - Use high voltage to represent 1
 - Use low voltage to represent 0



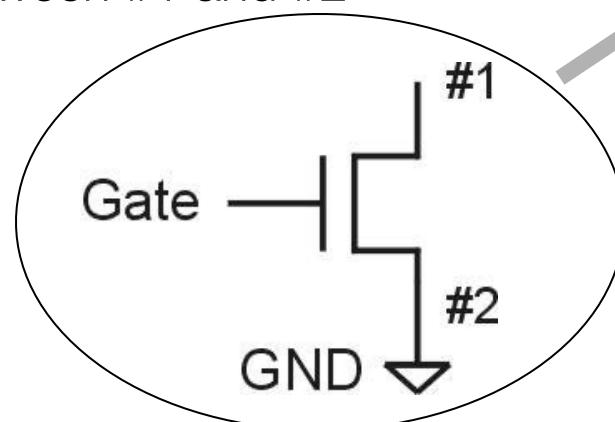
Why Bits?

Processors are made of transistors, which are Metal Oxide Semiconductor (MOS)

- two types: n-type and p-type

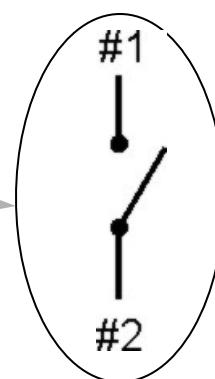
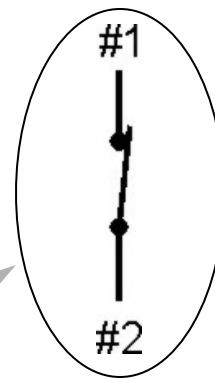
n-type (NMOS)

- when Gate has high voltage, short circuit between #1 and #2 (switch closed)
- when Gate has low voltage, open circuit between #1 and #2 (switch open)



Gate = 1

Gate = 0

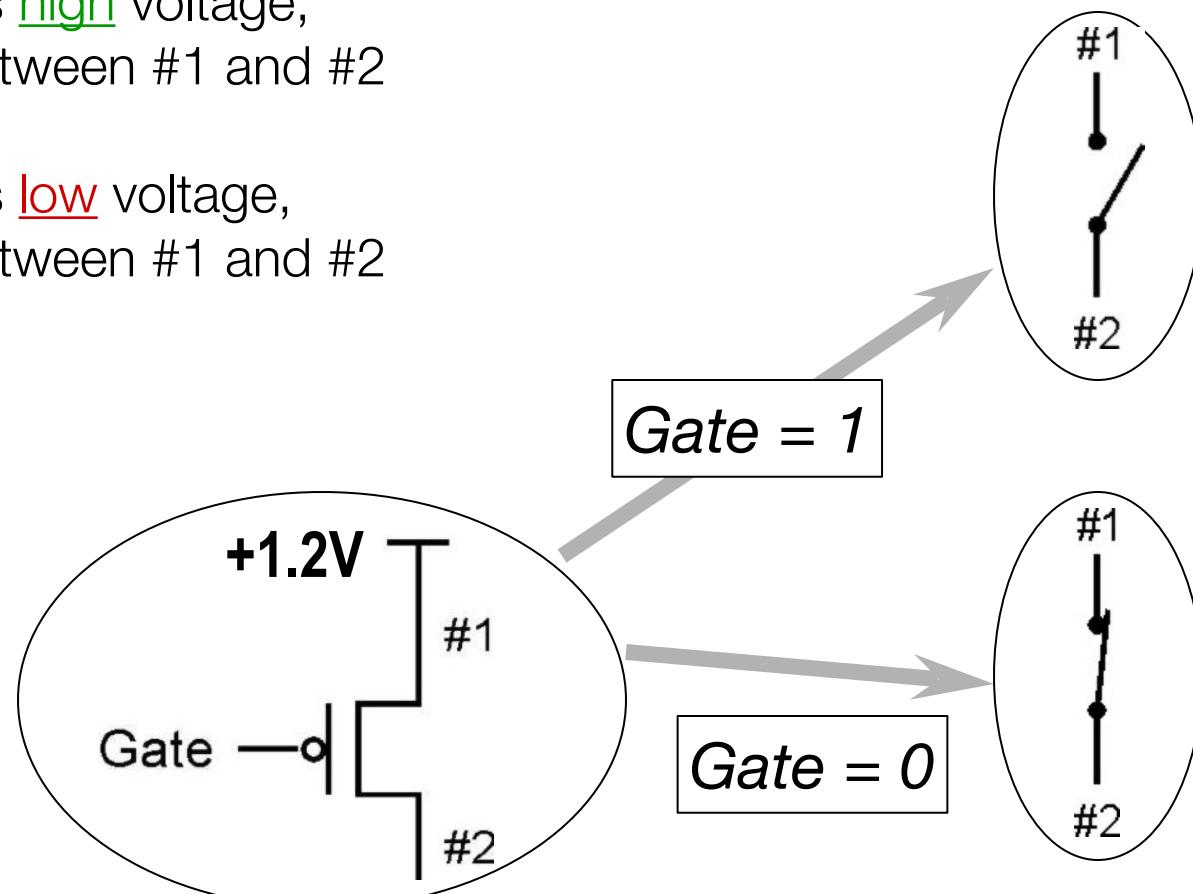


Terminal #2 must be connected to GND (0V).

Why Bits?

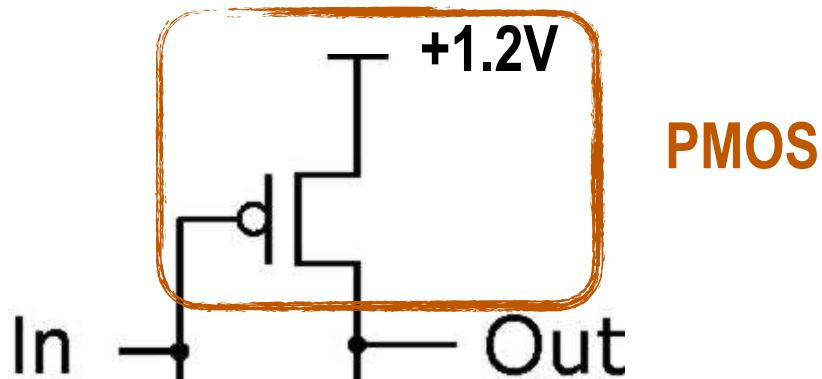
p-type is *complementary* to n-type (**PMOS**)

- when Gate has high voltage,
open circuit between #1 and #2
(switch open)
- when Gate has low voltage,
short circuit between #1 and #2
(switch closed)



Terminal #1 must be connected to +1.2V

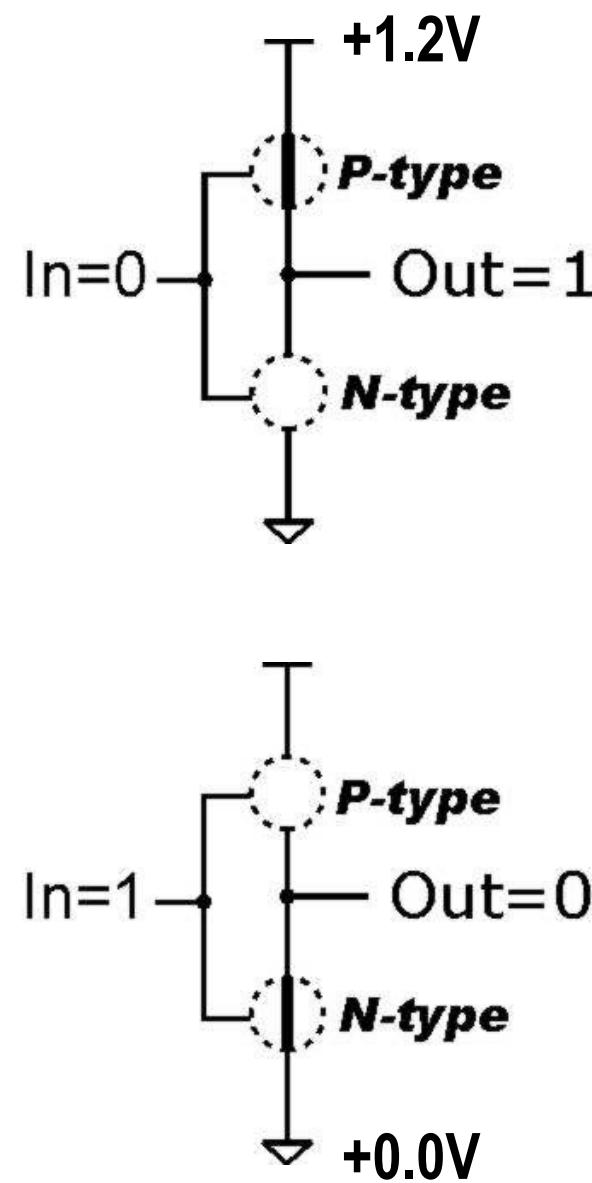
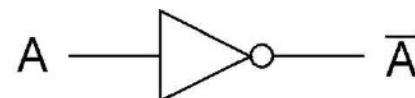
Inverter



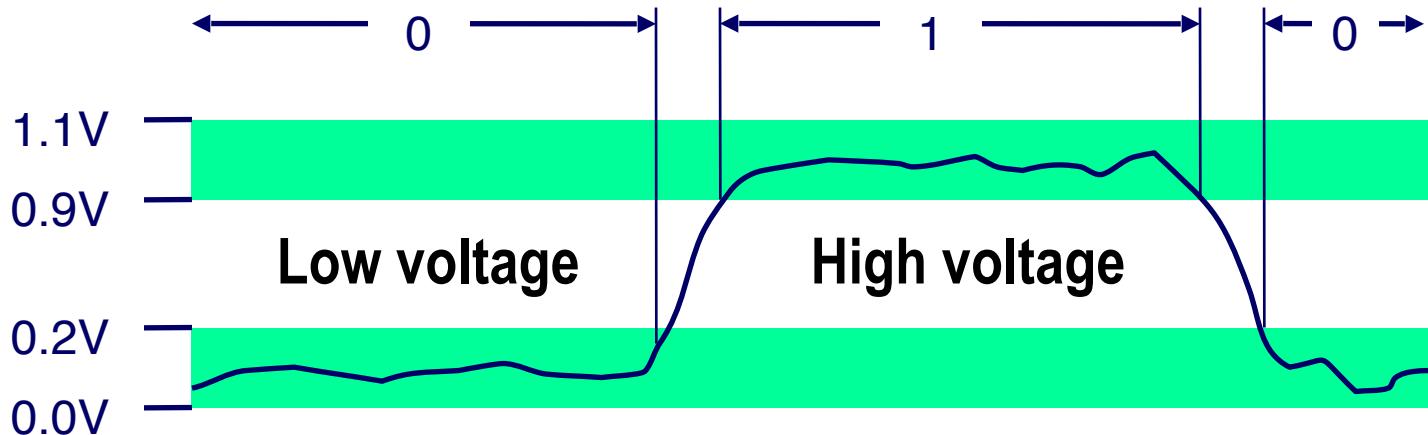
PMOS

NMOS

In	Out
0	1
1	0

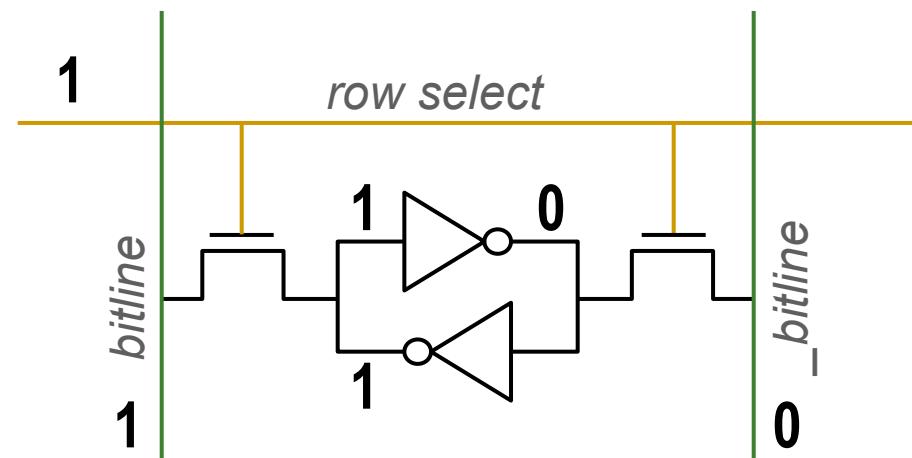


Store/Access Data



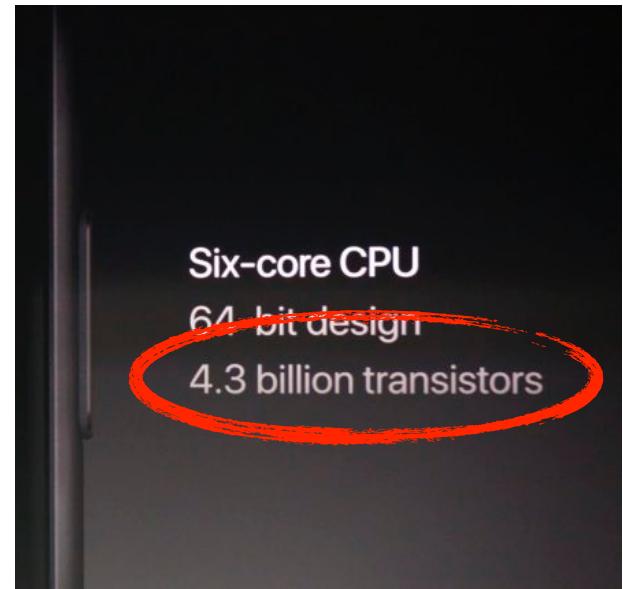
- Two cross coupled inverters store a single bit

- Feedback path persists the value in the “cell”
- 4 transistors for storage
- 2 transistors for access
- A “6T” cell



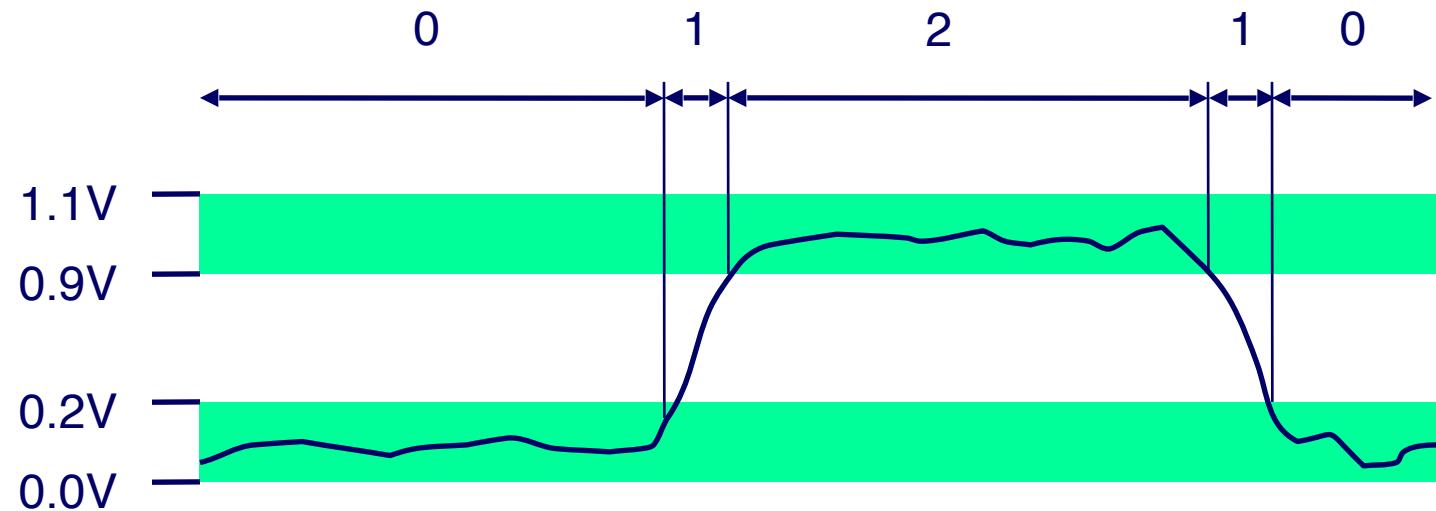
Transistors

- Computers are made of transistors
- Transistors have become smaller over the years
 - Not so much anymore...



Why Limit Ourselves Only to Binary?

- Voltage is continuous. Why interpret it only as 0s and 1s?
- Answer: Noise



Binary Notation

- Base 2 Number Representation (Binary)
- C.f., Base 10 number representation (Decimal)
- $21_{10} = 1*10^0 + 2*10^1 = 21$
- Weighted Positional Notation
 - Each bit has a weight depending on its position
- $1011_2 = 1*2^0 + 1*2^1 + 0*2^2 + 1*2^3 = 11_{10}$
- $b_3b_2b_1b_0 = b^0*2^0 + b^1*2^1 + b^2*2^2 + b^3*2^3$
- Binary Arithmetic

$$\begin{array}{r} 0110 \\ + 0101 \\ \hline 1011 \end{array} \qquad \begin{array}{r} 6 \\ + 5 \\ \hline 11 \end{array}$$

Decimal	Binary
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111
8	1000
9	1001
10	1010
11	1011
12	1100
13	1101
14	1110
15	1111

Hexdecimal (Hex) Notation

- **Base 16 Number Representation**
 - Use characters ‘0’ to ‘9’ and ‘A’ to ‘F’
 - Four bits per Hex digit
 - $11111110_2 = \text{FE}_{16}$
- Write FA1D37B_{16} in C as
 - `0xFA1D37B`
 - `0xfa1d37b`

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

Bit, Byte, Word

- **Byte = 8 bits**
 - Binary 0000000₂ to 1111111₂; Decimal: 0₁₀ to 255₁₀; Hex: 00₁₆ to FF₁₆
 - Least Significant Bit (LSb) vs. Most Significant Bit (MSb)



- **Word = 4 Bytes (32-bit machine) / 8 Bytes (64-bit machine)**
 - Least Significant Byte (LSB) vs. Most Significant Byte (MSB)

Today: Representing Information in Binary

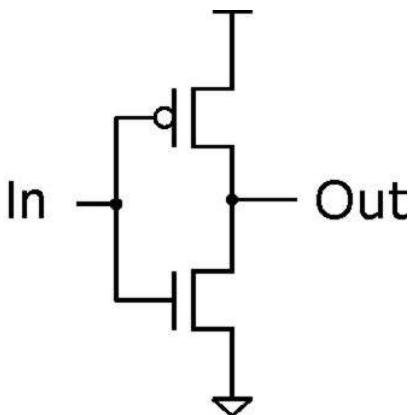
- Why Binary (bits)?
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting

Bit-level manipulations

Not

- $\sim A = 1$ when $A=0$

\sim	0	1
0	1	
1		0



Or

- $A \mid B = 1$ when either $A=1$ or $B=1$

I	0	1
0	0	1
1	1	1

And

- $A \& B = 1$ when both $A=1$ and $B=1$

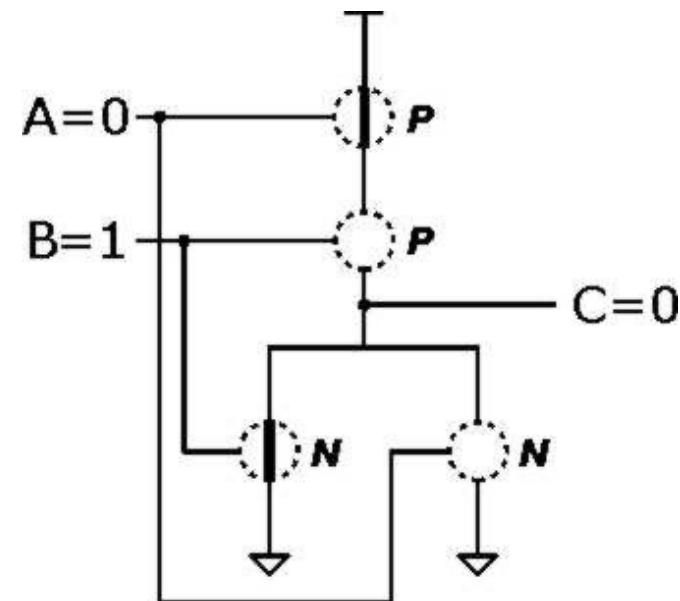
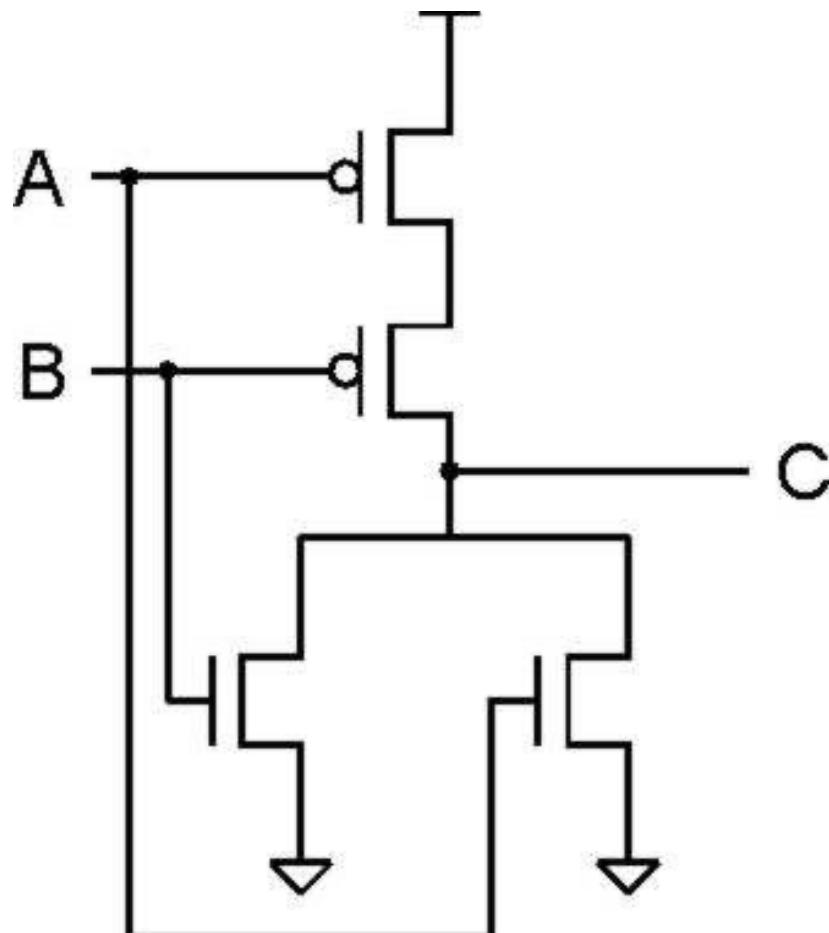
&	0	1
0	0	0
1	0	1

Exclusive-Or (Xor)

- $A \wedge B = 1$ when either $A=1$ or $B=1$, but not both

^	0	1
0	0	1
1	1	0

NOR (OR + NOT)



A	B	C
0	0	1
0	1	0
1	0	0
1	1	0

Bit Vector Operations

- Operate on Bit Vectors
 - Operations applied bitwise

$$\begin{array}{rcl} \begin{array}{r} 01101001 \\ \& 01010101 \\ \hline 01000001 \end{array} & \begin{array}{r} 01101001 \\ | 01010101 \\ \hline 01111101 \end{array} & \begin{array}{r} 01101001 \\ ^ 01010101 \\ \hline 00111100 \end{array} \\ & & \begin{array}{r} ~ 01010101 \\ \hline 10101010 \end{array} \end{array}$$

Bit-Level Operations in C

- Operations &, |, ~, ^ Available in C
 - Apply to any “integral” data type
 - long, int, short, char, unsigned
 - View arguments as bit vectors
 - Arguments applied bit-wise
- Examples (Char data type)
 - $\sim 0x41 \rightarrow 0xBE$
 - $\sim 01000001_2 \rightarrow 10111110_2$
 - $\sim 0x00 \rightarrow 0xFF$
 - $\sim 00000000_2 \rightarrow 11111111_2$
 - $0x69 \& 0x55 \rightarrow 0x41$
 - $01101001_2 \& 01010101_2 \rightarrow 01000001_2$
 - $0x69 | 0x55 \rightarrow 0x7D$
 - $01101001_2 | 01010101_2 \rightarrow 01111101_2$

Aside: Logic Operations in C

- Contrast to Logical Operators
 - `&&`, `||`, `!`
 - View 0 as “False”
 - Anything nonzero as “True”
 - Always return 0 or 1
 - Early termination (e.g., `0 && 1 && 1`)
- Examples (char data type)
 - `!0x41` → `0x00`
 - `!0x00` → `0x01`
 - `!!0x41` → `0x01`
 - `0x69 && 0x55` → `0x01`
 - `0x69 || 0x55` → `0x01`
 - `p && *p` (avoids null pointer access)

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector **x** left **y** positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector **x** right **y** positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq total amount of bits

Argument x	01100010
<< 3	00010000
Log. >> 2	00011000
Arith. >> 2	00011000

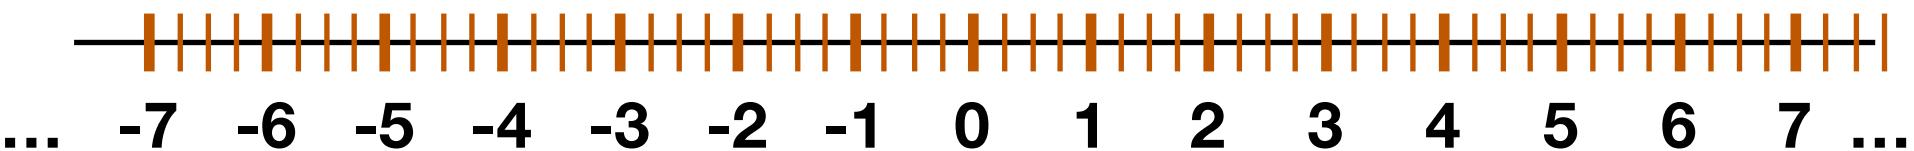
Argument x	10100010
<< 3	00010000
Log. >> 2	00101000
Arith. >> 2	11101000

Today: Representing Information in Binary

- Why Binary (bits)?
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting

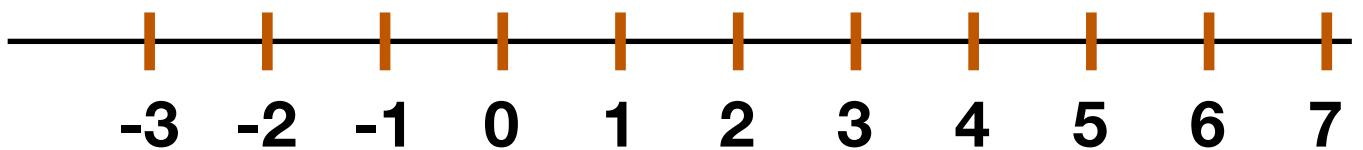
Representing Numbers in Binary

- Different types of number
 - Integer (Negative and Non-negative)
 - Fractions
 - Irrationals



Encoding Negative Numbers

- So far we have been discussing non-negative numbers: so called **unsigned**. How about negative numbers?
- Solution 1: Sign-magnitude
 - First bit represents sign; 0 for positive; 1 for negative
 - The rest represents magnitude



111	110	101	000	001	010	011	100	101	110	111
			100							

Sign-Magnitude Implications

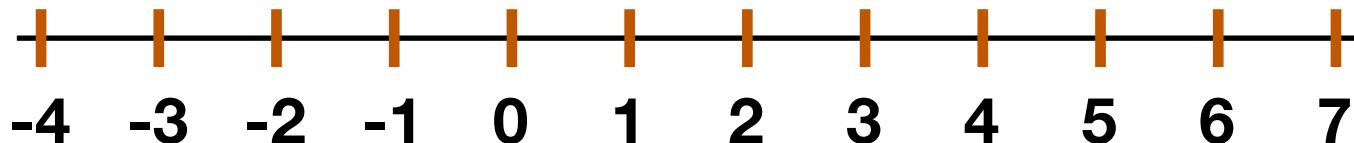
- Bits have different semantics
 - Two zeros...
 - Normal arithmetic doesn't work
 - Make hardware design harder

$$\begin{array}{r} \cancel{010} \\ +) \ 101 \\ \hline \cancel{111} \end{array}$$
$$\begin{array}{r} \cancel{2} \\ +) \ -1 \\ \hline \cancel{-3} \end{array}$$

Signed Value	Binary
0	000
1	001
2	010
3	011
-0	100
-1	101
-2	110
-3	111

Encoding Negative Numbers

- Solution 2: Two's Complement



Signed Weight	Unsigned Weight	Bit Position
2^0	2^0	0
2^1	2^1	1
-2^2	2^2	2

Signed	Unsigned	Binary
0	0	000
1	1	001
2	2	010
3	3	011
-4	4	100
-3	5	101
-2	6	110
-1	7	111

$$101_2 = 1 * 2^0 + 0 * 2^1 - 1 * 2^2 = -3_{10}$$

Two-Complement Encoding Example

x =	15213:	00111011	01101101
y =	-15213:	11000100	10010011

Weight	15213		-15213	
1	1	1	1	1
2	0	0	1	2
4	1	4	0	0
8	1	8	0	0
16	0	0	1	16
32	1	32	0	0
64	1	64	0	0
128	0	0	1	128
256	1	256	0	0
512	1	512	0	0
1024	0	0	1	1024
2048	1	2048	0	0
4096	1	4096	0	0
8192	1	8192	0	0
16384	0	0	1	16384
-32768	0	0	1	-32768
	Sum		15213	
			-15213	

Two-Complement Implications

- Only 1 zero
- Usual arithmetic still works
- There is a bit that represents the sign!
- Most widely used in today's machines

$$\begin{array}{r} 010 \\ +) 101 \\ \hline 111 \end{array}$$
$$\begin{array}{r} 2 \\ +) -3 \\ \hline -1 \end{array}$$

Signed	Binary
0	000
1	001
2	010
3	011
-4	100
-3	101
-2	110
-1	111

Numeric Ranges

- Unsigned Values
 - $UMin = 0$
000...0
 - $UMax = 2^w - 1$
111...1
- Two's Complement Values
 - $TMin = -2^{w-1}$
100...0
 - $TMax = 2^{w-1} - 1$
011...1
- Other Values
 - Minus 1
111...1

Values for $W = 16$

	Decimal	Hex	Binary
UMax	65535	FF FF	11111111 11111111
TMax	32767	7F FF	01111111 11111111
TMin	-32768	80 00	10000000 00000000
-1	-1	FF FF	11111111 11111111
0	0	00 00	00000000 00000000

Data Representations in C (in Bytes)

- By default variables are signed
- Unless explicitly declared as unsigned (e.g., `unsigned int`)
- Signed variables use two-complement encoding

C Data Type	32-bit	64-bit
<code>char</code>	1	1
<code>short</code>	2	2
<code>int</code>	4	4
<code>long</code>	4	8

Data Representations in C (in Bytes)

	W			
	8	16	32	64
UMax	255	65,535	4,294,967,295	18,446,744,073,709,551,615
TMax	127	32,767	2,147,483,647	9,223,372,036,854,775,807
TMin	-128	-32,768	-2,147,483,648	-9,223,372,036,854,775,808

C Data Type	32-bit	64-bit
char	1	1
short	2	2
int	4	4
long	4	8

- C Language
 - #include <limits.h>
 - Declares constants, e.g.,
 - ULONG_MAX
 - LONG_MAX
 - LONG_MIN
 - Values platform specific