

CSC 252: Computer Organization

Spring 2020: Lecture 4

Instructor: Yuhao Zhu

Department of Computer Science
University of Rochester

Announcement

- Programming Assignment 1 is out
 - Details: <https://www.cs.rochester.edu/courses/252/spring2020/labs/assignment1.html>
 - Due on Jan. 31, 11:59 PM
 - You have 3 slip days

19	20	21	22	23	24	25
26	27	28	29	30	31	Feb 1

Today

Due

Announcement

- Programming assignment 1 is in C language. Seek help from TAs.
- TAs are best positioned to answer your questions about programming assignments!!!
- Programming assignments do NOT repeat the lecture materials. They ask you to synthesize what you have learned from the lectures and work out something new.

Mapping Signed \leftrightarrow Unsigned

Bits	Signed	Unsigned
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	-8	8
1001	-7	9
1010	-6	10
1011	-5	11
1100	-4	12
1101	-3	13
1110	-2	14
1111	-1	15

Mapping Signed \leftrightarrow Unsigned

Bits	Signed	Unsigned
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	-8	8
1001	-7	9
1010	-6	10
1011	-5	11
1100	-4	12
1101	-3	13
1110	-2	14
1111	-1	15

Mapping Signed \leftrightarrow Unsigned



Today: Floating Point

- Background: Fractional binary numbers and fixed-point
- Floating point representation
- IEEE 754 standard
- Rounding, addition, multiplication
- Floating point in C
- Summary

Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
 - C.f., Decimal

Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
 - C.f., Decimal

$$12.45 = 1 * 10^1 + 2 * 10^0 + 4 * 10^{-1} + 5 * 10^{-2}$$

Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
 - C.f., Decimal

$$12.45 = \boxed{1 * 10^1} + 2 * 10^0 + 4 * 10^{-1} + 5 * 10^{-2}$$


Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
 - C.f., Decimal

$$12.45 = 1 * 10^1 + 2 * 10^0 + 4 * 10^{-1} + 5 * 10^{-2}$$


Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
 - C.f., Decimal

$$12.45 = 1 * 10^1 + 2 * 10^0 + \boxed{4 * 10^{-1}} + 5 * 10^{-2}$$


Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
 - C.f., Decimal

$$12.45 = 1 * 10^1 + 2 * 10^0 + 4 * 10^{-1} + \boxed{5 * 10^{-2}}$$


Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
 - C.f., Decimal

$$12.45 = 1 * 10^1 + 2 * 10^0 + 4 * 10^{-1} + 5 * 10^{-2}$$

$$10.01_2 = 1 * 2^1 + 0 * 2^0 + 0 * 2^{-1} + 1 * 2^{-2}$$

Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
 - C.f., Decimal

$$12.45 = 1 * 10^1 + 2 * 10^0 + 4 * 10^{-1} + 5 * 10^{-2}$$

$$10.01_2 = \boxed{1 * 2^1} + 0 * 2^0 + 0 * 2^{-1} + 1 * 2^{-2}$$


Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
 - C.f., Decimal

$$12.45 = 1 * 10^1 + 2 * 10^0 + 4 * 10^{-1} + 5 * 10^{-2}$$

$$10.01_2 = 1 * 2^1 + \boxed{0 * 2^0} + 0 * 2^{-1} + 1 * 2^{-2}$$


Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
 - C.f., Decimal

$$12.45 = 1 * 10^1 + 2 * 10^0 + 4 * 10^{-1} + 5 * 10^{-2}$$

$$10.01_2 = 1 * 2^1 + 0 * 2^0 + \boxed{0 * 2^{-1}} + 1 * 2^{-2}$$


Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
 - C.f., Decimal

$$12.45 = 1 * 10^1 + 2 * 10^0 + 4 * 10^{-1} + 5 * 10^{-2}$$

$$10.01_2 = 1 * 2^1 + 0 * 2^0 + 0 * 2^{-1} + \boxed{1 * 2^{-2}}$$

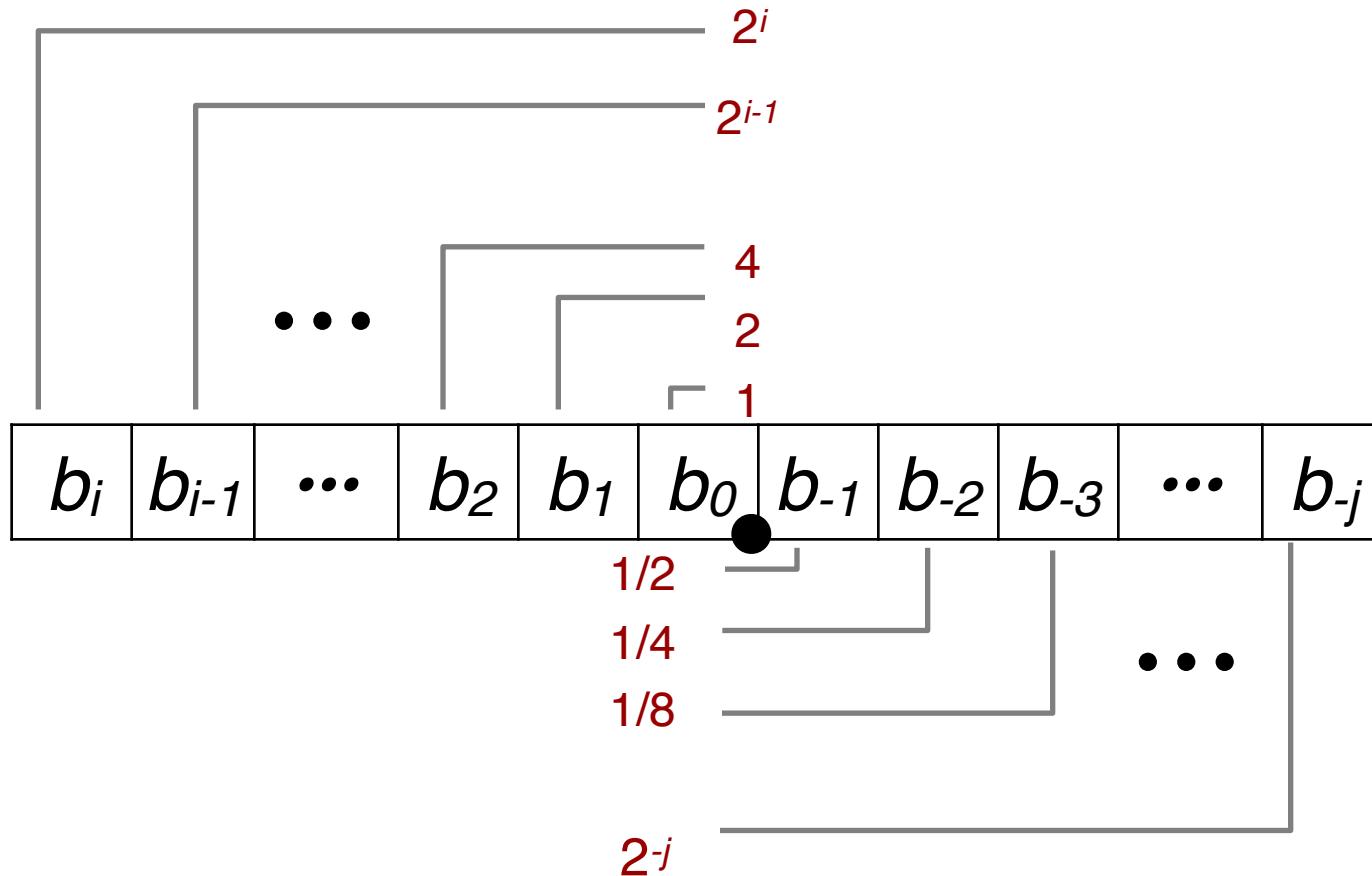

Can We Represent Fractions in Binary?

- What does 10.01_2 mean?
 - C.f., Decimal

$$12.45 = 1 * 10^1 + 2 * 10^0 + 4 * 10^{-1} + 5 * 10^{-2}$$

$$\begin{aligned}10.01_2 &= 1 * 2^1 + 0 * 2^0 + 0 * 2^{-1} + 1 * 2^{-2} \\&= 2.25_{10}\end{aligned}$$

Fractional Binary Numbers



Fractional Binary Numbers: Examples

Decimal Value	Binary Representation
5 3/4	101.11
2 7/8	10.111
1 7/16	1.0111

Fractional Binary Numbers: Examples

Decimal Value	Binary Representation	
5 3/4	101.11	
2 7/8	10.111	
1 7/16	1.0111	Exact Same Raw Bit Stream!

Fractional Binary Numbers: Examples

Decimal Value	Binary Representation
5 3/4	101.11
2 7/8	10.111
1 7/16	1.0111

Exact Same Raw Bit Stream!

- If we have a weighted positional notation system that has 5 bits, can the three numbers above all be represented in this notation system? How to do calculations?

Fractional Binary Numbers: Examples

Decimal Value	Binary Representation
5 3/4	101.11
2 7/8	10.111
1 7/16	1.0111

Exact Same Raw Bit Stream!

- If we have a weighted positional notation system that has 5 bits, can the three numbers above all be represented in this notation system? How to do calculations?
- We would need to remember:
 - The raw bit stream (5 bits)
 - Where the binary point is (potentially another 3 bits for 6 positions)

Fractional Binary Numbers: Examples

Decimal Value	Binary Representation
5 3/4	101.11
2 7/8	10.111
1 7/16	1.0111

Exact Same Raw Bit Stream!

- If we have a weighted positional notation system that has 5 bits, can the three numbers above all be represented in this notation system? How to do calculations?
- We would need to remember:
 - The raw bit stream (5 bits)
 - Where the binary point is (potentially another 3 bits for 6 positions)
- Makes calculations (e.g. addition) hard
 - Need to first align numbers according to the binary point

Fixed-Point Representation

Fixed-Point Representation

- Binary point stays fixed

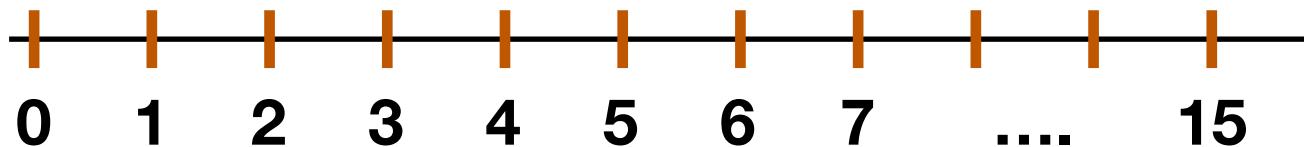
Fixed-Point Representation

- Binary point stays fixed

Decimal	Binary
0	0000.
1	0001.
2	0010.
3	0011.
4	0100.
5	0101.
6	0110.
7	0111.
8	1000.
9	1001.
10	1010.
11	1011.
12	1100.
13	1101.
14	1110.
15	1111.

Fixed-Point Representation

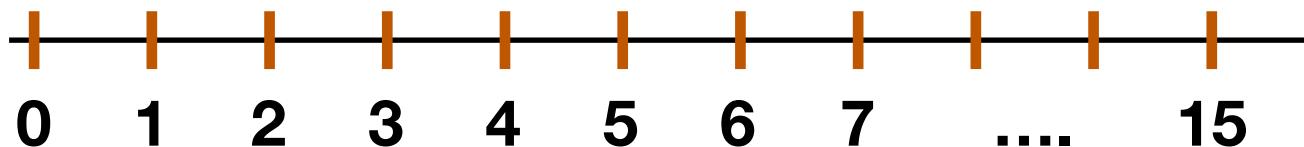
- Binary point stays fixed



Decimal	Binary
0	0000.
1	0001.
2	0010.
3	0011.
4	0100.
5	0101.
6	0110.
7	0111.
8	1000.
9	1001.
10	1010.
11	1011.
12	1100.
13	1101.
14	1110.
15	1111.

Fixed-Point Representation

- Binary point stays fixed



Decimal	Binary
0	00.00
0.25	00.01
0.5	00.10
0.75	00.11
1	01.00
1.25	01.01
1.5	01.10
1.75	01.11
2	10.00
2.25	10.01
2.5	10.10
2.75	10.11
3	11.00
3.25	11.01
3.5	11.10
3.75	11.11

Fixed-Point Representation

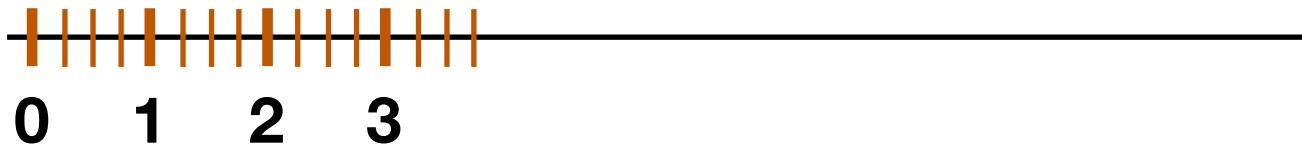
- Binary point stays fixed



Decimal	Binary
0	00.00
0.25	00.01
0.5	00.10
0.75	00.11
1	01.00
1.25	01.01
1.5	01.10
1.75	01.11
2	10.00
2.25	10.01
2.5	10.10
2.75	10.11
3	11.00
3.25	11.01
3.5	11.10
3.75	11.11

Fixed-Point Representation

- Binary point stays fixed
- Fixed interval between representable numbers
 - The interval in this example is 0.25_{10}



Decimal	Binary
0	00.00
0.25	00.01
0.5	00.10
0.75	00.11
1	01.00
1.25	01.01
1.5	01.10
1.75	01.11
2	10.00
2.25	10.01
2.5	10.10
2.75	10.11
3	11.00
3.25	11.01
3.5	11.10
3.75	11.11

Fixed-Point Representation

- Binary point stays fixed
- Fixed interval between representable numbers
 - The interval in this example is 0.25_{10}

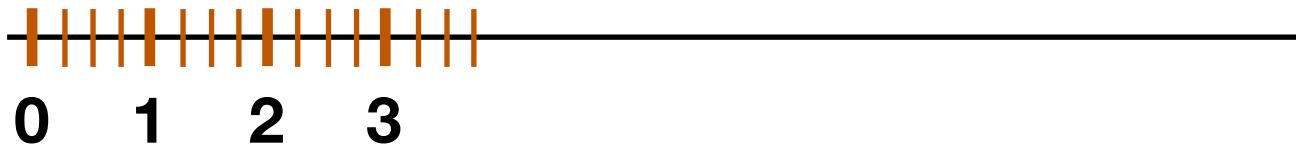


- Still need to remember the binary point, but just once for all numbers

Decimal	Binary
0	00.00
0.25	00.01
0.5	00.10
0.75	00.11
1	01.00
1.25	01.01
1.5	01.10
1.75	01.11
2	10.00
2.25	10.01
2.5	10.10
2.75	10.11
3	11.00
3.25	11.01
3.5	11.10
3.75	11.11

Fixed-Point Representation

- Binary point stays fixed
- Fixed interval between representable numbers
 - The interval in this example is 0.25_{10}

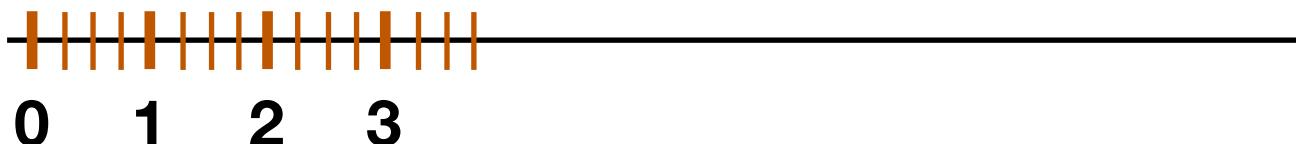


- Still need to remember the binary point, but just once for all numbers
- No need to align (already aligned)

Decimal	Binary
0	00.00
0.25	00.01
0.5	00.10
0.75	00.11
1	01.00
1.25	01.01
1.5	01.10
1.75	01.11
2	10.00
2.25	10.01
2.5	10.10
2.75	10.11
3	11.00
3.25	11.01
3.5	11.10
3.75	11.11

Fixed-Point Representation

- Binary point stays fixed
- Fixed interval between representable numbers
 - The interval in this example is 0.25_{10}



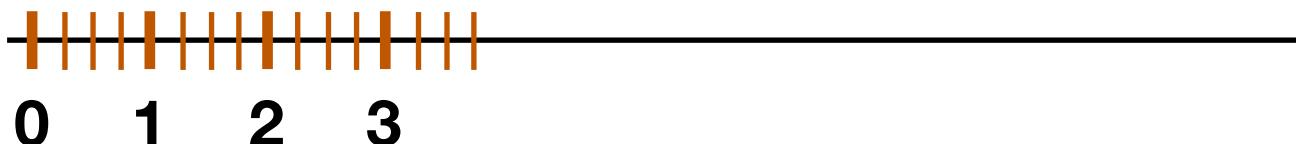
- Still need to remember the binary point, but just once for all numbers
- No need to align (already aligned)

$$\begin{array}{r} 01.10 \\ + 01.01 \\ \hline 10.11 \end{array}$$

Decimal	Binary
0	00.00
0.25	00.01
0.5	00.10
0.75	00.11
1	01.00
1.25	01.01
1.5	01.10
1.75	01.11
2	10.00
2.25	10.01
2.5	10.10
2.75	10.11
3	11.00
3.25	11.01
3.5	11.10
3.75	11.11

Fixed-Point Representation

- Binary point stays fixed
- Fixed interval between representable numbers
 - The interval in this example is 0.25_{10}



- Still need to remember the binary point, but just once for all numbers
- No need to align (already aligned)

$$\begin{array}{r} 01.10 \\ + 01.01 \\ \hline 10.11 \end{array} \qquad \begin{array}{r} 1.50 \\ + 1.25 \\ \hline 2.75 \end{array}$$

Decimal	Binary
0	00.00
0.25	00.01
0.5	00.10
0.75	00.11
1	01.00
1.25	01.01
1.5	01.10
1.75	01.11
2	10.00
2.25	10.01
2.5	10.10
2.75	10.11
3	11.00
3.25	11.01
3.5	11.10
3.75	11.11

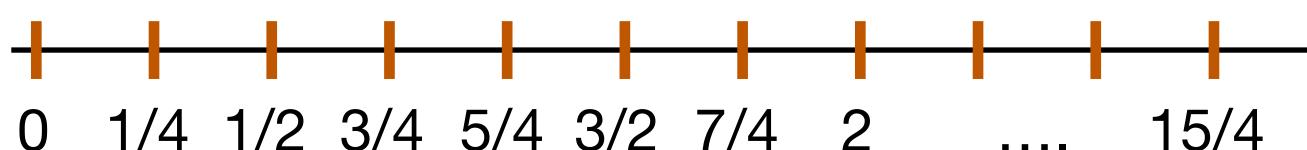
Limitations of Fixed-Point (#1)

Limitations of Fixed-Point (#1)

- Can exactly represent numbers only of the form $x/2^k$

Limitations of Fixed-Point (#1)

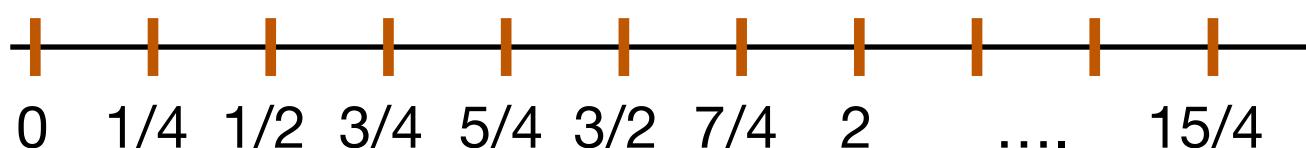
- Can exactly represent numbers only of the form $x/2^k$



$b_3b_2.b_1b_0$

Limitations of Fixed-Point (#1)

- Can exactly represent numbers only of the form $x/2^k$
 - Other rational numbers have repeating bit representations

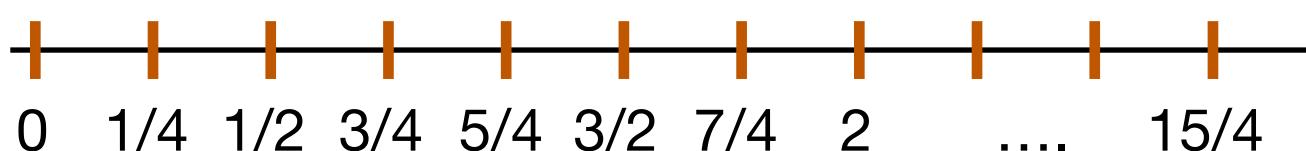


$b_3b_2.b_1b_0$

Limitations of Fixed-Point (#1)

- Can exactly represent numbers only of the form $x/2^k$
 - Other rational numbers have repeating bit representations

Decimal Value	Binary Representation
1/3	0.0101010101[01]...
1/5	0.001100110011[0011]...
1/10	0.0001100110011[0011]...



$b_3.b_2.b_1.b_0$

Limitations of Fixed-Point (#2)

Limitations of Fixed-Point (#2)

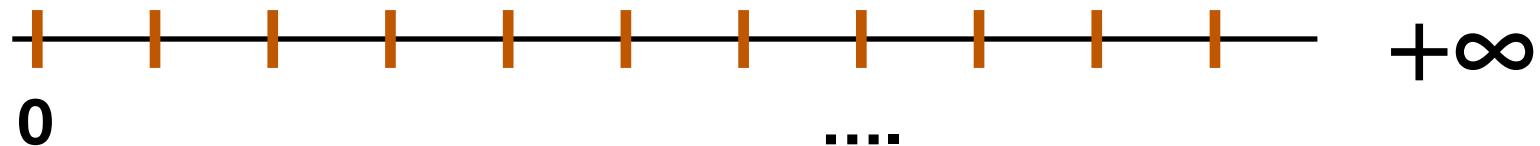
- Can't represent very small and very large numbers at the same time

Limitations of Fixed-Point (#2)

- Can't represent very small and very large numbers at the same time
 - To represent very large numbers, the (fixed) interval needs to be large, making it hard to represent small numbers

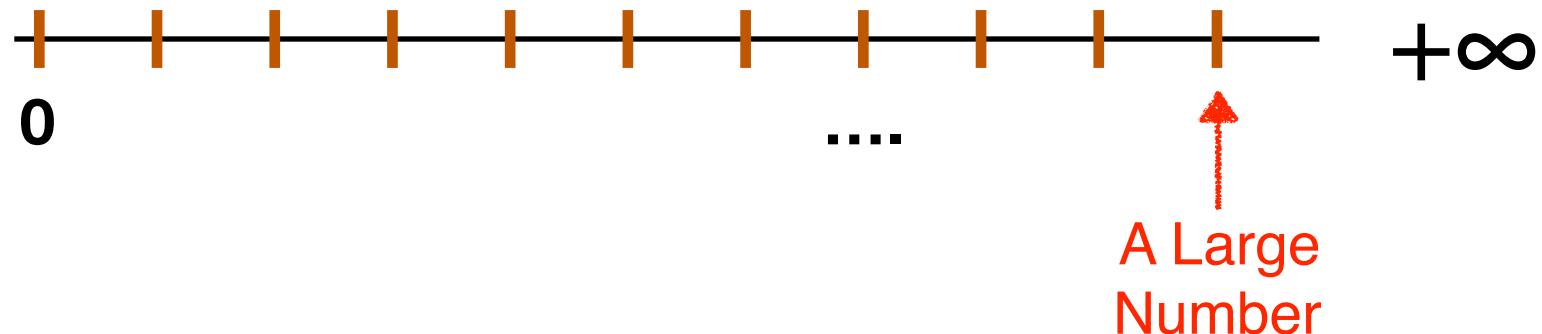
Limitations of Fixed-Point (#2)

- Can't represent very small and very large numbers at the same time
 - To represent very large numbers, the (fixed) interval needs to be large, making it hard to represent small numbers



Limitations of Fixed-Point (#2)

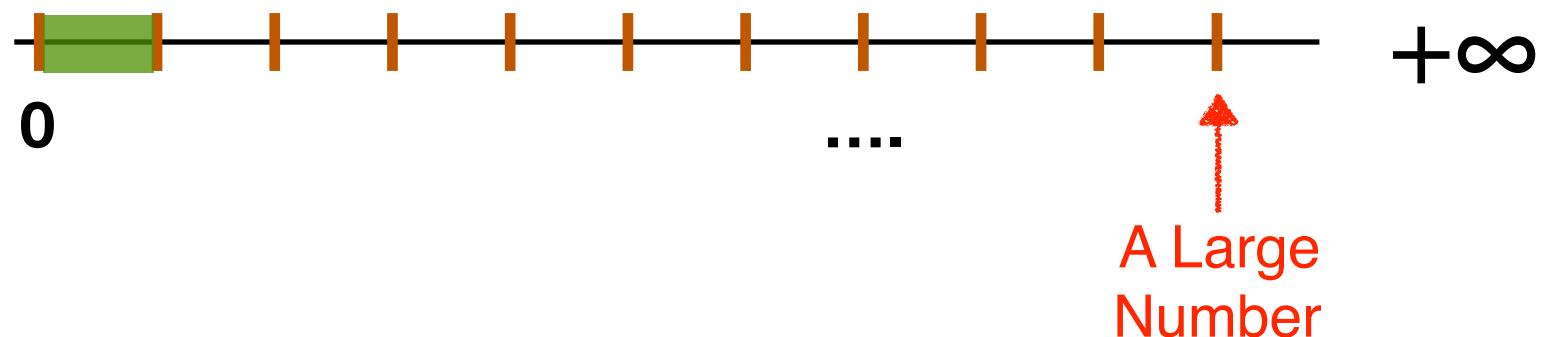
- Can't represent very small and very large numbers at the same time
 - To represent very large numbers, the (fixed) interval needs to be large, making it hard to represent small numbers



Limitations of Fixed-Point (#2)

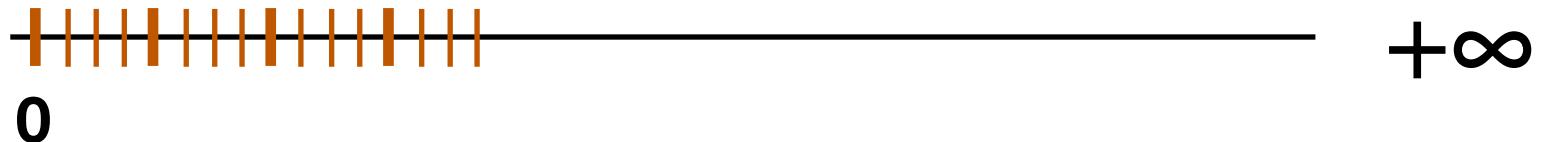
- Can't represent very small and very large numbers at the same time
 - To represent very large numbers, the (fixed) interval needs to be large, making it hard to represent small numbers

Unrepresentable
small numbers



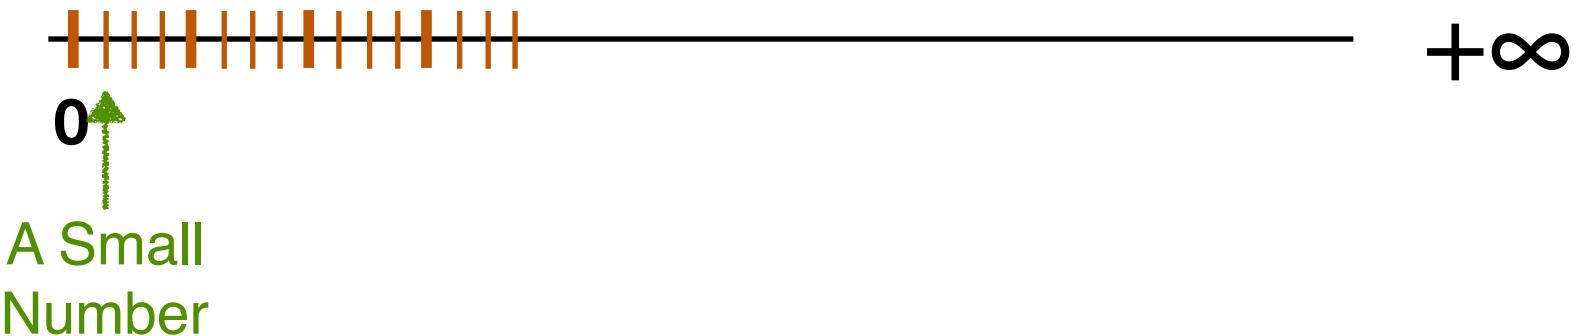
Limitations of Fixed-Point (#2)

- Can't represent very small and very large numbers at the same time
 - To represent very large numbers, the (fixed) interval needs to be large, making it hard to represent small numbers
 - To represent very small numbers, the (fixed) interval needs to be small, making it hard to represent large numbers



Limitations of Fixed-Point (#2)

- Can't represent very small and very large numbers at the same time
 - To represent very large numbers, the (fixed) interval needs to be large, making it hard to represent small numbers
 - To represent very small numbers, the (fixed) interval needs to be small, making it hard to represent large numbers



Limitations of Fixed-Point (#2)

- Can't represent very small and very large numbers at the same time
 - To represent very large numbers, the (fixed) interval needs to be large, making it hard to represent small numbers
 - To represent very small numbers, the (fixed) interval needs to be small, making it hard to represent large numbers

Unrepresentable
large numbers



Today: Floating Point

- Background: Fractional binary numbers and fixed-point
- **Floating point representation**
- IEEE 754 standard
- Rounding, addition, multiplication
- Floating point in C
- Summary

Primer: (Normalized) Scientific Notation

Primer: (Normalized) Scientific Notation

Decimal Value	Scientific Notation
2	2×10^0
-4,321.768	-4.321768×10^3
0.000 000 007 51	7.51×10^{-9}

Primer: (Normalized) Scientific Notation

- In decimal: $M \times 10^E$
 - E is an integer
 - Normalized form: $1 \leq |M| < 10$

Decimal Value	Scientific Notation
2	2×10^0
-4,321.768	-4.321768×10^3
0.000 000 007 51	7.51×10^{-9}

Primer: (Normalized) Scientific Notation

- In decimal: $M \times 10^E$
 - E is an integer
 - Normalized form: $1 \leq |M| < 10$

M × 10^E

Decimal Value	Scientific Notation
2	2×10^0
-4,321.768	-4.321768×10^3
0.000 000 007 51	7.51×10^{-9}

Primer: (Normalized) Scientific Notation

- In decimal: $M \times 10^E$
 - E is an integer
 - Normalized form: $1 \leq |M| < 10$

M × 10^E

↑

Significand

Decimal Value	Scientific Notation
2	2×10^0
-4,321.768	-4.321768×10^3
0.000 000 007 51	7.51×10^{-9}

Primer: (Normalized) Scientific Notation

- In decimal: $M \times 10^E$
 - E is an integer
 - Normalized form: $1 \leq |M| < 10$

M × 10^E

Significand Base

Decimal Value	Scientific Notation
2	2×10^0
-4,321.768	-4.321768×10^3
0.000 000 007 51	7.51×10^{-9}

Primer: (Normalized) Scientific Notation

- In decimal: $M \times 10^E$
 - E is an integer
 - Normalized form: $1 \leq |M| < 10$

M × 10^E ← Exponent

↑ ↑

Significand Base

Decimal Value	Scientific Notation
2	2×10^0
-4,321.768	-4.321768×10^3
0.000 000 007 51	7.51×10^{-9}

Primer: (Normalized) Scientific Notation

Binary Value	Scientific Notation
1110110110110	$(-1)^0 1.110110110110 \times 2^{12}$
-101.11	$(-1)^1 1.0111 \times 2^2$
0.00101	$(-1)^0 1.01 \times 2^{-3}$

Primer: (Normalized) Scientific Notation

$$(-1)^S M \times 2^E$$

Binary Value	Scientific Notation
1110110110110	$(-1)^0 1.110110110110 \times 2^{12}$
-101.11	$(-1)^1 1.0111 \times 2^2$
0.00101	$(-1)^0 1.01 \times 2^{-3}$

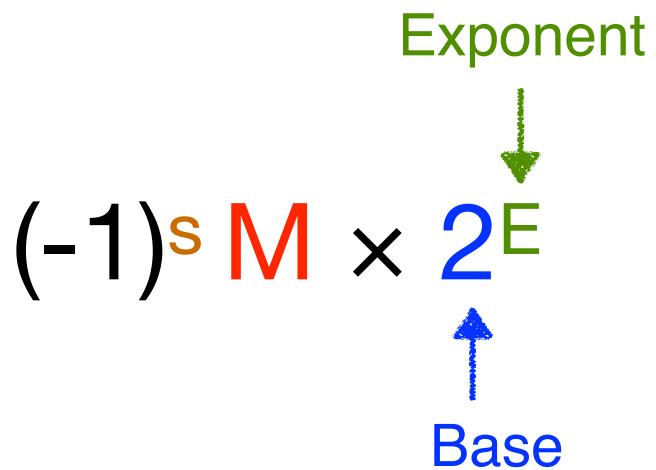
Primer: (Normalized) Scientific Notation

$$(-1)^S M \times 2^E$$

↑
Base

Binary Value	Scientific Notation
1110110110110	$(-1)^0 1.110110110110 \times 2^{12}$
-101.11	$(-1)^1 1.0111 \times 2^2$
0.00101	$(-1)^0 1.01 \times 2^{-3}$

Primer: (Normalized) Scientific Notation

(-1)^S M × 2^E

Exponent
↓
Base
↑

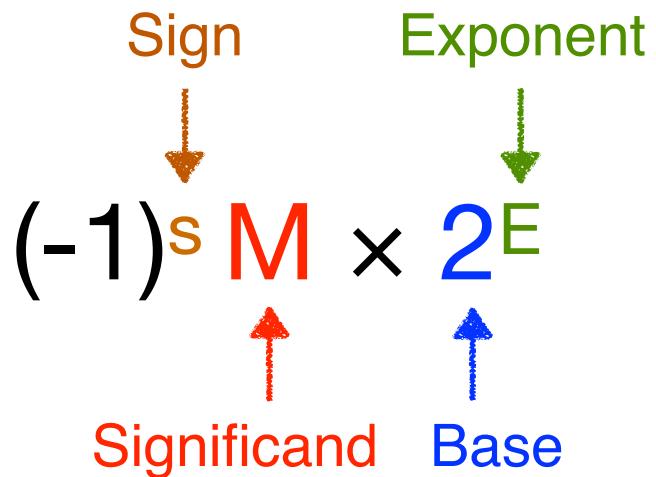
Binary Value	Scientific Notation
1110110110110	(-1) ⁰ 1.110110110110 × 2 ¹²
-101.11	(-1) ¹ 1.0111 × 2 ²
0.00101	(-1) ⁰ 1.01 × 2 ⁻³

Primer: (Normalized) Scientific Notation

Exponent
↓
 $(-1)^s M \times 2^E$
↑ ↑
Significand **Base**

Binary Value	Scientific Notation
1110110110110	$(-1)^0 1.110110110110 \times 2^{12}$
-101.11	$(-1)^1 1.0111 \times 2^2$
0.00101	$(-1)^0 1.01 \times 2^{-3}$

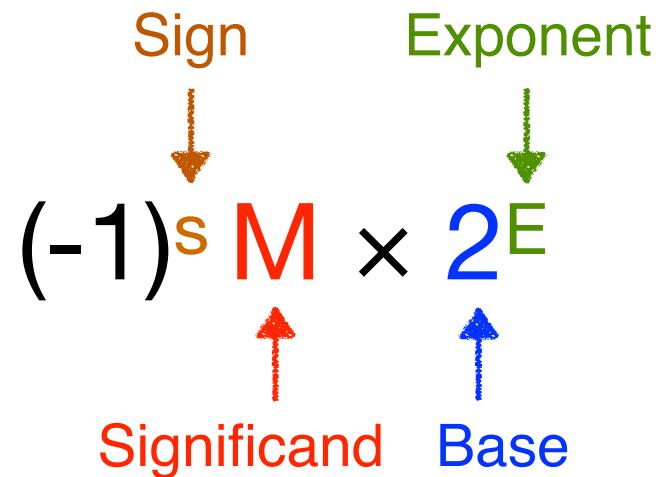
Primer: (Normalized) Scientific Notation



Binary Value	Scientific Notation
1110110110110	$(-1)^0 1.110110110110 \times 2^{12}$
-101.11	$(-1)^1 1.0111 \times 2^2$
0.00101	$(-1)^0 1.01 \times 2^{-3}$

Primer: (Normalized) Scientific Notation

- In binary: $(-1)^s M 2^E$
- Normalized form:
 - $1 \leq M < 2$
 - $M = 1.b_0b_1b_2b_3\dots$

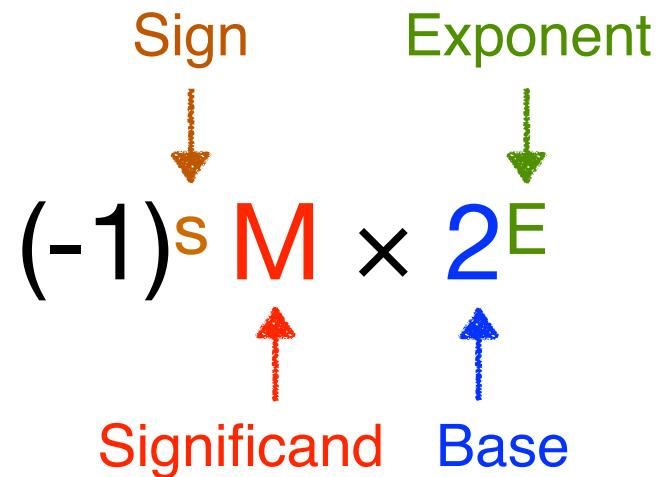


Binary Value	Scientific Notation
1110110110110	$(-1)^0 1.110110110110 \times 2^{12}$
-101.11	$(-1)^1 1.0111 \times 2^2$
0.00101	$(-1)^0 1.01 \times 2^{-3}$

Primer: (Normalized) Scientific Notation

- In binary: $(-1)^s M 2^E$
- Normalized form:
 - $1 \leq M < 2$
 - $M = 1.b_0b_1b_2b_3\dots$

Fraction



Binary Value	Scientific Notation
1110110110110	$(-1)^0 1.110110110110 \times 2^{12}$
-101.11	$(-1)^1 1.0111 \times 2^2$
0.00101	$(-1)^0 1.01 \times 2^{-3}$

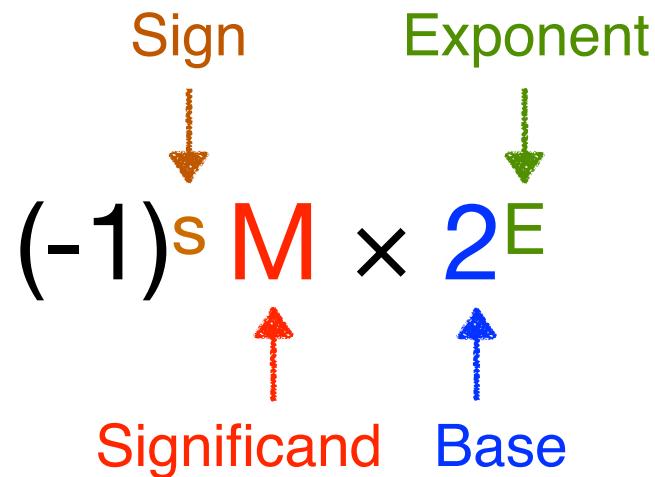
Primer: (Normalized) Scientific Notation

- In binary: $(-1)^s M 2^E$

- Normalized form:

- $1 \leq M < 2$
- $M = 1.b_0b_1b_2b_3\dots$

Fraction

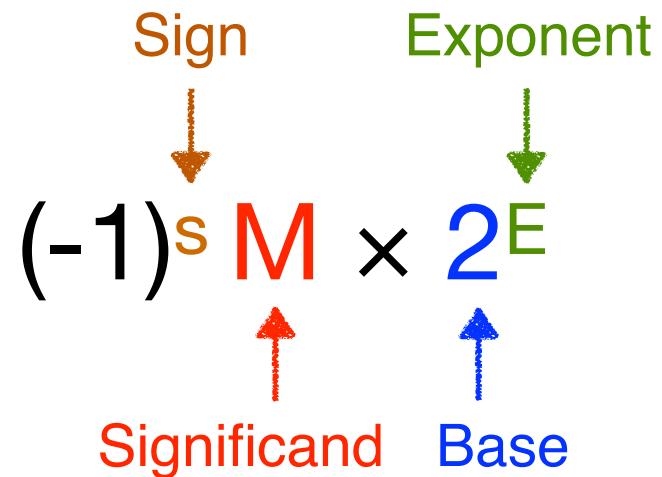


- If I tell you that there is a number where:

- Fraction = 0101
- $s = 1$
- $E = 10$
- You could reconstruct the number as $(-1)^1 1.0101 \times 2^{10}$

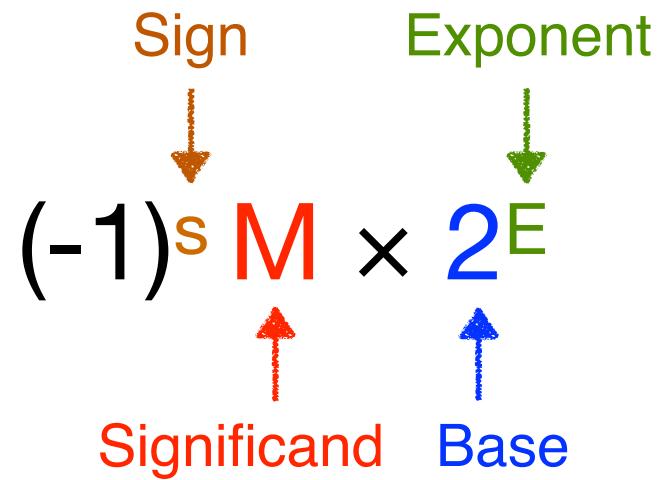
Primer: Floating Point Representation

- In binary: $(-1)^s M 2^E$
 - Normalized form:
 - $1 \leq M < 2$
 - $M = 1.b_0b_1b_2b_3\dots$
- Fraction



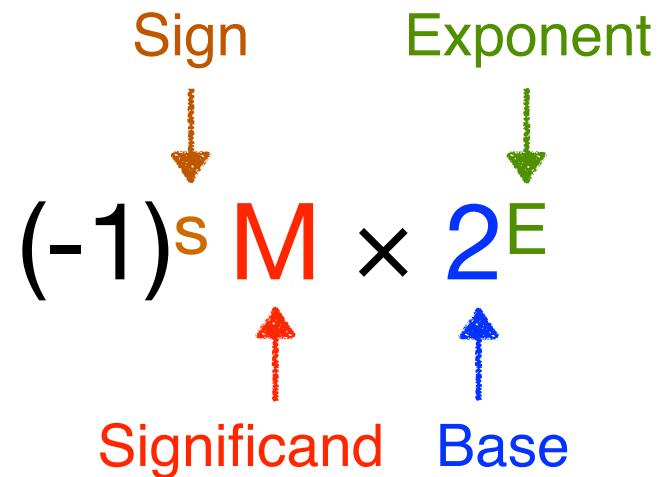
Primer: Floating Point Representation

- In binary: $(-1)^s M 2^E$
- Normalized form:
 - $1 \leq M < 2$
 - $M = 1.b_0b_1b_2b_3\dots$
Fraction
- Encoding



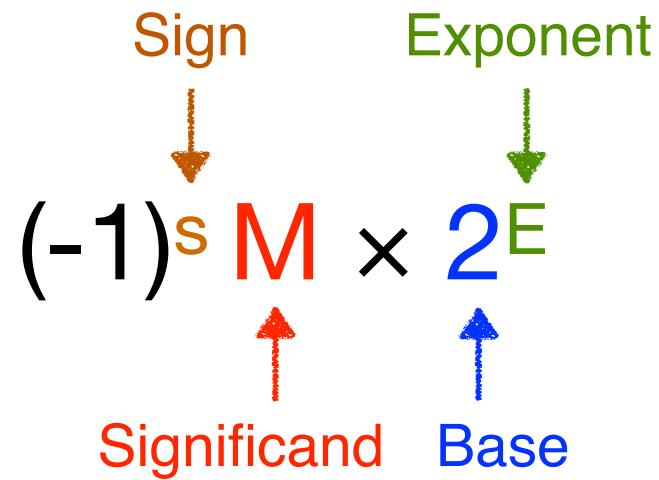
Primer: Floating Point Representation

- In binary: $(-1)^s M 2^E$
- Normalized form:
 - $1 \leq M < 2$
 - $M = 1.b_0b_1b_2b_3\dots$
Fraction
- Encoding



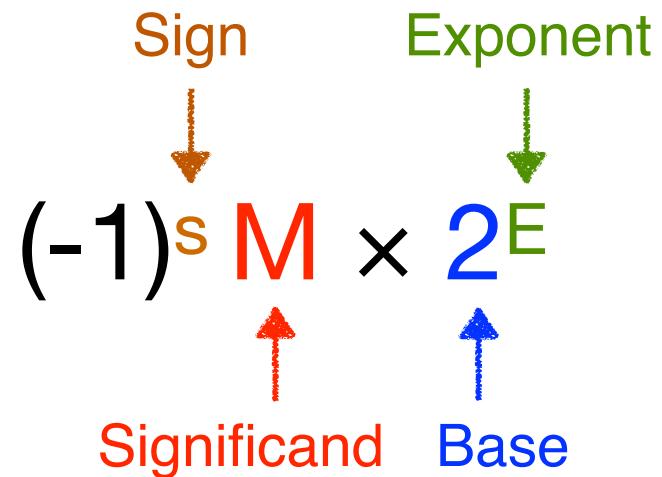
Primer: Floating Point Representation

- In binary: $(-1)^s M 2^E$
- Normalized form:
 - $1 \leq M < 2$
 - $M = 1.b_0b_1b_2b_3\dots$
Fraction
- Encoding
 - MSB s is sign bit s



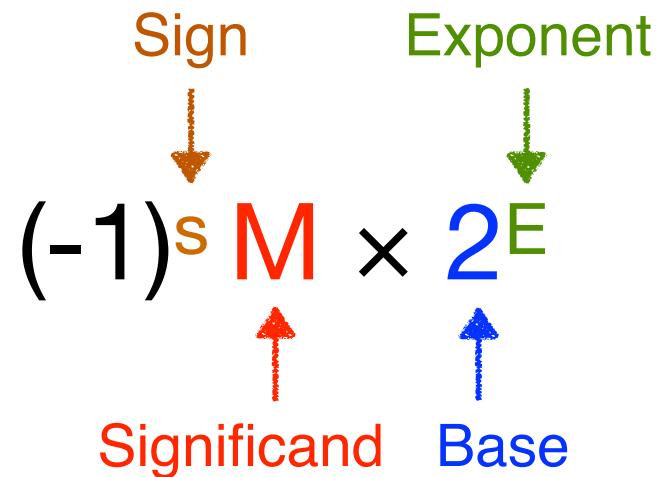
Primer: Floating Point Representation

- In binary: $(-1)^s M 2^E$
- Normalized form:
 - $1 \leq M < 2$
 - $M = 1.b_0b_1b_2b_3\dots$
Fraction
- Encoding
 - MSB s is sign bit s
 - exp field encodes Exponent (but not exactly the same, more later)



Primer: Floating Point Representation

- In binary: $(-1)^s M 2^E$
- Normalized form:
 - $1 \leq M < 2$
 - $M = 1.b_0b_1b_2b_3\dots$
- Encoding
 - MSB s is sign bit s
 - exp field encodes Exponent (but not exactly the same, more later)
 - $frac$ field encodes Fraction (but not exactly the same, more later)

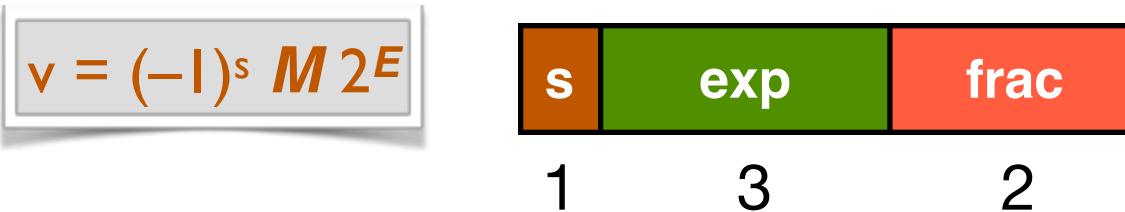


6-bit Floating Point Example

$$v = (-1)^s M \cdot 2^E$$

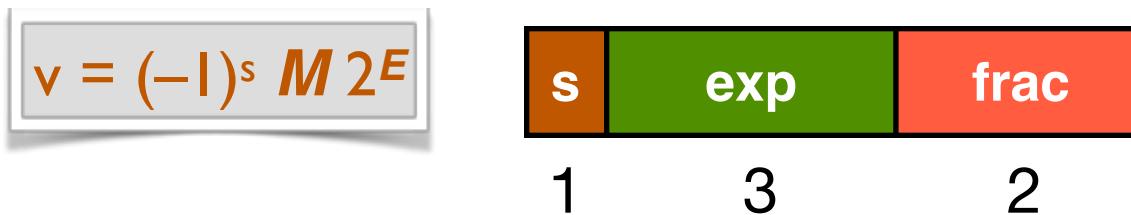


6-bit Floating Point Example



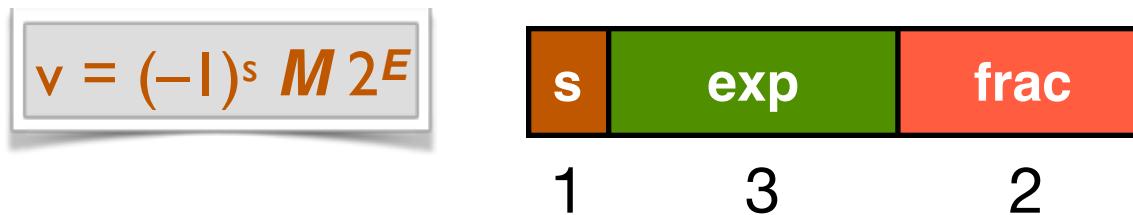
- exp has 3 bits, interpreted as an unsigned value

6-bit Floating Point Example



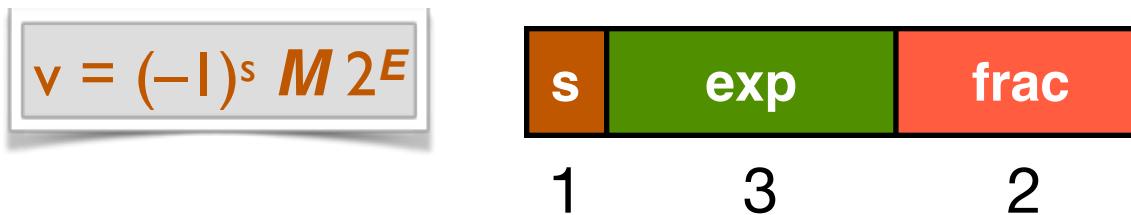
- exp has 3 bits, interpreted as an unsigned value
 - If exp were E , we could represent exponents from **0 to 7**

6-bit Floating Point Example



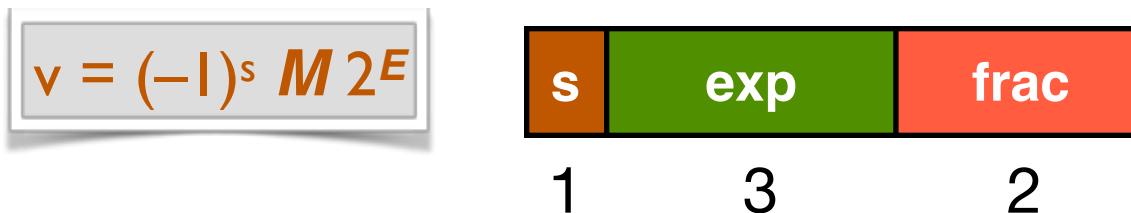
- *exp* has 3 bits, interpreted as an unsigned value
 - If *exp* were *E*, we could represent exponents from **0 to 7**
 - How about negative exponent?

6-bit Floating Point Example



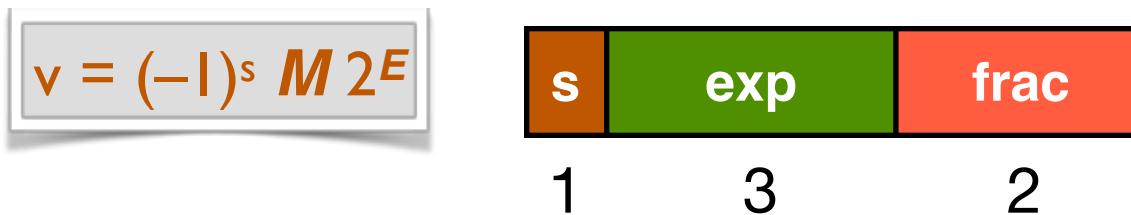
- exp has 3 bits, interpreted as an unsigned value
 - If exp were E , we could represent exponents from **0 to 7**
 - How about negative exponent?
 - Add a bias term: $E = exp - bias$ (i.e., $exp = E + bias$)

6-bit Floating Point Example



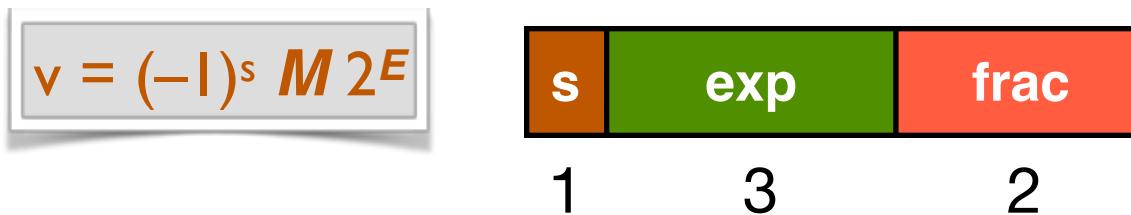
- **exp** has 3 bits, interpreted as an unsigned value
 - If **exp** were E , we could represent exponents from **0 to 7**
 - How about negative exponent?
 - Add a bias term: $E = \text{exp} - \text{bias}$ (i.e., $\text{exp} = E + \text{bias}$)
 - bias is always $2^{k-1} - 1$, where k is number of exponent bits

6-bit Floating Point Example



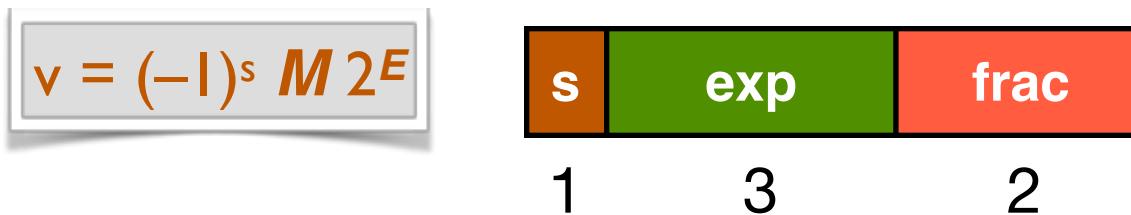
- **exp** has 3 bits, interpreted as an unsigned value
 - If **exp** were **E**, we could represent exponents from **0 to 7**
 - How about negative exponent?
 - Add a bias term: $E = exp - bias$ (i.e., $exp = E + bias$)
 - bias is always $2^{k-1} - 1$, where k is number of exponent bits
- Example when we use 3 bits for **exp** (i.e., $k = 3$):

6-bit Floating Point Example



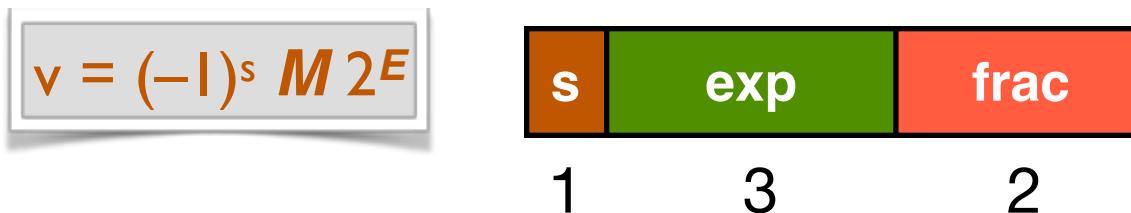
- exp has 3 bits, interpreted as an unsigned value
 - If exp were E , we could represent exponents from **0 to 7**
 - How about negative exponent?
 - Add a bias term: $E = \text{exp} - \text{bias}$ (i.e., $\text{exp} = E + \text{bias}$)
 - bias is always $2^{k-1} - 1$, where k is number of exponent bits
- Example when we use 3 bits for exp (i.e., $k = 3$):
 - bias = 3

6-bit Floating Point Example



- **exp** has 3 bits, interpreted as an unsigned value
 - If **exp** were **E**, we could represent exponents from **0 to 7**
 - How about negative exponent?
 - Add a bias term: $E = \text{exp} - \text{bias}$ (i.e., $\text{exp} = E + \text{bias}$)
 - bias is always $2^{k-1} - 1$, where k is number of exponent bits
- Example when we use 3 bits for **exp** (i.e., $k = 3$):
 - bias = 3
 - If $E = -2$, **exp** is 1 (001_2)

6-bit Floating Point Example



- exp has 3 bits, interpreted as an unsigned value
 - If exp were E , we could represent exponents from **0 to 7**
 - How about negative exponent?
 - Add a bias term: $E = \text{exp} - \text{bias}$ (i.e., $\text{exp} = E + \text{bias}$)
 - bias is always $2^{k-1} - 1$, where k is number of exponent bits
- Example when we use 3 bits for exp (i.e., $k = 3$):
 - bias = 3
 - If $E = -2$, exp is 1 (001_2)

E	exp
-3	000
-2	001
-1	010
0	011
1	100
2	101
3	110
4	111

6-bit Floating Point Example

$$v = (-1)^s M 2^E$$



1 3 2

- **exp** has 3 bits, interpreted as an unsigned value
 - If **exp** were E , we could represent exponents from **0 to 7**
 - How about negative exponent?
 - Add a bias term: $E = \text{exp} - \text{bias}$ (i.e., $\text{exp} = E + \text{bias}$)
 - bias is always $2^{k-1} - 1$, where k is number of exponent bits
- Example when we use 3 bits for exp (i.e., $k = 3$):
 - bias = 3
 - If $E = -2$, exp is 1 (001_2)
 - Reserve 000 and 111 for other purposes (more on this later)

E	exp
-3	000
-2	001
-1	010
0	011
1	100
2	101
3	110
4	111

6-bit Floating Point Example

$$v = (-1)^s M 2^E$$



- **exp** has 3 bits, interpreted as an unsigned value
 - If **exp** were E , we could represent exponents from **0 to 7**
 - How about negative exponent?
 - Add a bias term: $E = \text{exp} - \text{bias}$ (i.e., $\text{exp} = E + \text{bias}$)
 - bias is always $2^{k-1} - 1$, where k is number of exponent bits
- Example when we use 3 bits for **exp** (i.e., $k = 3$):
 - bias = 3
 - If $E = -2$, **exp** is 1 (001_2)
 - Reserve 000 and 111 for other purposes (more on this later)
 - We can now represent exponents from **-2 (exp 001)** to **3 (exp 110)**

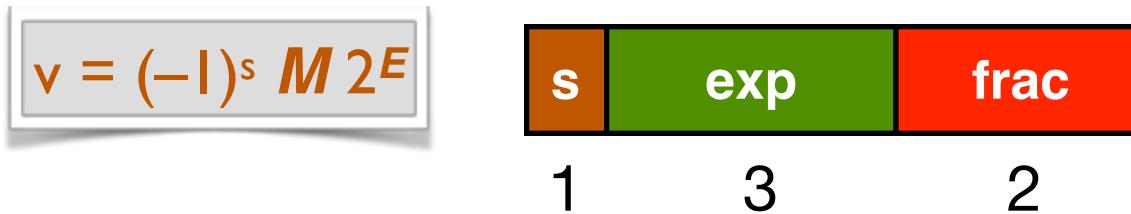
E	exp
-3	000
-2	001
-1	010
0	011
1	100
2	101
3	110
4	111

6-bit Floating Point Example

$$v = (-1)^s M \cdot 2^E$$

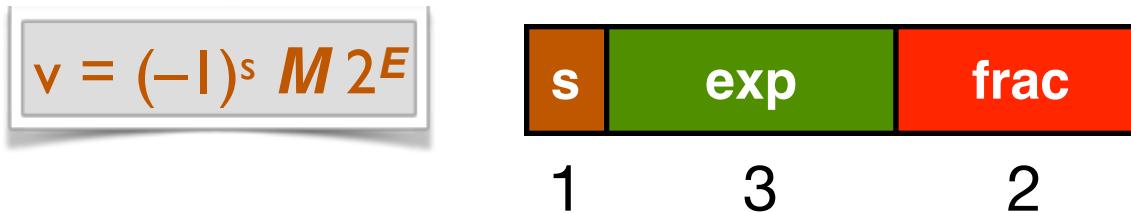


6-bit Floating Point Example



- *frac* has 2 bits, append them after “1.” to form M
 - *frac* = 10 implies M = 1.10

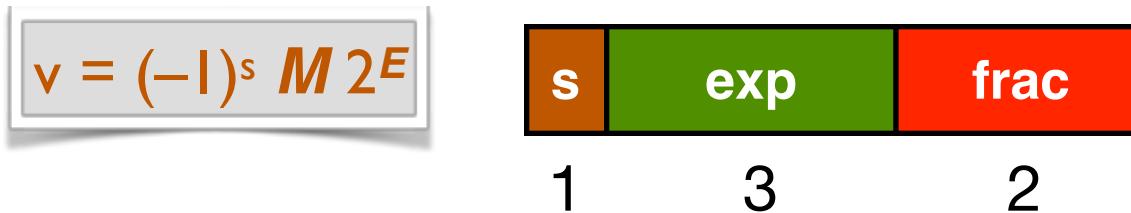
6-bit Floating Point Example



- *frac* has 2 bits, append them after “1.” to form M
 - *frac* = 10 implies M = 1.10
- Putting it Together: An Example:

$$-10.1_2 = (-1)^1 \ 1.01 \times 2^1$$

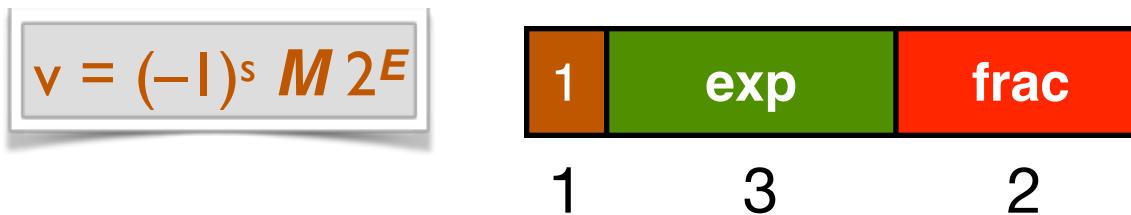
6-bit Floating Point Example



- *frac* has 2 bits, append them after “1.” to form M
 - *frac* = 10 implies M = 1.10
- Putting it Together: An Example:

$$-10.1_2 = (-1)^1 \downarrow 1.01 \times 2^1$$

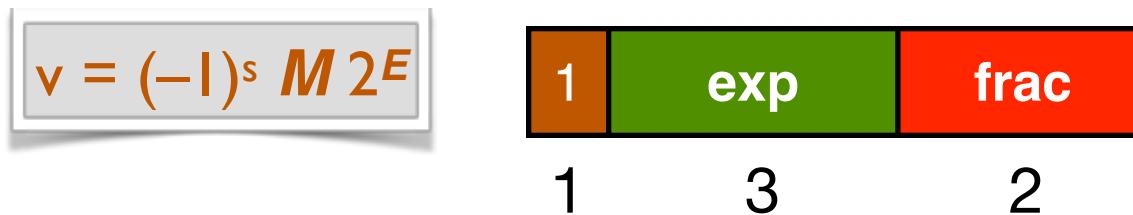
6-bit Floating Point Example



- *frac* has 2 bits, append them after “1.” to form M
 - *frac* = 10 implies M = 1.10
- Putting it Together: An Example:

$$-10.1_2 = (-1)^1 \downarrow 1.01 \times 2^1$$

6-bit Floating Point Example



- *frac* has 2 bits, append them after “1.” to form M
 - *frac* = 10 implies M = 1.10
- Putting it Together: An Example:

$$-10.1_2 = (-1)^1 \downarrow 1.01 \times 2^1$$

6-bit Floating Point Example

$$v = (-1)^s M 2^E$$



1 3 2

- *frac* has 2 bits, append them after “1.” to form M
 - *frac* = 10 implies M = 1.10
- Putting it Together: An Example:

$$-10.1_2 = (-1)^1 \ 1.01 \times 2^1$$



E	exp
-3	000
-2	001
-1	010
0	011
1	100
2	101
3	110
4	111

6-bit Floating Point Example

$$v = (-1)^s M 2^E$$



1 3 2

- *frac* has 2 bits, append them after “1.” to form M
 - *frac* = 10 implies M = 1.10
- Putting it Together: An Example:

$$-10.1_2 = (-1)^1 1.01 \times 2^1$$

E	exp
-3	000
-2	001
-1	010
0	011
1	100
2	101
3	110
4	111

6-bit Floating Point Example

$$v = (-1)^s M 2^E$$



1 3 2

- *frac* has 2 bits, append them after “1.” to form M
 - *frac* = 10 implies M = 1.10
- Putting it Together: An Example:

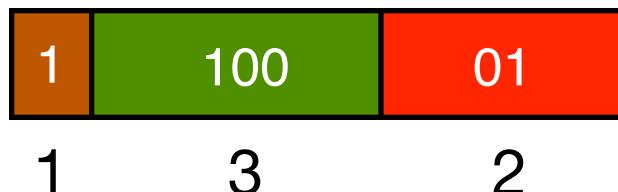


$$-10.1_2 = (-1)^1 1.01 \times 2^1$$

E	exp
-3	000
-2	001
-1	010
0	011
1	100
2	101
3	110
4	111

6-bit Floating Point Example

$$v = (-1)^s M 2^E$$



- *frac* has 2 bits, append them after “1.” to form M
 - *frac* = 10 implies M = 1.10
- Putting it Together: An Example:



$$-10.1_2 = (-1)^1 1.01 \times 2^1$$

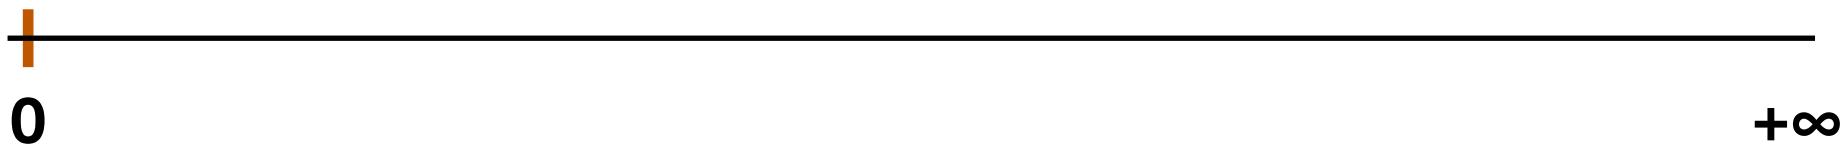
E	exp
-3	000
-2	001
-1	010
0	011
1	100
2	101
3	110
4	111

Representable Numbers (Positive Only)

$$v = (-1)^s M \cdot 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111



Representable Numbers (Positive Only)

$$v = (-1)^s M \cdot 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111



Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111



Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

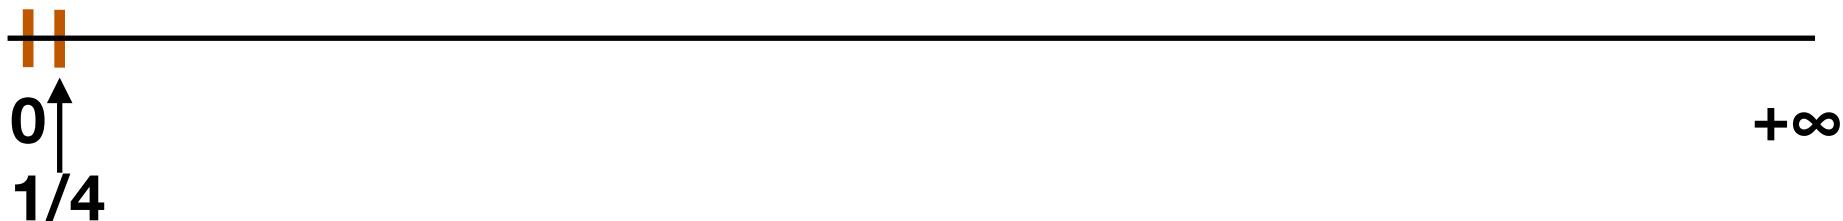


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

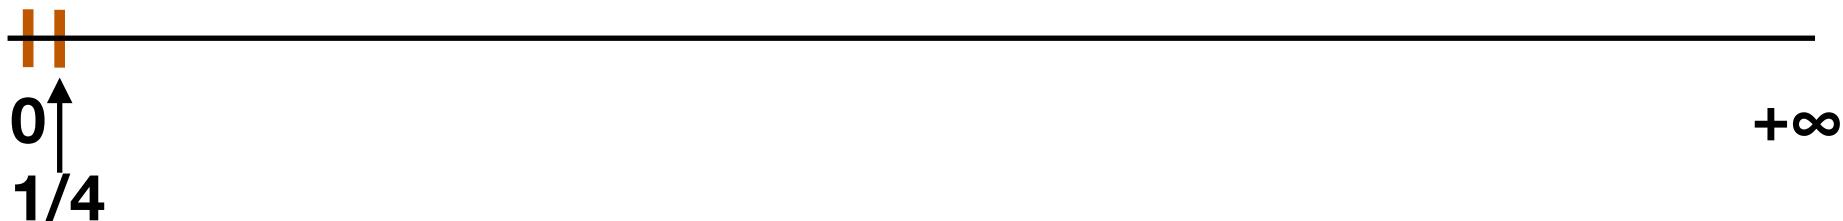


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

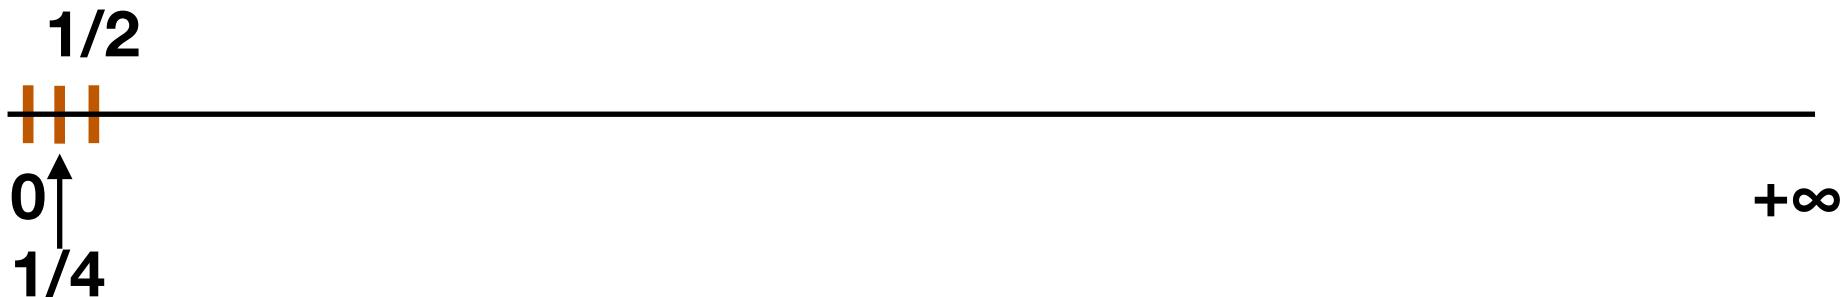


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

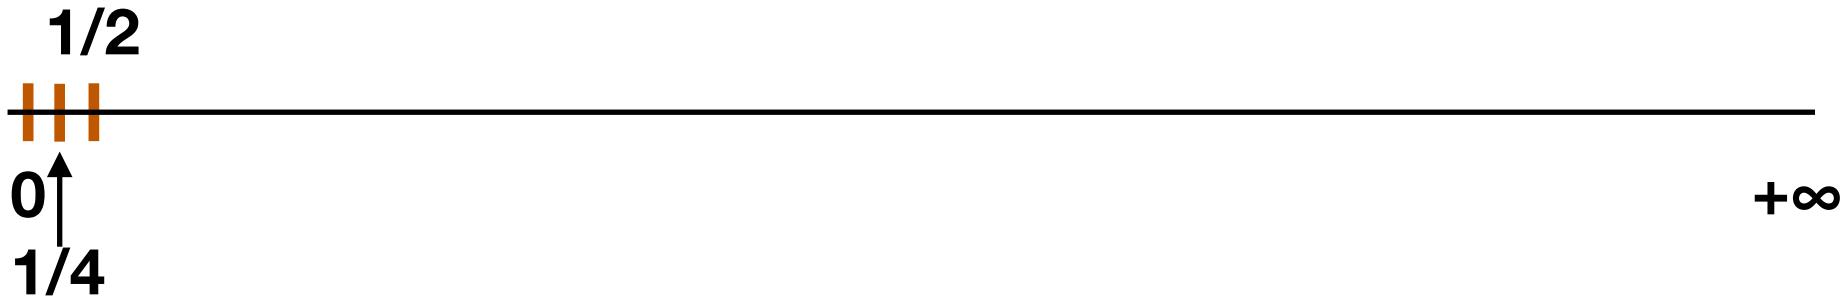


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

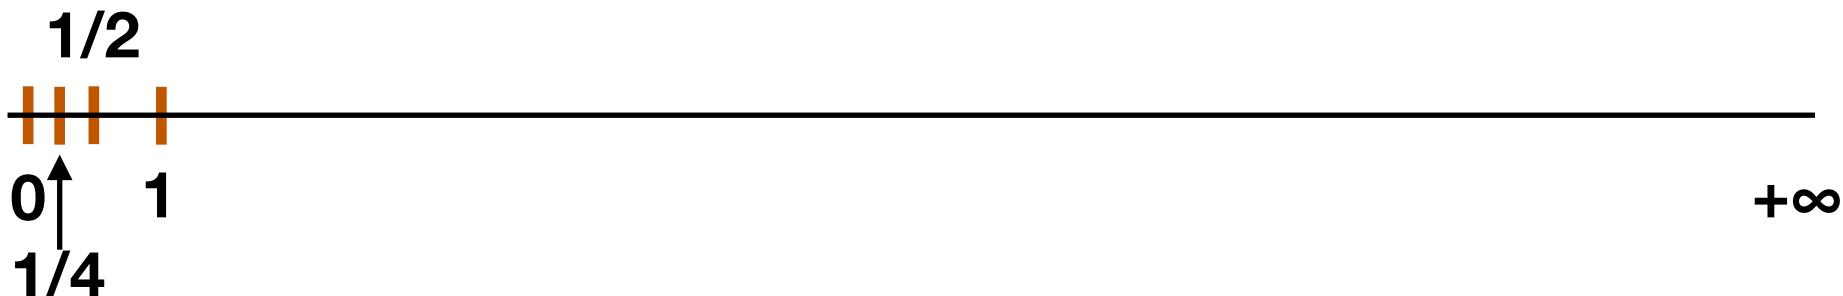


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

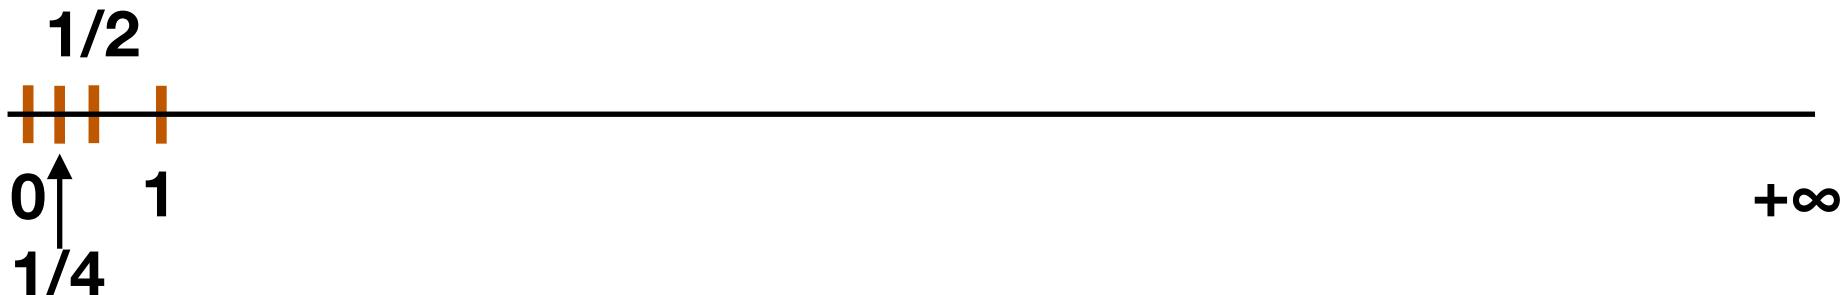


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

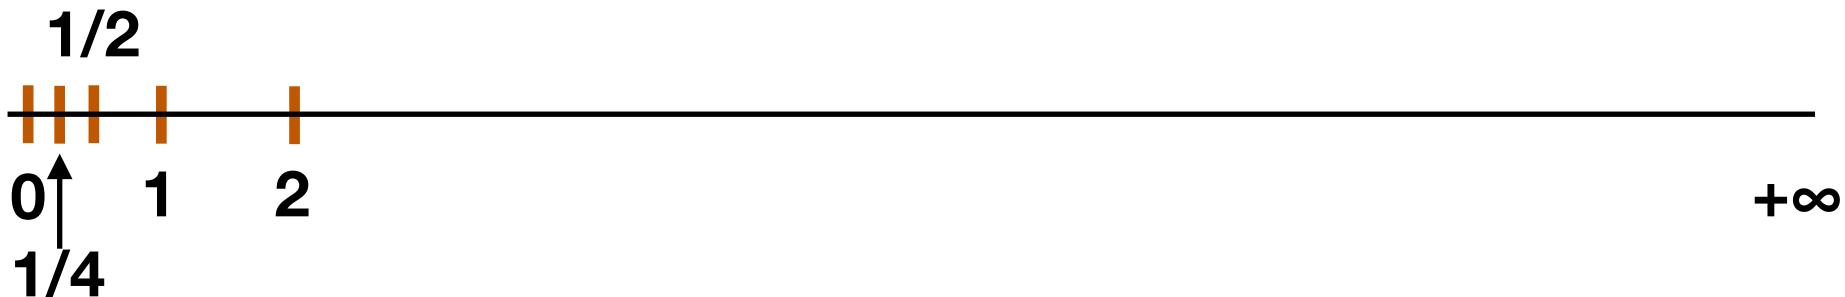


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

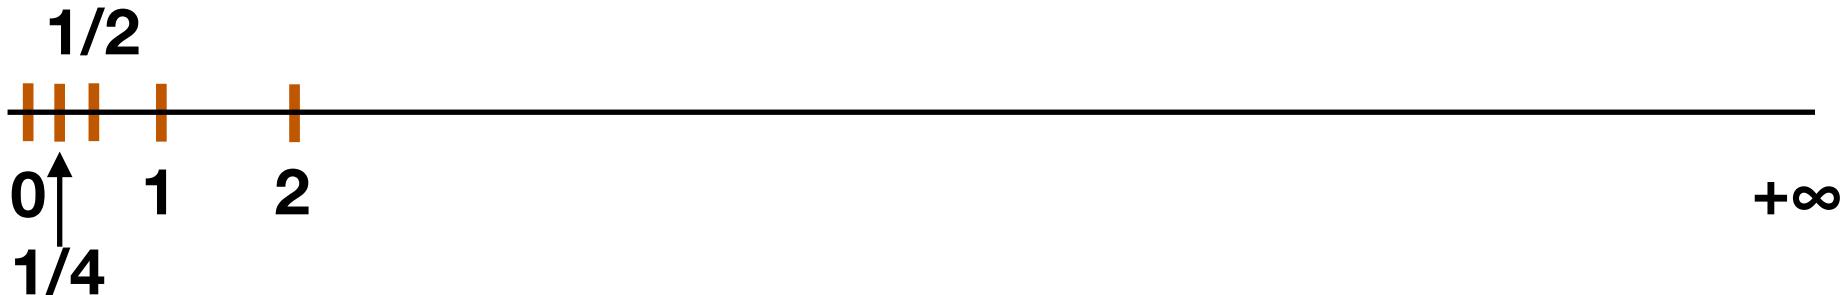


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

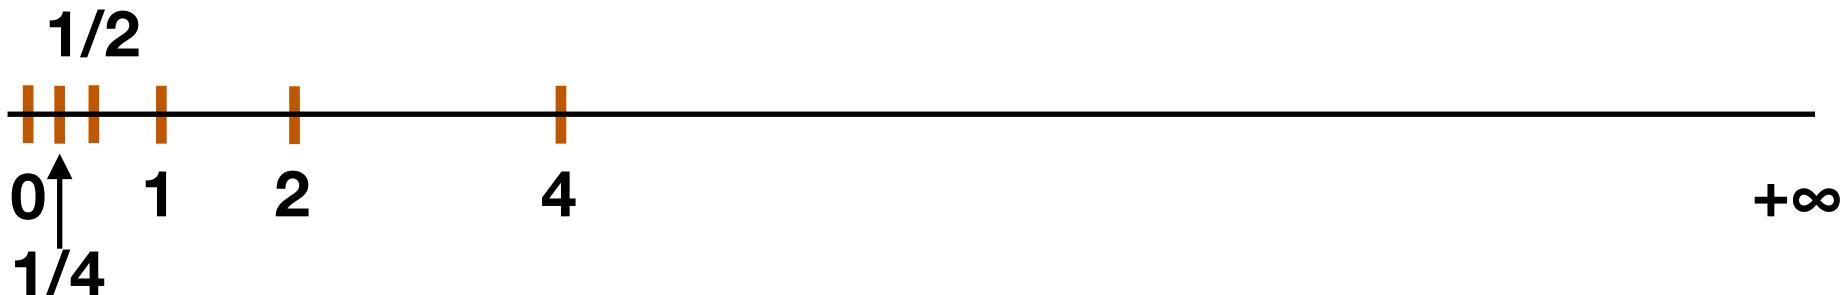


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

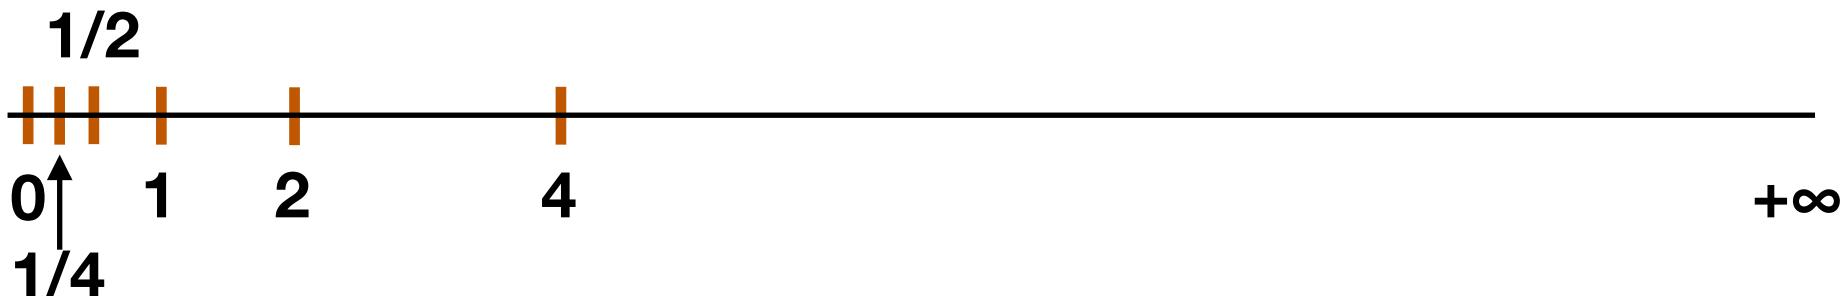


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

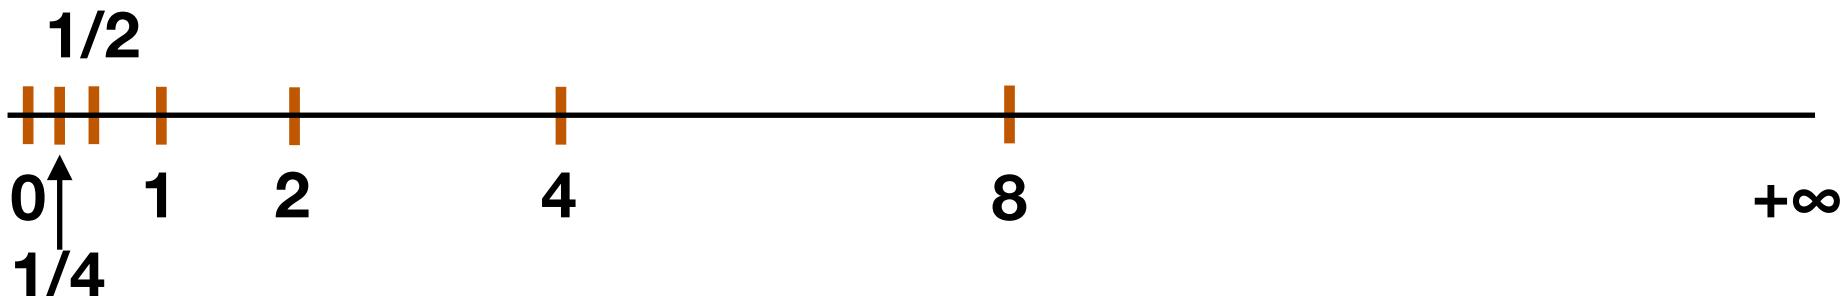


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

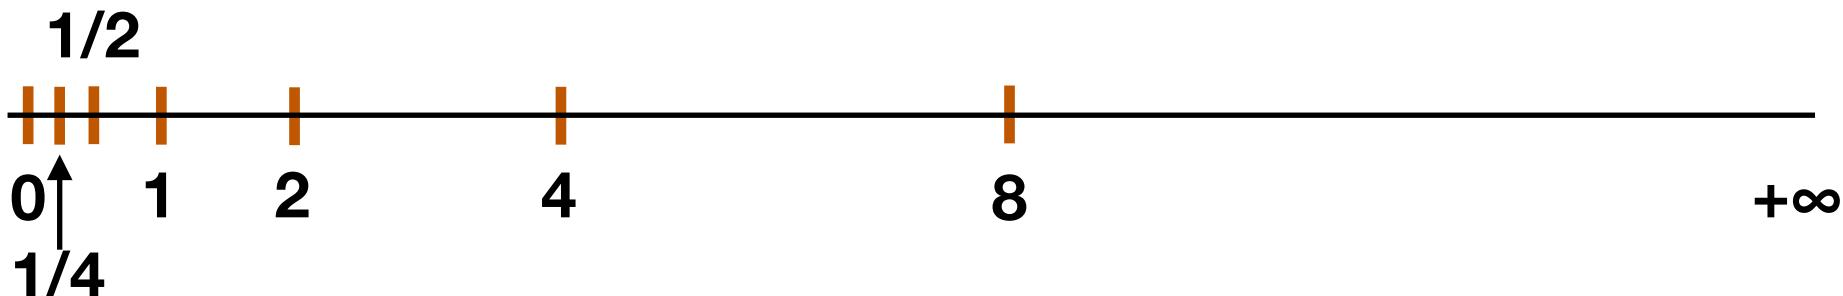


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

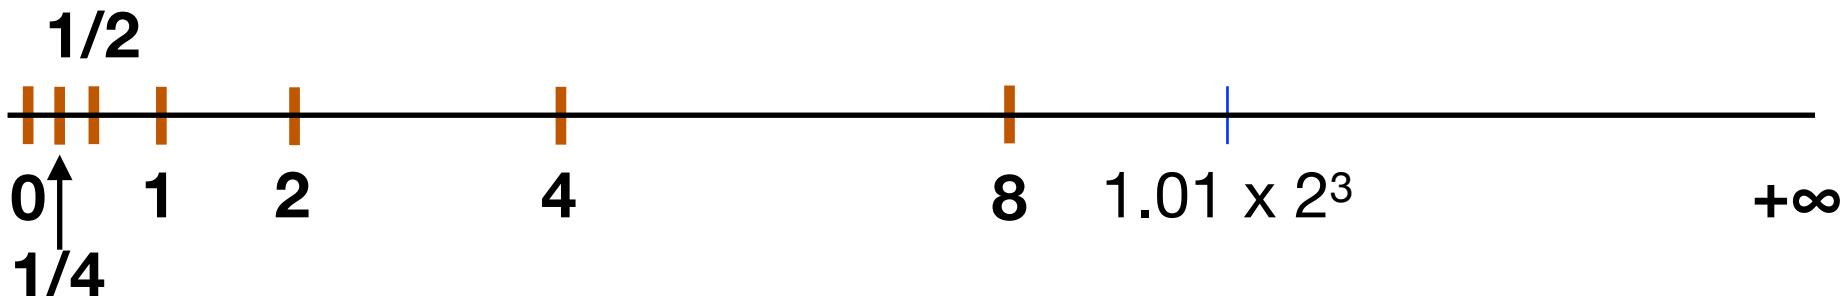


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

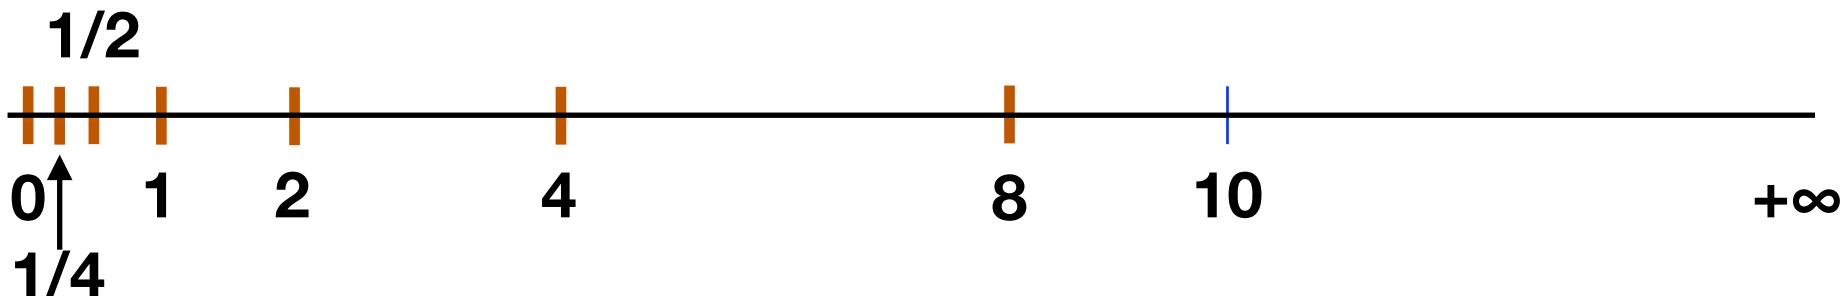


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

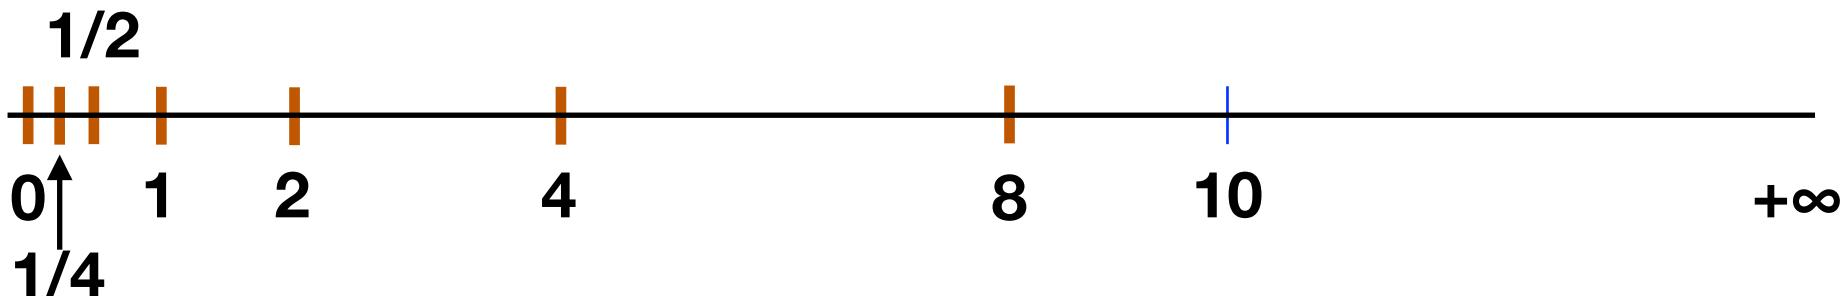


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

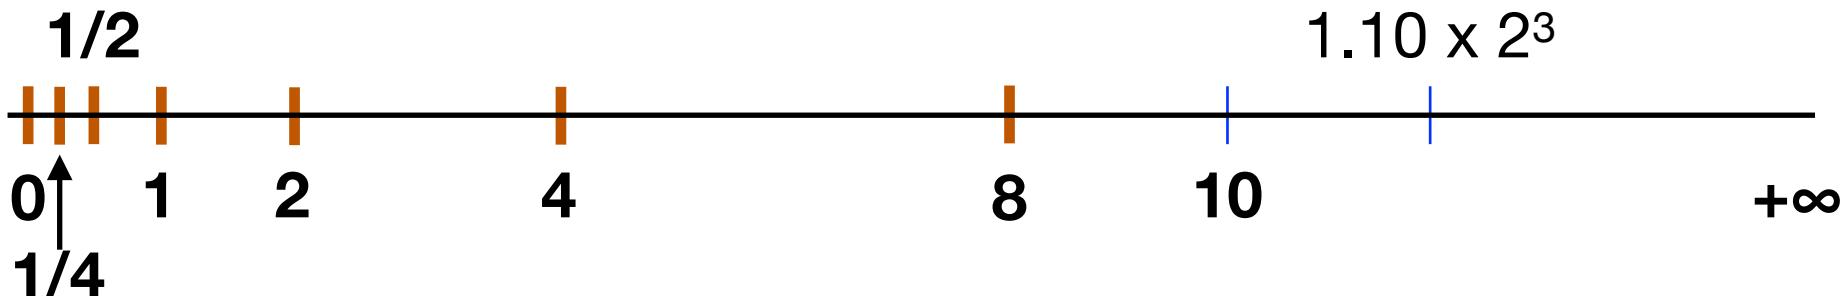


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

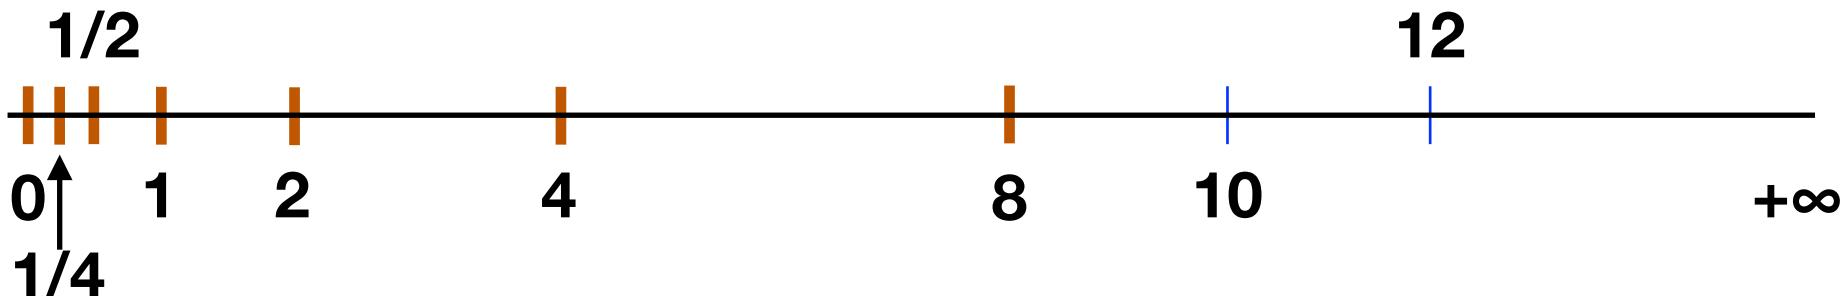


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

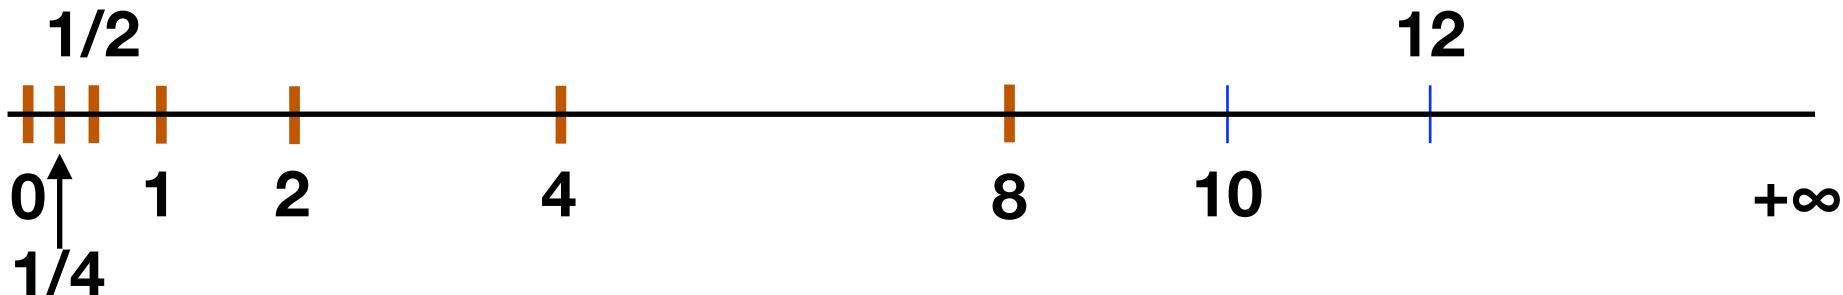


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

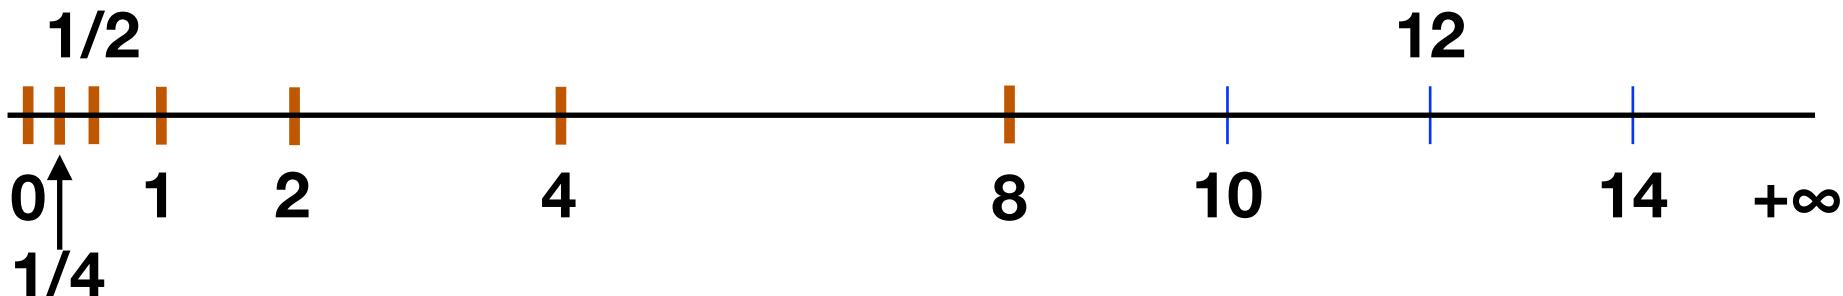


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$

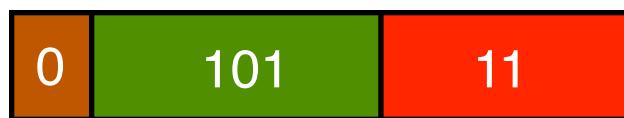


E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

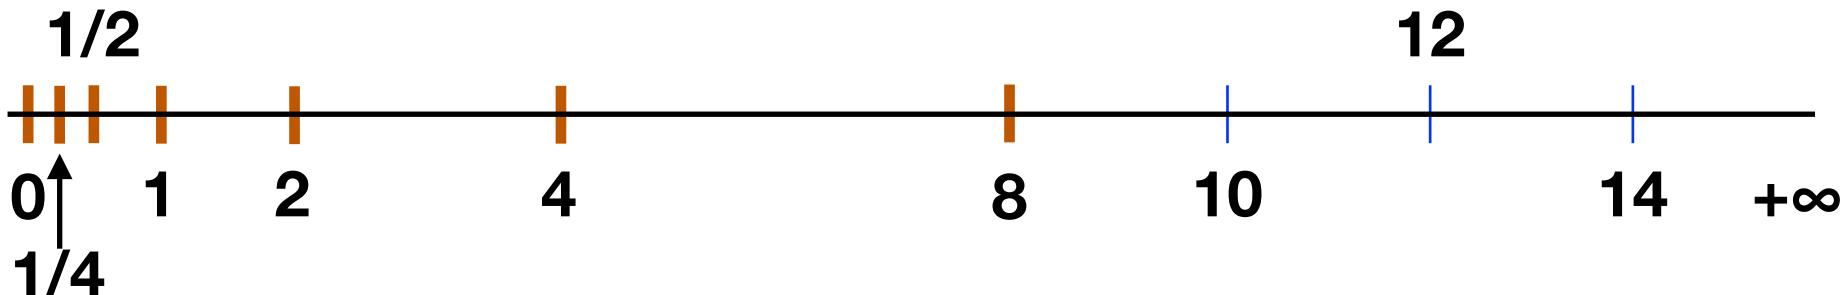


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

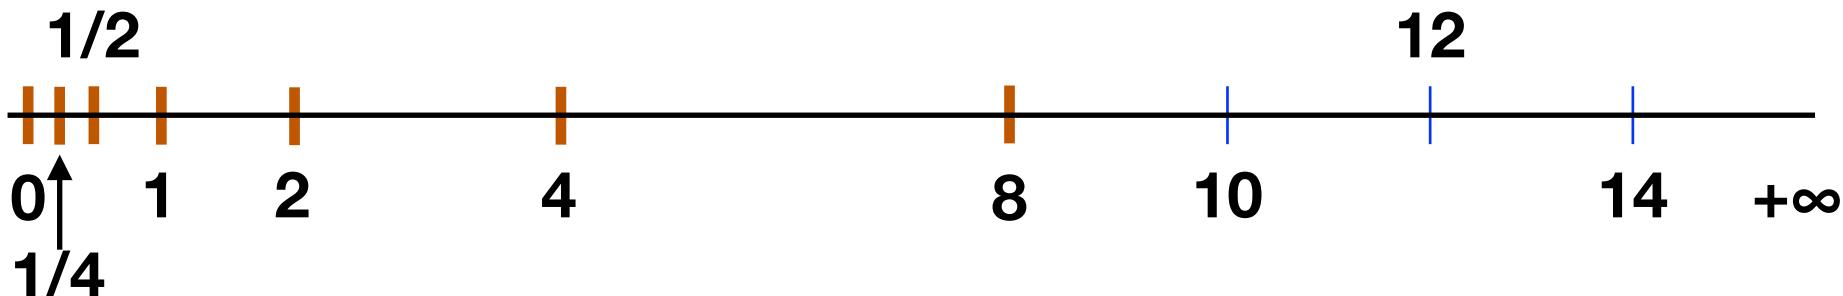


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

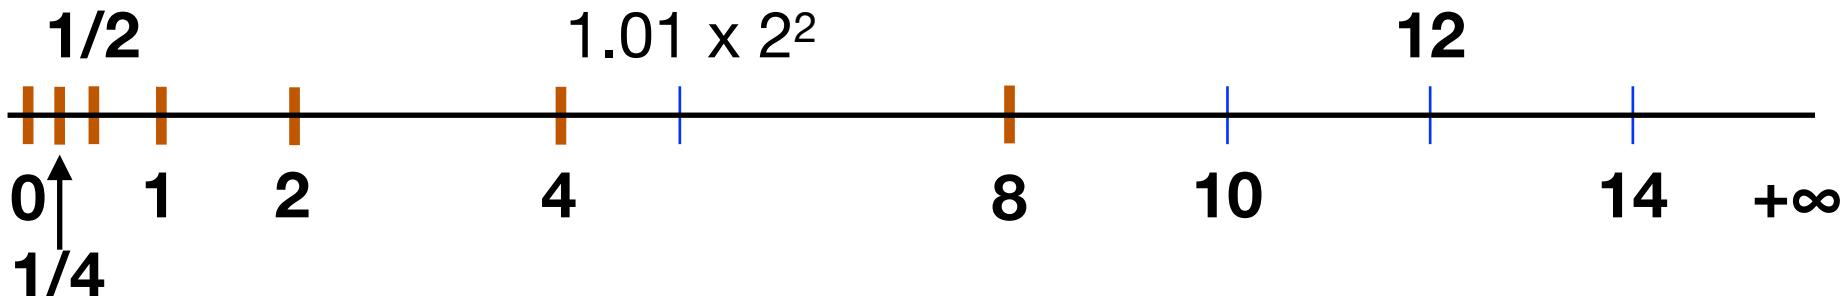


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$

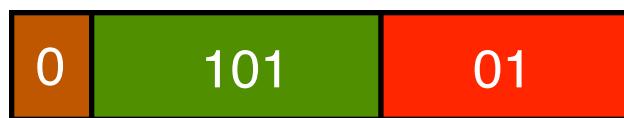


E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

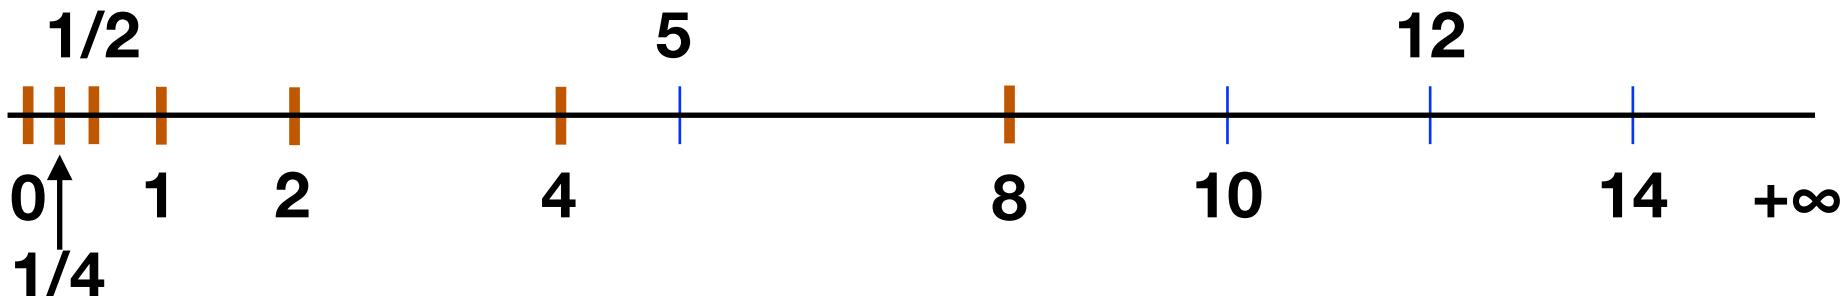


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$

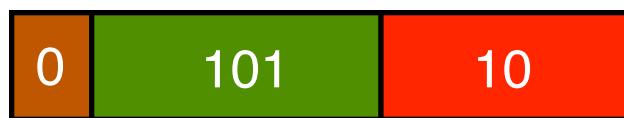


E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

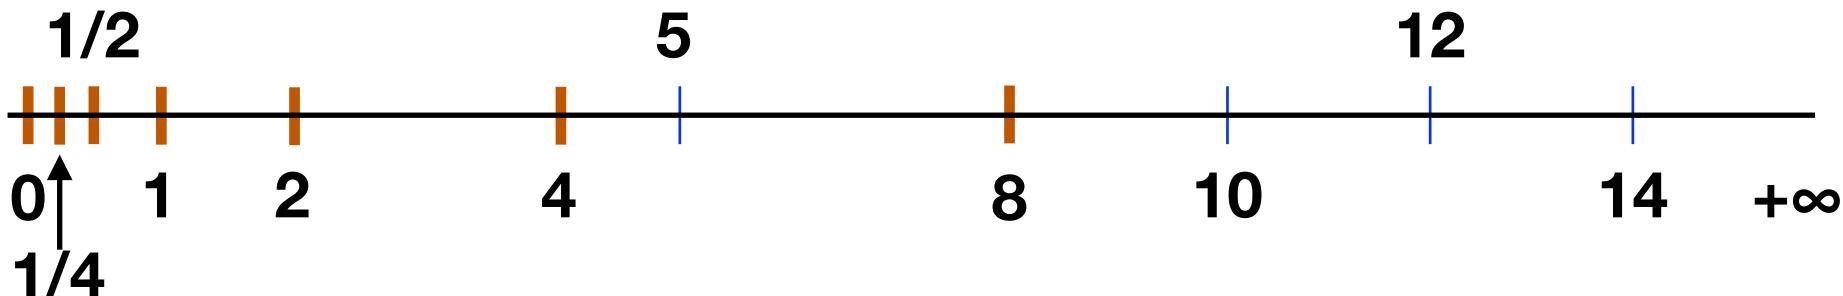


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$

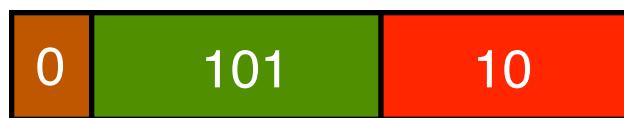


E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

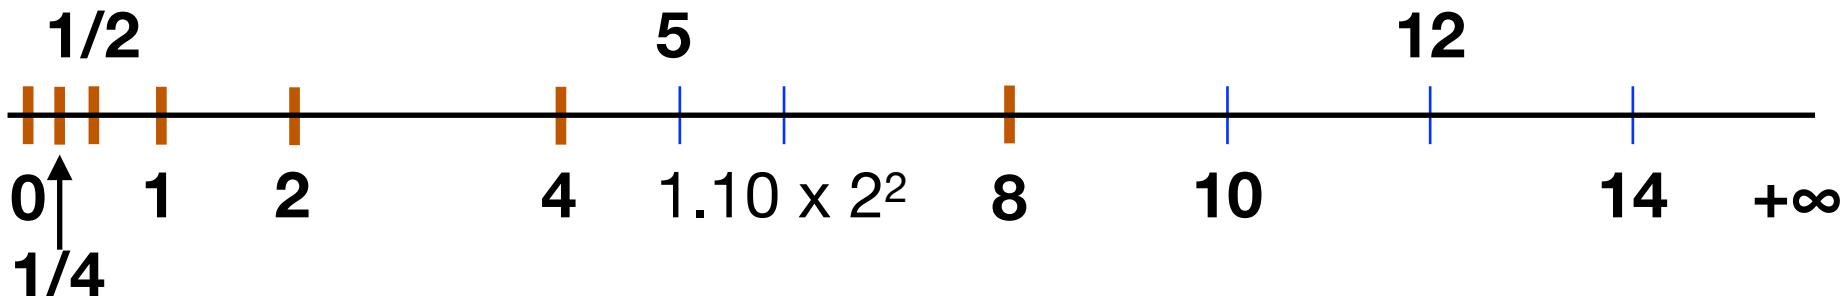


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

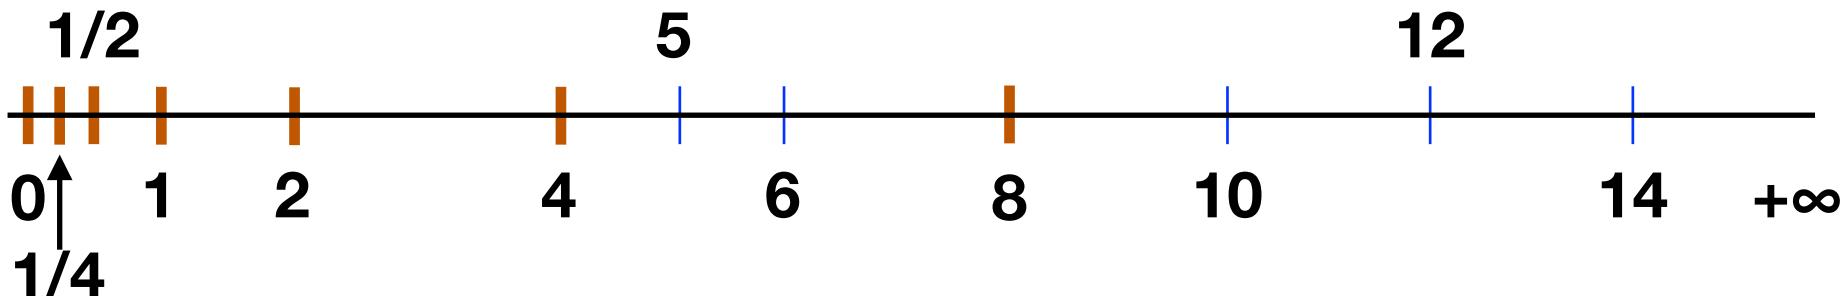


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$

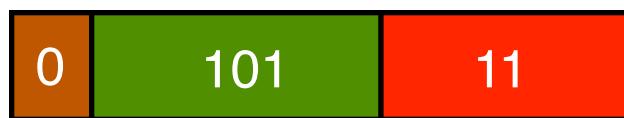


E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

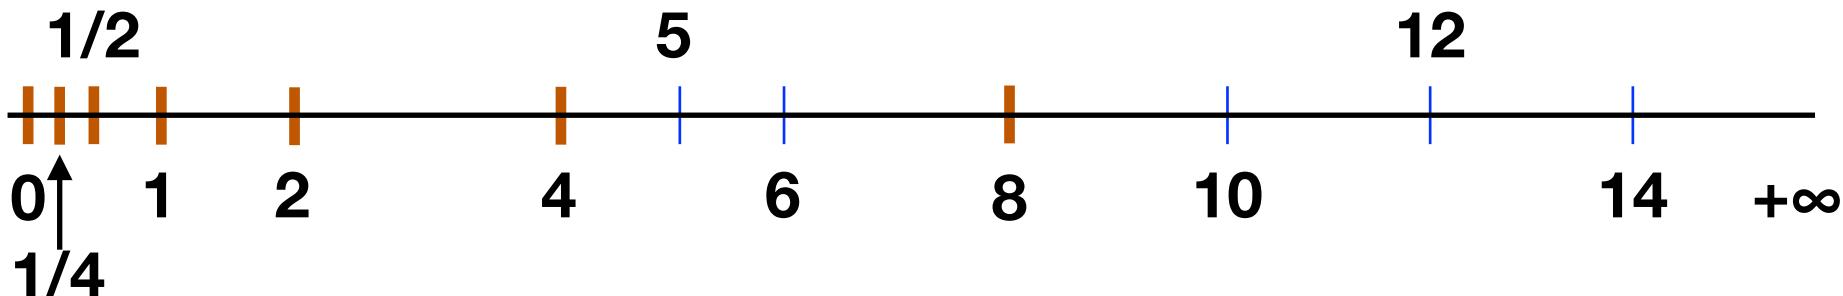


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$

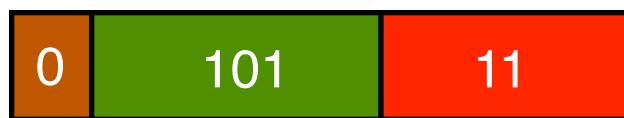


E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

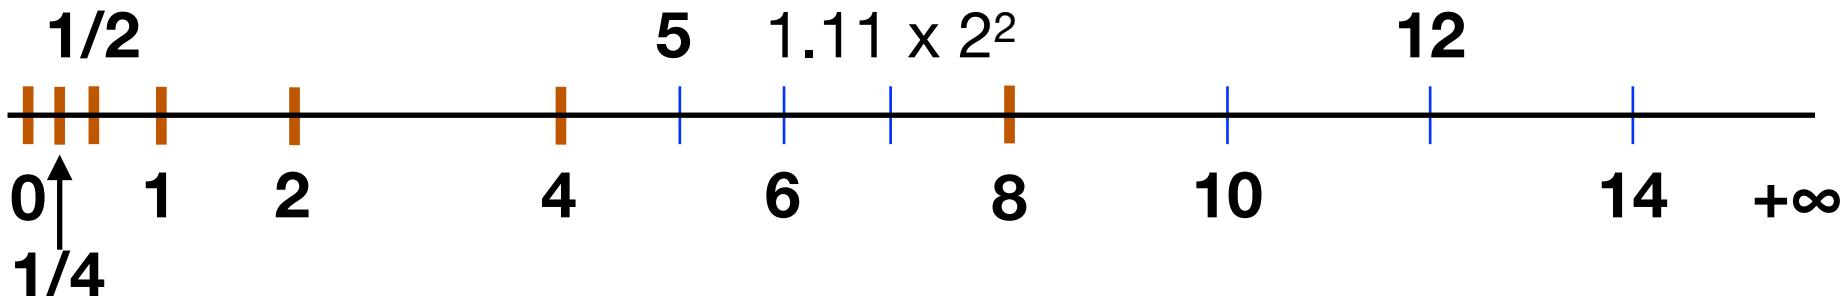


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

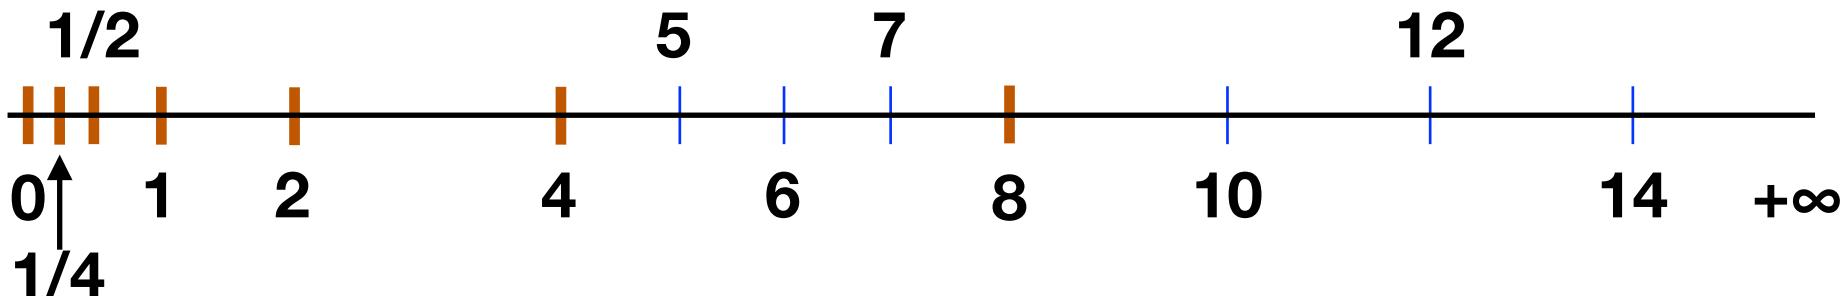


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

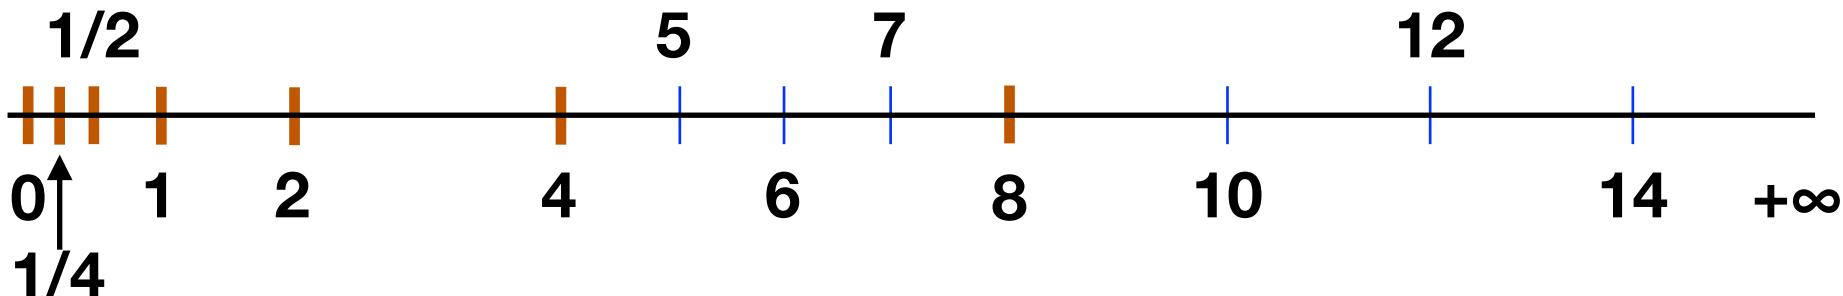


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

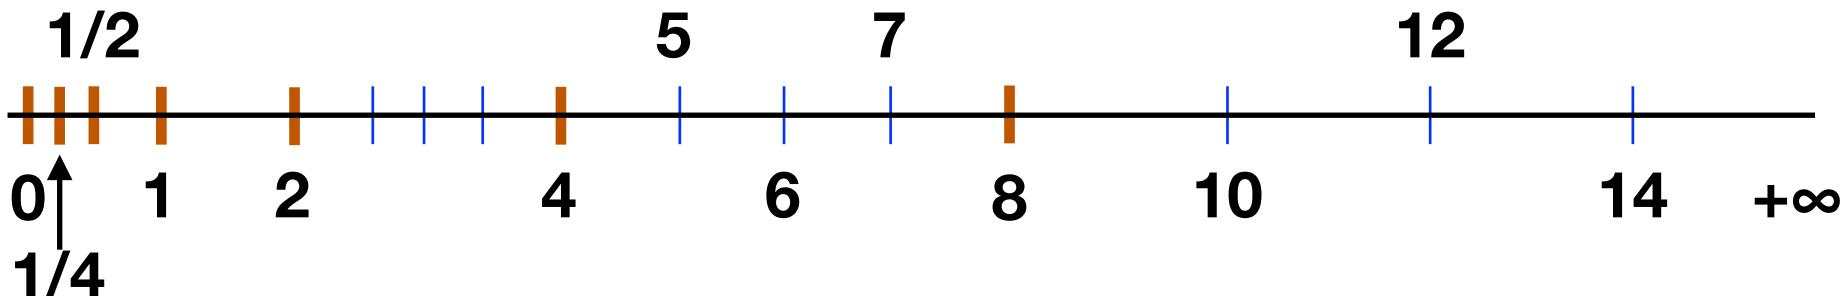


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

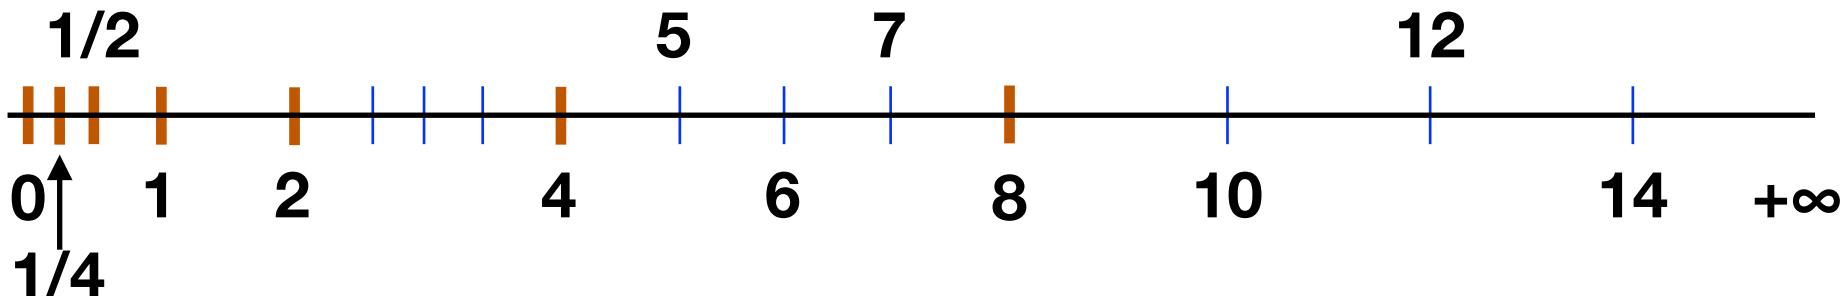


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

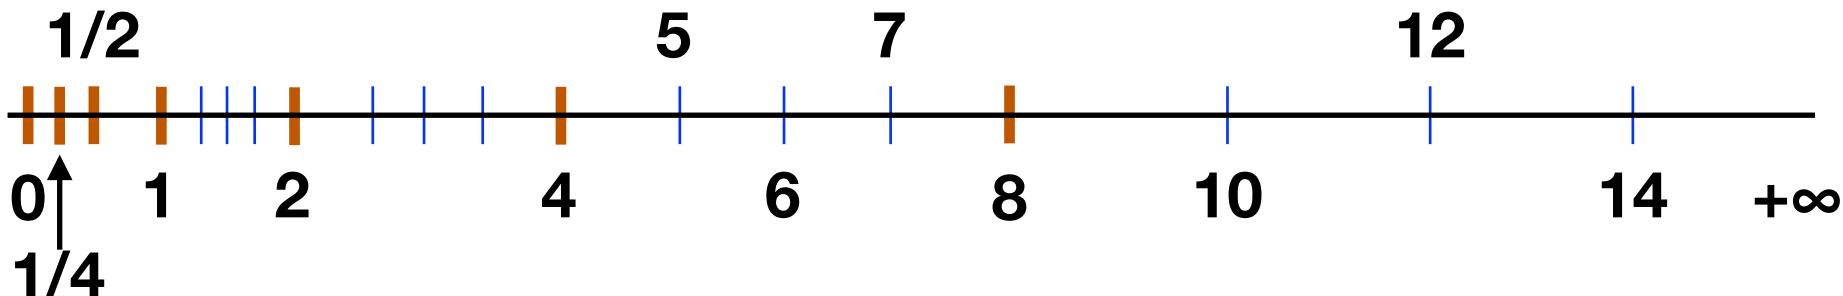


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

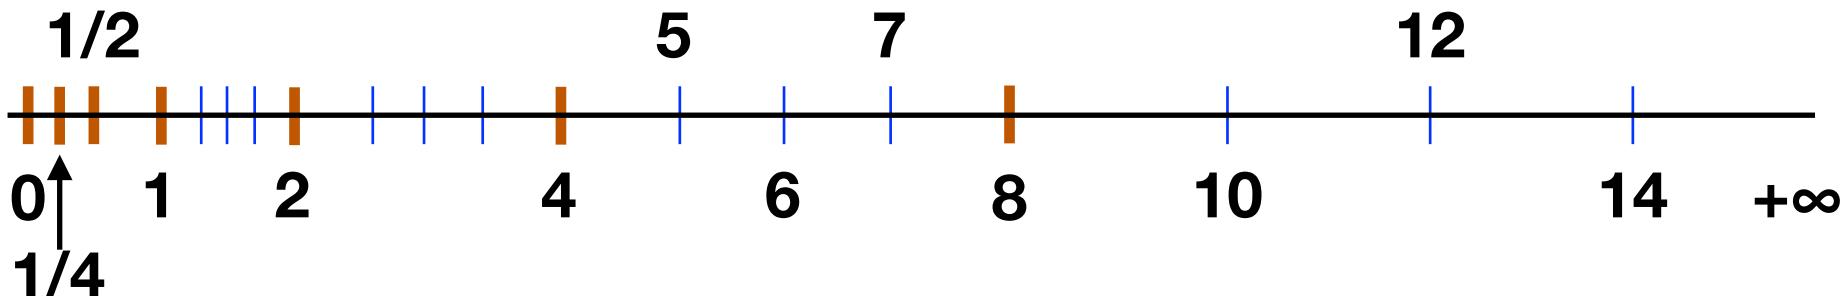


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

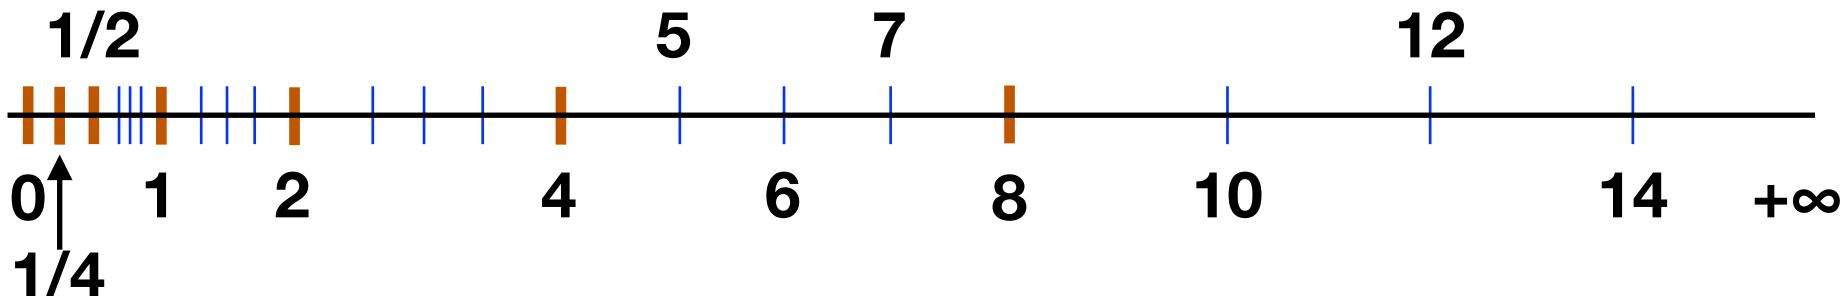


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$

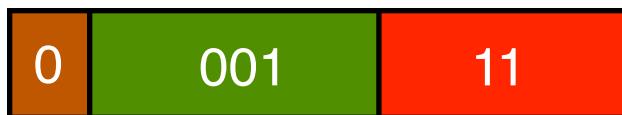


E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

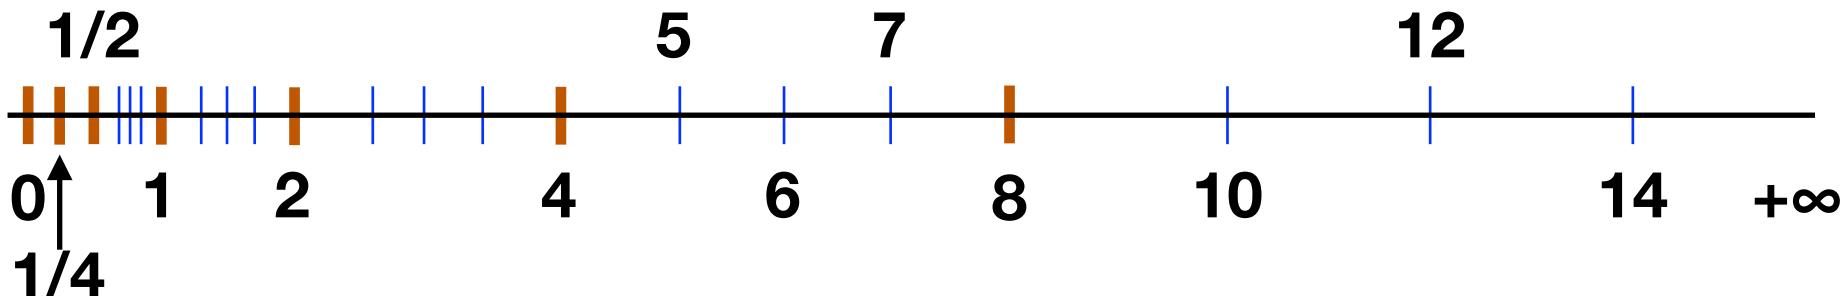


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

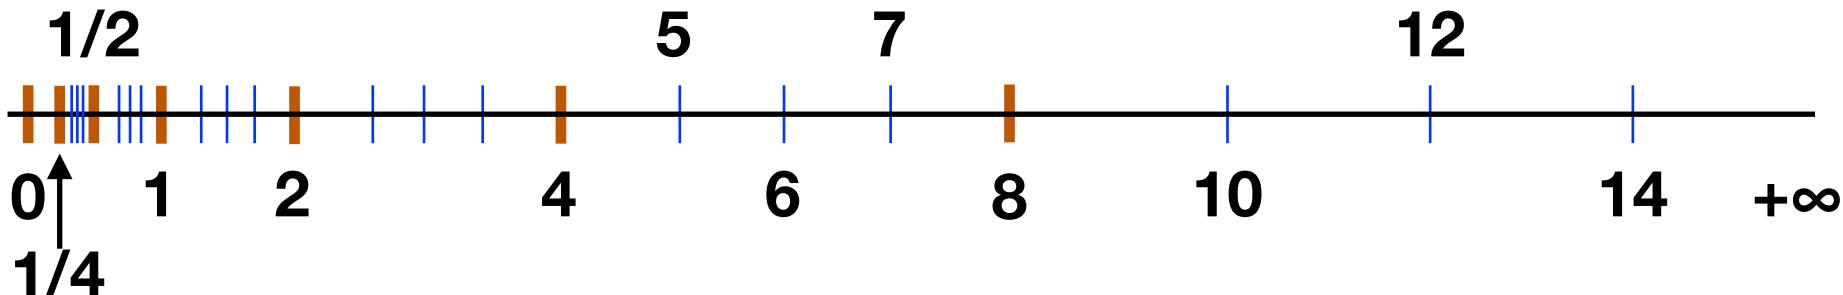


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111



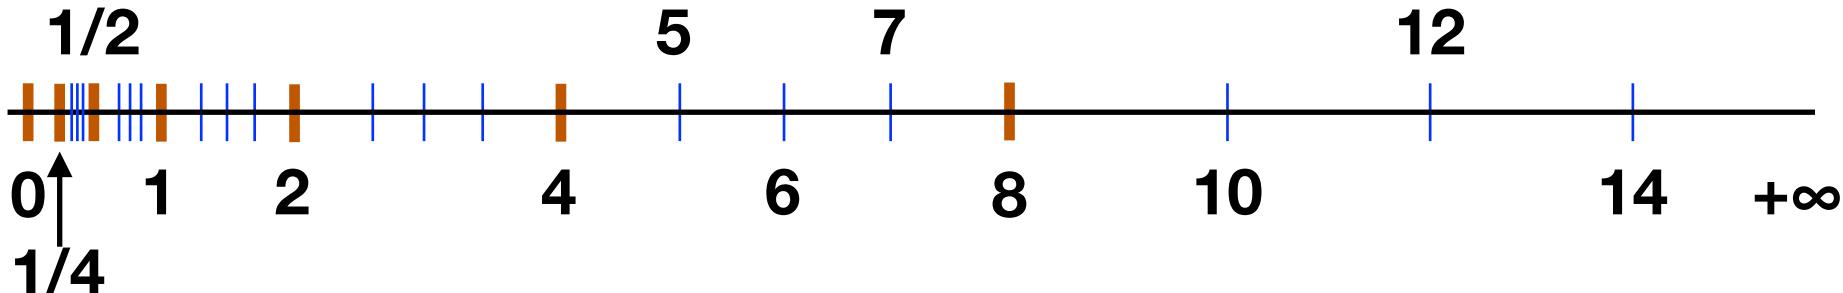
Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

- Uneven interval (c.f., fixed interval in fixed-point)
 - More dense toward 0, sparser toward infinite
 - Allow encoding small and large numbers at the same time



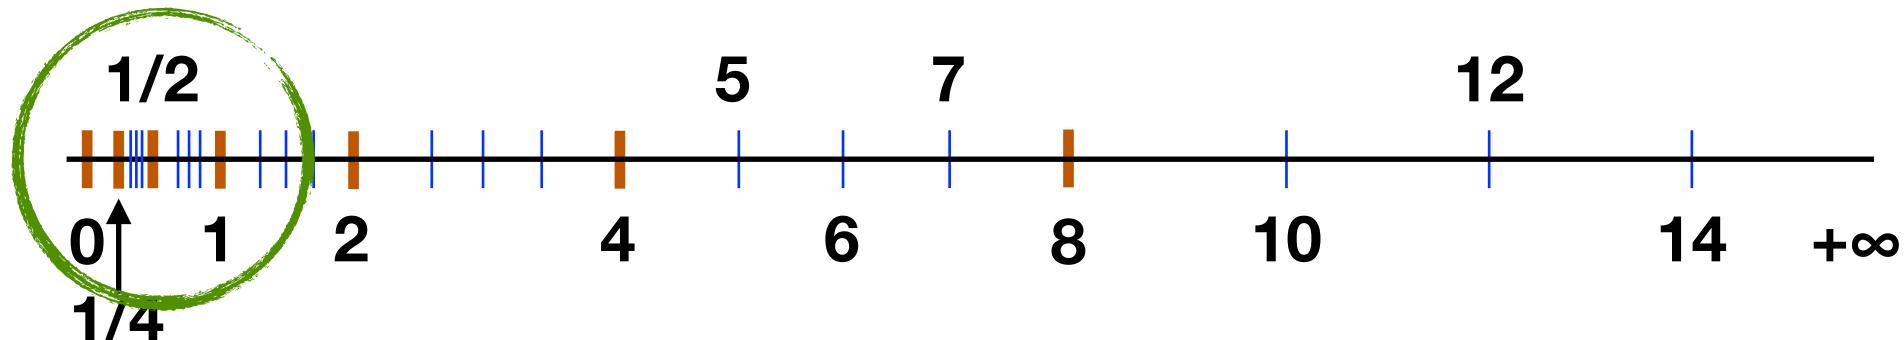
Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

- Uneven interval (c.f., fixed interval in fixed-point)
 - More dense toward 0, sparser toward infinite
 - Allow encoding small and large numbers at the same time

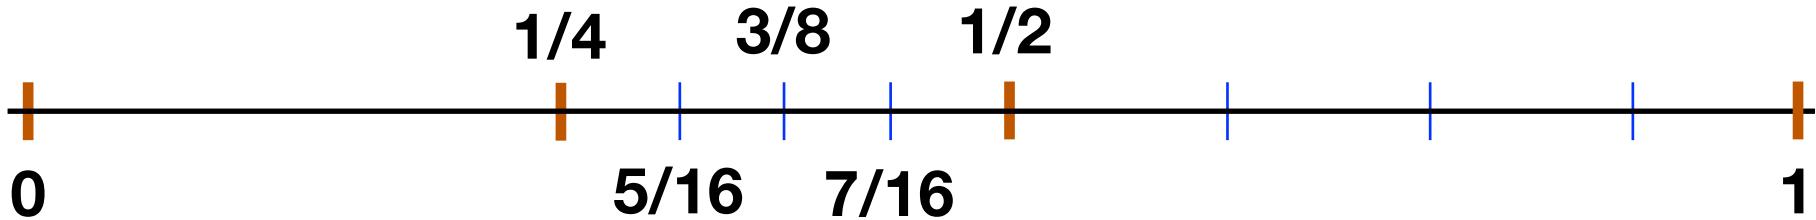


Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111



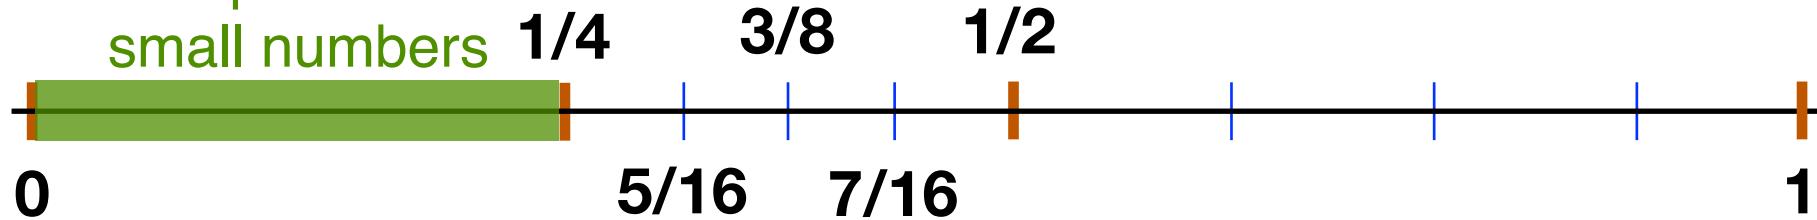
Representable Numbers (Positive Only)

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

Unrepresented
small numbers



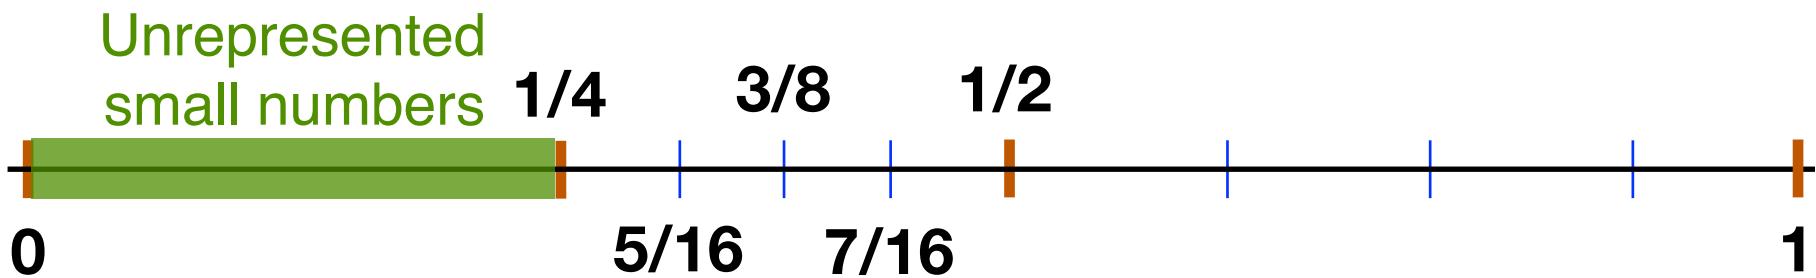
Representable Numbers (Positive Only)

$$v = (-1)^s M \cdot 2^E$$



- Underflow: always round to 0 is inelegant

E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111



Representable Numbers (Positive Only)

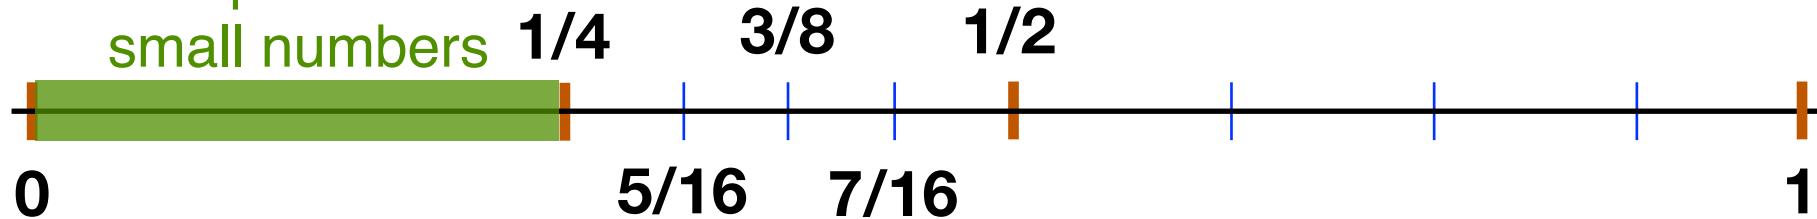
$$v = (-1)^s M \cdot 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

- Underflow: always round to 0 is inelegant

Unrepresented
small numbers



Representable Numbers (Positive Only)

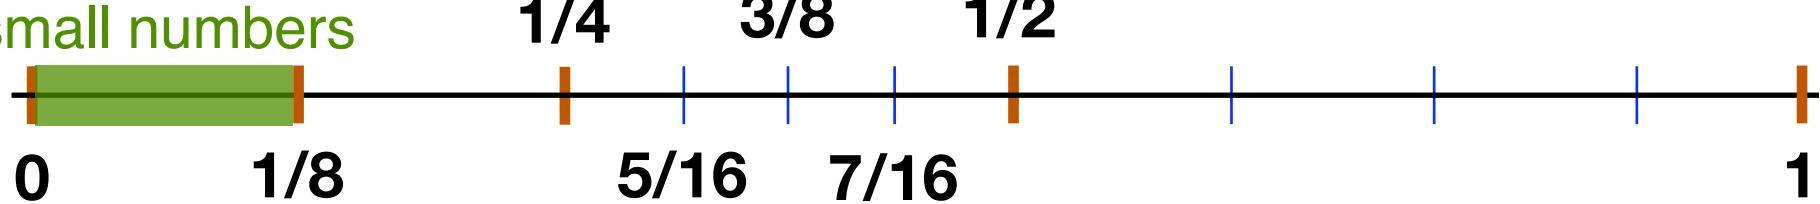
$$v = (-1)^s M \cdot 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

- Underflow: always round to 0 is inelegant

Unrepresented
small numbers



Representable Numbers (Positive Only)

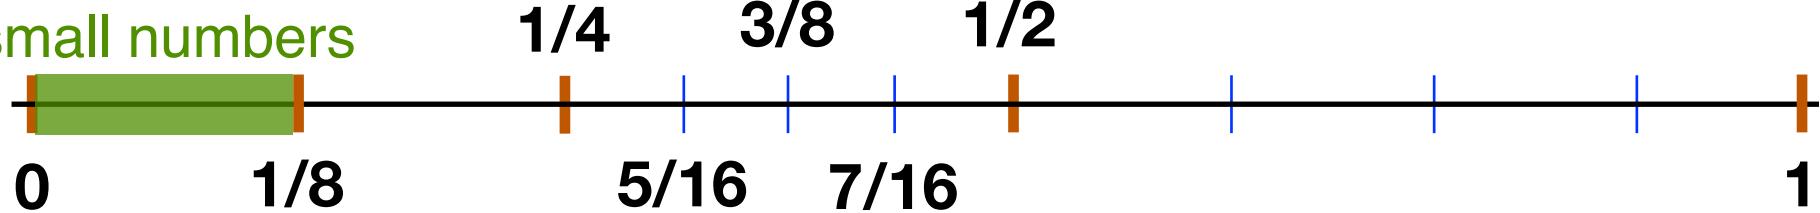
$$v = (-1)^s M \cdot 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

- Underflow: always round to 0 is inelegant
- Using 000 for exp would only postpone the problem rather than solving it

Unrepresented
small numbers



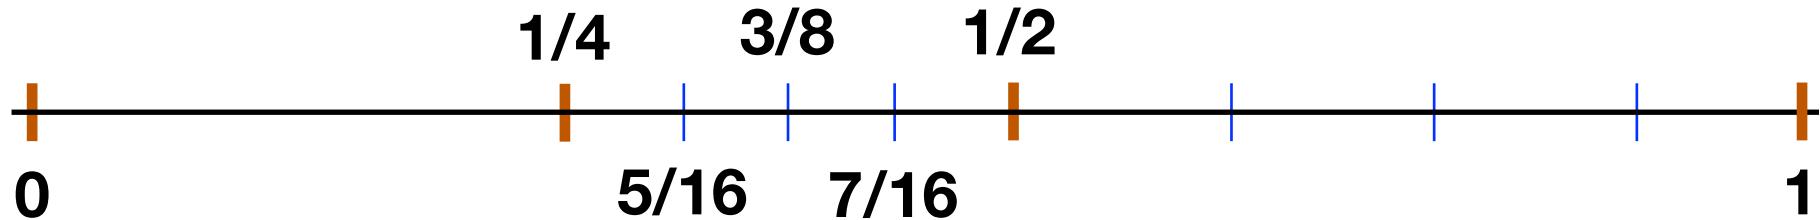
Subnormal (De-normalized) Numbers

$$v = (-1)^s M \cdot 2^E$$



- Idea: Evenly divide between 0 and 1/4 rather than exponentially decreasing **when exp = 0** (subnormal/denormalized numbers)

E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111



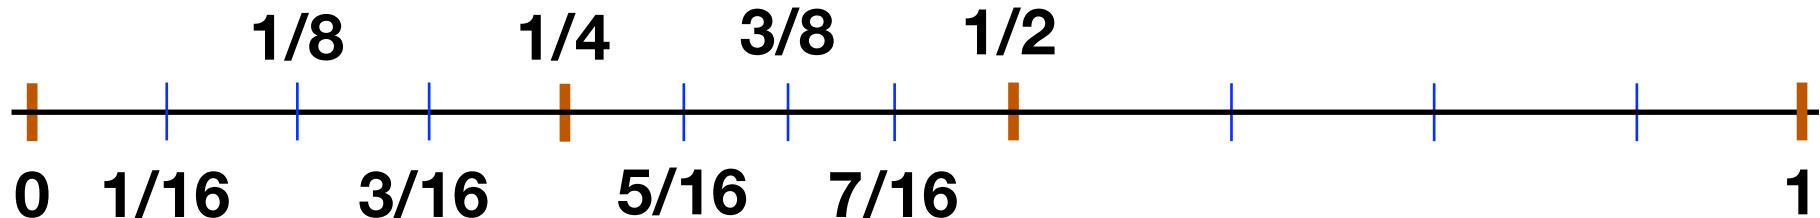
Subnormal (De-normalized) Numbers

$$v = (-1)^s M \cdot 2^E$$



- Idea: Evenly divide between 0 and 1/4 rather than exponentially decreasing **when exp = 0** (subnormal/denormalized numbers)

E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111



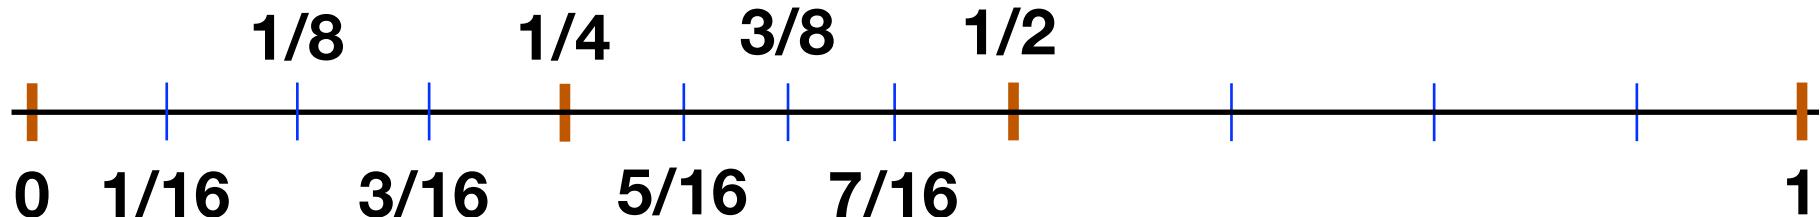
Subnormal (De-normalized) Numbers

$$v = (-1)^s M 2^E$$



- Idea: Evenly divide between 0 and 1/4 rather than exponentially decreasing **when exp = 0** (subnormal/denormalized numbers)
- $E = (\text{exp} + 1) - \text{bias}$ (instead of $\text{exp} - \text{bias}$)
- $M = 0.\text{frac}$ (instead of $1.\text{frac}$)

E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111



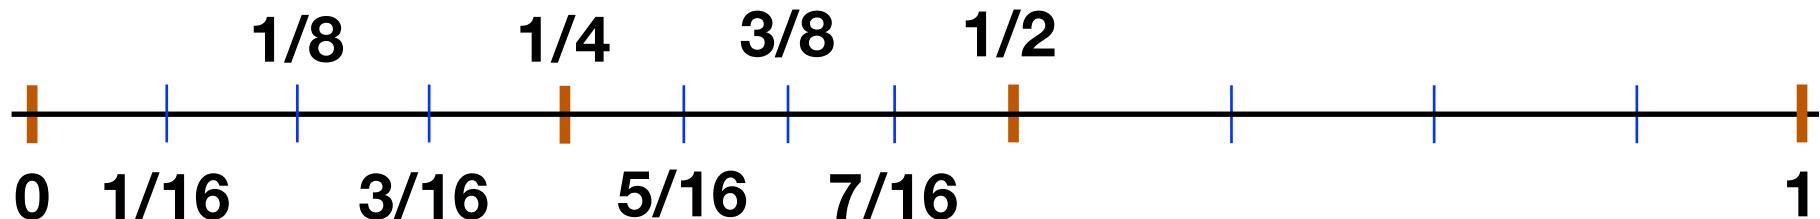
Subnormal (De-normalized) Numbers

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

- Idea: Evenly divide between 0 and 1/4 rather than exponentially decreasing **when exp = 0** (subnormal/denormalized numbers)
- $E = (\text{exp} + 1) - \text{bias}$ (instead of $\text{exp} - \text{bias}$)
- $M = 0.\text{frac}$ (instead of $1.\text{frac}$)



$$= (-1)^0 0.01 \times 2^{(0+1-3)} = 1/16$$

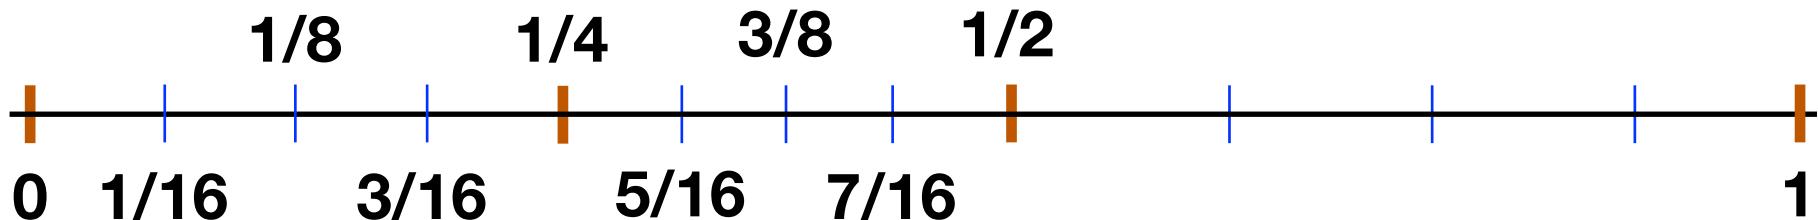
Subnormal (De-normalized) Numbers

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-3	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

- Idea: Evenly divide between 0 and 1/4 rather than exponentially decreasing **when exp = 0** (subnormal/denormalized numbers)
- $E = (\text{exp} + 1) - \text{bias}$ (instead of $\text{exp} - \text{bias}$)
- $M = 0.\text{frac}$ (instead of $1.\text{frac}$)
- Subnormal numbers allow graceful underflow




$$= (-1)^0 0.01 \times 2^{(0+1-3)} = 1/16$$

Special Values

$$v = (-1)^s M 2^E$$



E	exp	E	exp
-2	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

Special Values

$$v = (-1)^s M \cdot 2^E$$



E	exp	E	exp
-2	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

- There are many special values in scientific computing
 - $+\/-\infty$, Not-a-Numbers (NaNs) e.g., $0/0$, $0/\infty$, ∞/∞ , $\sqrt{-1}$, $\infty - \infty$, $\infty \times 0$, etc.

Special Values

$$v = (-1)^s M \cdot 2^E$$



E	exp	E	exp
-2	000	1	100
-2	001	2	101
-1	010	3	110
0	011	4	111

- There are many special values in scientific computing
 - $+\/-\infty$, Not-a-Numbers (NaNs) e.g., $0/0$, $0/\infty$, ∞/∞ , $\sqrt{-1}$, $\infty - \infty$, $\infty \times 0$, etc.
- $\text{exp} = 111$ is reserved to represent these numbers

Special Values

$$v = (-1)^s M \cdot 2^E$$



E	exp	E	exp
-2	000	1	100
-2	001	2	101
-1	010	3	110
0	011		111

- There are many special values in scientific computing
 - $+\/-\infty$, Not-a-Numbers (NaNs) e.g., $0/0$, $0/\infty$, ∞/∞ , $\sqrt{-1}$, $\infty - \infty$, $\infty \times 0$, etc.
- $\text{exp} = 111$ is reserved to represent these numbers

Special Values

$$v = (-1)^s M \cdot 2^E$$



E	exp	E	exp
-2	000	1	100
-2	001	2	101
-1	010	3	110
0	011		111

- There are many special values in scientific computing
 - $+\/-\infty$, Not-a-Numbers (NaNs) e.g., $0/0$, $0/\infty$, ∞/∞ , $\sqrt{-1}$, $\infty - \infty$, $\infty \times 0$, etc.
- $\text{exp} = 111$ is reserved to represent these numbers
- $\text{exp} = 111$, $\text{frac} = 00$
 - $+\/-\infty$ (depending on the s bit). Overflow results.
 - Arithmetic on ∞ is exact: $1.0/0.0 = -1.0/-0.0 = +\infty$, $1.0/-0.0 = -\infty$

Special Values

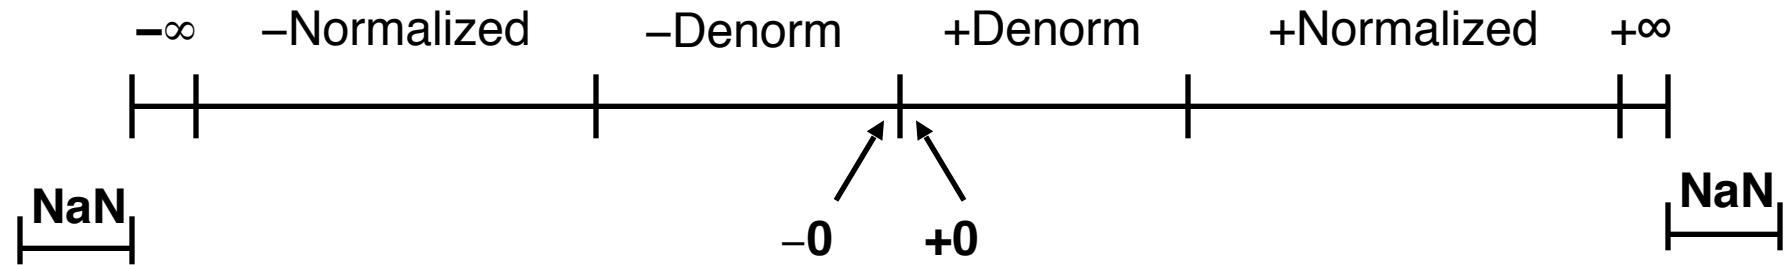
$$v = (-1)^s M \cdot 2^E$$



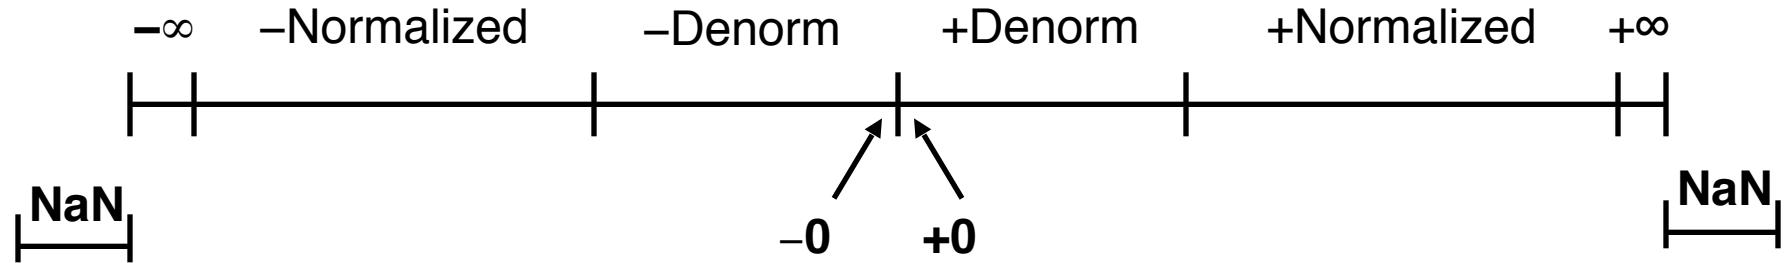
E	exp	E	exp
-2	000	1	100
-2	001	2	101
-1	010	3	110
0	011		111

- There are many special values in scientific computing
 - $+\/-\infty$, Not-a-Numbers (NaNs) e.g., $0/0$, $0/\infty$, ∞/∞ , $\sqrt{-1}$, $\infty - \infty$, $\infty \times 0$, etc.
- $\text{exp} = 111$ is reserved to represent these numbers
- $\text{exp} = 111$, $\text{frac} = 00$
 - $+\/-\infty$ (depending on the s bit). Overflow results.
 - Arithmetic on ∞ is exact: $1.0/0.0 = -1.0/-0.0 = +\infty$, $1.0/-0.0 = -\infty$
- $\text{exp} = 111$, $\text{frac} \neq 00$
 - Represent NaNs

Visualization: Floating Point Encodings



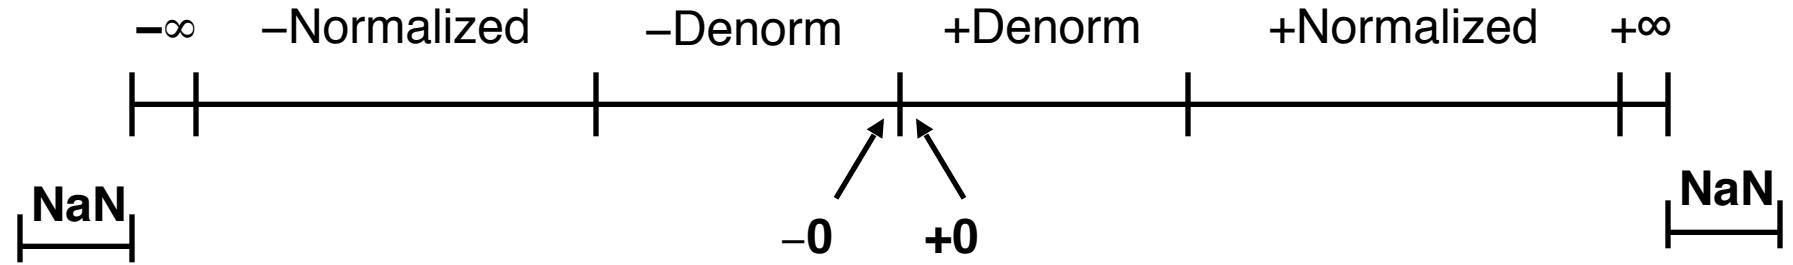
Visualization: Floating Point Encodings



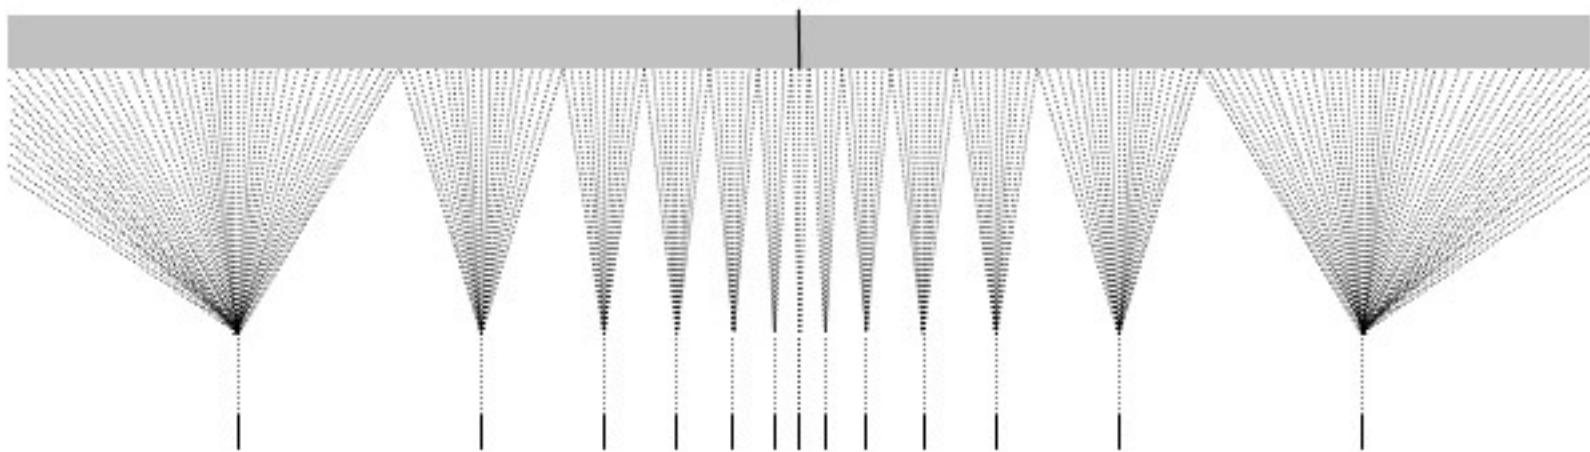
Infinite Amount of Real Numbers



Visualization: Floating Point Encodings

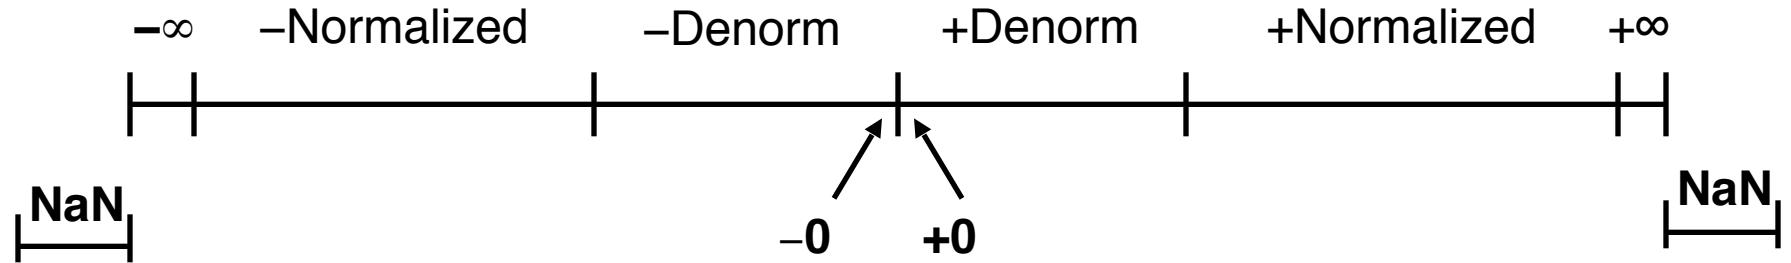


Infinite Amount of Real Numbers

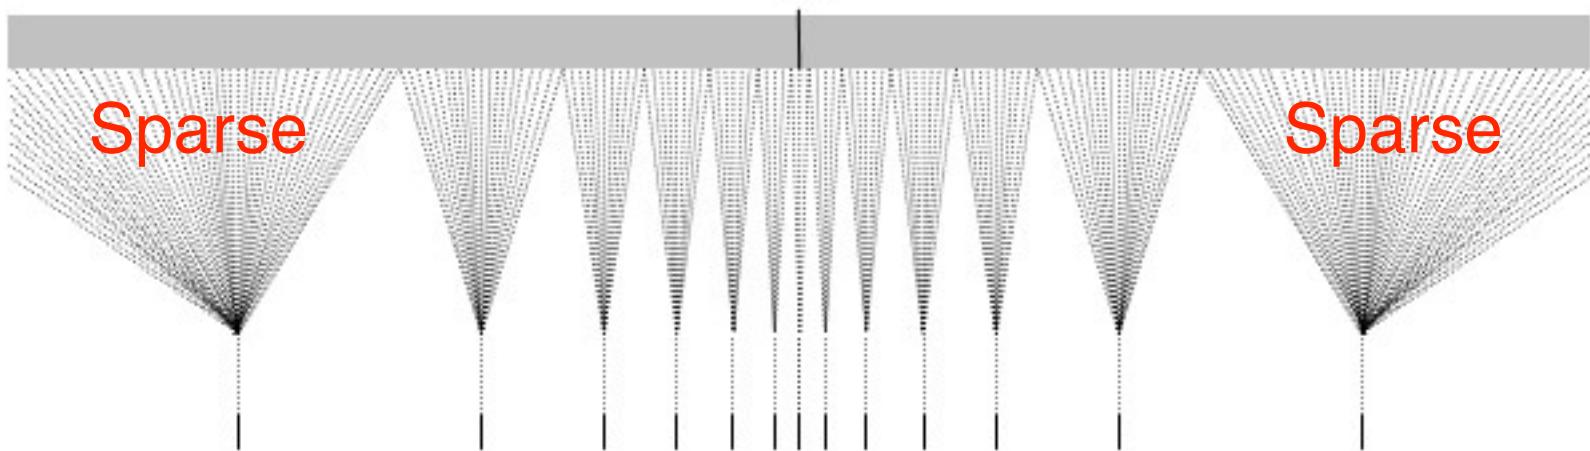


Finite Amount of Floating Point Numbers

Visualization: Floating Point Encodings

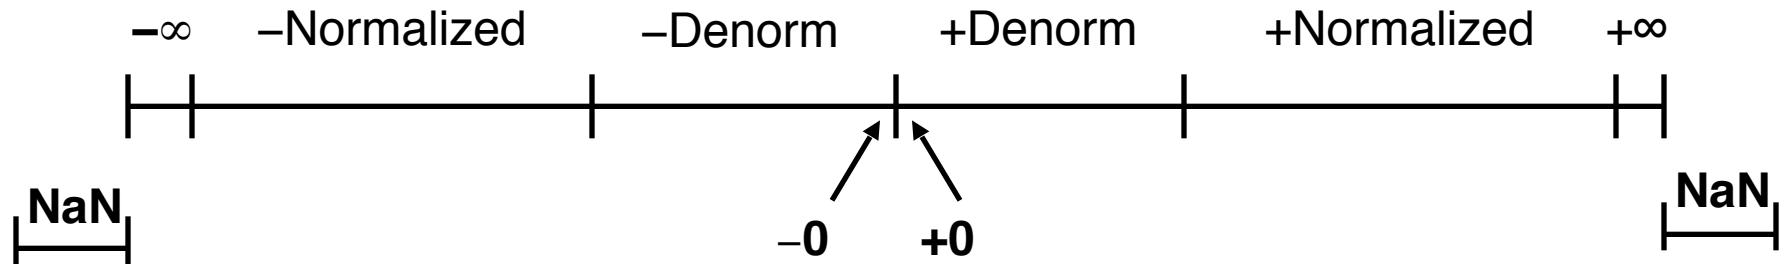


Infinite Amount of Real Numbers

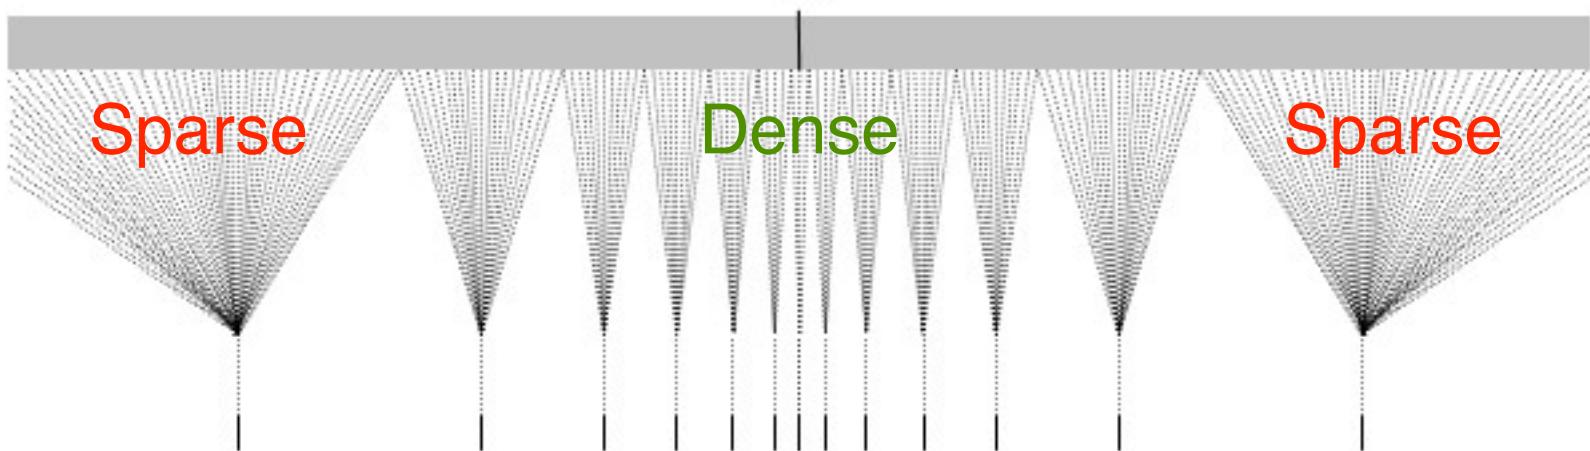


Finite Amount of Floating Point Numbers

Visualization: Floating Point Encodings



Infinite Amount of Real Numbers



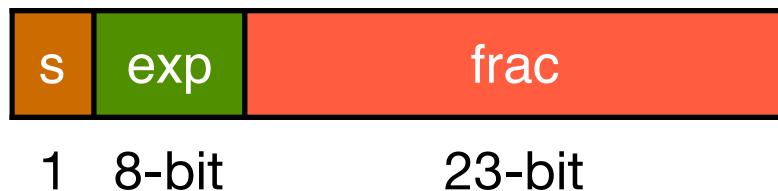
Finite Amount of Floating Point Numbers

Today: Floating Point

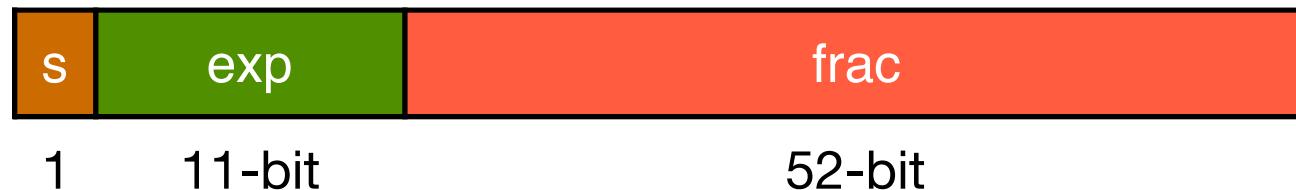
- Background: Fractional binary numbers and fixed-point
- Floating point representation
- IEEE 754 standard
- Rounding, addition, multiplication
- Floating point in C
- Summary

IEEE 754 Floating Point Standard

- Single precision: 32 bits



- Double precision: 64 bits



IEEE Floating Point

- IEEE Standard 754
 - Established in 1985 as uniform standard for floating point arithmetic
 - Before that, many idiosyncratic formats
 - Supported by all major CPUs (and even GPUs and other processors)
- Driven by numerical concerns
 - Nice standards for rounding, overflow, underflow
 - Hard to make fast in hardware
 - Numerical analysts predominated over hardware designers in defining standard

Single Precision (32-bit) Example

$$v = (-1)^s M 2^E$$

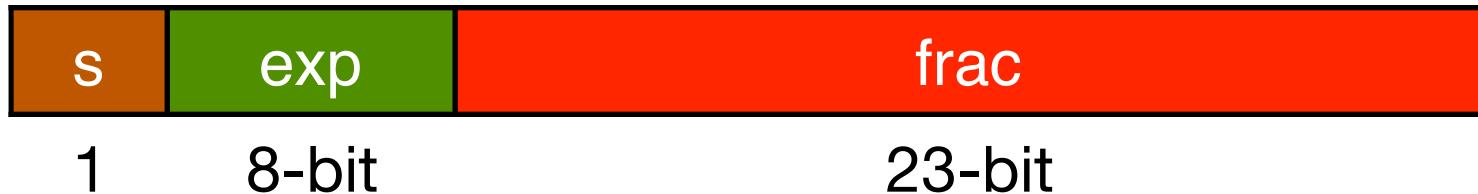
$$\text{bias} = 2^{(8-1)-1} = 127$$



Single Precision (32-bit) Example

$$v = (-1)^s M 2^E$$

$$\text{bias} = 2^{(8-1)-1} = 127$$

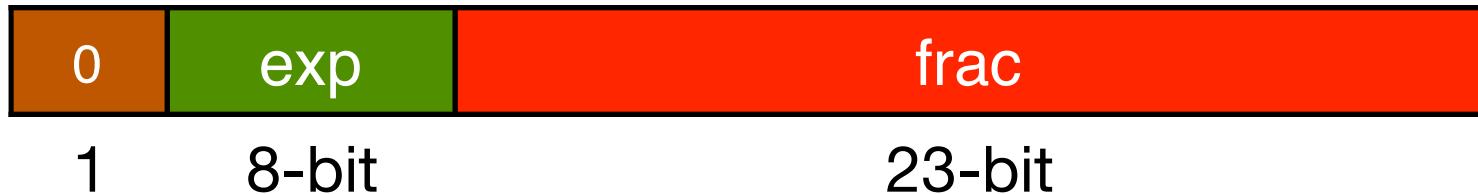


$$\begin{aligned}15213_{10} &= 11101101101101_2 \\&= (-1)^0 1.1101101101101_2 \times 2^{13}\end{aligned}$$

Single Precision (32-bit) Example

$$v = (-1)^s M 2^E$$

$$\text{bias} = 2^{(8-1)-1} = 127$$



$$\begin{aligned} 15213_{10} &= 11101101101101_2 \\ &= (-1)^0 1.1101101101101_2 \times 2^{13} \end{aligned}$$

Single Precision (32-bit) Example

$$v = (-1)^s M 2^E$$

$$\text{bias} = 2^{(8-1)-1} = 127$$



$$\begin{aligned} 15213_{10} &= 11101101101101_2 \\ &= (-1)^0 1.1101101101101_2 \times 2^{13} \end{aligned}$$

Single Precision (32-bit) Example

$$v = (-1)^s M 2^E$$

$$\text{bias} = 2^{(8-1)-1} = 127$$



$$\begin{aligned}15213_{10} &= 11101101101101_2 \\&= (-1)^0 \cdot 1.1101101101101_2 \times 2^{13}\end{aligned}$$

$$\text{exp} = E + \text{bias} = 140_{10}$$

Single Precision (32-bit) Example

$$v = (-1)^s M 2^E$$

$$\text{bias} = 2^{(8-1)-1} = 127$$



$$\begin{aligned} 15213_{10} &= 11101101101101_2 \\ &= (-1)^0 1.1101101101101_2 \times 2^{13} \end{aligned}$$

$$\text{exp} = E + \text{bias} = 140_{10}$$

Single Precision (32-bit) Example

$$v = (-1)^s M 2^E$$

$$\text{bias} = 2^{(8-1)-1} = 127$$



$$\begin{aligned}15213_{10} &= 11101101101101_2 \\&= (-1)^0 1.\underline{1101101101101}_2 \times 2^{13}\end{aligned}$$

$$\text{exp} = E + \text{bias} = 140_{10}$$

Single Precision (32-bit) Example

$$v = (-1)^s M 2^E$$

$$\text{bias} = 2^{(8-1)-1} = 127$$



$$\begin{aligned}15213_{10} &= 11101101101101_2 \\&= (-1)^0 1.\underline{1101101101101}_2 \times 2^{13}\end{aligned}$$

$$\text{exp} = E + \text{bias} = 140_{10}$$

Today: Floating Point

- Background: Fractional binary numbers and fixed-point
- Floating point representation
- IEEE 754 standard
- Rounding, addition, multiplication
- Floating point in C
- Summary

Floating Point Computations

- The problem: Computing on floating point numbers might produce a result that can't be precisely represented
- Basic idea
 - We perform the operation & produce the infinitely **precise** result
 - Make it fit into desired precision
 - Possibly **overflow** if exponent too large
 - Possibly **round** to fit into frac

Rounding Modes (Decimal)

Rounding Modes (Decimal)

- Common ones:
 - Towards zero (chop)
 - Round down ($-\infty$)
 - Round up ($+\infty$)

Rounding Modes (Decimal)

- Common ones:
 - Towards zero (chop)
 - Round down ($-\infty$)
 - Round up ($+\infty$)
- Better: Nearest Even —> To nearest; if equally near, then to the one having an even least significant digit (bit)

Rounding Modes (Decimal)

- Common ones:
 - Towards zero (chop)
 - Round down ($-\infty$)
 - Round up ($+\infty$)
- Better: Nearest Even —> To nearest; if equally near, then to the one having an even least significant digit (bit)

Rounding Mode	1.40	1.60	1.50	2.50	-1.50
Towards zero	1	1	1	2	-1
Round down ($-\infty$)	1	1	1	2	-2
Round up ($+\infty$)	2	2	2	3	-1
Nearest even (default)	1	2	2	2	-2

Rounding Modes (Binary Example)

- Nearest Even; if equally near, then to the one having an even least significant digit (bit)
- Assuming 3 bits for *frac*

Rounding Modes (Binary Example)

- Nearest Even; if equally near, then to the one having an even least significant digit (bit)
- Assuming 3 bits for *frac*

Precise Value	Rounded Value	Notes
1.000 ⁰¹¹	1.000	1.000 is the nearest (down)
1.000 ¹¹⁰	1.001	1.001 is the nearest (up)
1.000 ¹⁰⁰	1.000	1.000 is the nearest even (down)
1.001 ¹⁰⁰	1.010	1.010 is the nearest even (up)

Rounding Modes (Binary Example)

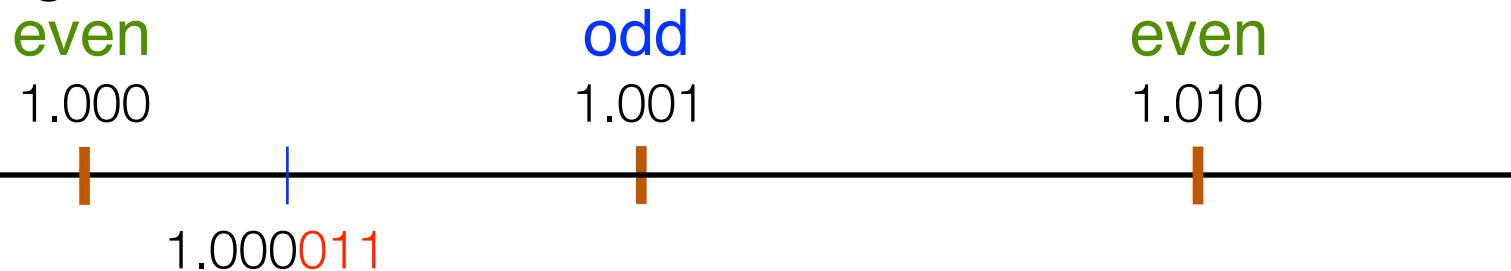
- Nearest Even; if equally near, then to the one having an even least significant digit (bit)
- Assuming 3 bits for *frac*



Precise Value	Rounded Value	Notes
1.000 ₁₁₀	1.000	1.000 is the nearest (down)
1.000 ₁₁₀	1.001	1.001 is the nearest (up)
1.000 ₁₀₀	1.000	1.000 is the nearest even (down)
1.001 ₁₀₀	1.010	1.010 is the nearest even (up)

Rounding Modes (Binary Example)

- Nearest Even; if equally near, then to the one having an even least significant digit (bit)
- Assuming 3 bits for *frac*

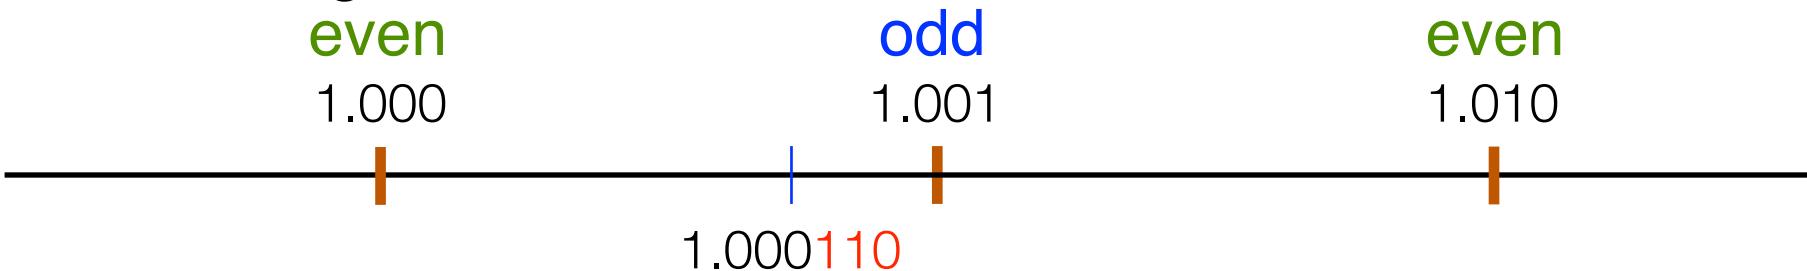


Precise Value	Rounded Value	Notes
1.000011	1.000	1.000 is the nearest (down)
1.000110	1.001	1.001 is the nearest (up)
1.000100	1.000	1.000 is the nearest even (down)
1.001100	1.010	1.010 is the nearest even (up)



Rounding Modes (Binary Example)

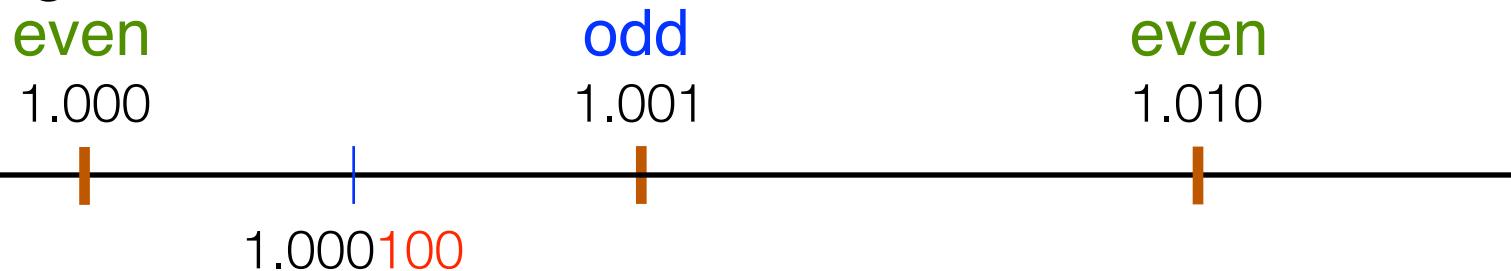
- Nearest Even; if equally near, then to the one having an even least significant digit (bit)
- Assuming 3 bits for *frac*



Precise Value	Rounded Value	Notes
1.000011	1.000	1.000 is the nearest (down)
1.000110	1.001	1.001 is the nearest (up)
1.000100	1.000	1.000 is the nearest even (down)
1.001100	1.010	1.010 is the nearest even (up)

Rounding Modes (Binary Example)

- Nearest Even; if equally near, then to the one having an even least significant digit (bit)
- Assuming 3 bits for *frac*

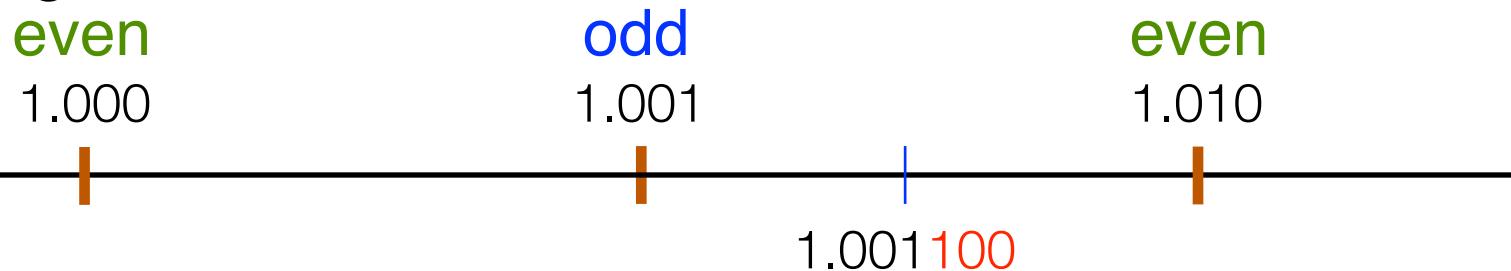


Precise Value	Rounded Value	Notes
1.000011	1.000	1.000 is the nearest (down)
1.000110	1.001	1.001 is the nearest (up)
1.000100	1.000	1.000 is the nearest even (down)
1.001100	1.010	1.010 is the nearest even (up)



Rounding Modes (Binary Example)

- Nearest Even; if equally near, then to the one having an even least significant digit (bit)
- Assuming 3 bits for *frac*



Precise Value	Rounded Value	Notes
1.000011	1.000	1.000 is the nearest (down)
1.000110	1.001	1.001 is the nearest (up)
1.000100	1.000	1.000 is the nearest even (down)
1.001100	1.010	1.010 is the nearest even (up)



Floating Point Addition

Floating Point Addition

$$\bullet (-1)^{s_1} M_1 2^{E_1} + (-1)^{s_2} M_2 2^{E_2}$$

$$1.000 \times 2^{-1} + 11.10 \times 2^{-3}$$

Floating Point Addition

- $(-1)^{s_1} M_1 2^{E_1} + (-1)^{s_2} M_2 2^{E_2}$

$$1.000 \times 2^{-1} + 11.10 \times 2^{-3}$$



align $1.000 \times 2^{-1} + 0.111 \times 2^{-1}$

Floating Point Addition

- $(-1)^{s1} M1 2^{E1} + (-1)^{s2} M2 2^{E2}$

$$1.000 \times 2^{-1} + 11.10 \times 2^{-3}$$



$$1.000 \times 2^{-1} + 0.111 \times 2^{-1}$$



add

$$1.111 \times 2^{-1}$$

Floating Point Addition

- $(-1)^{s_1} M_1 2^{E_1} + (-1)^{s_2} M_2 2^{E_2}$

- Exact Result: $(-1)^s M 2^E$

- Sign s , significand M :

- Result of signed align & add

- Exponent E: E_1

- Assume $E_1 > E_2$

$$1.000 \times 2^{-1} + 11.10 \times 2^{-3}$$

$$1.000 \times 2^{-1} + 0.111 \times 2^{-1}$$

$$1.111 \times 2^{-1}$$

Floating Point Addition

- $(-1)^{s_1} M_1 2^{E_1} + (-1)^{s_2} M_2 2^{E_2}$

- Exact Result: $(-1)^s M 2^E$

- Sign s , significand M :

- Result of signed align & add

- Exponent E: E_1

- Assume $E_1 > E_2$

- Fixing

- If $M \geq 2$, shift M right, increment E

- If $M < 1$, shift M left k positions, decrement E by k

- Overflow if E out of range

- Round M to fit frac precision

$$1.000 \times 2^{-1} + 11.10 \times 2^{-3}$$

$$1.000 \times 2^{-1} + 0.111 \times 2^{-1}$$

$$1.111 \times 2^{-1}$$

Mathematical Properties of FP Add

Mathematical Properties of FP Add

- Commutative?

Mathematical Properties of FP Add

- Commutative? Yes

Mathematical Properties of FP Add

- Commutative? **Yes**
- Associative?

Mathematical Properties of FP Add

- Commutative? **Yes**
- Associative?
 - Overflow and inexactness of rounding
 - $(3.14 + 1e10) - 1e10 = 0, 3.14 + (1e10 - 1e10) = 3.14$

Mathematical Properties of FP Add

- Commutative? **Yes**
- Associative? **No**
 - Overflow and inexactness of rounding
 - $(3.14 + 1e10) - 1e10 = 0, 3.14 + (1e10 - 1e10) = 3.14$
- 0 is additive identity?

Mathematical Properties of FP Add

- Commutative? **Yes**
- Associative?
 - Overflow and inexactness of rounding
 - $(3.14 + 1e10) - 1e10 = 0, 3.14 + (1e10 - 1e10) = 3.14$**No**
- 0 is additive identity? **Yes**

Mathematical Properties of FP Add

- Commutative? **Yes**
- Associative? **No**
 - Overflow and inexactness of rounding
 - $(3.14 + 1e10) - 1e10 = 0, 3.14 + (1e10 - 1e10) = 3.14$
- 0 is additive identity? **Yes**
- Every element has additive inverse (negation)?

Mathematical Properties of FP Add

- Commutative? **Yes**
- Associative?
 - Overflow and inexactness of rounding
 - $(3.14 + 1e10) - 1e10 = 0, 3.14 + (1e10 - 1e10) = 3.14$**No**
- 0 is additive identity? **Yes**
- Every element has additive inverse (negation)? **Almost**
 - Except for infinities & NaNs

Mathematical Properties of FP Add

- Commutative? **Yes**
- Associative? **No**
 - Overflow and inexactness of rounding
 - $(3.14 + 1e10) - 1e10 = 0, 3.14 + (1e10 - 1e10) = 3.14$
- 0 is additive identity? **Yes**
- Every element has additive inverse (negation)? **Almost**
 - Except for infinities & NaNs
- Monotonicity: $a \geq b \Rightarrow a+c \geq b+c?$

Mathematical Properties of FP Add

- Commutative? **Yes**
- Associative? **No**
 - Overflow and inexactness of rounding
 - $(3.14 + 1e10) - 1e10 = 0, 3.14 + (1e10 - 1e10) = 3.14$
- 0 is additive identity? **Yes**
- Every element has additive inverse (negation)? **Almost**
 - Except for infinities & NaNs
- Monotonicity: $a \geq b \Rightarrow a+c \geq b+c?$ **Almost**
 - Except for infinities & NaNs

Floating Point Multiplication

Floating Point Multiplication

- $(-1)^{s_1} M_1 2^{E_1} \times (-1)^{s_2} M_2 2^{E_2}$

Floating Point Multiplication

- $(-1)^{s_1} M_1 2^{E_1} \times (-1)^{s_2} M_2 2^{E_2}$
- Exact Result: $(-1)^s M 2^E$
 - Sign s : $s_1 \wedge s_2$
 - Significand M : $M_1 \times M_2$
 - Exponent E : $E_1 + E_2$

Floating Point Multiplication

- $(-1)^{s_1} M_1 2^{E_1} \times (-1)^{s_2} M_2 2^{E_2}$

- Exact Result: $(-1)^s M 2^E$

- Sign s : $s_1 \wedge s_2$
- Significand M : $M_1 \times M_2$
- Exponent E : $E_1 + E_2$

- Fixing

- If $M \geq 2$, shift M right, increment E
- If E out of range, overflow
- Round M to fit frac precision

Floating Point Multiplication

- $(-1)^{s_1} M_1 2^{E_1} \times (-1)^{s_2} M_2 2^{E_2}$
- Exact Result: $(-1)^s M 2^E$
 - Sign s : $s_1 \wedge s_2$
 - Significand M : $M_1 \times M_2$
 - Exponent E : $E_1 + E_2$
- Fixing
 - If $M \geq 2$, shift M right, increment E
 - If E out of range, overflow
 - Round M to fit frac precision
- Implementation
 - Biggest chore is multiplying significands

Mathematical Properties of FP Mult

Mathematical Properties of FP Mult

- Multiplication Commutative?

Mathematical Properties of FP Mult

- Multiplication Commutative? **Yes**

Mathematical Properties of FP Mult

- Multiplication Commutative? **Yes**
- Multiplication is Associative?

Mathematical Properties of FP Mult

- Multiplication Commutative? **Yes**
- Multiplication is Associative?
 - Possibility of overflow, inexactness of rounding
 - Ex: $(1e20 * 1e20) * 1e-20 = \text{inf}$, $1e20 * (1e20 * 1e-20) = 1e20$

Mathematical Properties of FP Mult

- Multiplication Commutative? **Yes**
- Multiplication is Associative?
 - Possibility of overflow, inexactness of rounding
 - Ex: $(1e20 * 1e20) * 1e-20 = \text{inf}$, $1e20 * (1e20 * 1e-20) = 1e20$
- 1 is multiplicative identity?

Mathematical Properties of FP Mult

- Multiplication Commutative? **Yes**
- Multiplication is Associative?
 - Possibility of overflow, inexactness of rounding
 - Ex: $(1e20 * 1e20) * 1e-20 = \text{inf}$, $1e20 * (1e20 * 1e-20) = 1e20$
- 1 is multiplicative identity? **Yes**

Mathematical Properties of FP Mult

- Multiplication Commutative? **Yes**
- Multiplication is Associative?
 - Possibility of overflow, inexactness of rounding
 - Ex: $(1e20 * 1e20) * 1e-20 = \text{inf}$, $1e20 * (1e20 * 1e-20) = 1e20$**No**
- 1 is multiplicative identity? **Yes**
- Multiplication distributes over addition?

Mathematical Properties of FP Mult

- Multiplication Commutative? **Yes**
- Multiplication is Associative?
 - Possibility of overflow, inexactness of rounding
 - Ex: $(1e20 * 1e20) * 1e-20 = \text{inf}$, $1e20 * (1e20 * 1e-20) = 1e20$**No**
- 1 is multiplicative identity? **Yes**
- Multiplication distributes over addition?
 - Possibility of overflow, inexactness of rounding
 - $1e20 * (1e20 - 1e20) = 0.0$, $1e20 * 1e20 - 1e20 * 1e20 = \text{NaN}$**No**

Mathematical Properties of FP Mult

- Multiplication Commutative? **Yes**
- Multiplication is Associative?
 - Possibility of overflow, inexactness of rounding
 - Ex: $(1e20 * 1e20) * 1e-20 = \text{inf}$, $1e20 * (1e20 * 1e-20) = 1e20$**No**
- 1 is multiplicative identity? **Yes**
- Multiplication distributes over addition?
 - Possibility of overflow, inexactness of rounding
 - $1e20 * (1e20 - 1e20) = 0.0$, $1e20 * 1e20 - 1e20 * 1e20 = \text{NaN}$**No**
- Monotonicity: $a \geq b \ \& \ c \geq 0 \Rightarrow a * c \geq b * c?$

Mathematical Properties of FP Mult

- Multiplication Commutative? **Yes**
- Multiplication is Associative?
 - Possibility of overflow, inexactness of rounding
 - Ex: $(1e20 * 1e20) * 1e-20 = \inf$, $1e20 * (1e20 * 1e-20) = 1e20$**No**
- 1 is multiplicative identity? **Yes**
- Multiplication distributes over addition?
 - Possibility of overflow, inexactness of rounding
 - $1e20 * (1e20 - 1e20) = 0.0$, $1e20 * 1e20 - 1e20 * 1e20 = \text{NaN}$**No**
- Monotonicity: $a \geq b \ \& \ c \geq 0 \Rightarrow a * c \geq b * c?$ **Almost**
 - Except for infinities & NaNs