

CSC 252: Computer Organization

Spring 2021: Lecture 26

Instructor: Yuhao Zhu

Department of Computer Science
University of Rochester

Announcements

SUN 25	MON 26	TUE 27	WED 28	THU 29	FRI 30	SAT May 1
2	3	4	5	6 Today Last Lecture	7	8 A5 Due
9	10	11	12	13	14	15 Final

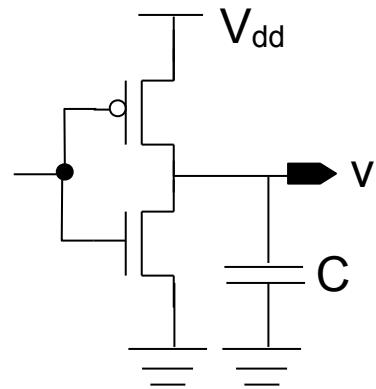
Announcements

- Final exam: May 12, 19:15 PM -- 22:15 PM; online.
- Past exam & Problem set: <https://www.cs.rochester.edu/courses/252/spring2021/handouts.html>
- Exam will be electronic using Gradescope, but we will send you an PDF version so that you can work offline in case
 - 1) you don't have Internet access at the exam time or
 - 2) you lose Internet access.
 - Write down the answers on a scratch paper, take pictures, and send us the pictures

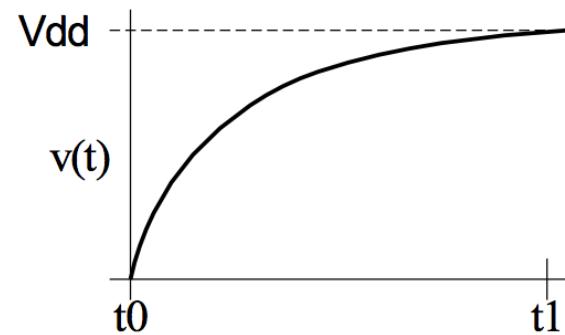
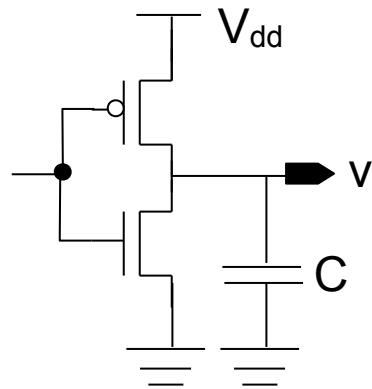
Announcements

- Open book test: any sort of paper-based product, e.g., book, **notes**, magazine, old tests.
- Exams are designed to test your ability to apply what you have learned and not your memory (though a good memory could help).
- **Nothing electronic (including laptop, cell phone, calculator, etc) other than the computer you use to take the exam.**
- **Nothing biological**, including your roommate, husband, wife, your hamster, another professor, etc.
- “**I don’t know**” gets 15% partial credit. Must erase everything else.

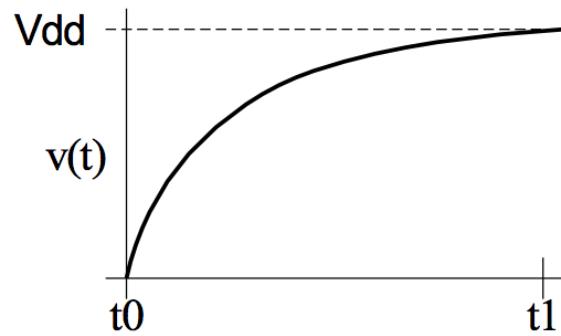
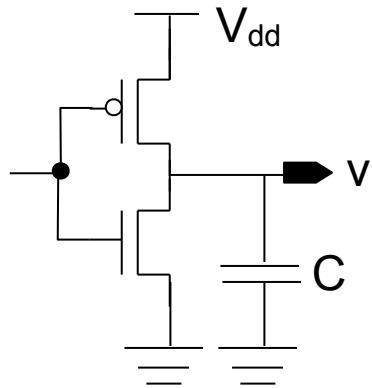
Dynamic Power



Dynamic Power

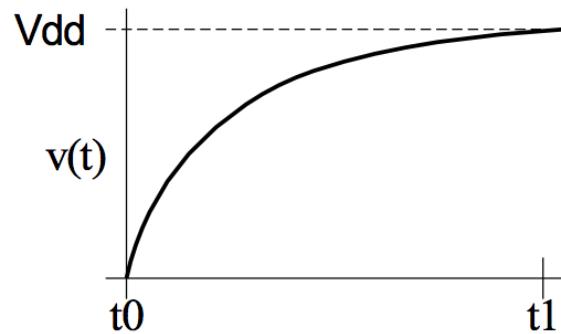
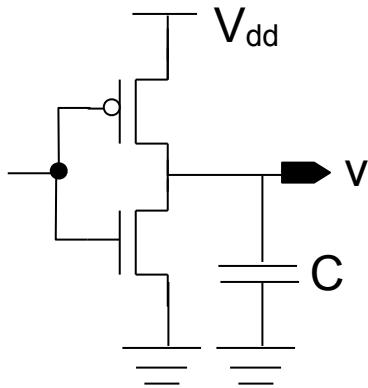


Dynamic Power



$$\begin{aligned} E_{sw} &= \int_{t_0}^{t_1} P(t) dt = \int_{t_0}^{t_1} (V_{dd} - v) \cdot i(t) dt = \int_{t_0}^{t_1} (V_{dd} - v) \cdot c (dv/dt) dt = \\ &= c V_{dd} \int_{t_0}^{t_1} dv - c \int_{t_0}^{t_1} v \cdot dv \xrightarrow{\text{_____}} c V_{dd}^2 - \underset{\uparrow}{1/2} c V_{dd}^2 = \boxed{1/2 c V_{dd}^2} \end{aligned}$$

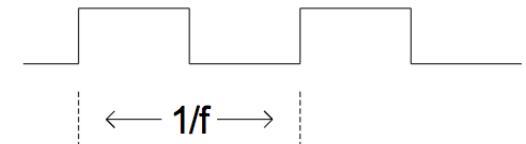
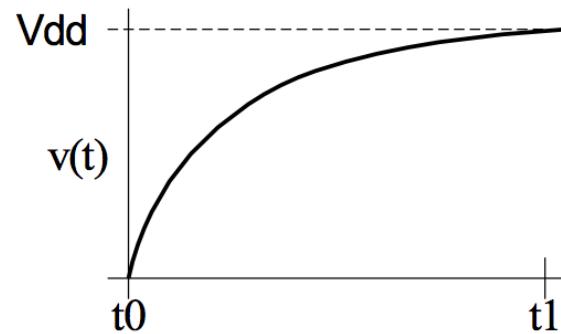
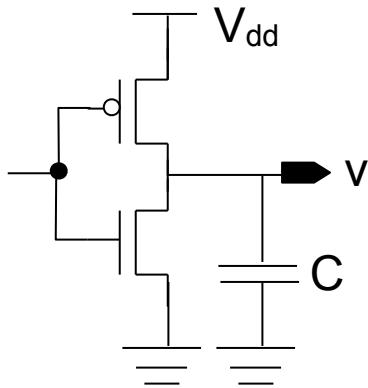
Dynamic Power



Energy dissipated for every transition ($0 \rightarrow 1$ or $1 \rightarrow 0$)

$$\begin{aligned} E_{sw} &= \int_{t_0}^{t_1} P(t) dt = \int_{t_0}^{t_1} (V_{dd} - v) \cdot i(t) dt = \int_{t_0}^{t_1} (V_{dd} - v) \cdot c (dv/dt) dt = \\ &= c V_{dd} \int_{t_0}^{t_1} dv - c \int_{t_0}^{t_1} v \cdot dv \xrightarrow{\text{_____}} c V_{dd}^2 - \underset{\uparrow}{1/2} c V_{dd}^2 = \boxed{1/2 c V_{dd}^2} \end{aligned}$$

Dynamic Power



Average dynamic power of a transistor:

$$P = \alpha \cdot (E / T) = \alpha \cdot E f = \frac{1}{2} \alpha C V_{dd}^2 f$$

α : switch activity factor. No switching, no dynamic power consumption

Dynamic Power

$$P = k C V^2 f$$

Dynamic Power

$$P = k C V^2 f$$

- Increasing f requires V to be increased proportionally

Dynamic Power

$$P = k C V^2 f$$

- Increasing f requires V to be increased proportionally
 - Intuitively: a higher frequency means a shorter cycle time, which means the critical path of your processor needs to be shorter, which requires faster transistors, which you get by increasing the voltage

Dynamic Power

$$P = k C V^2 f$$

- Increasing f requires V to be increased proportionally
 - Intuitively: a higher frequency means a shorter cycle time, which means the critical path of your processor needs to be shorter, which requires faster transistors, which you get by increasing the voltage
 - “Overclocking” just increases the clock speed without increasing voltage => machine might crash (cycle time shorter than the critical path delay)

Dynamic Power

$$P = k C V^2 f$$

- Increasing f requires V to be increased proportionally
 - Intuitively: a higher frequency means a shorter cycle time, which means the critical path of your processor needs to be shorter, which requires faster transistors, which you get by increasing the voltage
 - “Overclocking” just increases the clock speed without increasing voltage => machine might crash (cycle time shorter than the critical path delay)
 - Corollary: reducing voltage requires reducing frequency

Dynamic Power

$$P = k C V^2 f$$

- Increasing f requires V to be increased proportionally
 - Intuitively: a higher frequency means a shorter cycle time, which means the critical path of your processor needs to be shorter, which requires faster transistors, which you get by increasing the voltage
 - “Overclocking” just increases the clock speed without increasing voltage => machine might crash (cycle time shorter than the critical path delay)
 - Corollary: reducing voltage requires reducing frequency
 - 15% reduction in voltage requires about 15% slow down in frequency

Dynamic Power

$$P = k C V^2 f$$

- Increasing f requires V to be increased proportionally
 - Intuitively: a higher frequency means a shorter cycle time, which means the critical path of your processor needs to be shorter, which requires faster transistors, which you get by increasing the voltage
 - “Overclocking” just increases the clock speed without increasing voltage => machine might crash (cycle time shorter than the critical path delay)
 - Corollary: reducing voltage requires reducing frequency
 - 15% reduction in voltage requires about 15% slow down in frequency
 - What’s the impact on dynamic power? $0.85^3 \approx 60\%$ -> 40% dynamic power reduction.

Dynamic Power Favors Parallelisms

$$P = k C f^3$$

- Dynamic power favors parallel processing over higher clock rate

Dynamic Power Favors Parallelisms

$$P = k C f^3$$

- Dynamic power favors parallel processing over higher clock rate
 - Take a core and replicate it 4 times: 4x speedup & **4x power**

Dynamic Power Favors Parallelisms

$$P = k C f^3$$

- Dynamic power favors parallel processing over higher clock rate
 - Take a core and replicate it 4 times: 4x speedup & **4x power**
 - Take a core and clock it 4 times faster: 4x speedup but **64x dynamic power!**

Dynamic Power Favors Parallelisms

$$P = k C f^3$$

- Dynamic power favors parallel processing over higher clock rate
 - Take a core and replicate it 4 times: 4x speedup & **4x power**
 - Take a core and clock it 4 times faster: 4x speedup but **64x dynamic power!**
- Another way to think about this

Dynamic Power Favors Parallelisms

$$P = k C f^3$$

- Dynamic power favors parallel processing over higher clock rate
 - Take a core and replicate it 4 times: 4x speedup & **4x power**
 - Take a core and clock it 4 times faster: 4x speedup but **64x dynamic power!**
- Another way to think about this
 - If a task can be perfectly parallelized by 4 cores, we can reduce the clock frequency of each core to 1/4 while retaining the same performance

Dynamic Power Favors Parallelisms

$$P = k C f^3$$

- Dynamic power favors parallel processing over higher clock rate
 - Take a core and replicate it 4 times: 4x speedup & **4x power**
 - Take a core and clock it 4 times faster: 4x speedup but **64x dynamic power!**
- Another way to think about this
 - If a task can be perfectly parallelized by 4 cores, we can reduce the clock frequency of each core to 1/4 while retaining the same performance
 - Dynamic power becomes $4 \times (1/4)^3 = 1/16$

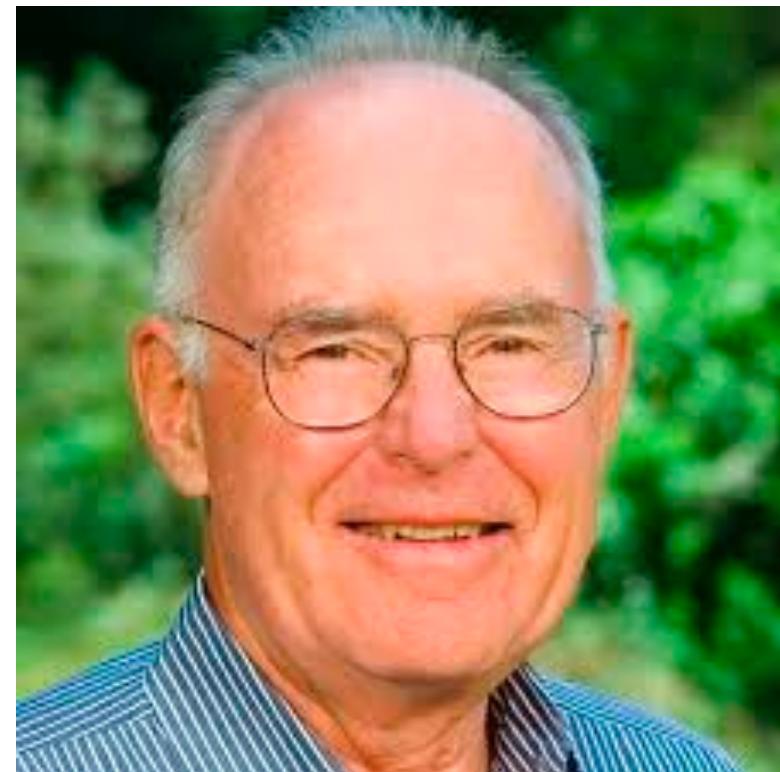
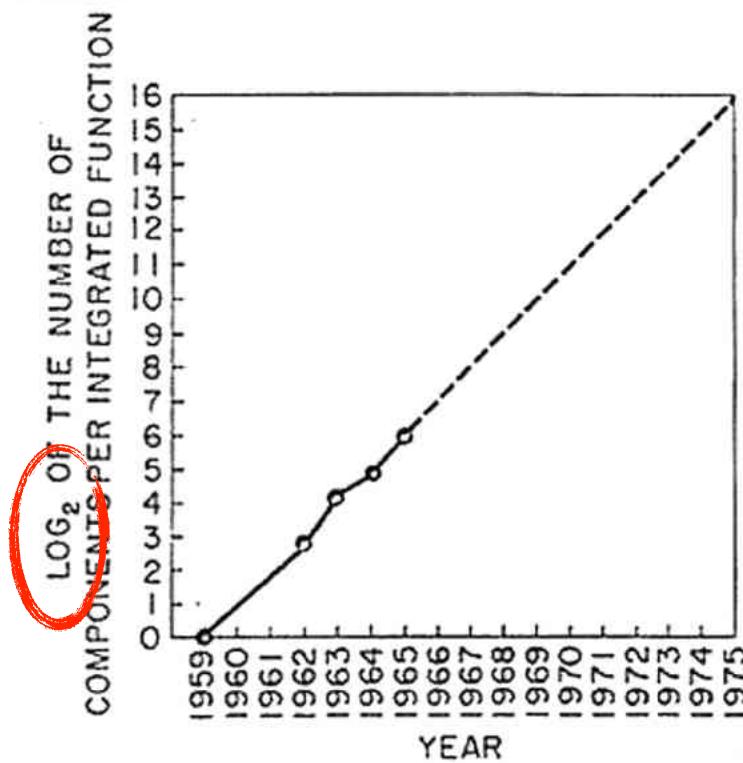
Moore's Law

- Gordon Moore in 1965 predicted that the number of transistors doubles every year



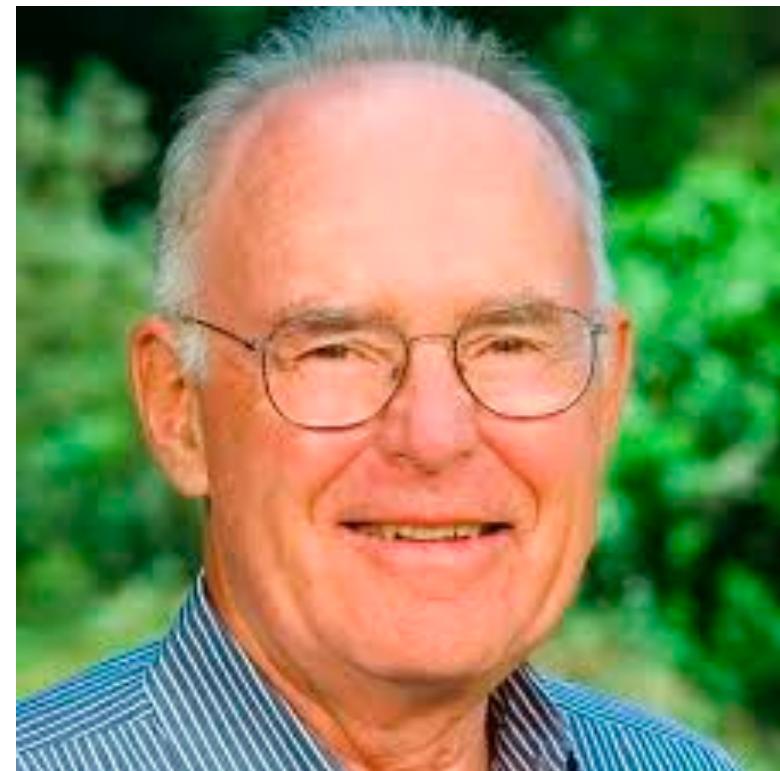
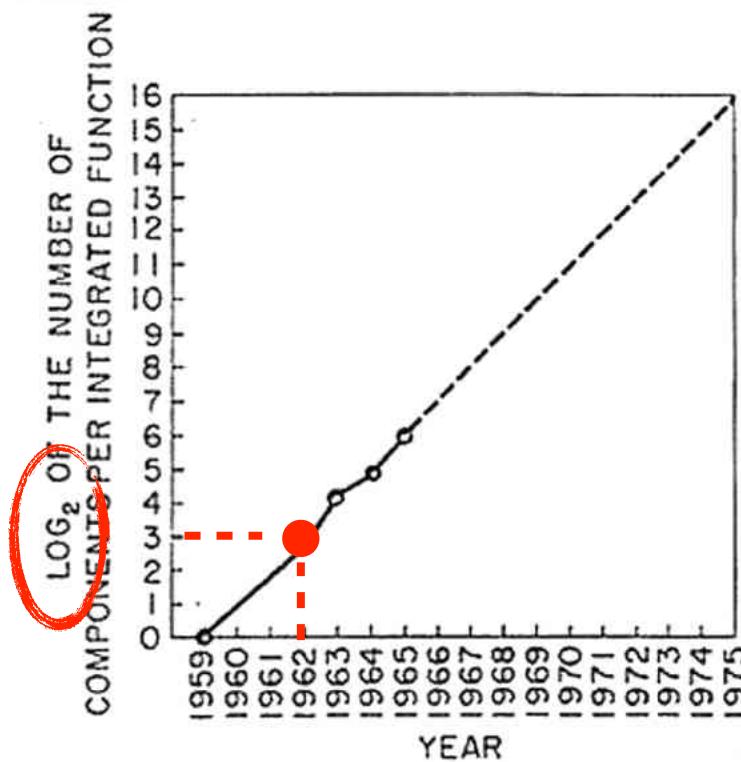
Moore's Law

- Gordon Moore in 1965 predicted that the number of transistors doubles every year



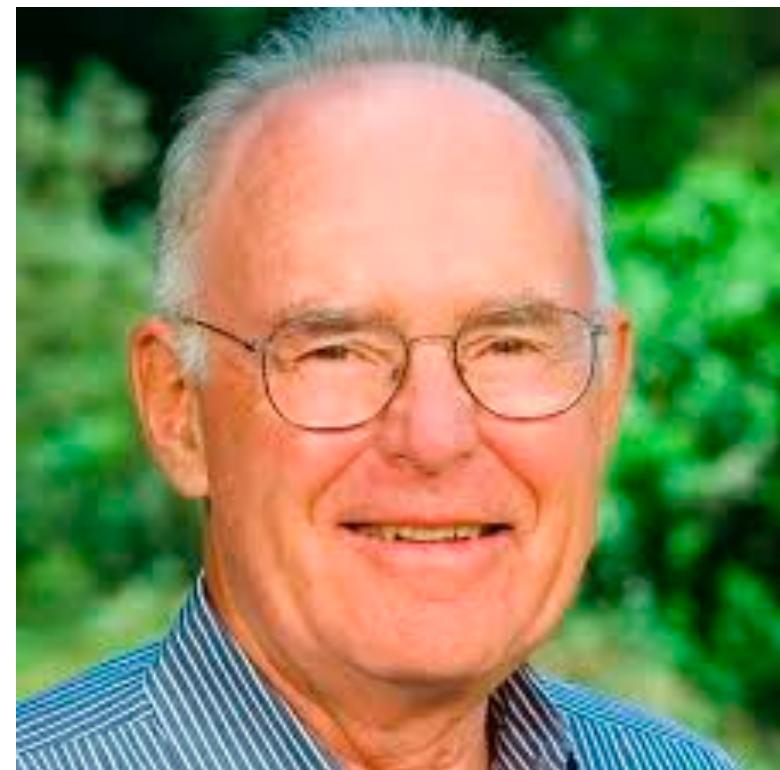
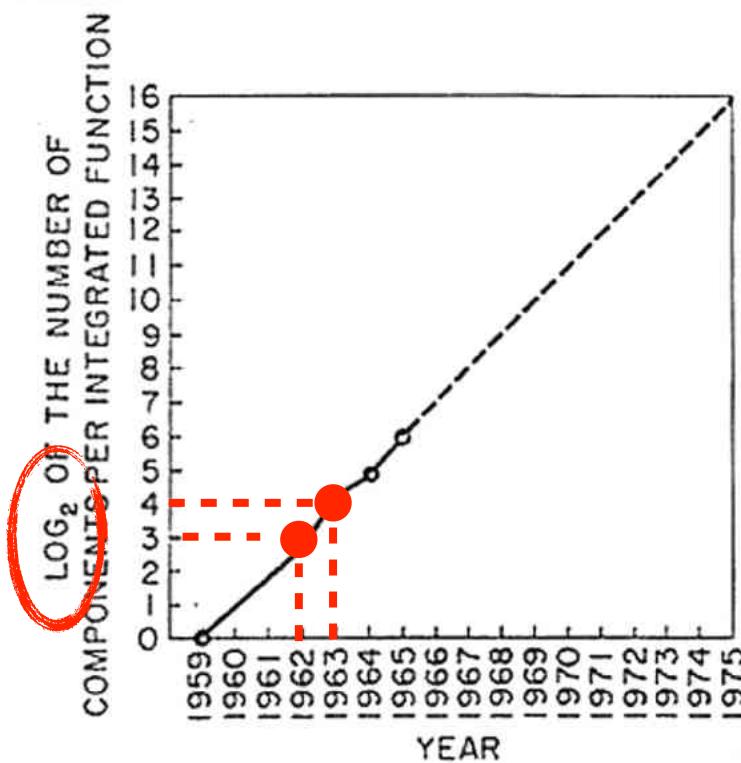
Moore's Law

- Gordon Moore in 1965 predicted that the number of transistors doubles every year



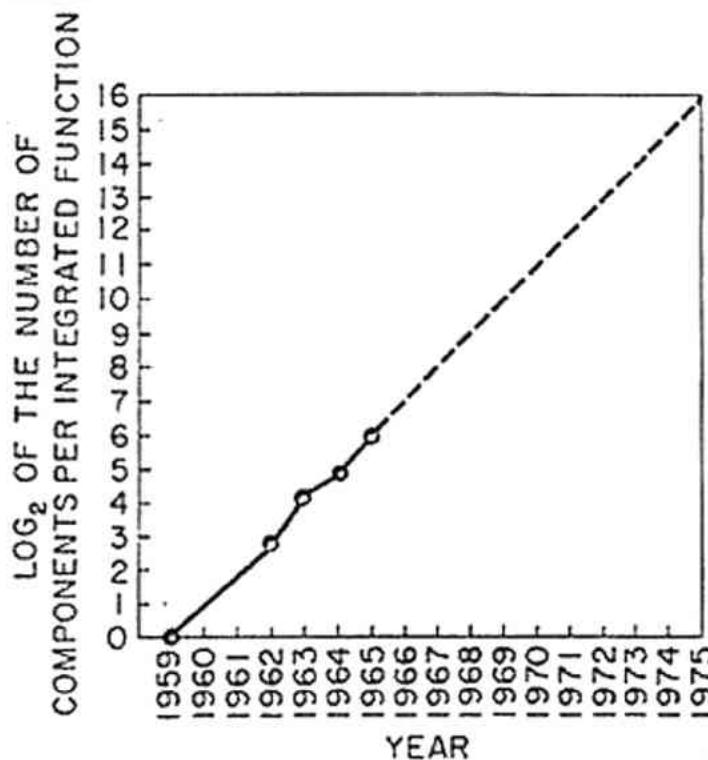
Moore's Law

- Gordon Moore in 1965 predicted that the number of transistors doubles every year



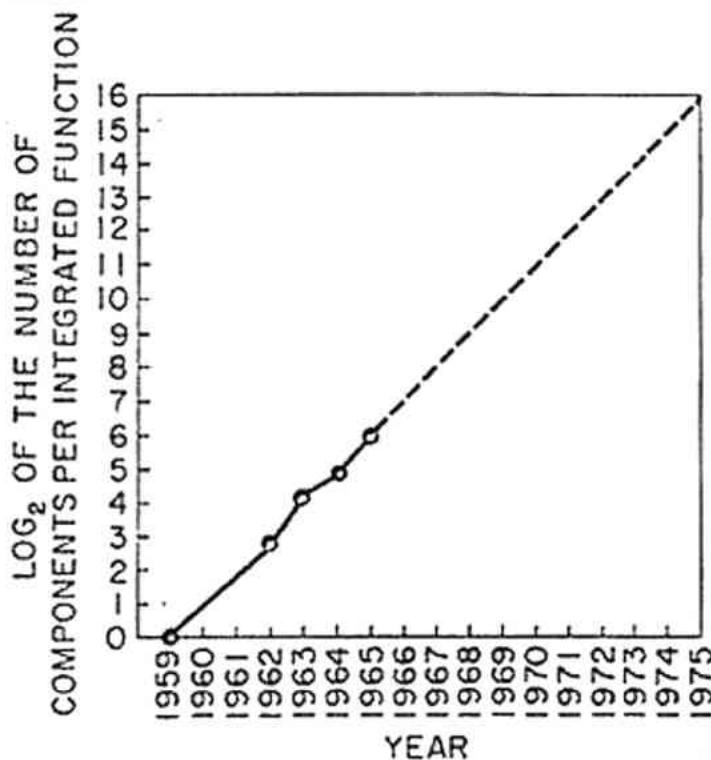
Moore's Law

- Gordon Moore in 1965 predicted that the number of transistors doubles every year
- In 1975 he revised the prediction to doubling every 2 years

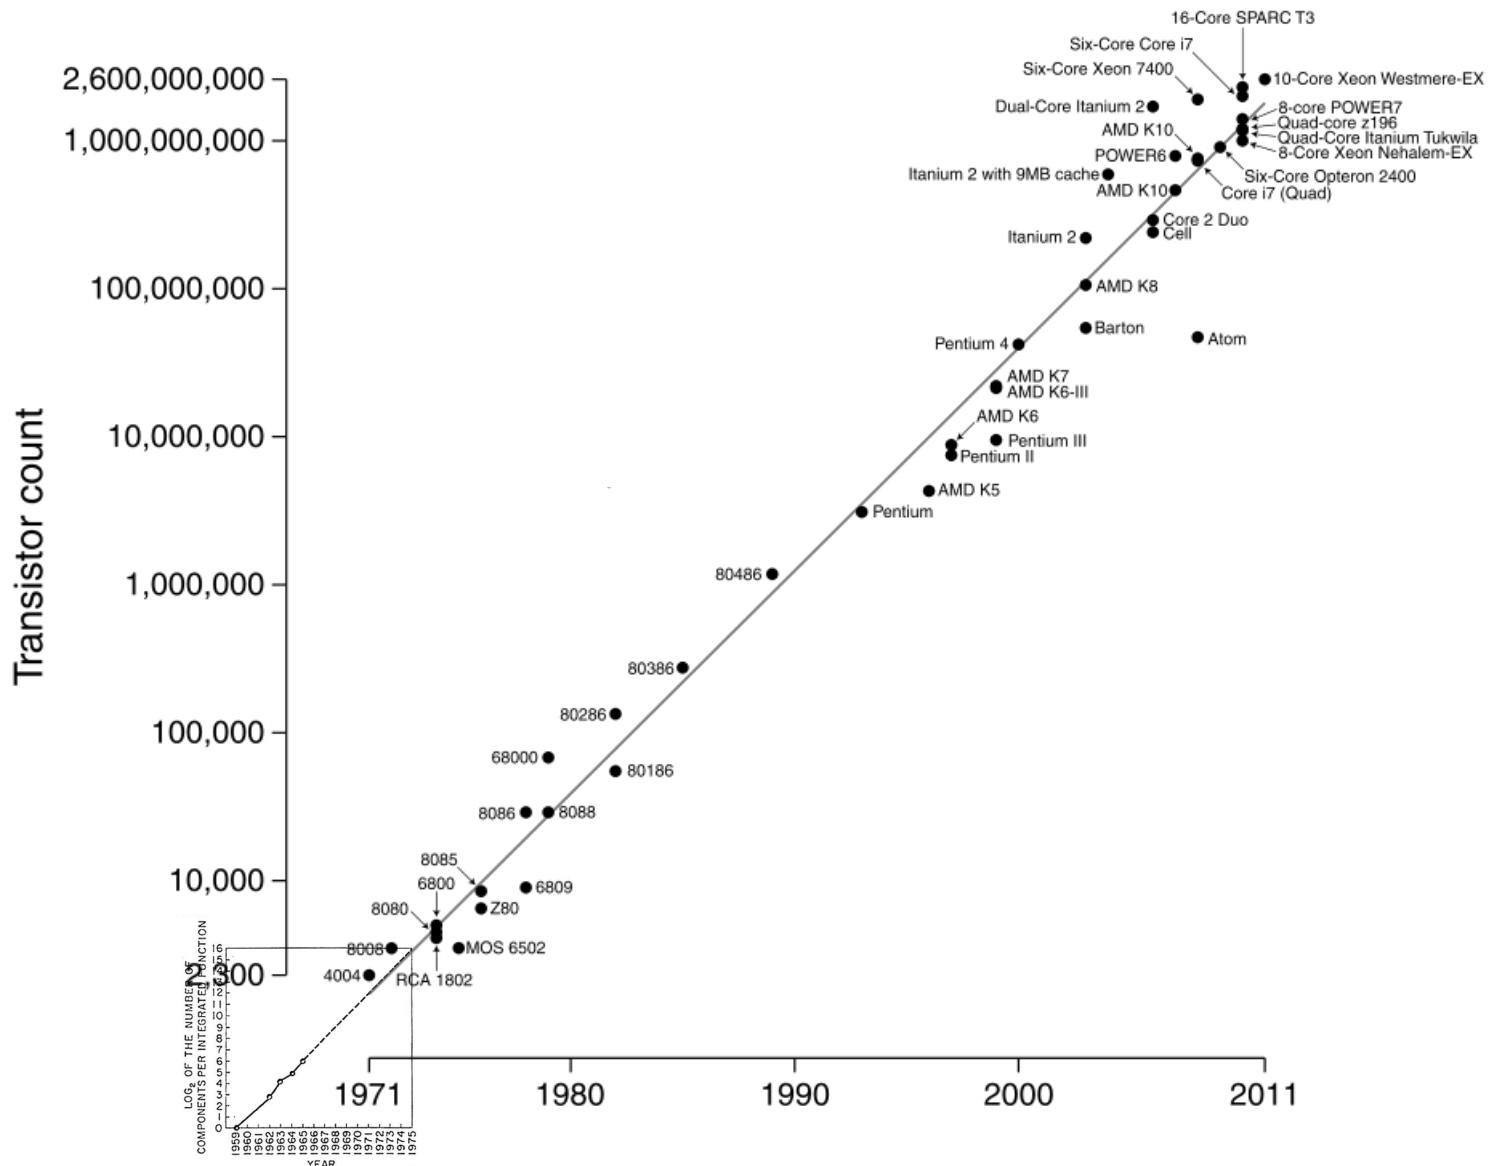


Moore's Law

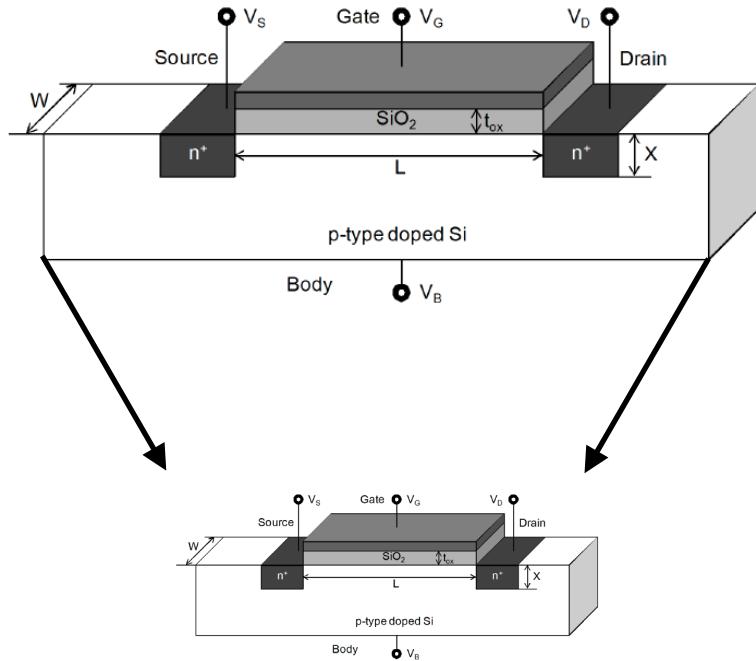
- Gordon Moore in 1965 predicted that the number of transistors doubles every year
- In 1975 he revised the prediction to doubling every 2 years
- Today's widely-known Moore's Law: number of transistors double about every 18 months (Moore never used the number 18...)



Moore's Law



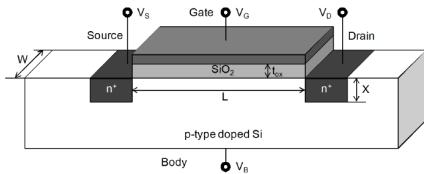
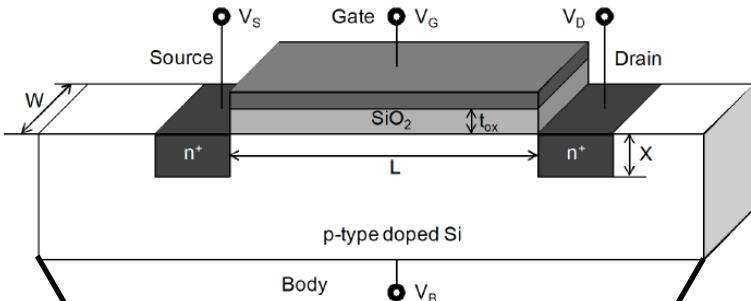
Dennard Scaling



Scale factor $\alpha < 1$

$\alpha = 0.7 \Rightarrow 2X$ more transistors!

Dennard Scaling

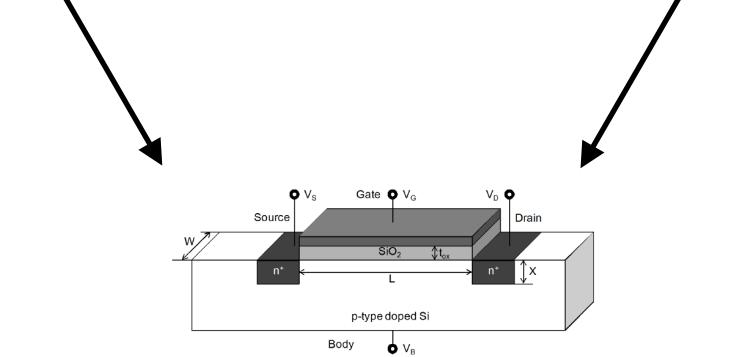
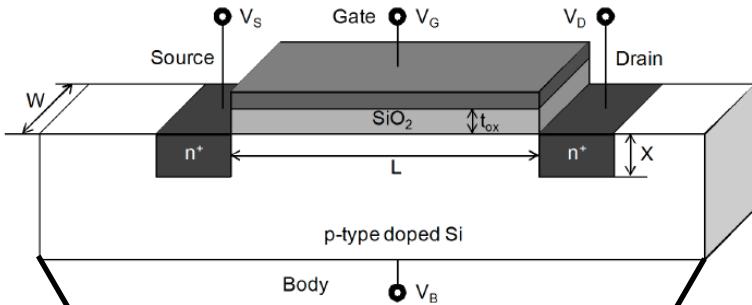


Parameter	Value	Scaled Value
Dopant concentrations	N_a, N_d	$N_a/a, N_d/a$
Dimensions	L, W, t_{ox}	aL, aW, aT_{ox}
Field	E	E
Voltage	V	aV
Capacitance	C	aC
Current	I	aI

Scale factor $\alpha < 1$

$\alpha = 0.7 \Rightarrow 2X$ more transistors!

Dennard Scaling

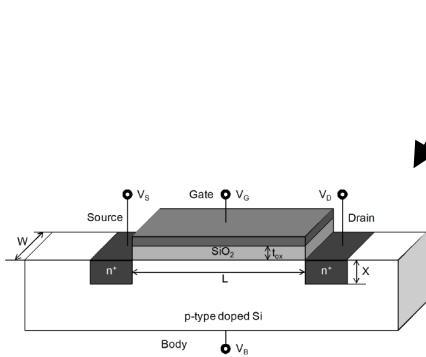
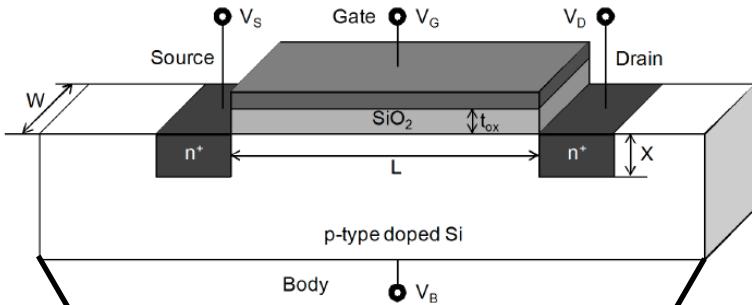


Parameter	Value	Scaled Value
Dopant concentrations	N_A, N_D	$N_A/\alpha, N_D/\alpha$
Dimensions	L, W, t_{ox}	$\alpha L, \alpha W, \alpha t_{ox}$
Field	E	E
Voltage	V	αV
Capacitance	C	αC
Current	I	αI
Transistors/Area	d	d/α^2

Scale factor $\alpha < 1$

$\alpha = 0.7 \Rightarrow 2X$ more transistors!

Dennard Scaling



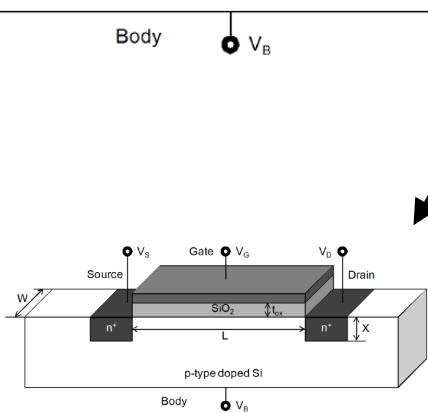
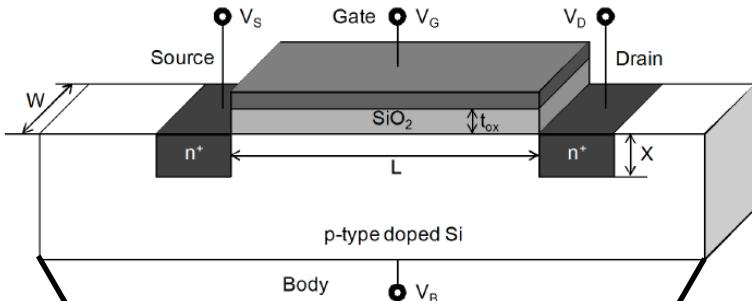
Parameter	Value	Scaled Value
Dopant concentrations	N_A, N_D	$N_A/a, N_D/a$
Dimensions	L, W, t_{ox}	aL, aW, aT_{ox}
Field	E	E
Voltage	V	aV
Capacitance	C	aC
Current	I	aI
Transistors/Area	d	d/a^2
Propagation time ($\sim CV/I$)	t	at
Frequency ($1/t$)	f	f/a



Scale factor $\alpha < 1$

$\alpha = 0.7 \Rightarrow 2X$ more transistors!

Dennard Scaling



Parameter	Value	Scaled Value
Dopant concentrations	N_A, N_D	$N_A/a, N_D/a$
Dimensions	L, W, t_{ox}	aL, aW, aT_{ox}
Field	E	E
Voltage	V	aV
Capacitance	C	aC
Current	I	aI
Transistors/Area	d	d/a^2

Propagation time ($\sim CV/I$)	t	at
Frequency ($1/t$)	f	f/a



Power (CV^2f)	P	a^2P
Power/area (Power density)	P_d	P_d

Scale factor $\alpha < 1$

$\alpha = 0.7 \Rightarrow 2X$ more transistors!

Implications of Dennard Scaling and Moore's Law

- Each new processor generation can have more transistors (Moore's Law), and will run at higher frequency but won't consume more power under the same area budget.

Implications of Dennard Scaling and Moore's Law

- Each new processor generation can have more transistors (Moore's Law), and will run at higher frequency but won't consume more power under the same area budget.
- More transistors means better microarchitecture, which leads to better performance even under the same frequency.

Implications of Dennard Scaling and Moore's Law

- Each new processor generation can have more transistors (Moore's Law), and will run at higher frequency but won't consume more power under the same area budget.
- More transistors means better microarchitecture, which leads to better performance even under the same frequency.
- Higher frequency means better performance even under the same microarchitecture.

Implications of Dennard Scaling and Moore's Law

- Each new processor generation can have more transistors (Moore's Law), and will run at higher frequency but won't consume more power under the same area budget.
- More transistors means better microarchitecture, which leads to better performance even under the same frequency.
- Higher frequency means better performance even under the same microarchitecture.
- Overall, software gets a free ride: wait for the next generation of hardware and performance will naturally increase without consuming more power.

Implications of Dennard Scaling and Moore's Law

- Each new processor generation can have more transistors (Moore's Law), and will run at higher frequency but won't consume more power under the same area budget.
- More transistors means better microarchitecture, which leads to better performance even under the same frequency.
- Higher frequency means better performance even under the same microarchitecture.
- Overall, software gets a free ride: wait for the next generation of hardware and performance will naturally increase without consuming more power.

Moore's law gave us more transistors;
Dennard scaling made them useful.

Bob Colwell, DAC 2013, June 4, 2013

2005: End of Dennard Scaling

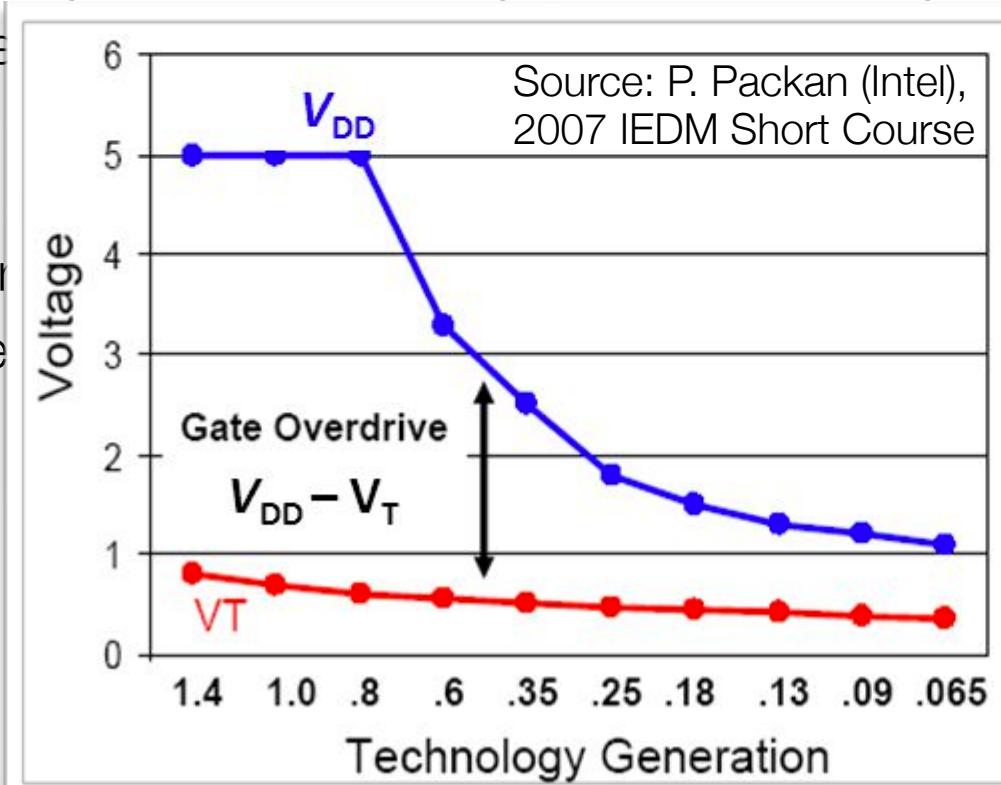
- What Happened?
 - Supply voltage V_{dd} stops scaling (Can't drop voltage below ~ 1 V)
 - Remember Power = CV^2f

2005: End of Dennard Scaling

- What Happened?
 - Supply voltage V_{dd} stops scaling (Can't drop voltage below ~ 1 V)
 - Remember Power = CV^2f
- Why?
 - There is a fundamental limit as to how much voltage we need to switch a transistor, called threshold voltage (V_{th}).
 - V_{th} stopped scaling because leakage power/reliability/variation becomes huge issues, and accordingly V_{dd} stops scaling

2005: End of Dennard Scaling

- What Happened?
 - Supply voltage V_{dd} stops scaling (Can't drop voltage below ~1 V)
 - Remember
- Why?
 - There is a switch a transistors
 - V_{th} stopped scaling becomes

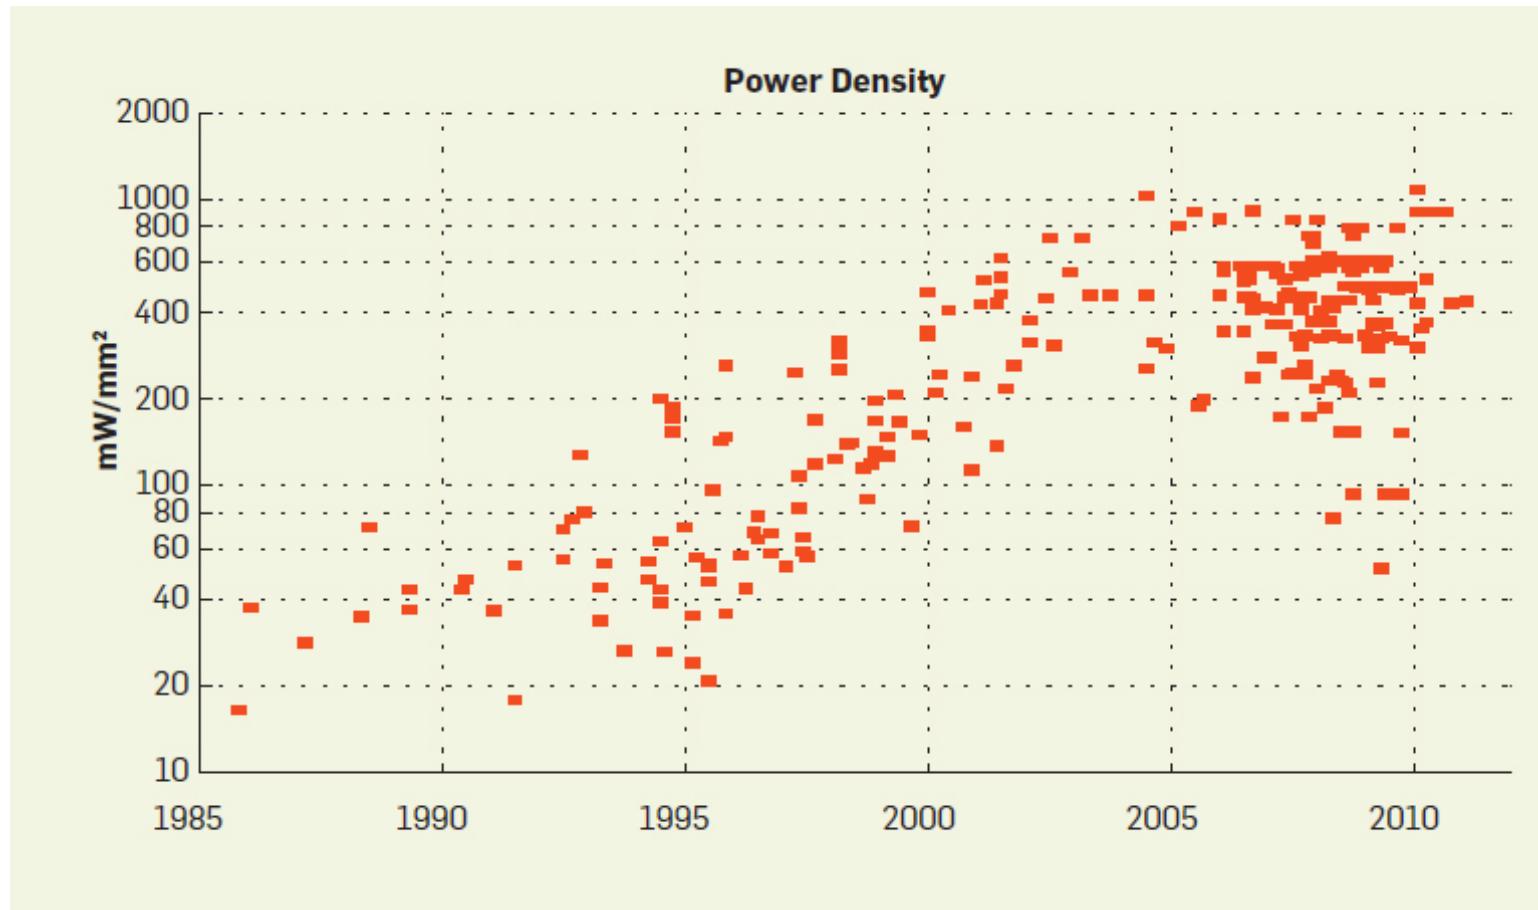


the need to
variation
g

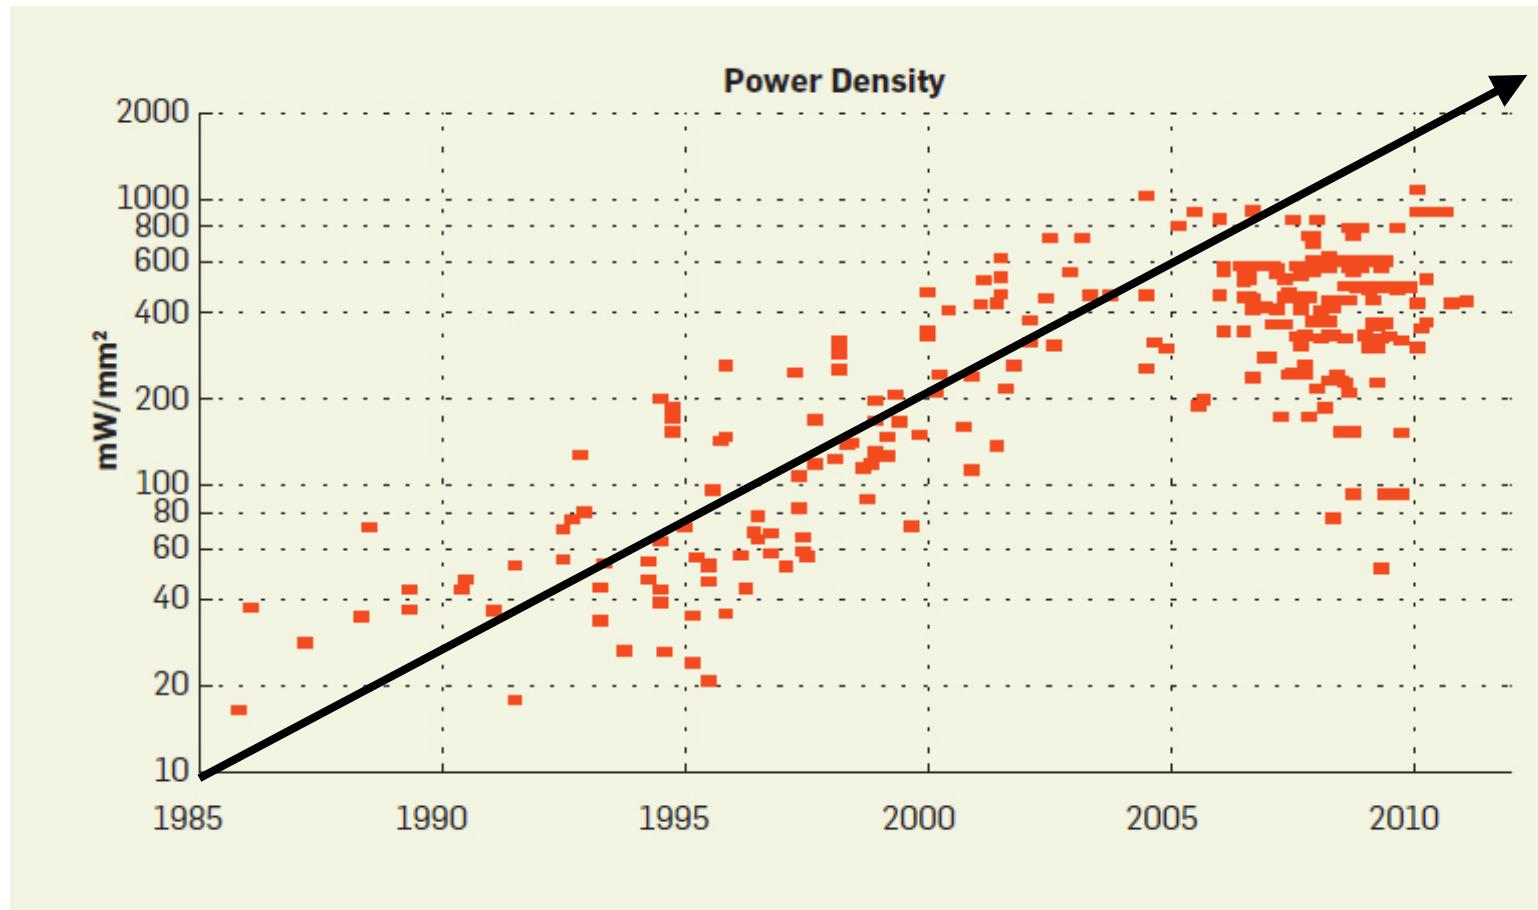
2005: End of Dennard Scaling

- What Happened?
 - Supply voltage V_{dd} stops scaling (Can't drop voltage below ~ 1 V)
 - Remember Power = CV^2f
- Why?
 - There is a fundamental limit as to how much voltage we need to switch a transistor, called threshold voltage (V_{th}).
 - V_{th} stopped scaling because leakage power/reliability/variation becomes huge issues, and accordingly V_{dd} stops scaling
- The demise of Dennard Scaling means the power density (power consumption per unit area) will increase rather than staying stable.

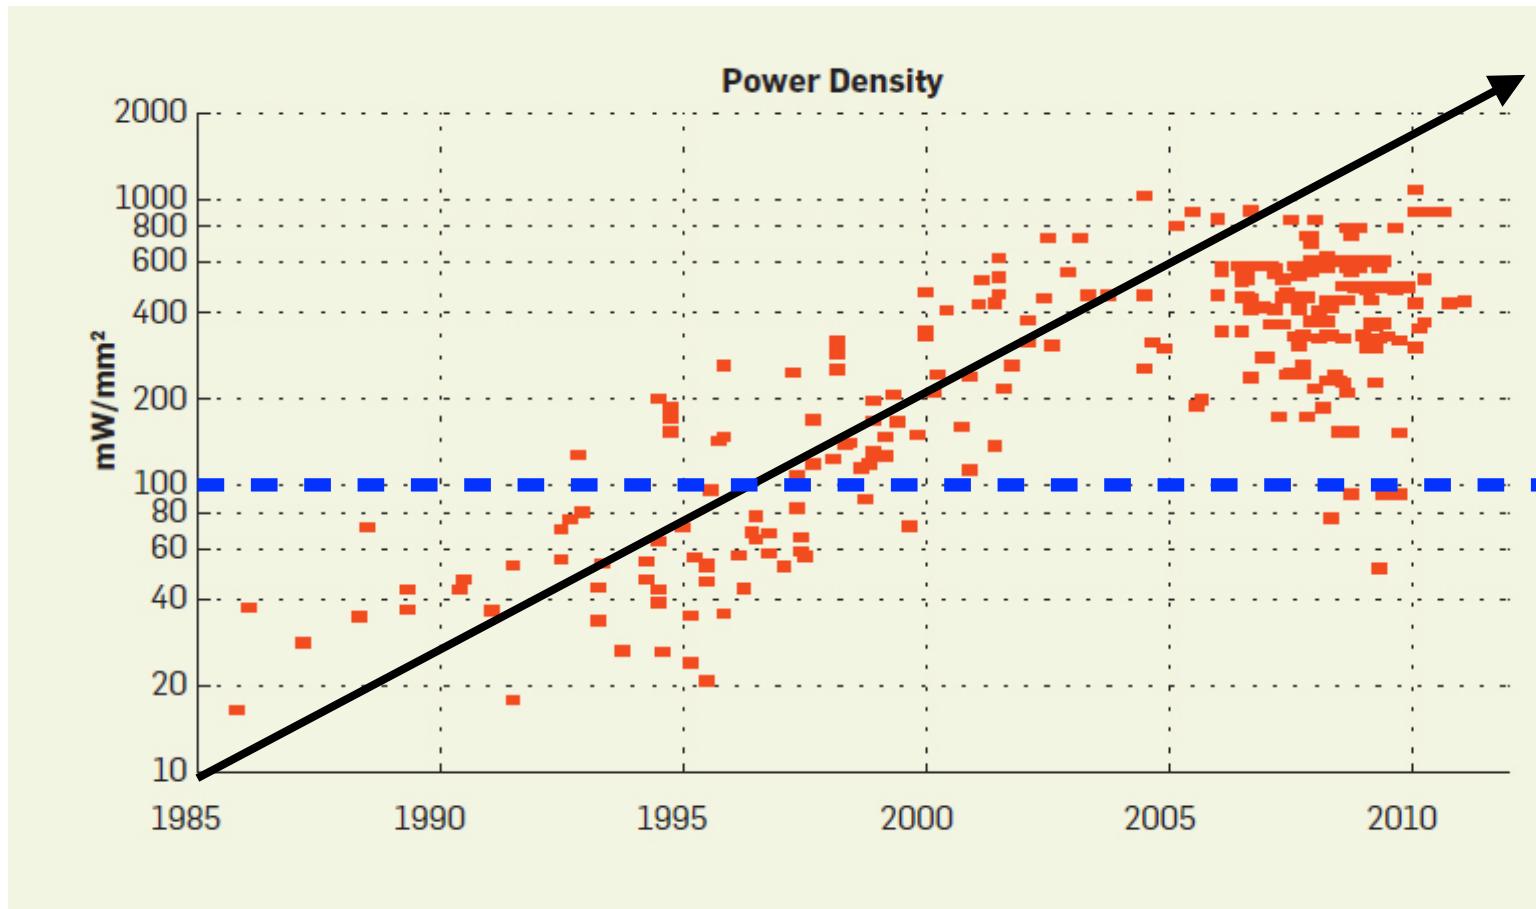
2005: End of Dennard Scaling



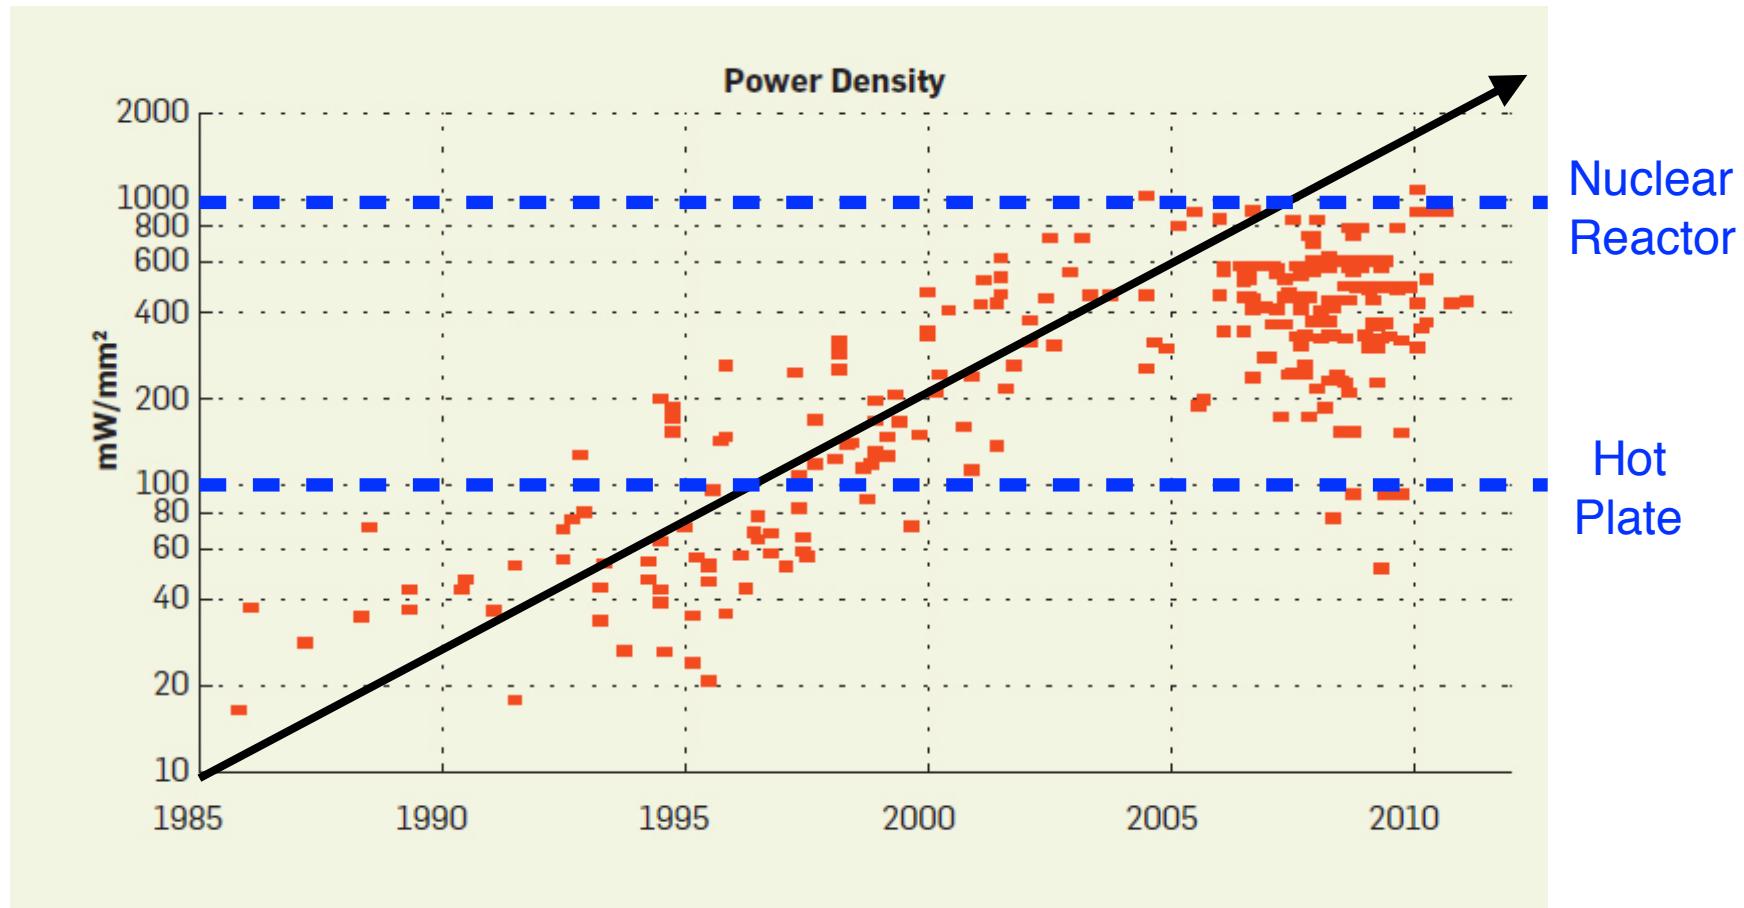
2005: End of Dennard Scaling



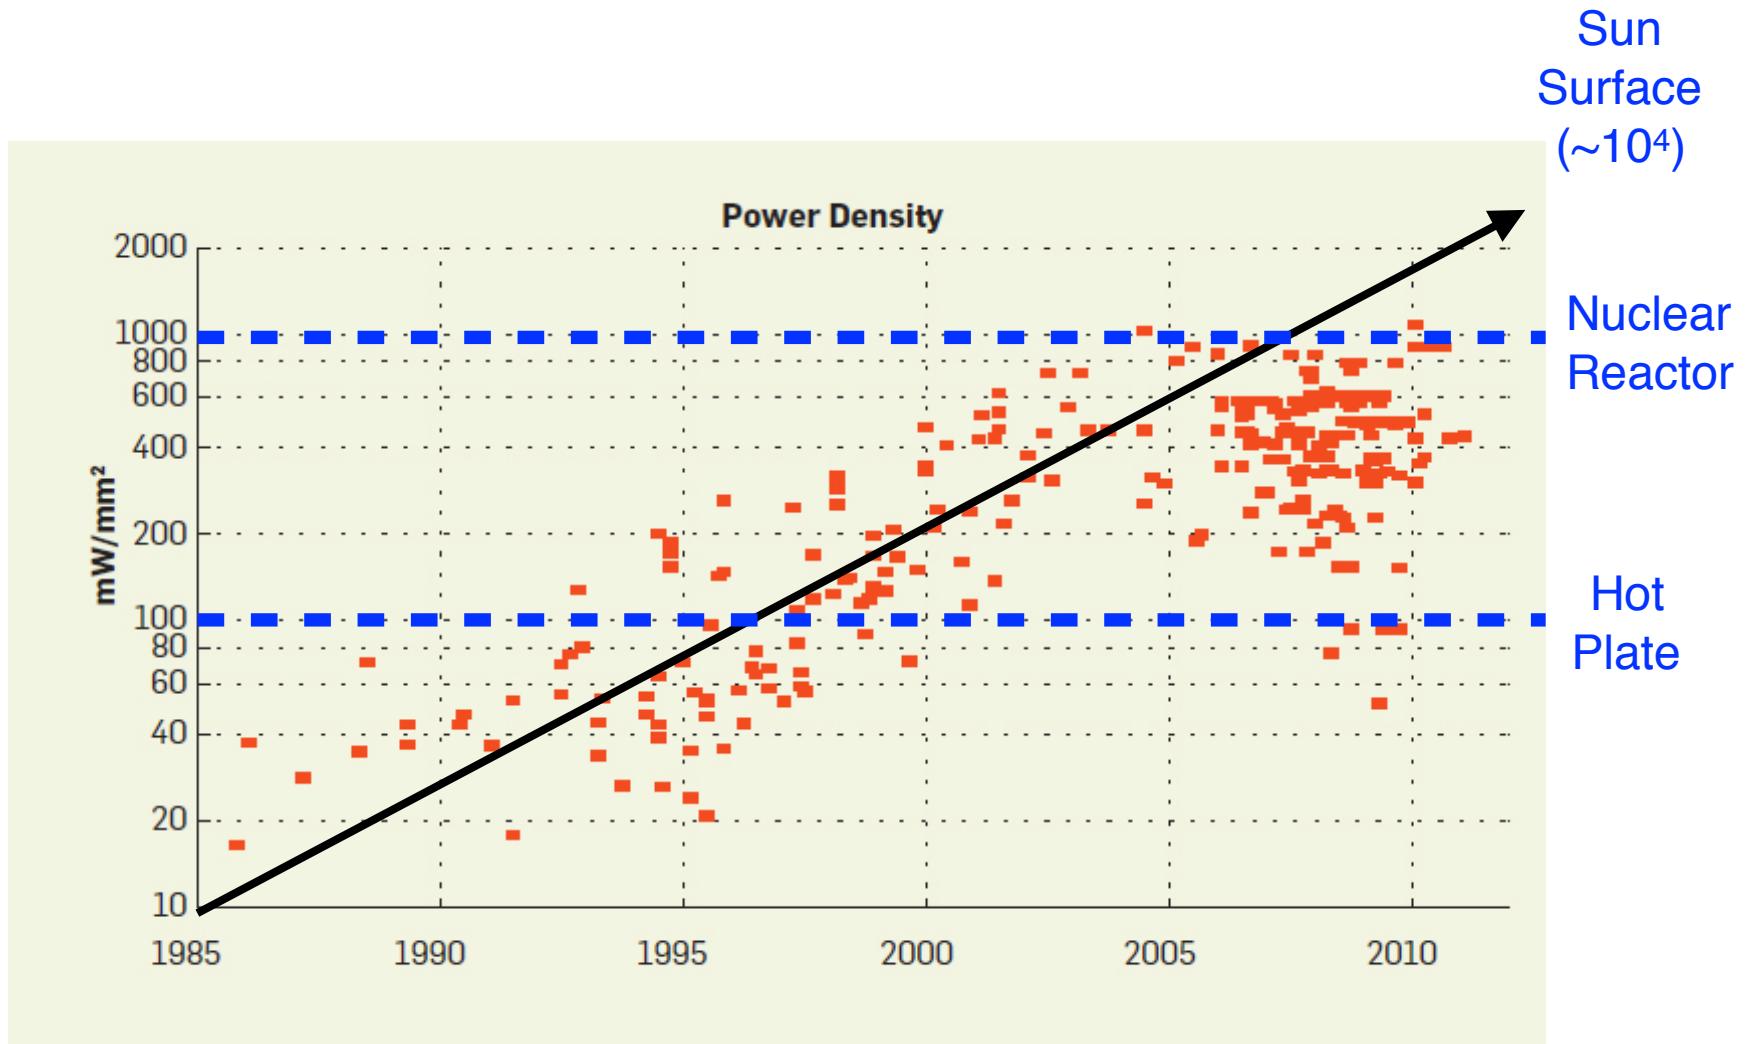
2005: End of Dennard Scaling



2005: End of Dennard Scaling



2005: End of Dennard Scaling



Dark Silicon

n. [därk, sǐl'i-kən, -kön']

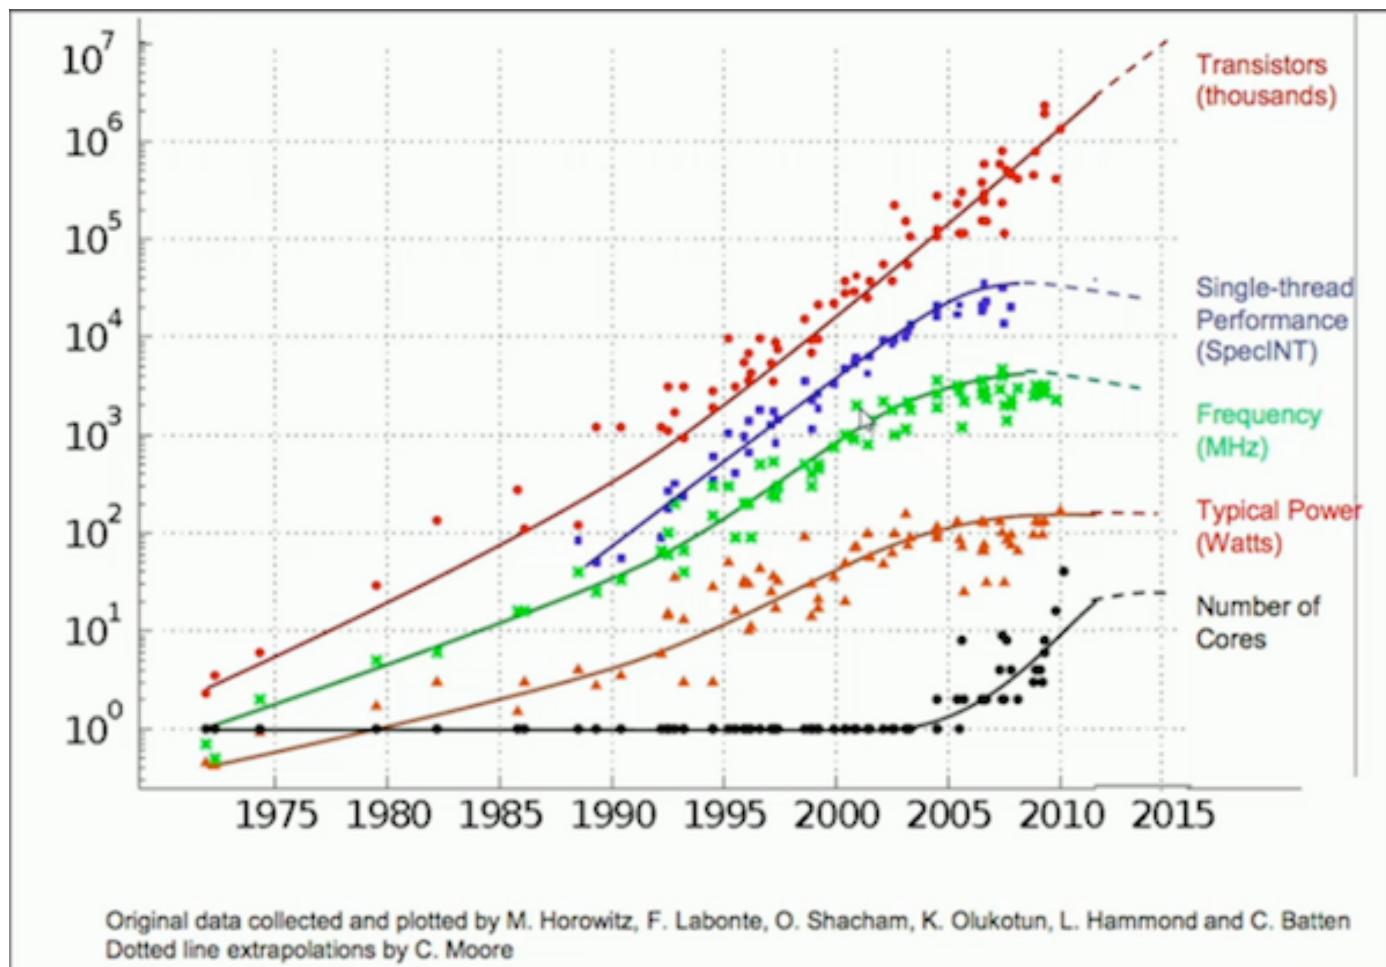
More transistors on chip (due to Moore's Law), but a growing fraction cannot actually be used due to power limits (due to the end of Dennard Scaling).

2005: End of Dennard Scaling

- Initial response has been to lower frequency and increase cores / chip

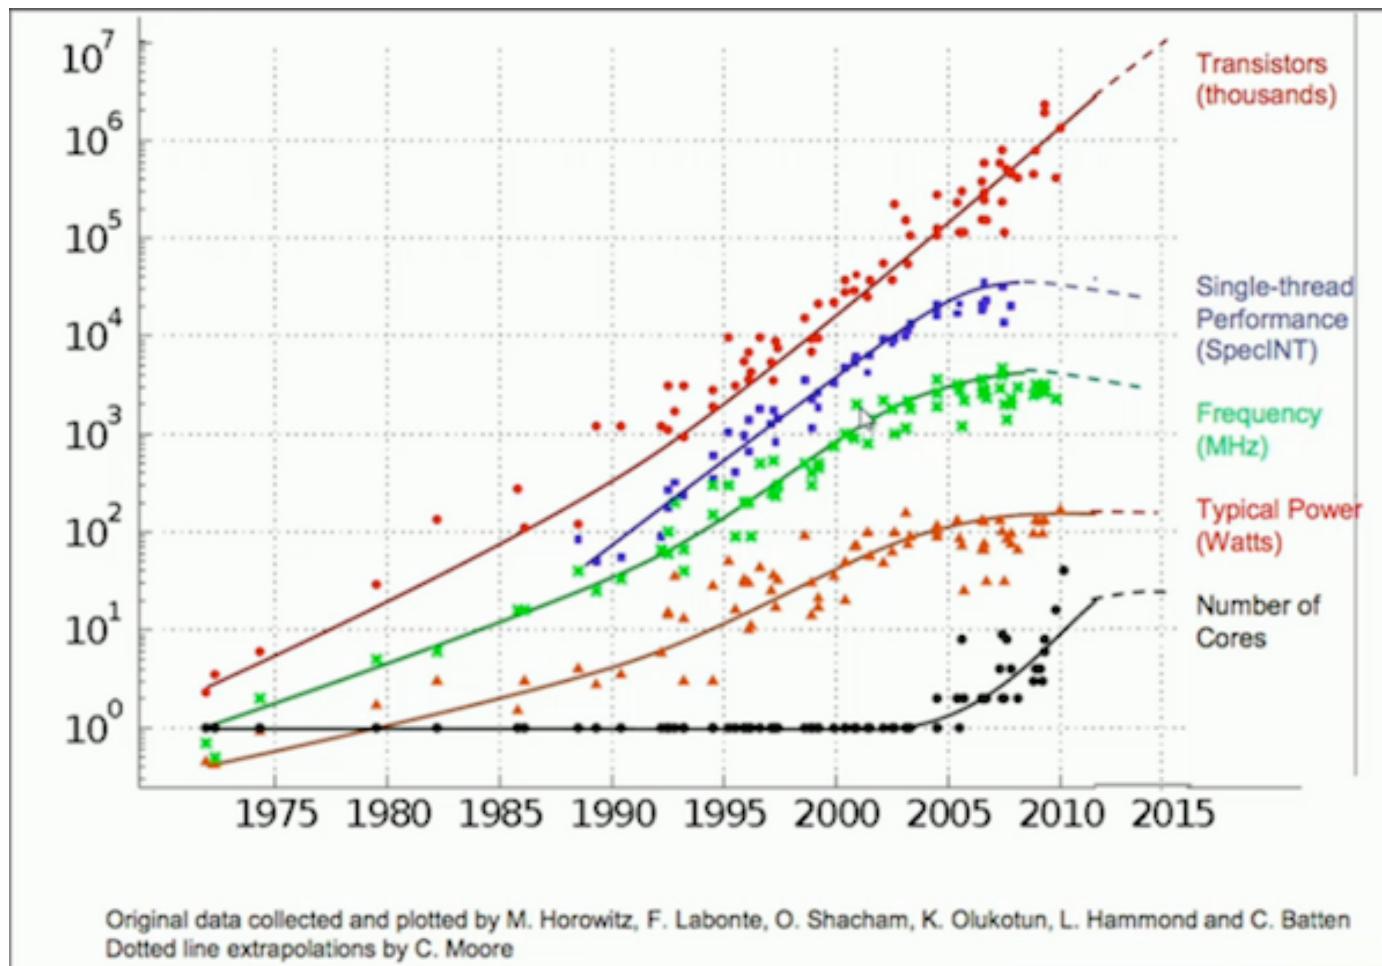
2005: End of Dennard Scaling

- Initial response has been to lower frequency and increase cores / chip



2005: End of Dennard Scaling

- Initial response has been to lower frequency and increase cores / chip
- There is a limit to core scaling. Why?



2007: A Revolutionary New Computer





No Moore's Law for batteries

Fred Schlachter¹

American Physical Society, Washington, DC 20045

The public has become accustomed to rapid progress in mobile phone technology, computers, and access to information; tablet computers, smart phones, and other powerful new devices are familiar to most people on the planet.

These developments are due in part to the ongoing exponential increase in computer processing power, doubling approximately every 2 years for the past several decades. This pattern is usually called Moore's Law and is named for Gordon Moore, a co-founder of Intel. The law is not a law like that for gravity; it is an empirical observation, which has become a self-fulfilling prophecy. Unfortunately, much of the public has come to expect that all technology does, will, or should follow such a law, which is not consistent with our everyday observations: For example, the maximum

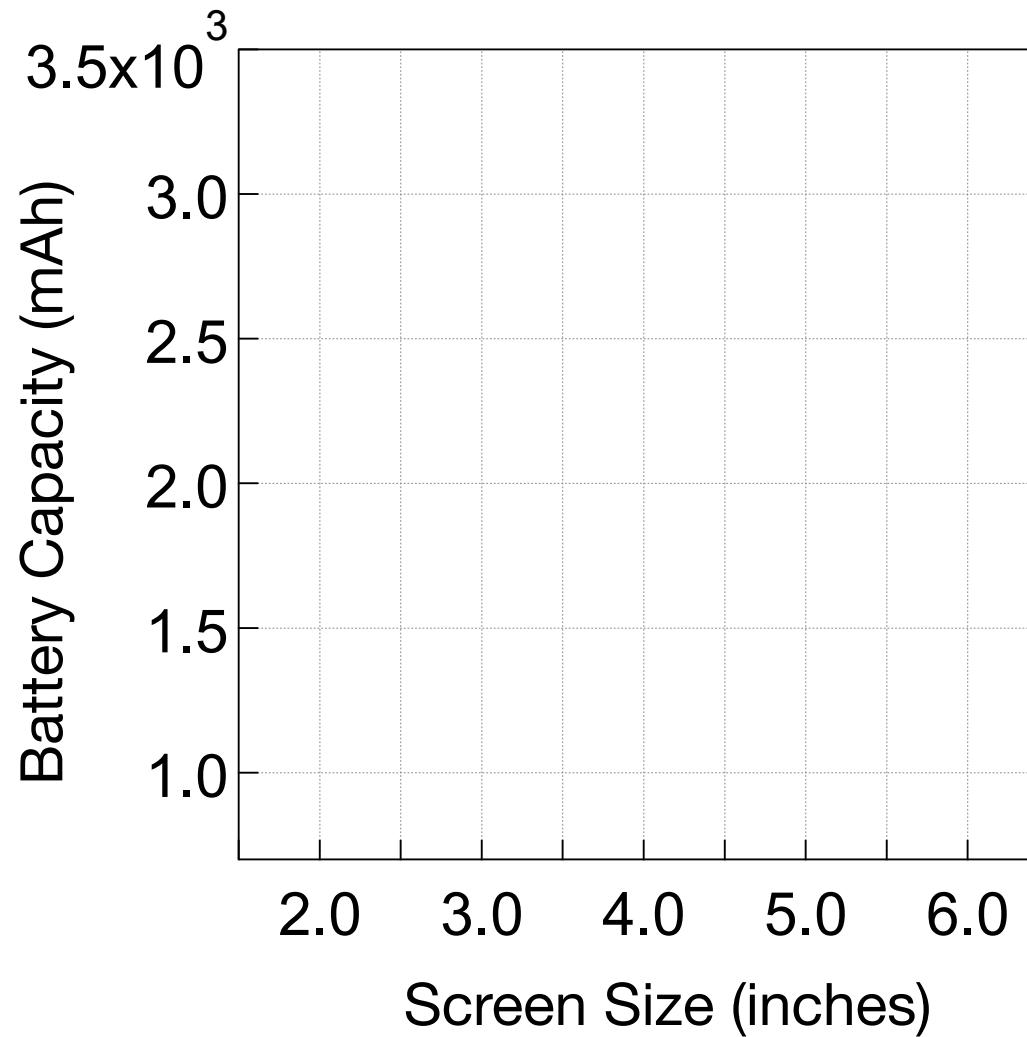
there is a Moore's Law for computer processors is that electrons are small and they do not take up space on a chip. Chip performance is limited by the lithography technology used to fabricate the chips; as lithography improves ever smaller features can be made on processors. Batteries are not like this. Ions, which transfer charge in batteries, are large, and they take up space, as do anodes, cathodes, and electrolytes. A D-cell battery stores more energy than an AA-cell. Potentials in a battery are dictated by the relevant chemical reactions, thus limiting eventual battery performance. Significant improvement in battery capacity can only be made by changing to a different chemistry.

Scientists and battery experts, who have been optimistic in the recent past about improving lithium-ion batteries and about de-

“Improving” Energy Capacity

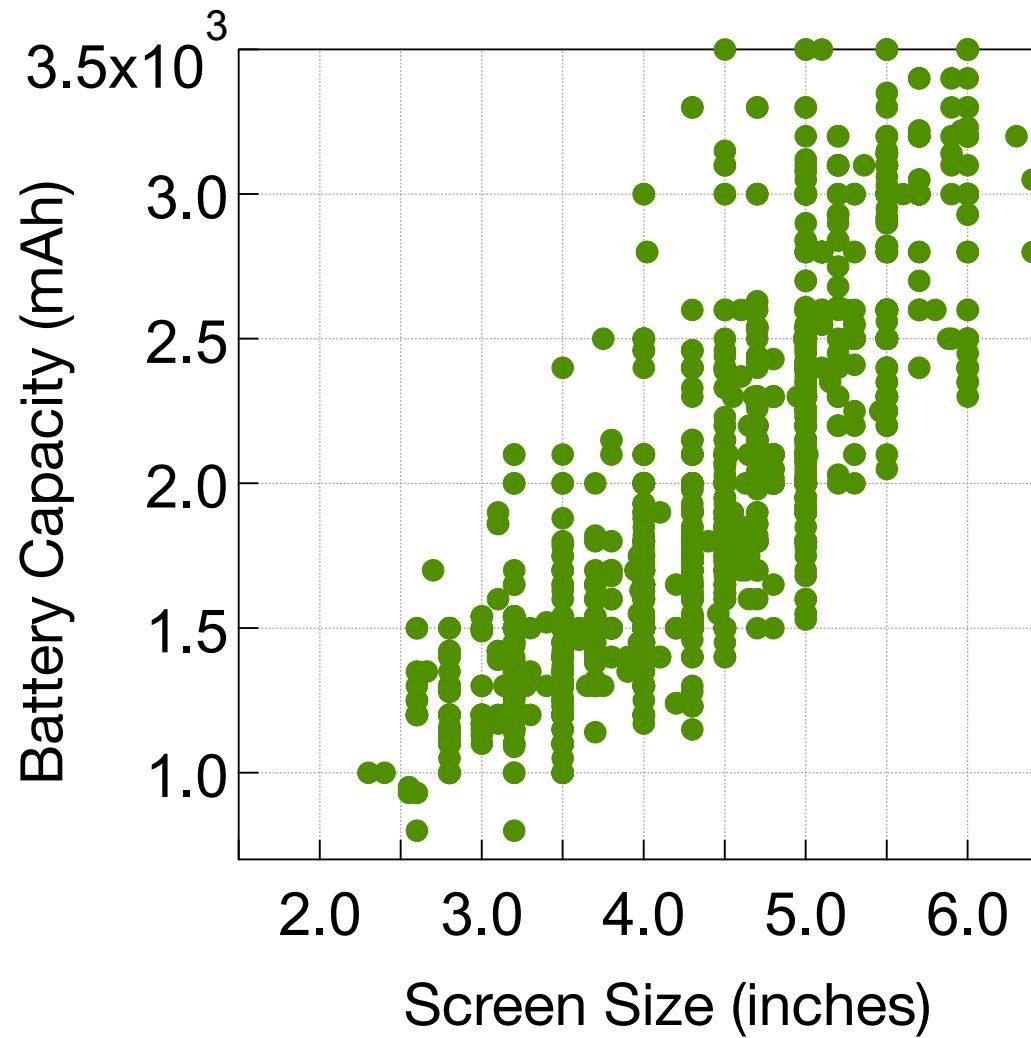
600 smartphone from 2006 to 2014 on <http://www.gsmarena.com/makers.php3>

“Improving” Energy Capacity



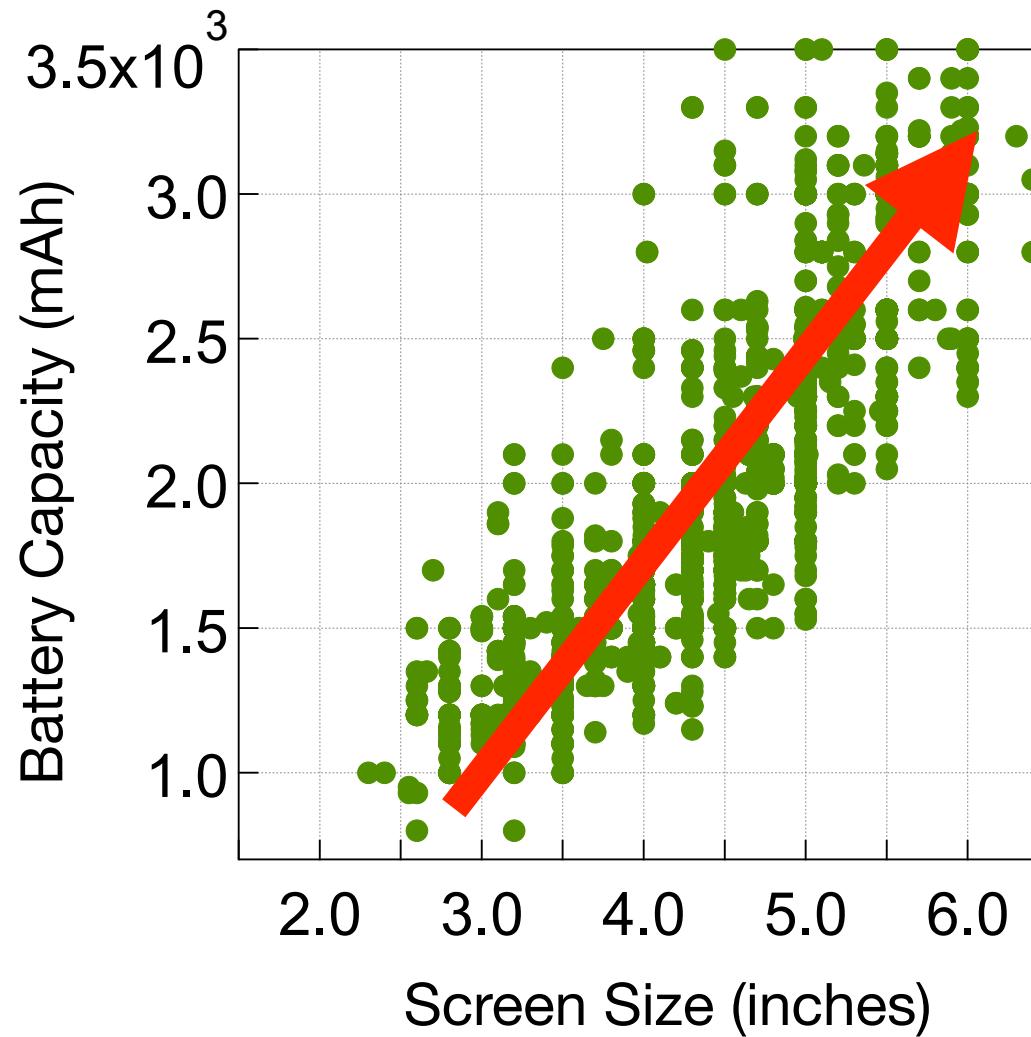
600 smartphone from 2006 to 2014 on <http://www.gsmarena.com/makers.php3>

“Improving” Energy Capacity



600 smartphone from 2006 to 2014 on <http://www.gsmarena.com/makers.php3>

“Improving” Energy Capacity



600 smartphone from 2006 to 2014 on <http://www.gsmarena.com/makers.php3>

“Improving” Energy Capacity

“Improving” Energy Capacity



“Improving” Energy Capacity



“Improving” Energy Capacity

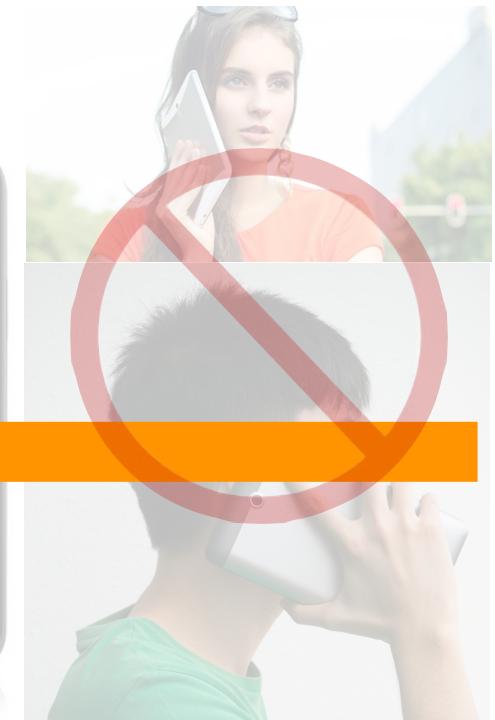
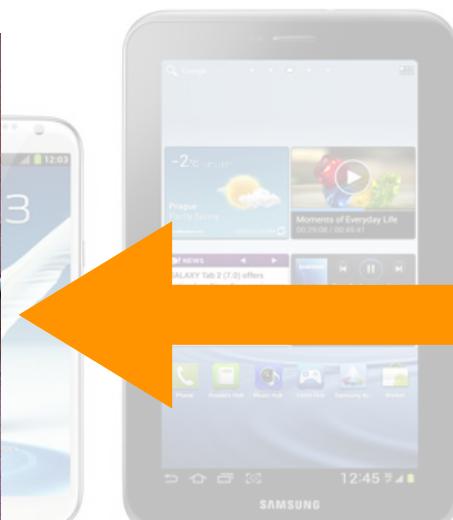


SMARTPHONE

PHABLET

TABLET

“Improving” Energy Capacity

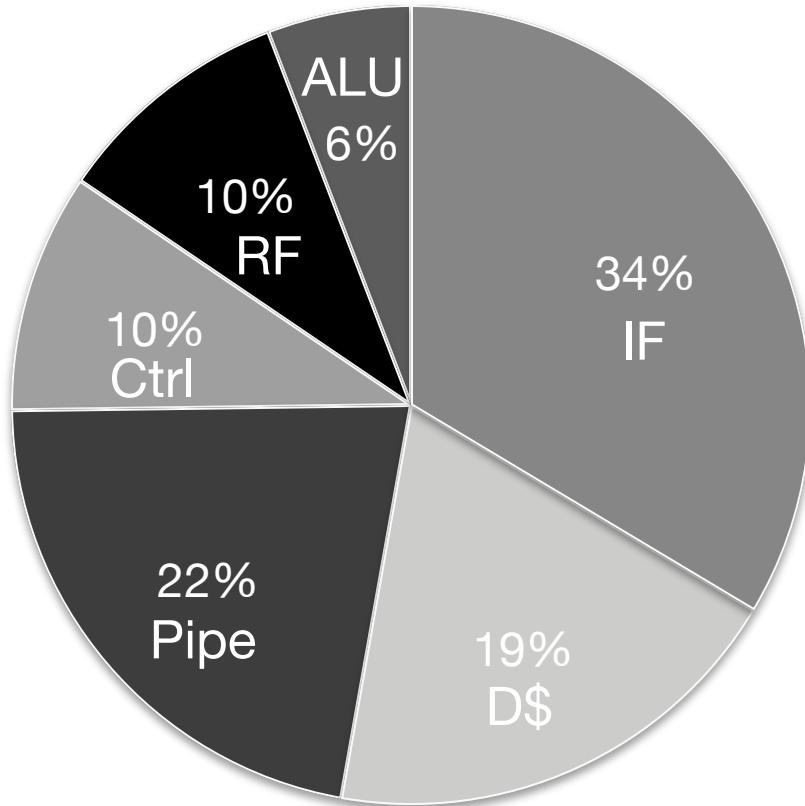


Sources of Energy-Inefficiencies

General-Purpose CPU = Instruction Delivery + Data Feeding + Execution + Control, where instruction delivery, data feeding & control are pure overhead

Sources of Energy-Inefficiencies

General-Purpose CPU = Instruction Delivery + Data Feeding + Execution + Control, where instruction delivery, data feeding & control are pure overhead



IF: Instruction fetch

Ctrl: Other control logics

Pipe: Pipeline reg, bus, clock

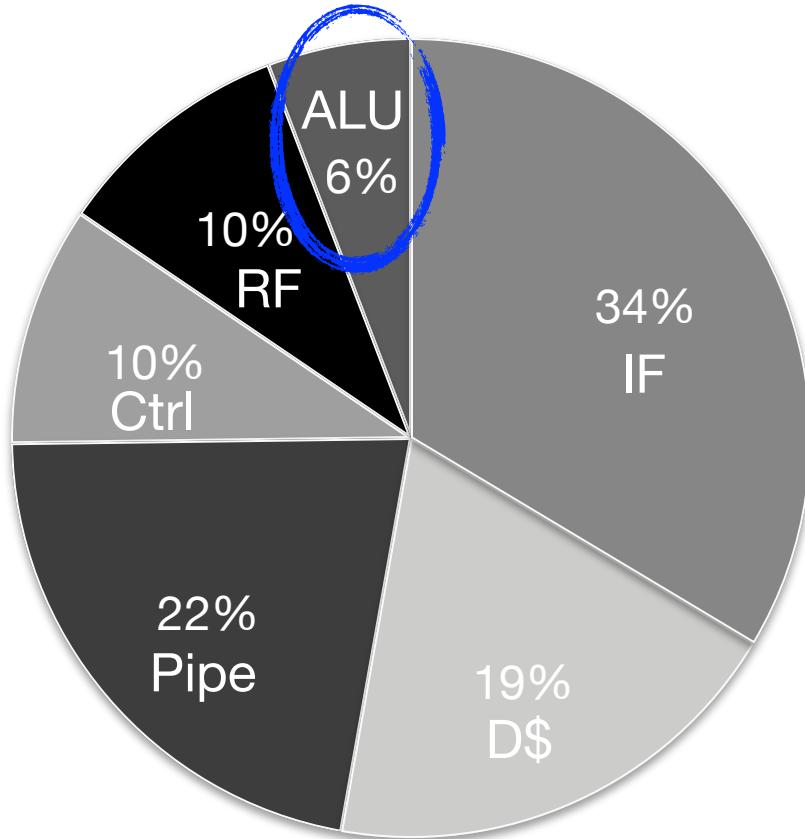
D\$: Data cache

RF: Register file

ALU: Functional units

Sources of Energy-Inefficiencies

General-Purpose CPU = Instruction Delivery + Data Feeding + Execution + Control, where instruction delivery, data feeding & control are pure overhead



IF: Instruction fetch

Ctrl: Other control logics

Pipe: Pipeline reg, bus, clock

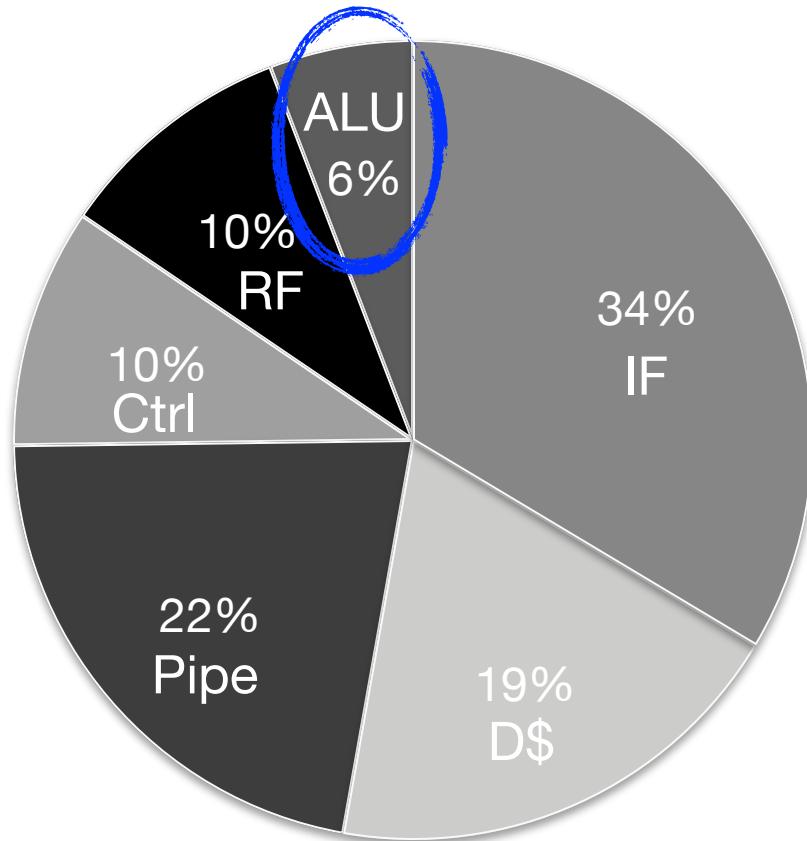
D\$: Data cache

RF: Register file

ALU: Functional units

Sources of Energy-Inefficiencies

General-Purpose CPU = Instruction Delivery + Data Feeding + Execution + Control, where instruction delivery, data feeding & control are pure overhead



Pure Overhead

IF: Instruction fetch

Ctrl: Other control logics

Pipe: Pipeline reg, bus, clock

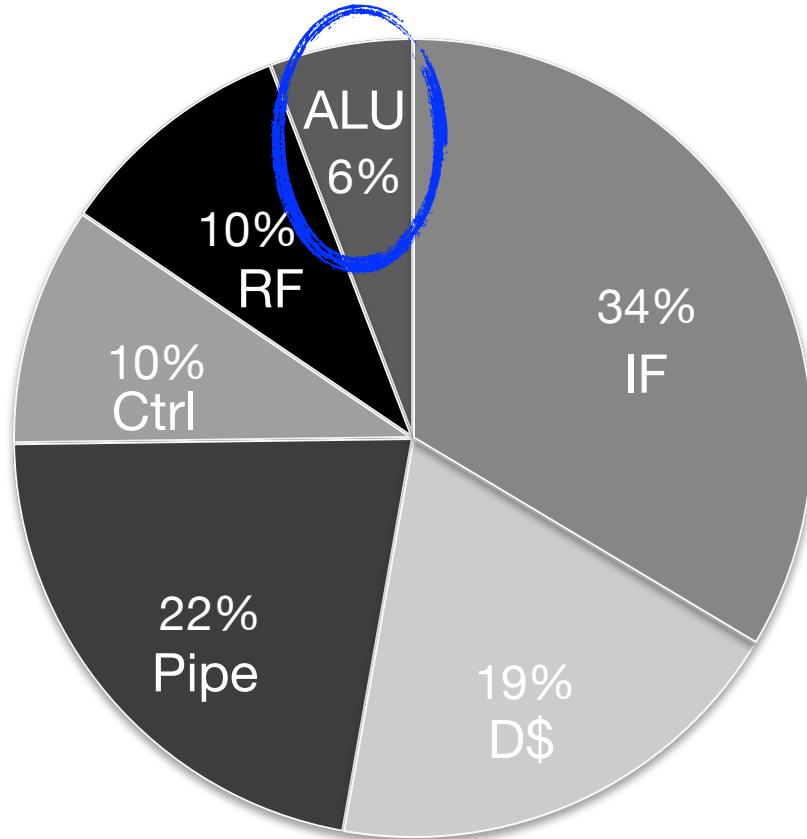
D\$: Data cache

RF: Register file

ALU: Functional units

Sources of Energy-Inefficiencies

General-Purpose CPU = Instruction Delivery + Data Feeding + Execution + Control, where instruction delivery, data feeding & control are pure overhead



Pure Overhead

IF: Instruction fetch

Ctrl: Other control logics

Pipe: Pipeline reg, bus, clock

D\$: Data cache

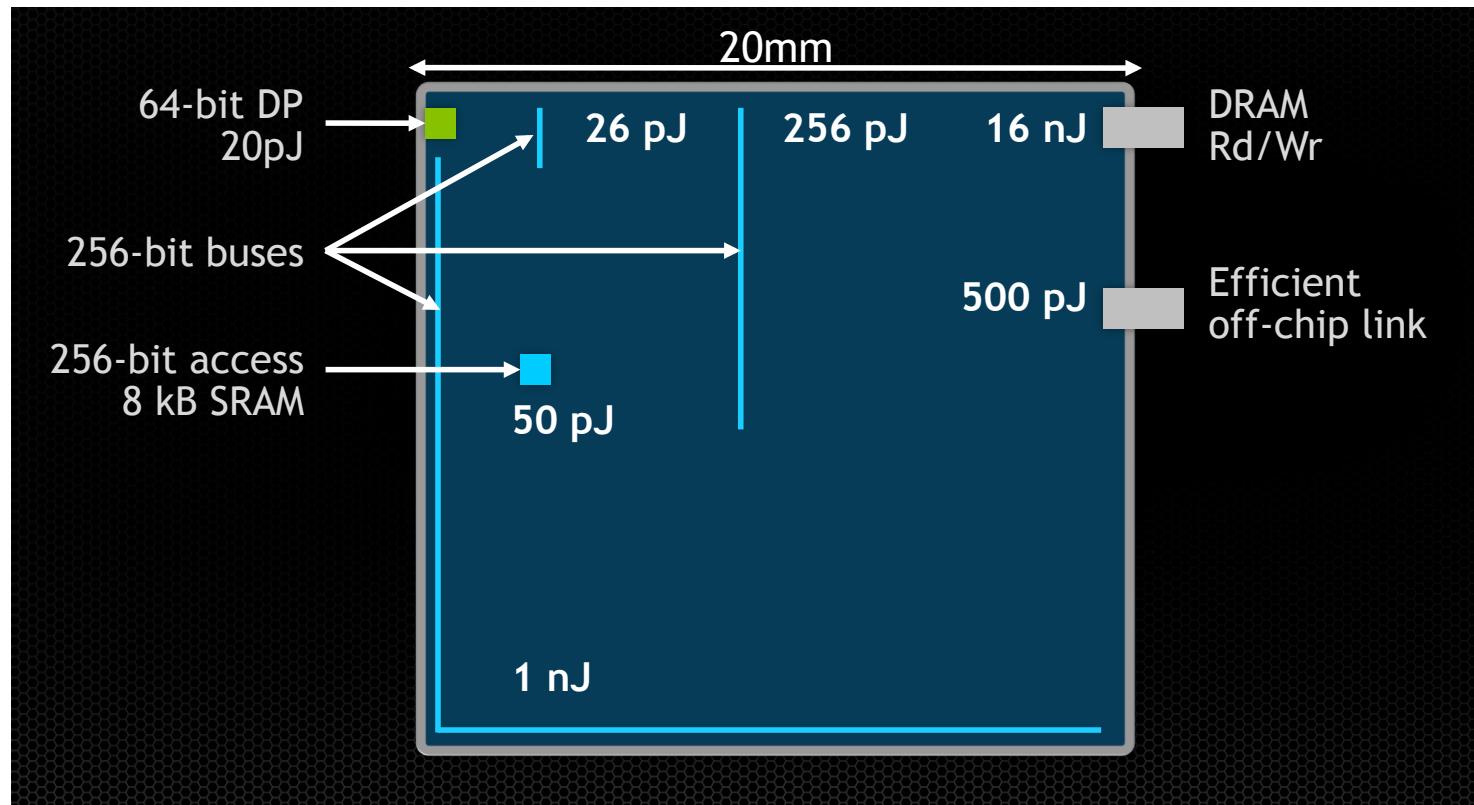
RF: Register file

ALU: Functional units

Doing Actual Work

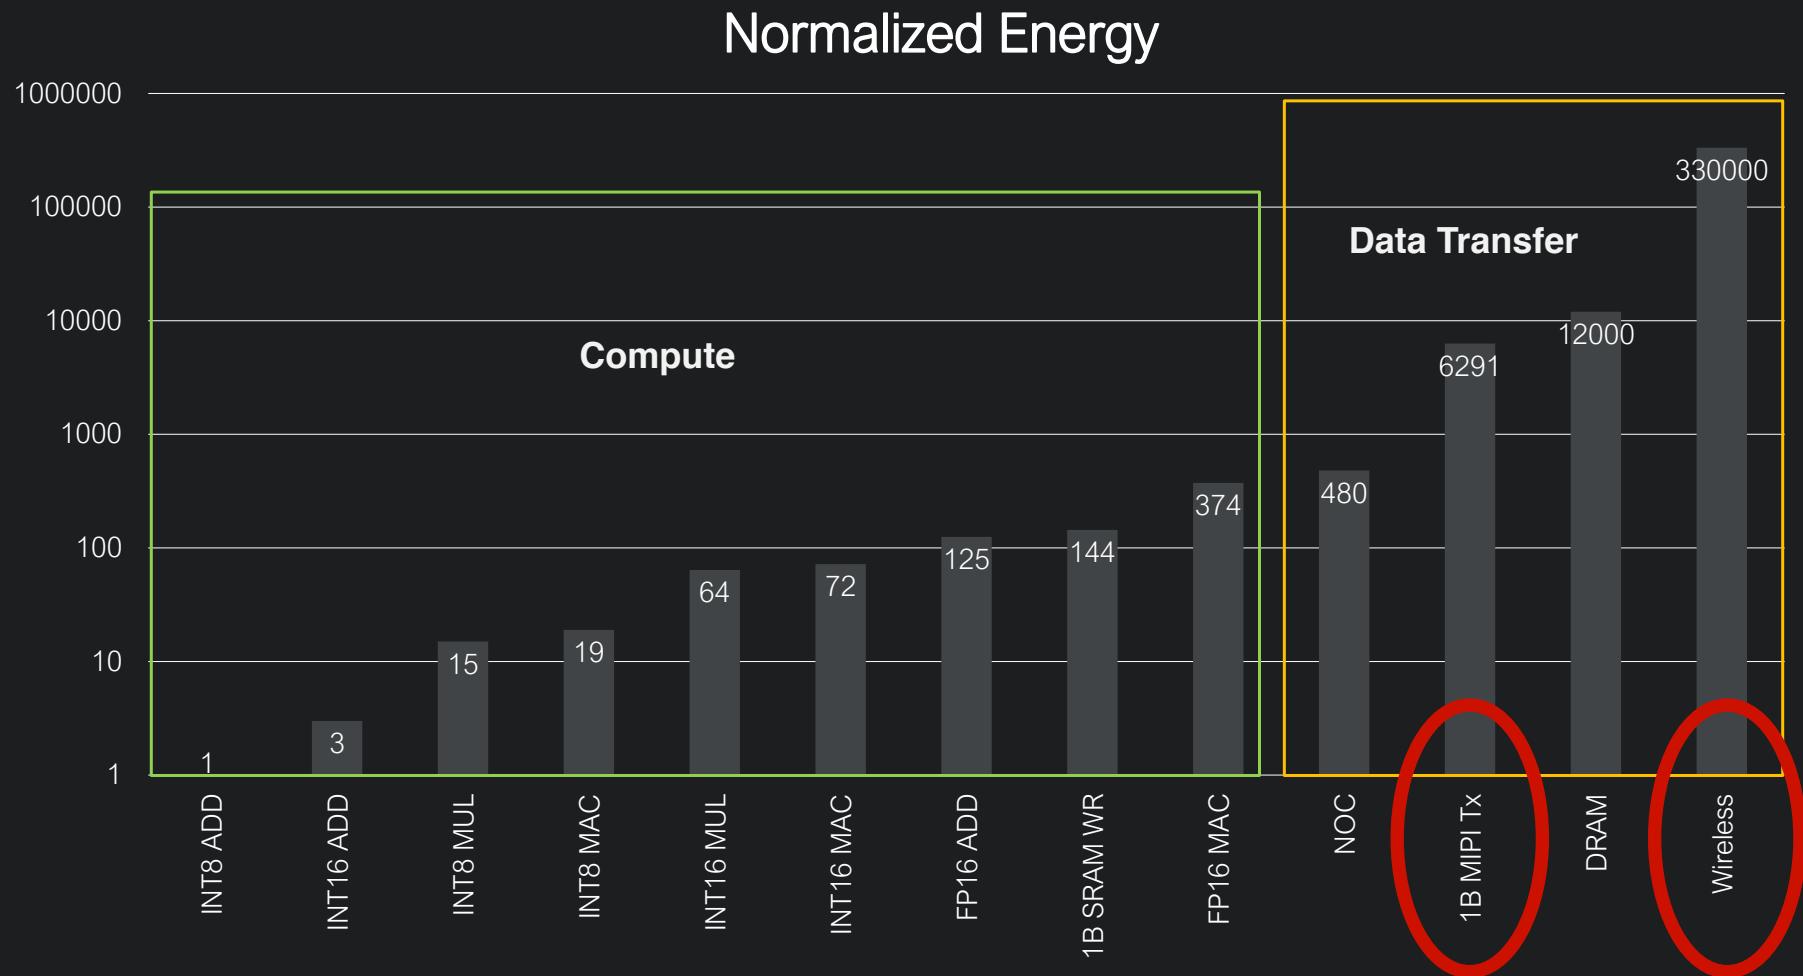
Computation vs. Data Movement

Data movement energy >> computation energy



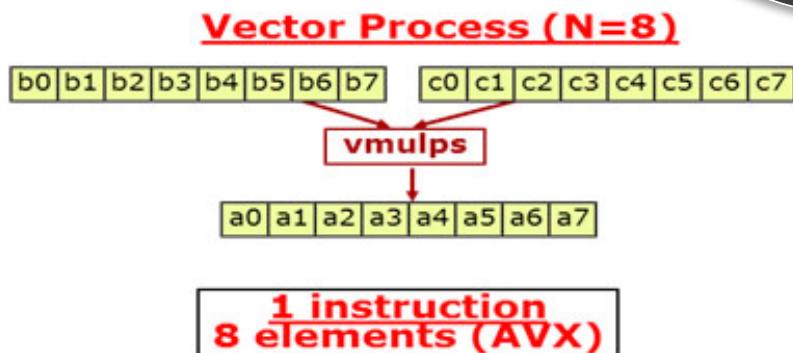
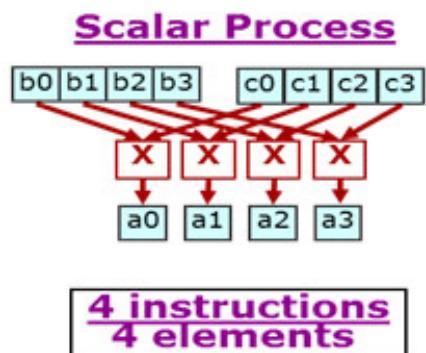
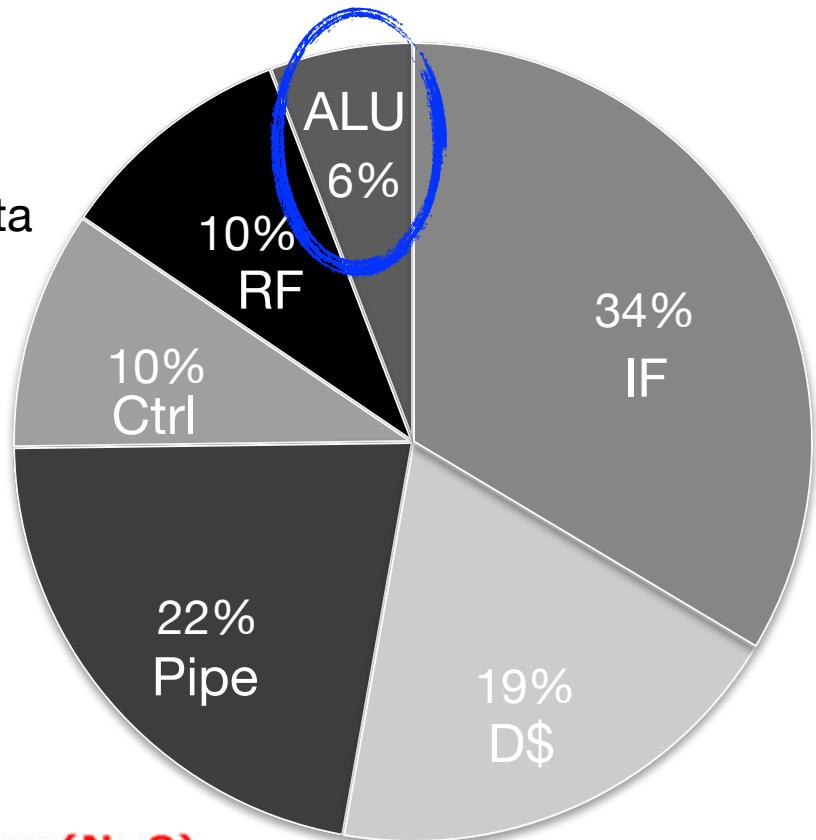
Computation vs. Data Movement

Data movement energy >> computation energy



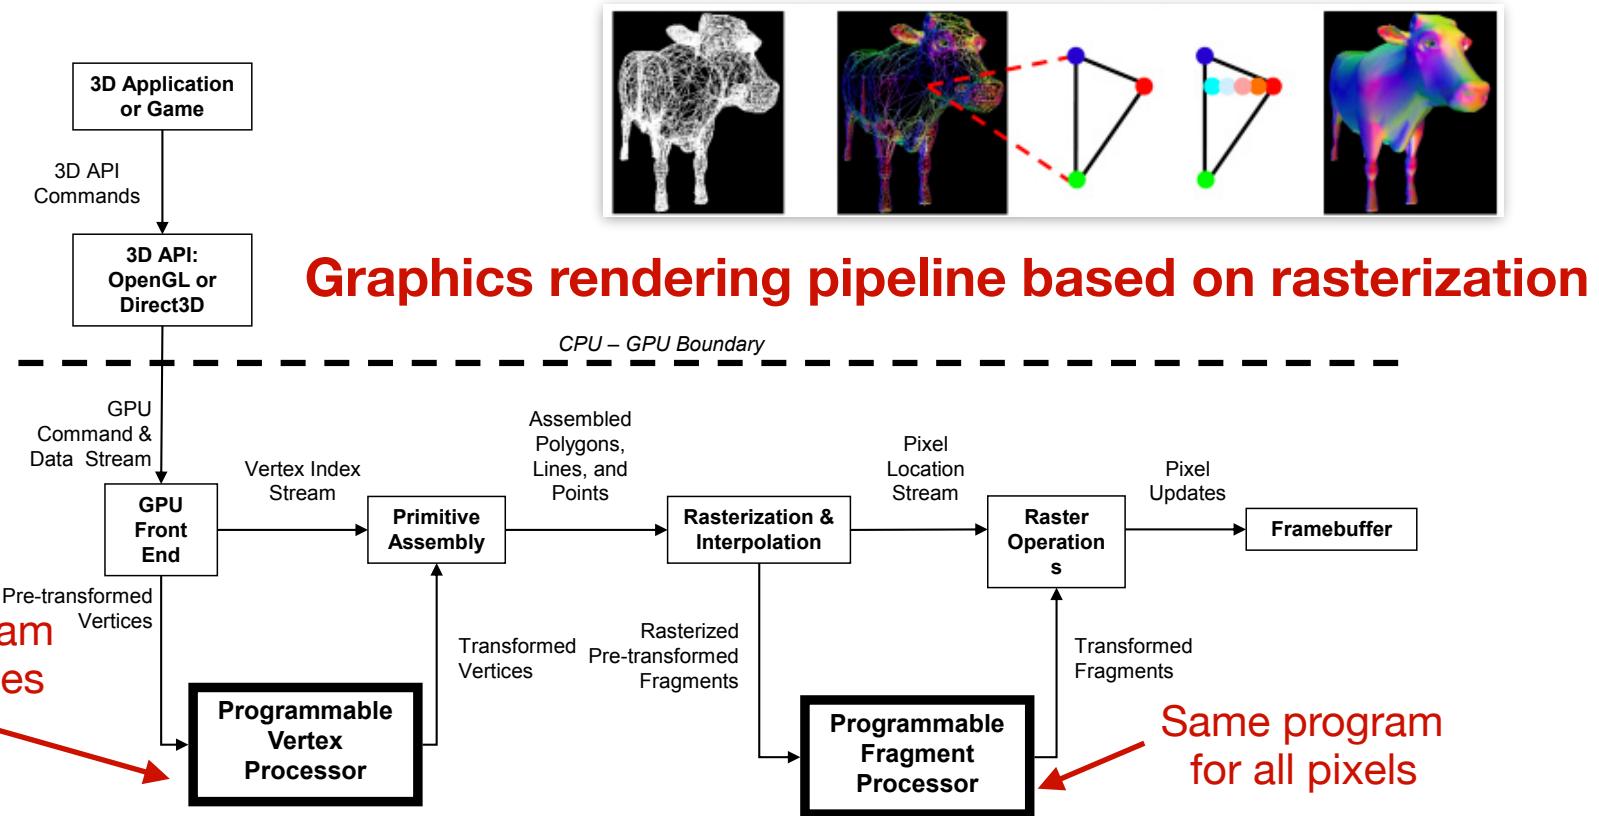
SIMD

- Single Instruction (operating on) Multiple Data
- Amortizing the cost of instruction delivery/ control across many execution units (even cores).
- Almost all modern ISAs provide such instructions:
 - x86: MMX/SSE/AVX
 - Arm: Neon



Graphics Processing Units/GPUs (SIMT)

- Designed for graphics rendering, which is massively parallel.

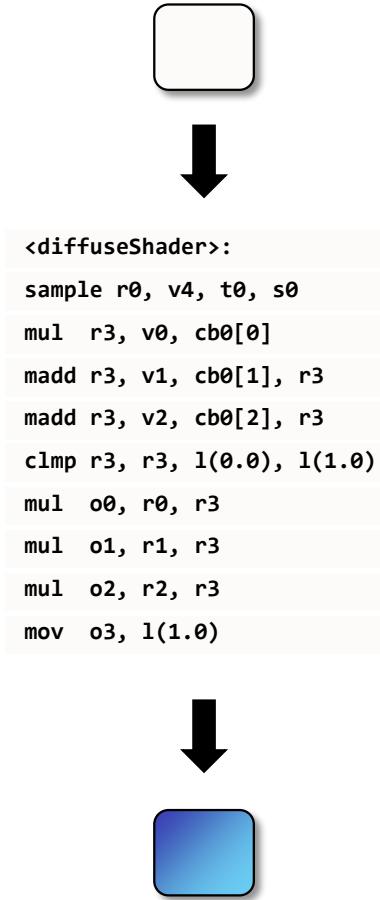
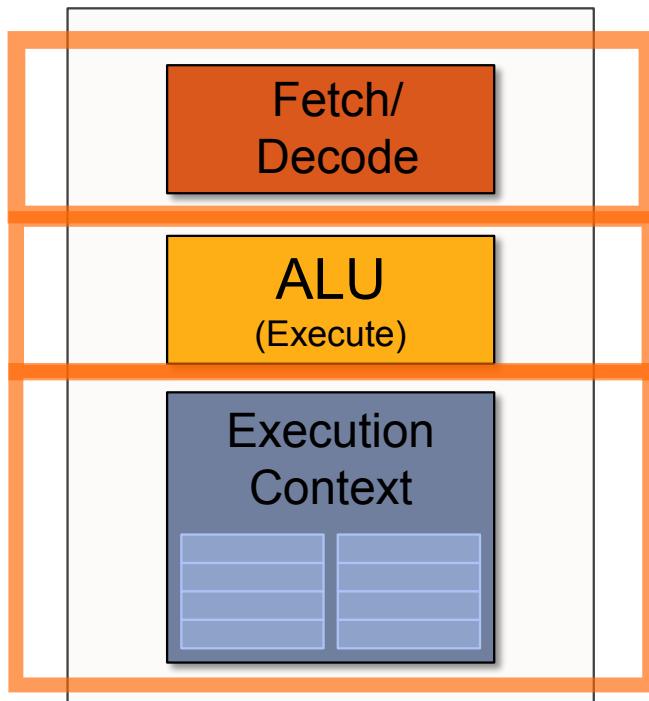


© David Kirk/NVIDIA and
Wen-mei W. Hwu, 2007

ECE 498AL, University of Illinois,
Urbana-Champaign

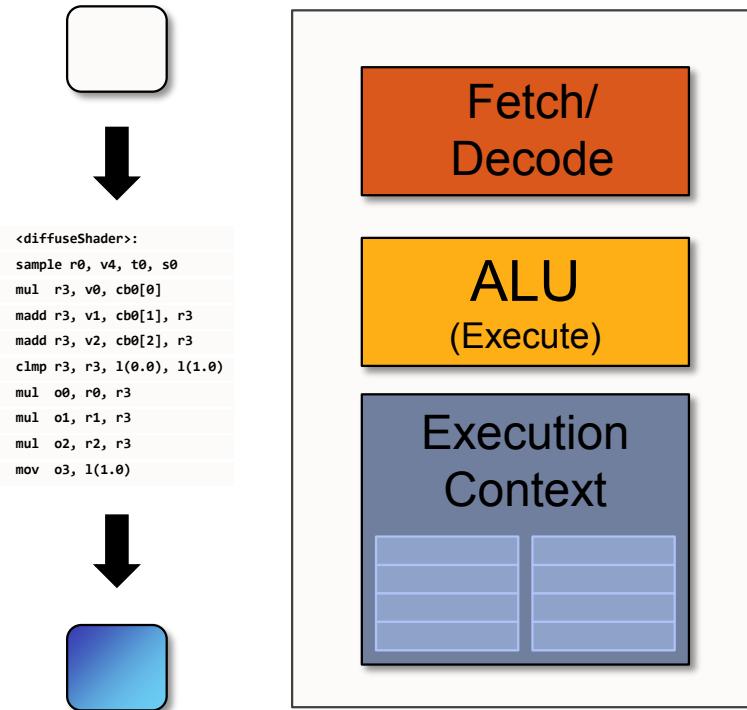
Computer Architecture

Execute shader

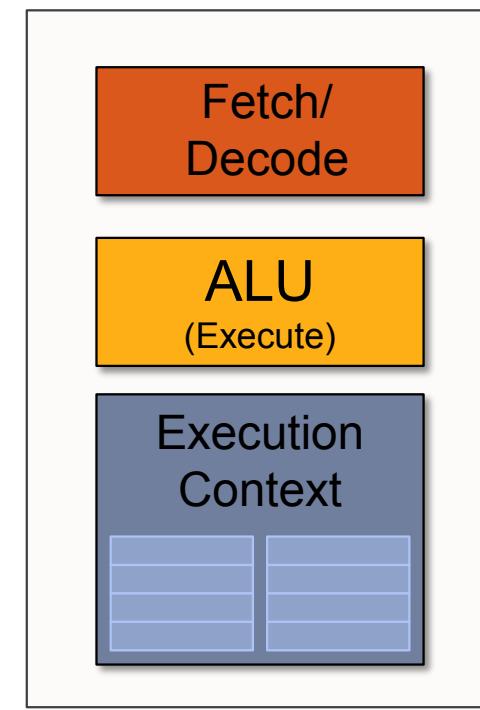


Two cores (two fragments in parallel)

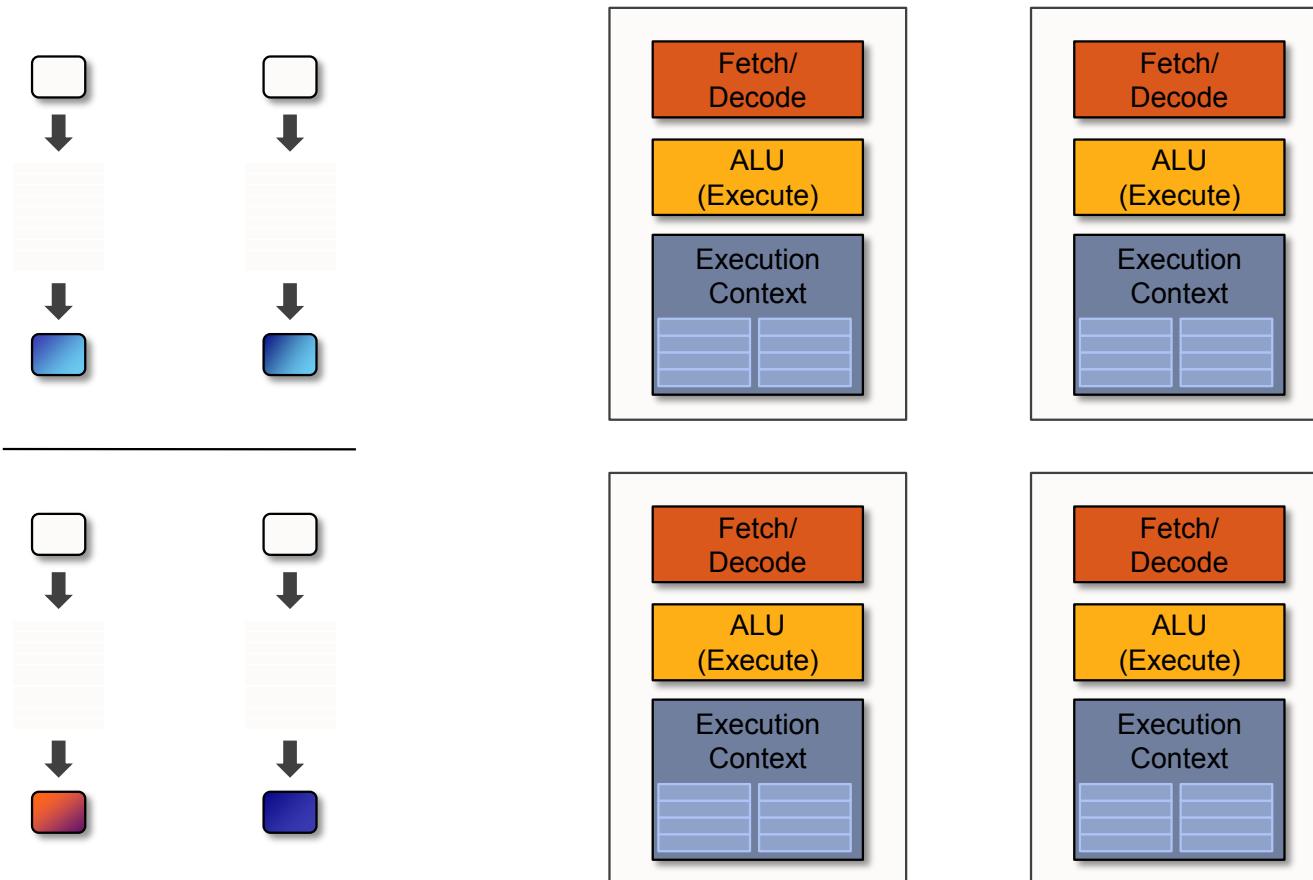
fragment 1



fragment 2



Four cores (four fragments in parallel)



Sixteen cores (sixteen fragments in parallel)

25

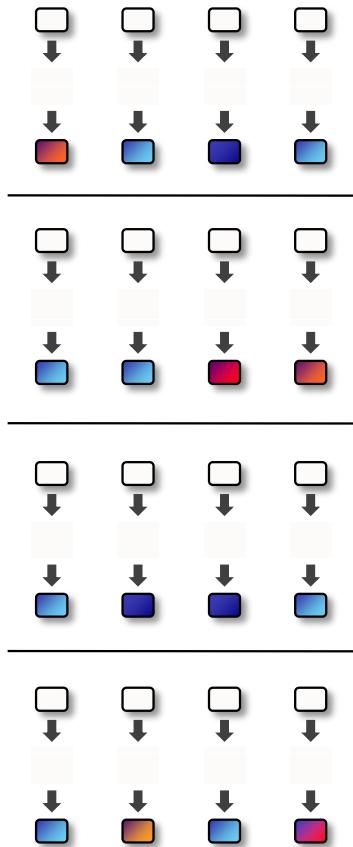


16 cores = 16 simultaneous instruction streams

29



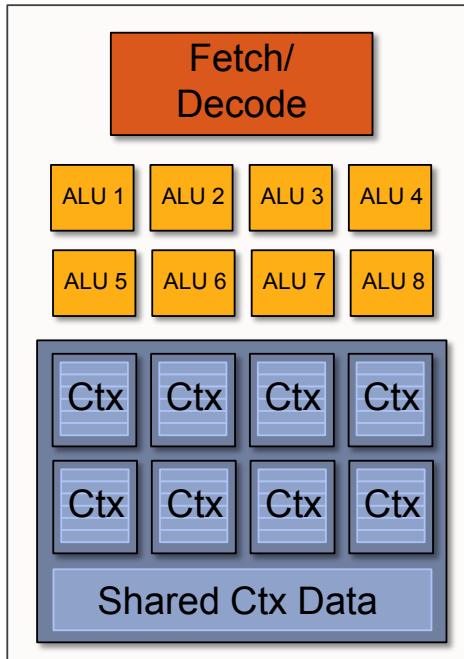
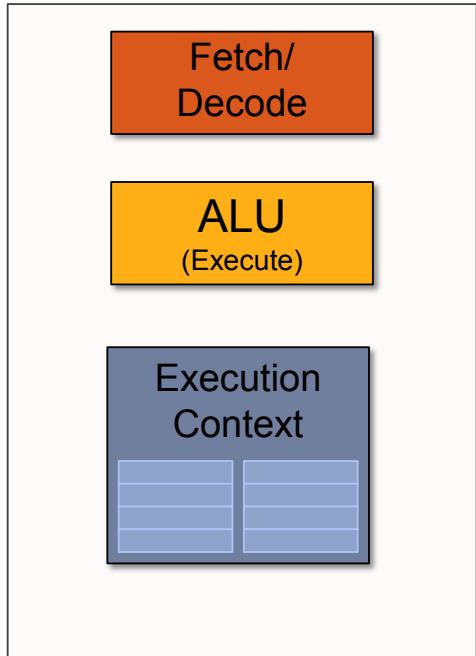
Instruction stream coherence



But... many fragments should be able to share an instruction stream!

```
<diffuseShader>:
sample r0, v4, t0, s0
mul r3, v0, cb0[0]
madd r3, v1, cb0[1], r3
madd r3, v2, cb0[2], r3
clmp r3, r3, 1(0.0), 1(1.0)
mul o0, r0, r3
mul o1, r1, r3
mul o2, r2, r3
mov o3, 1(1.0)
```

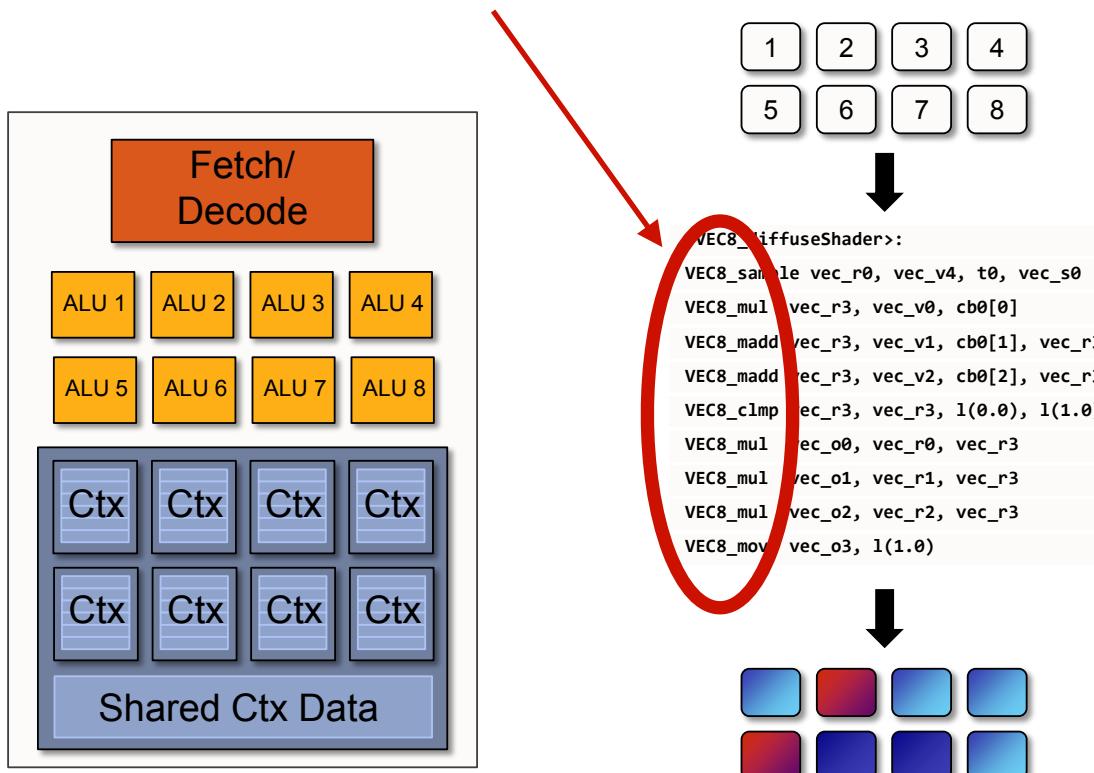




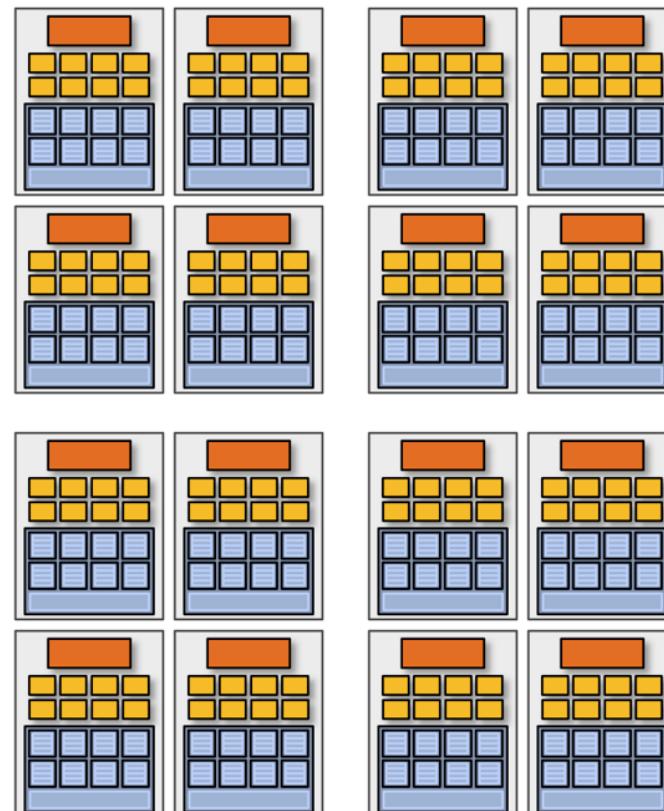
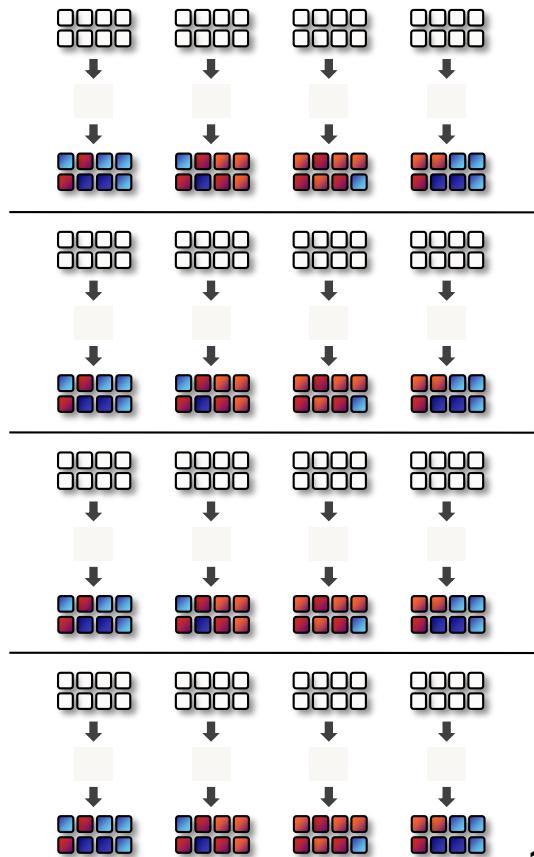
Amortize cost/complexity of managing an instruction stream across many ALUs

SIMD processing

SIMD/vector instructions, each operates on a vector of 8 elements here.



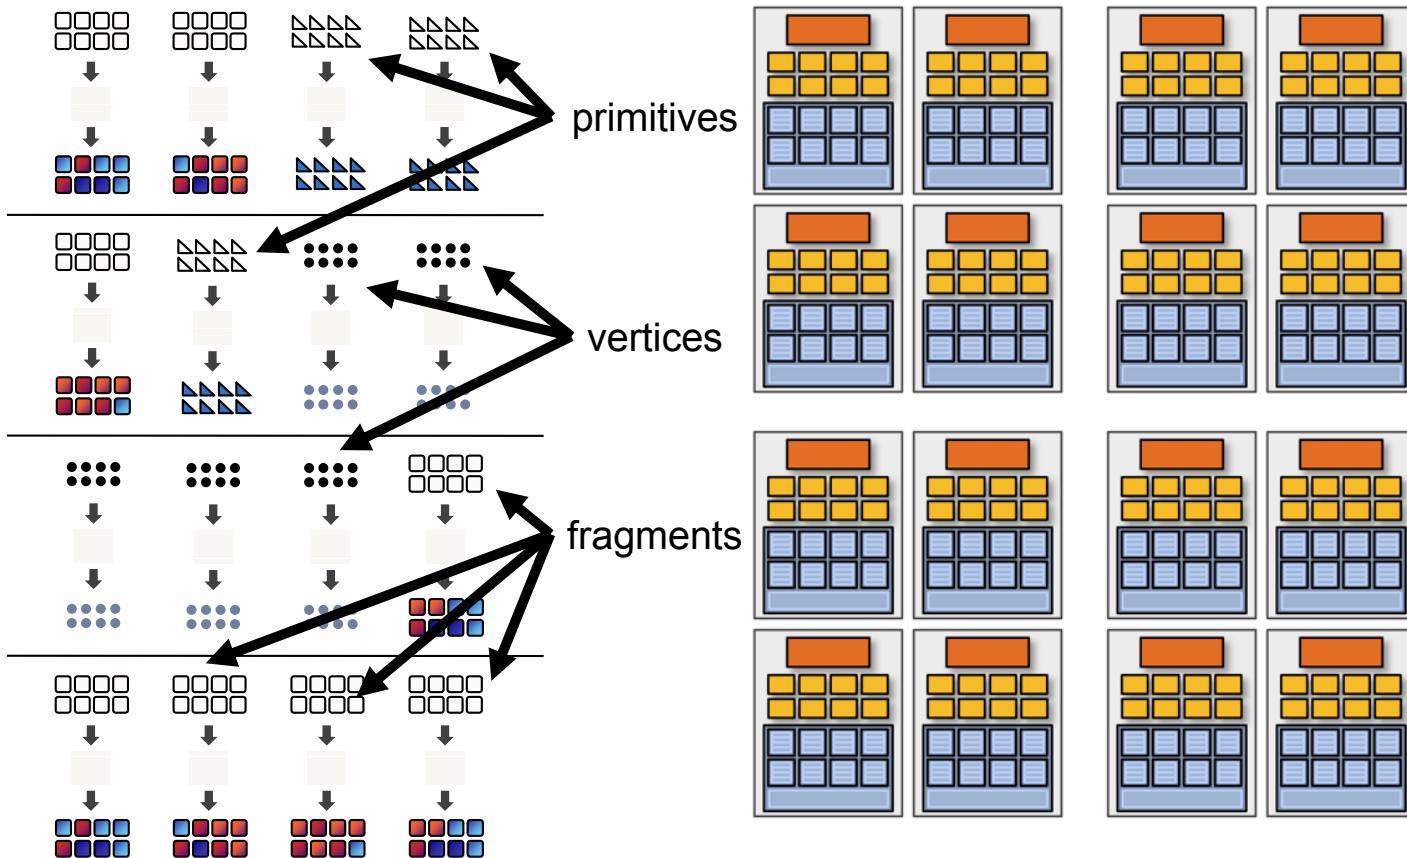
16 cores, each with 8 ALUs. Each core here runs the same program (fragment shader)



16 cores = 128 ALUs
= 16 simultaneous instruction streams

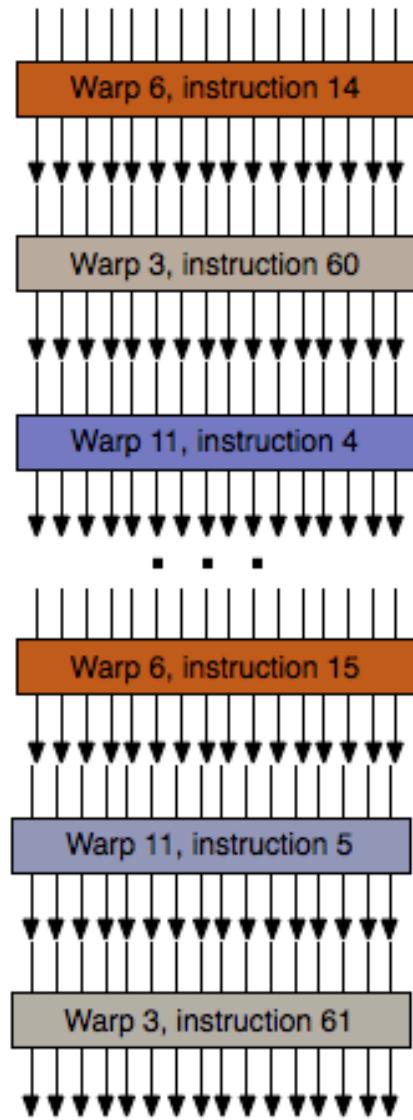


16 cores, each with 8 ALUs. Cores here run different programs
(some are processing vertices, some are processing fragments)



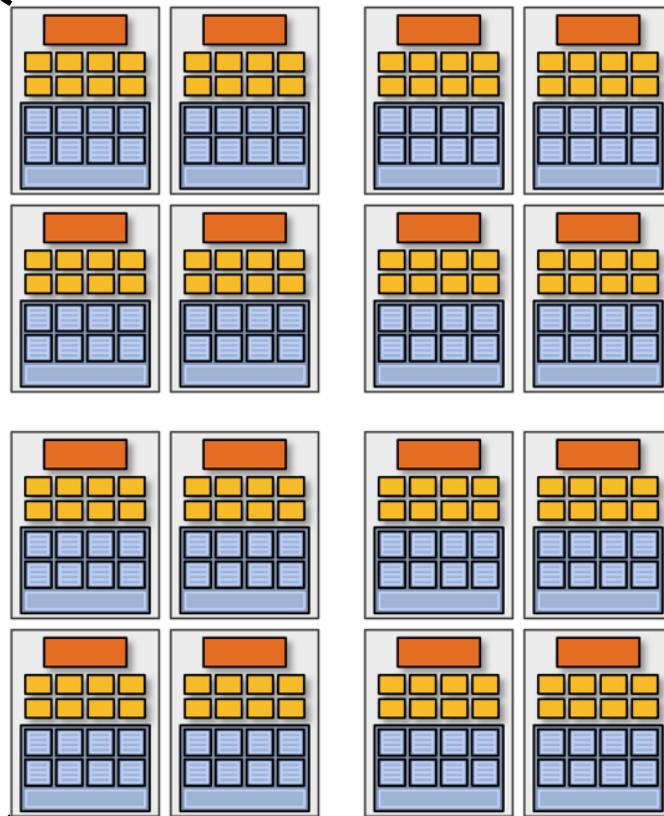
Each Core Does Fine-Grained Multi-threading

Wrap: a group of threads (8 here)



Time

No need for branch prediction and out-of-order execution. Simple core design.



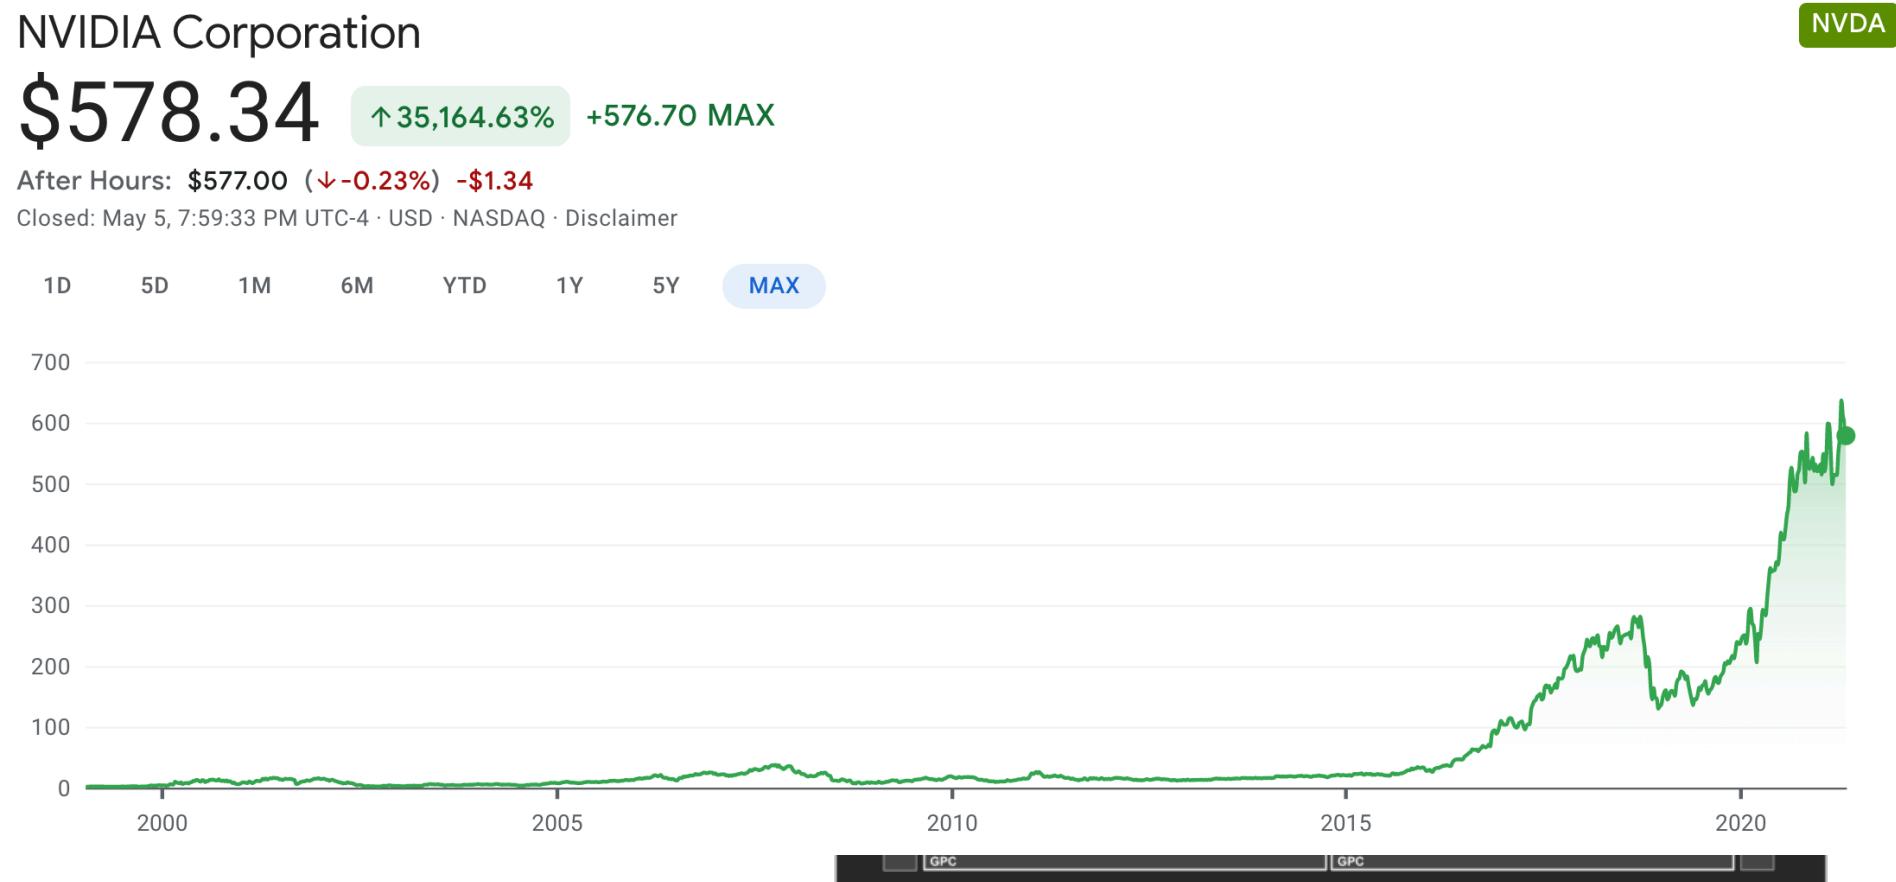
Nvidia Maxwell GPU (2014)

- Today: General Purpose GPU (GPGPU), used for any massive parallel applications:
 - Physics simulation
 - Deep learning
 - Computer vision



Nvidia Maxwell GPU (2014)

- Today: General Purpose GPU (GPGPU), used for any massive parallel applications:



Entering the Era of Specialization

Entering the Era of Specialization

- GPUs are very efficient for massively parallel program

Entering the Era of Specialization

- GPUs are very efficient for massively parallel program
- But are still fairly general, so there are still many inefficiencies
 - Still need to fetch and decode instructions
 - Still have (very large) caches, so data delivery isn't efficient

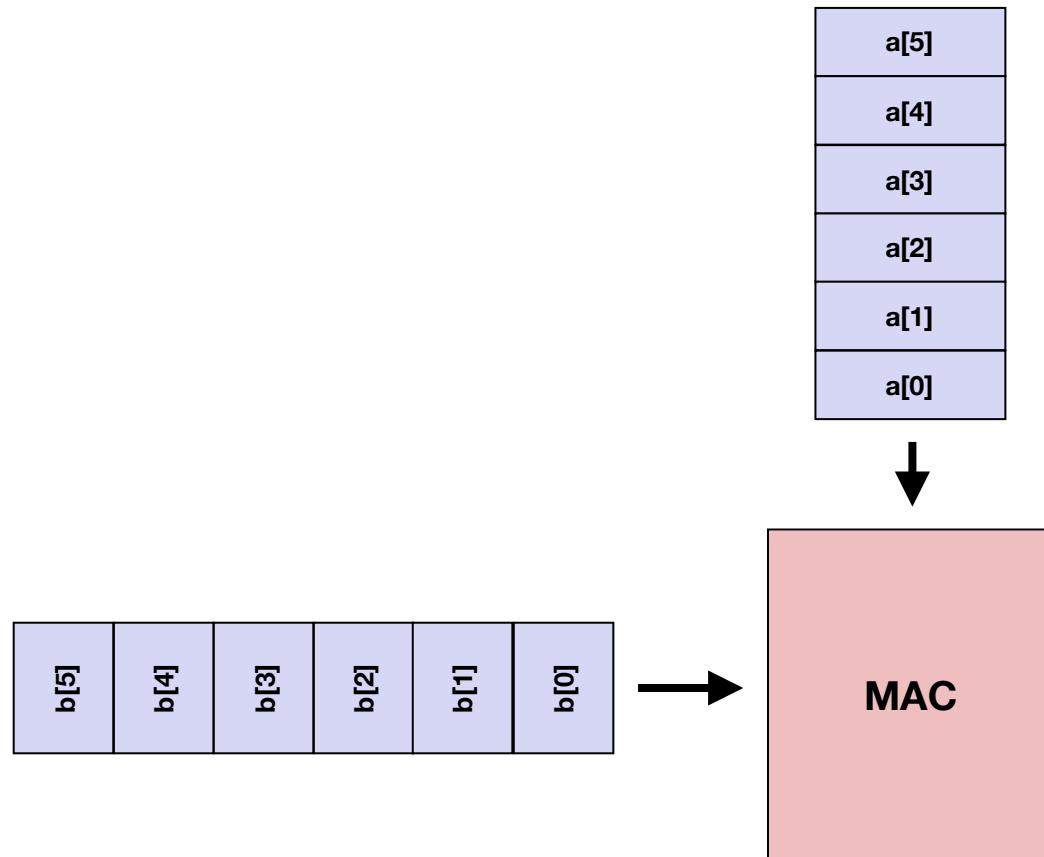
Entering the Era of Specialization

- GPUs are very efficient for massively parallel program
- But are still fairly general, so there are still many inefficiencies
 - Still need to fetch and decode instructions
 - Still have (very large) caches, so data delivery isn't efficient
- Idea: instead of building general-purpose processors that can do everything, but inefficiently, let's build specialized processors that can only do limited things, but extremely efficiently.

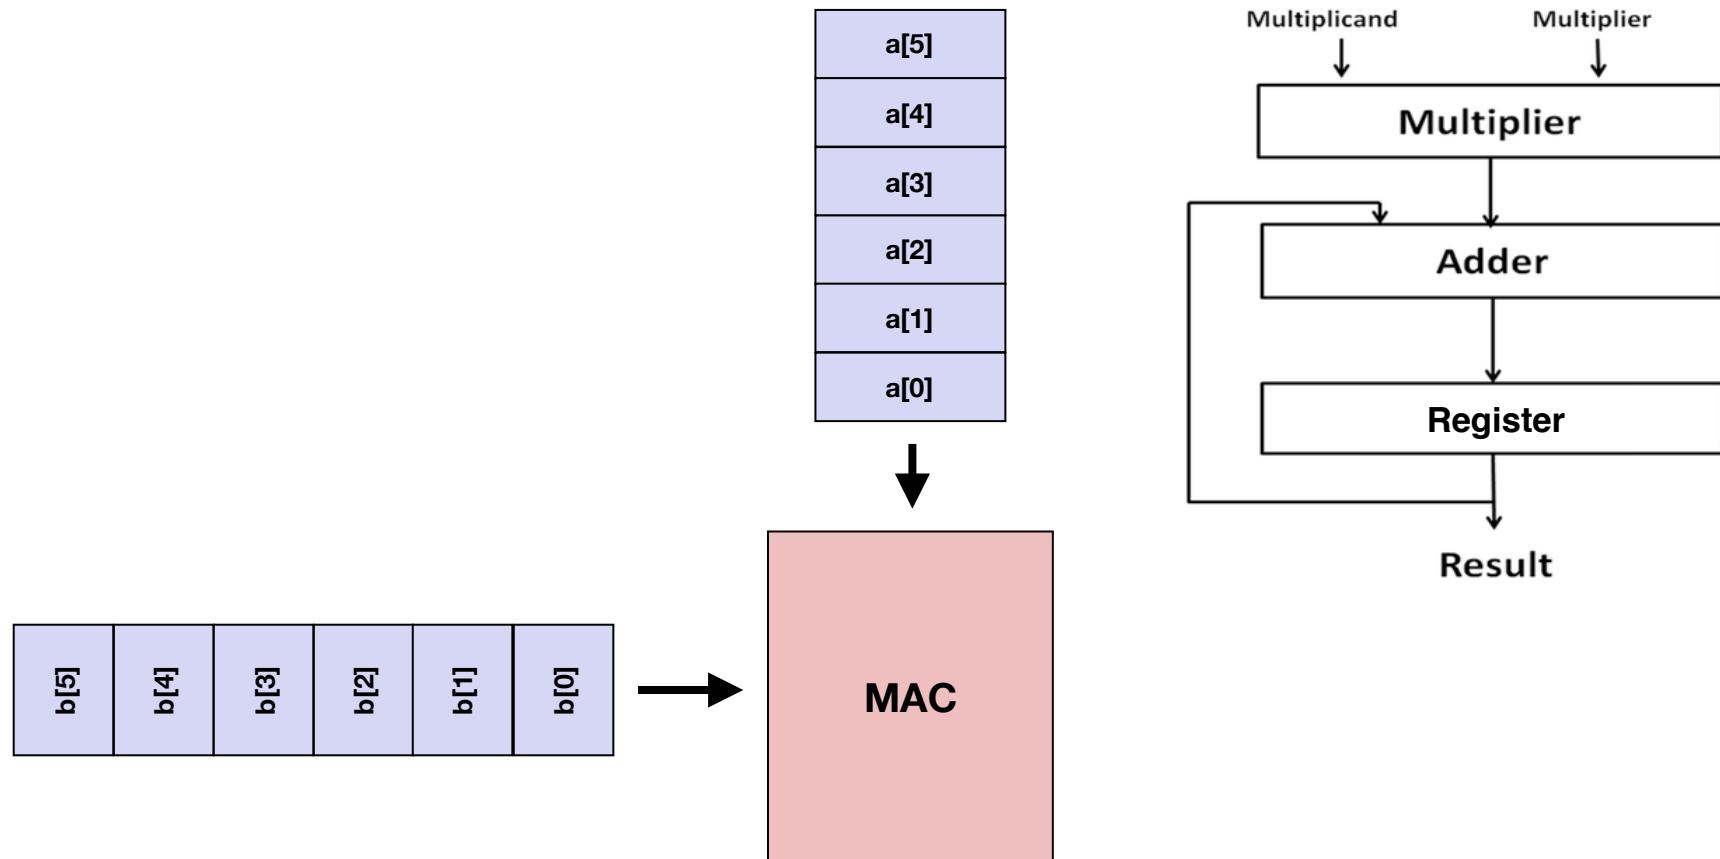
Entering the Era of Specialization

- GPUs are very efficient for massively parallel program
- But are still fairly general, so there are still many inefficiencies
 - Still need to fetch and decode instructions
 - Still have (very large) caches, so data delivery isn't efficient
- Idea: instead of building general-purpose processors that can do everything, but inefficiently, let's build specialized processors that can only do limited things, but extremely efficiently.
- A.k.a., domain-specific accelerators

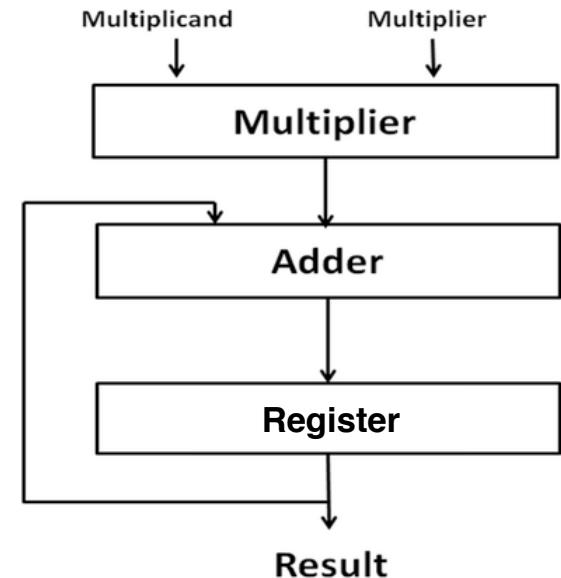
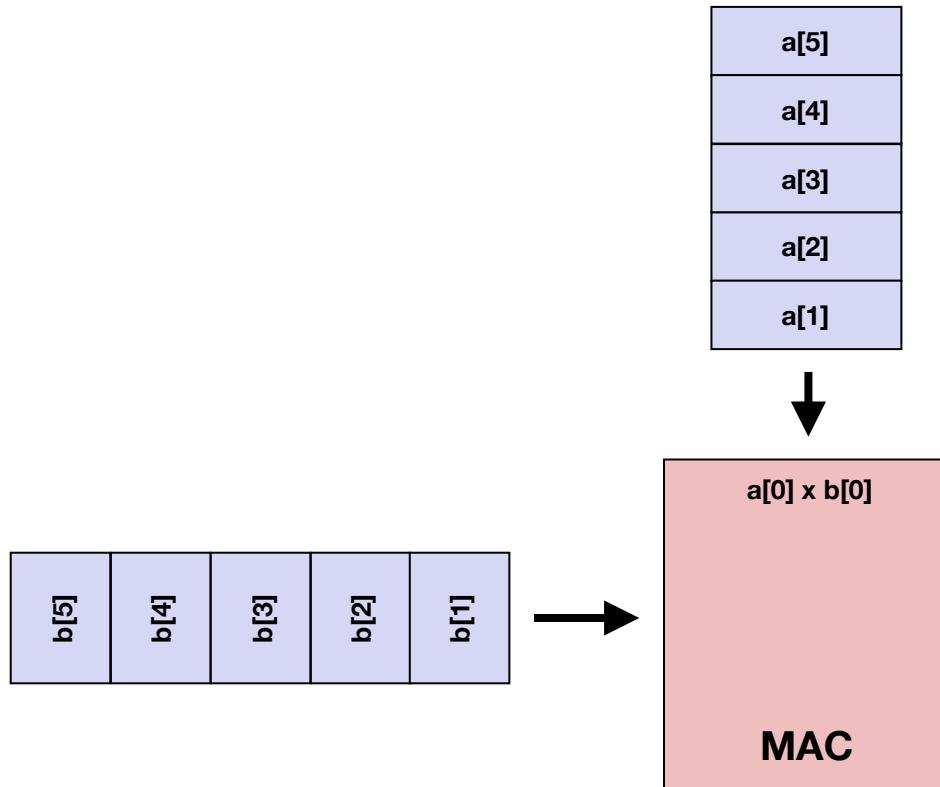
Example: Vector Dot Product



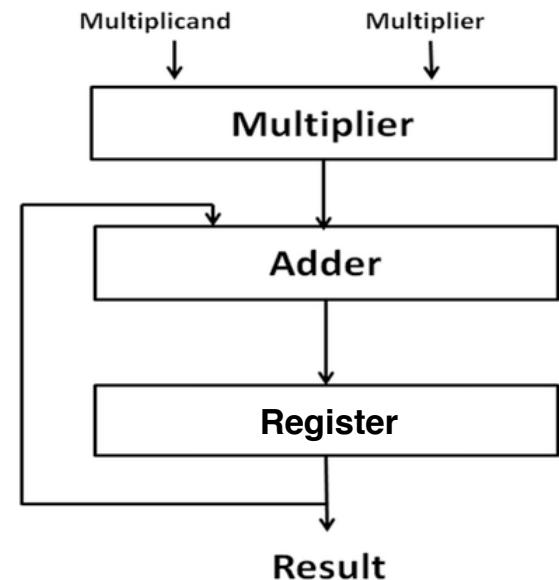
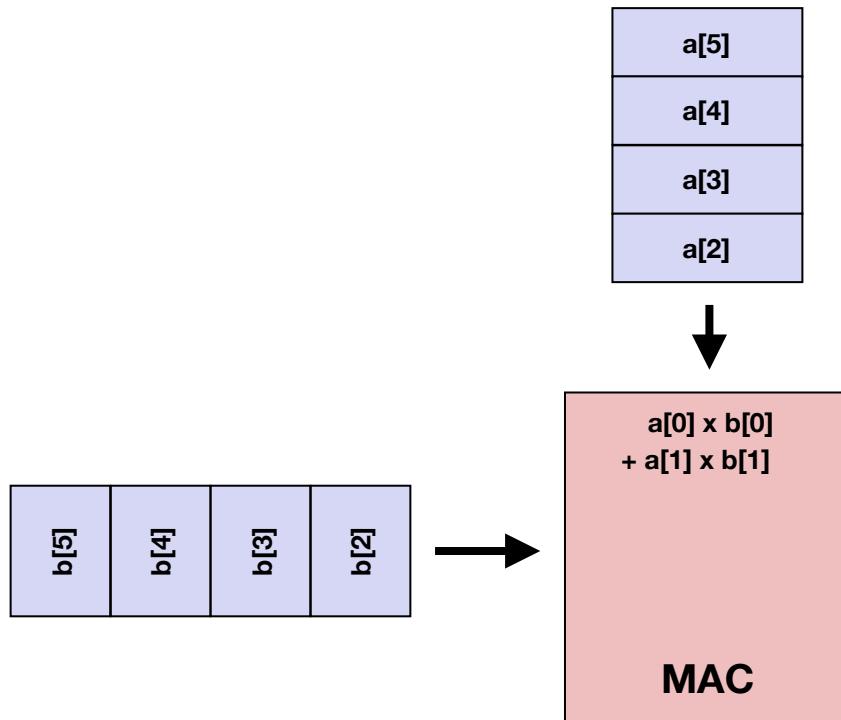
Example: Vector Dot Product



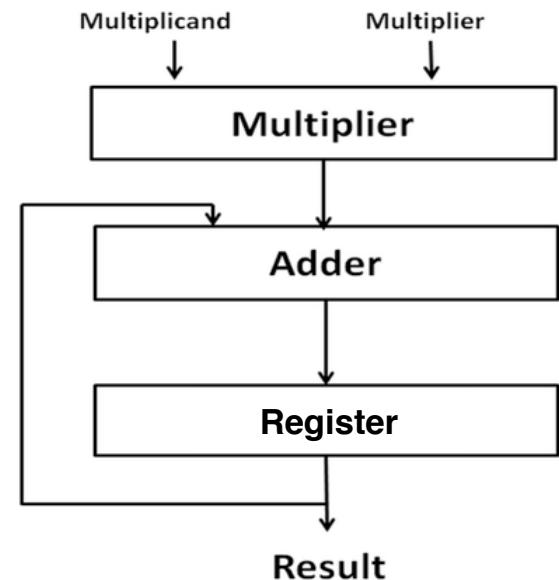
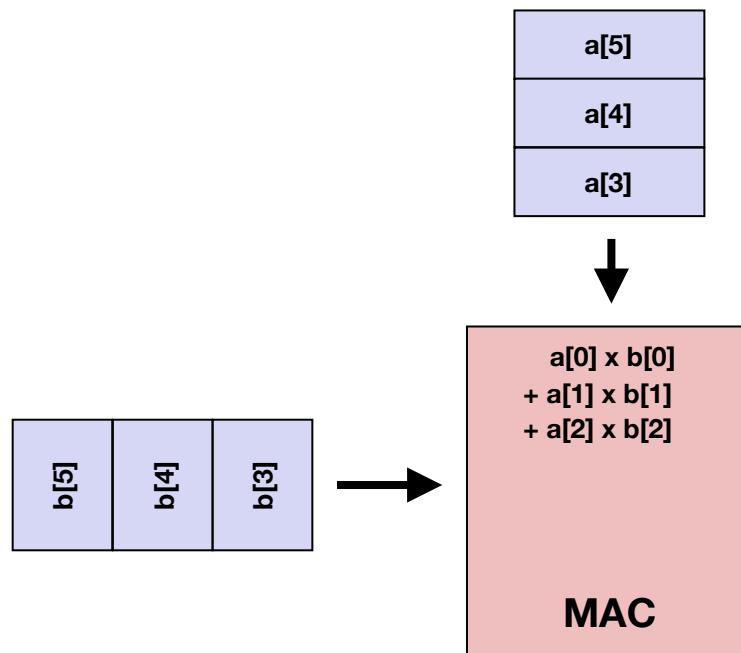
Example: Vector Dot Product



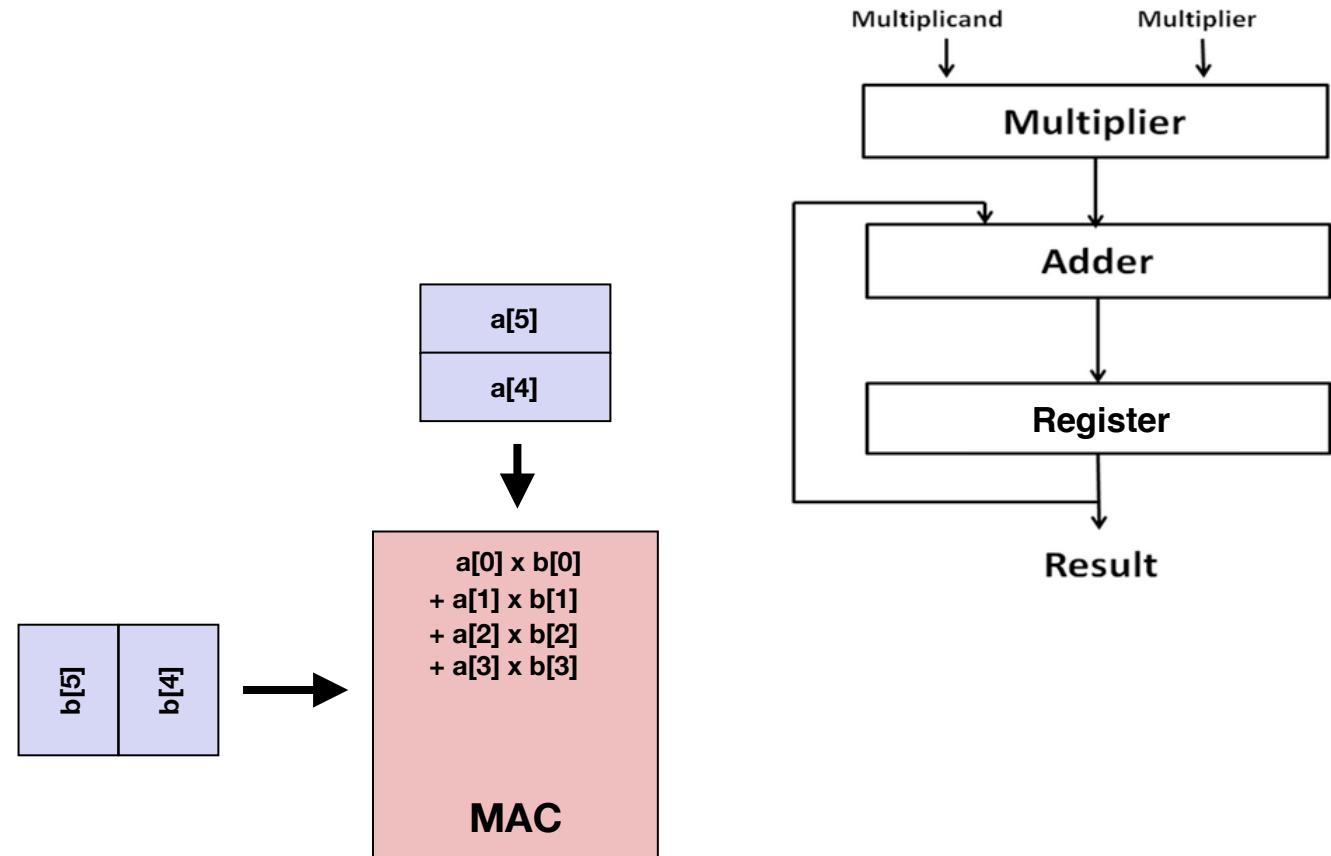
Example: Vector Dot Product



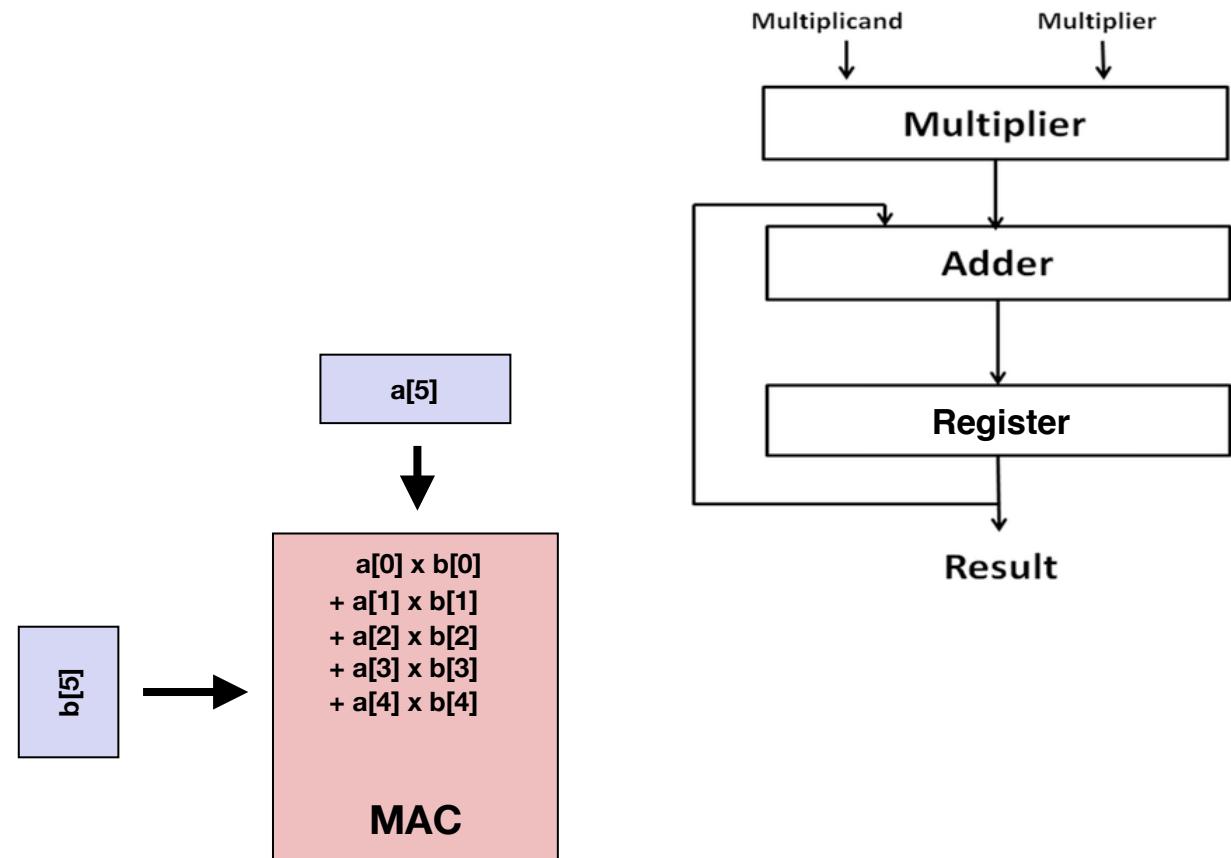
Example: Vector Dot Product



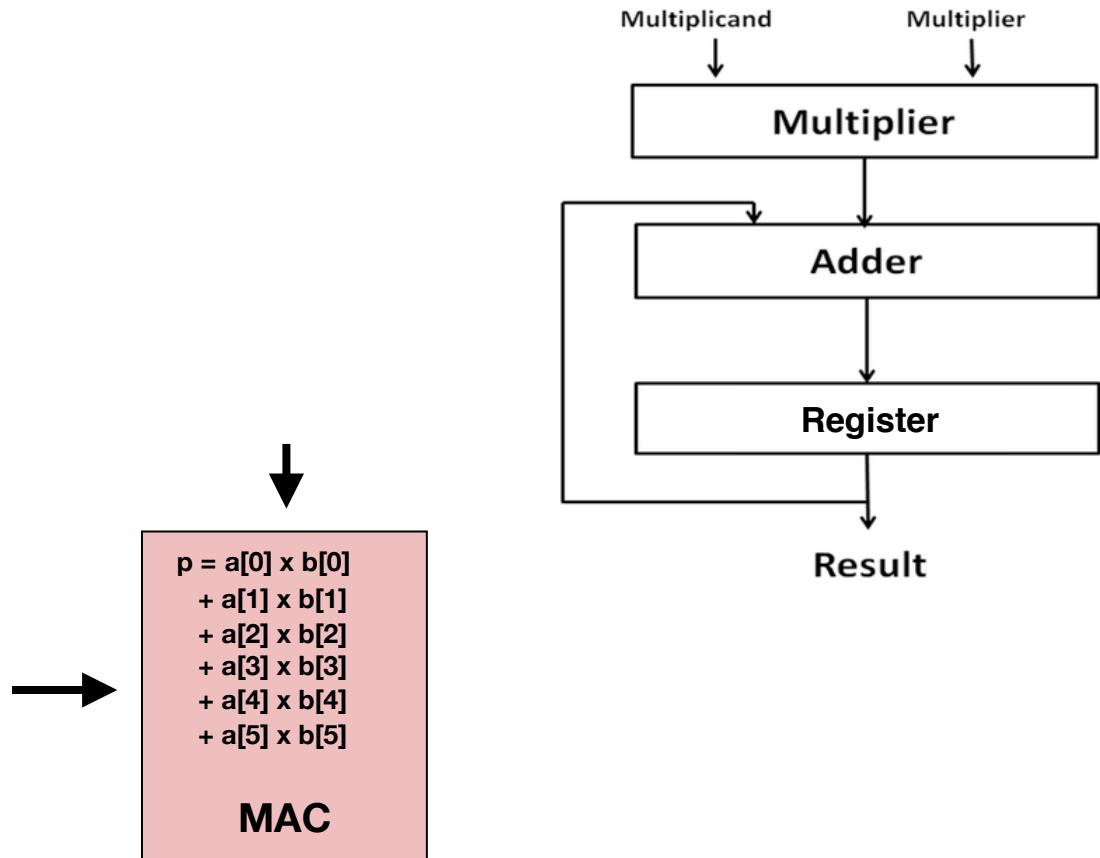
Example: Vector Dot Product



Example: Vector Dot Product

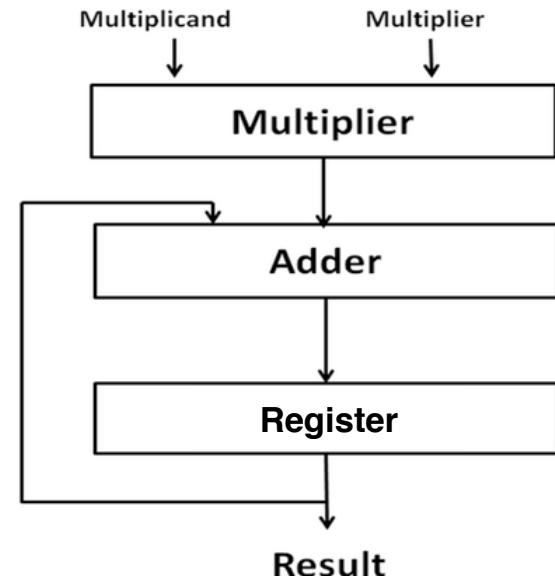
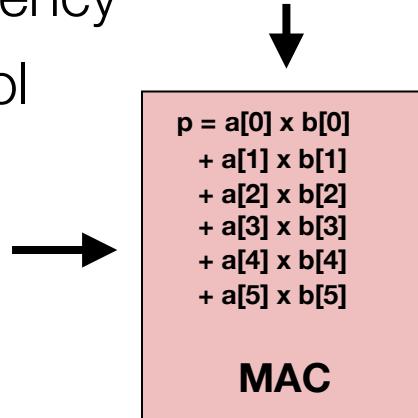


Example: Vector Dot Product

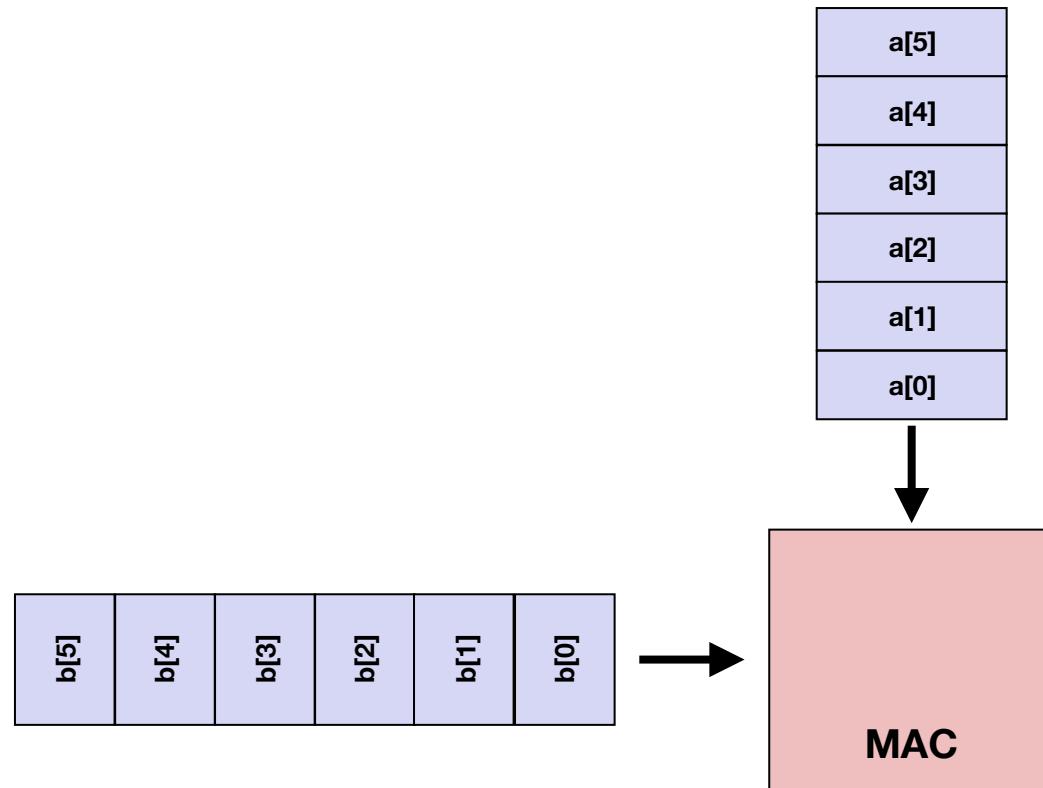


Example: Vector Dot Product

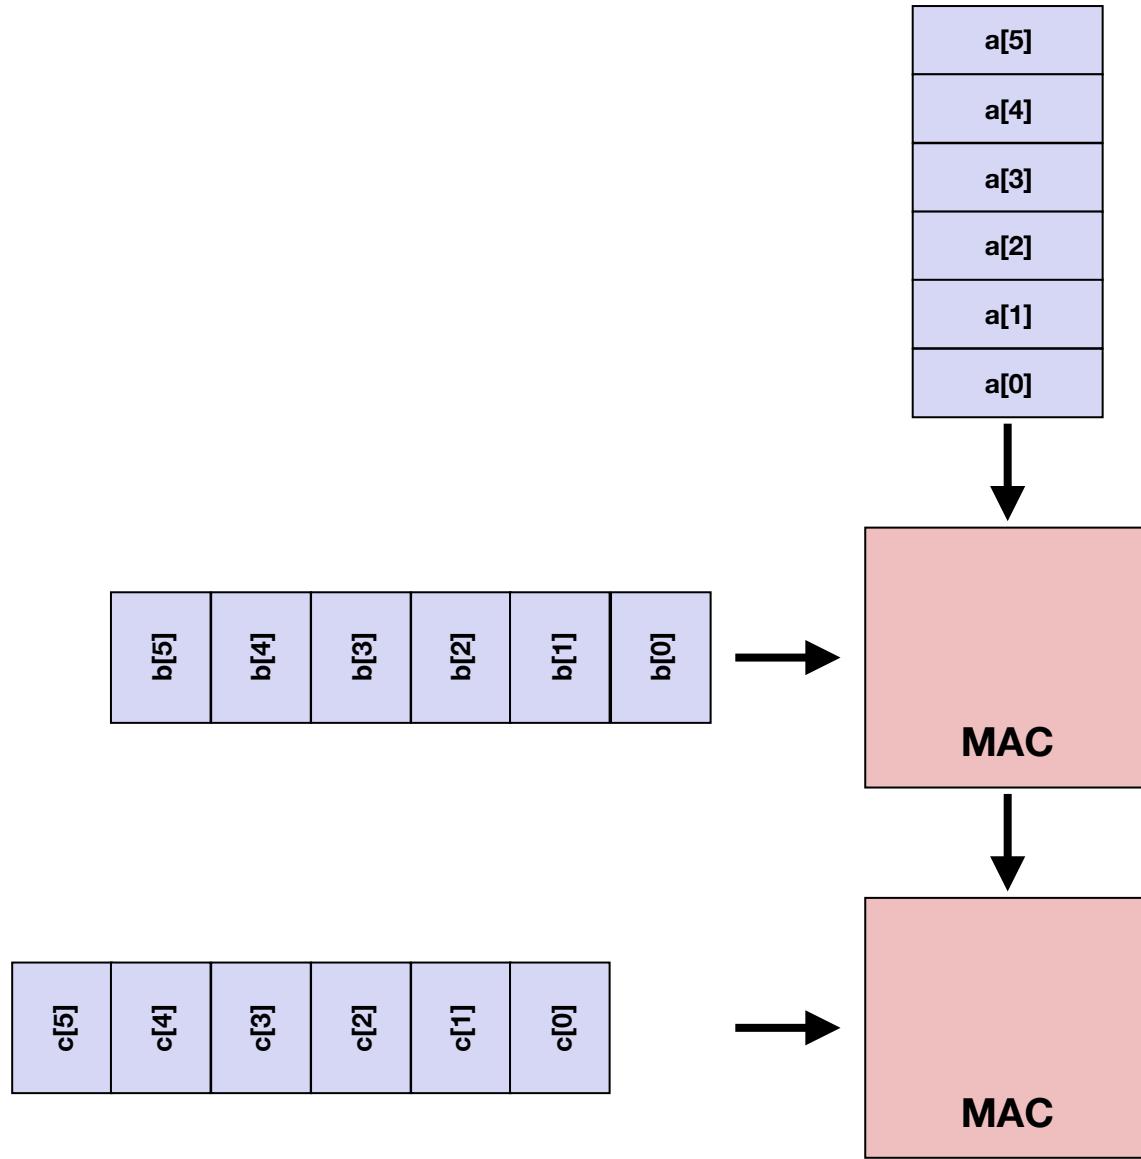
- Does nothing but vector dot product
 - No instruction fetch and decode
(there is no instruction)
 - The register is close to the ALU and gets reused over and over: good data delivery efficiency
 - Very simple control



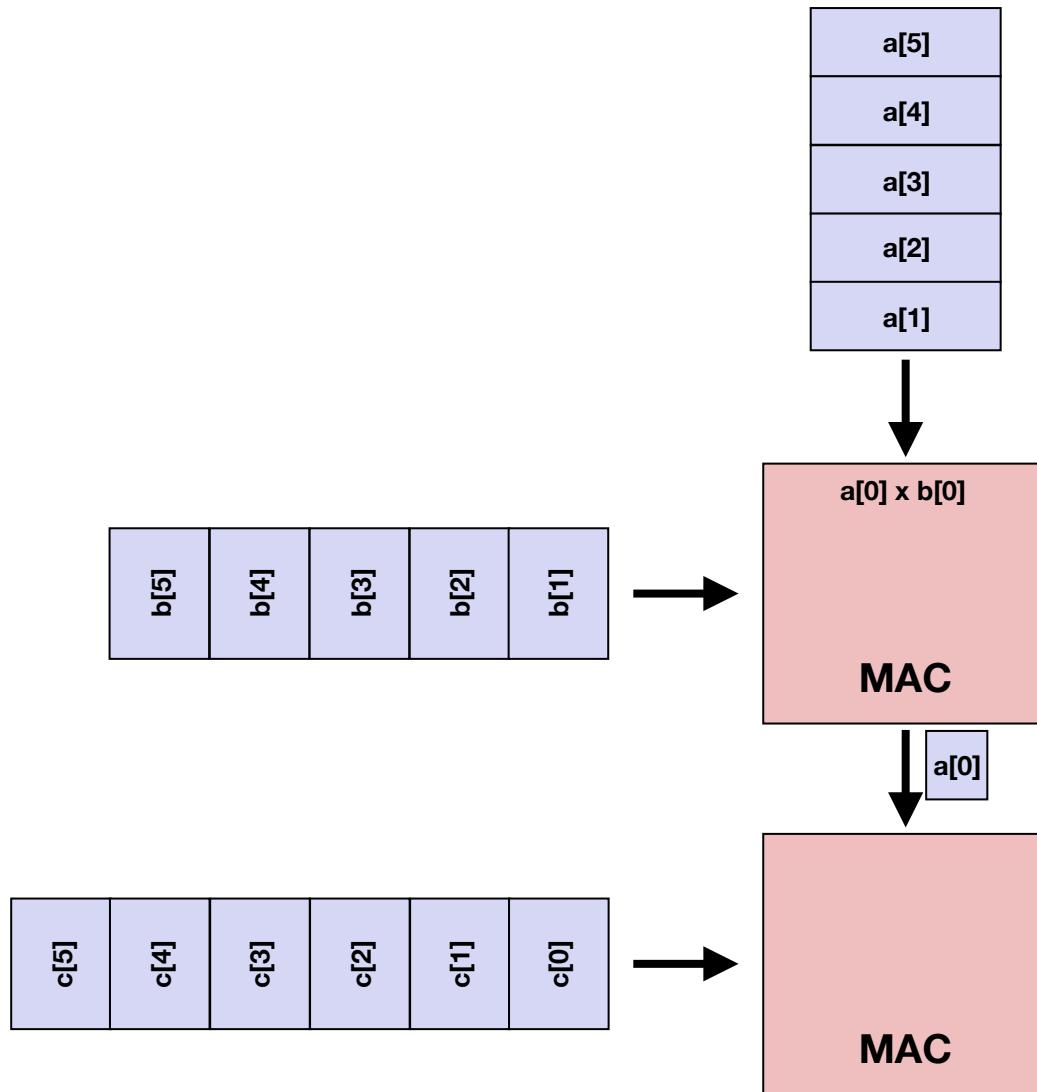
Matrix Vector Multiplication



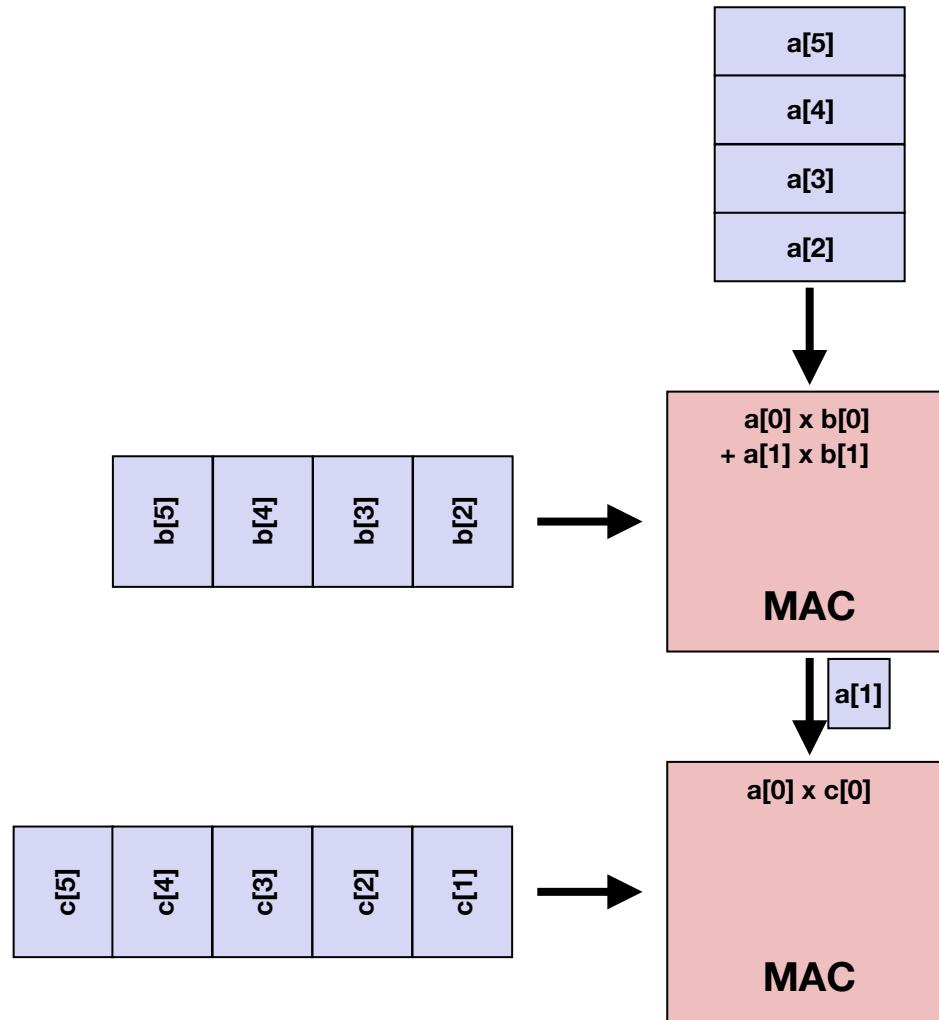
Matrix Vector Multiplication



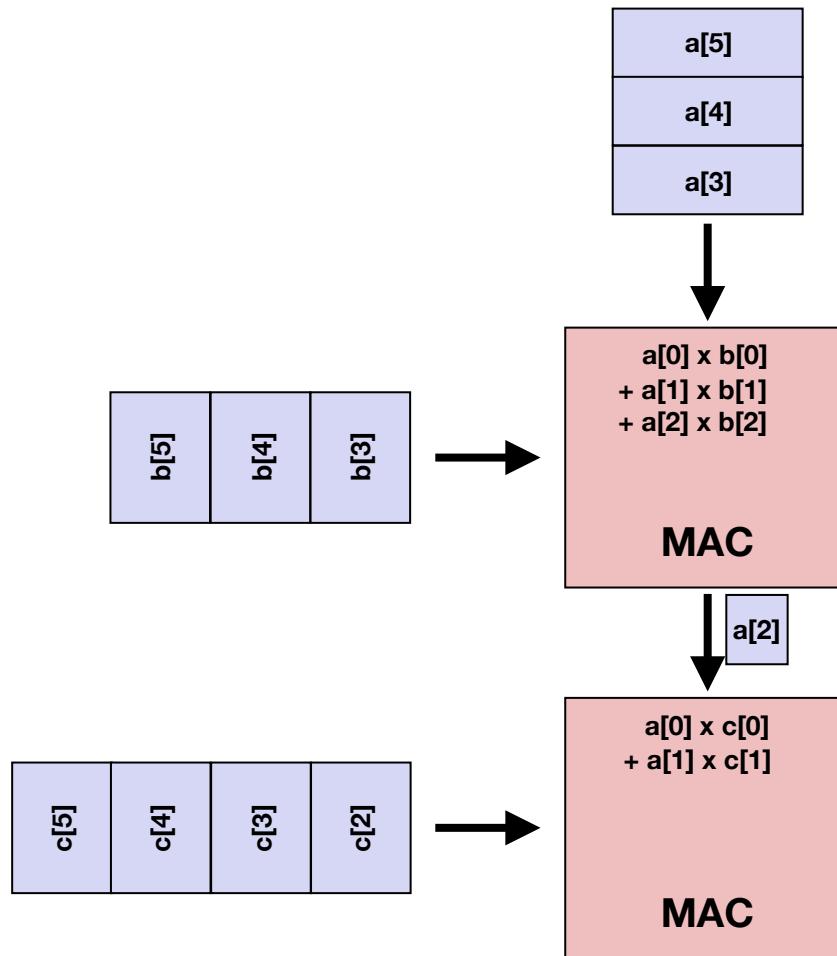
Matrix Vector Multiplication



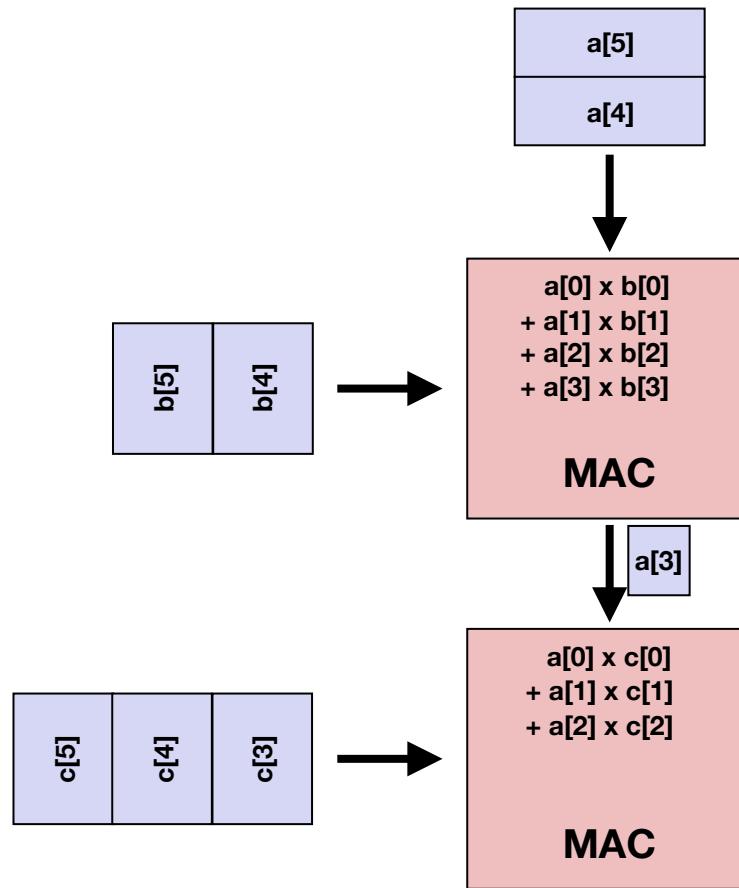
Matrix Vector Multiplication



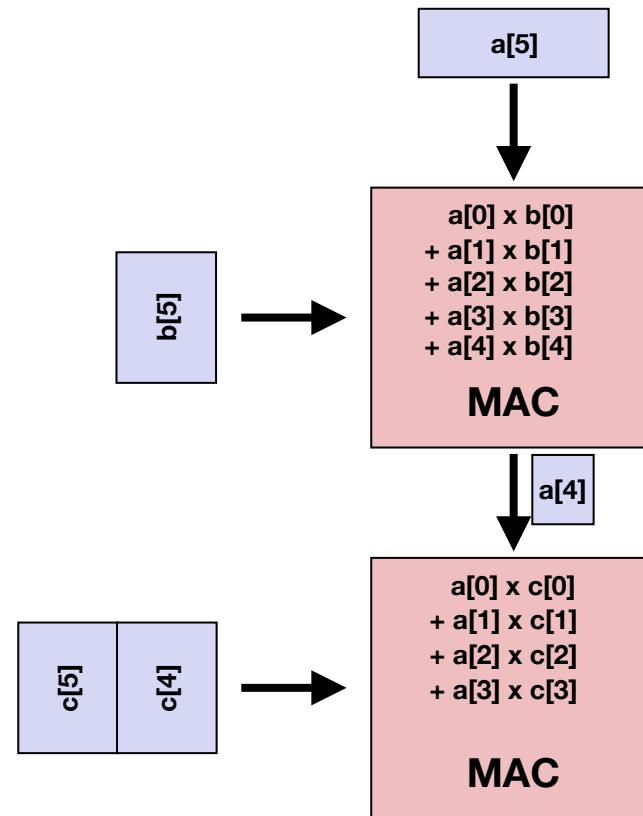
Matrix Vector Multiplication



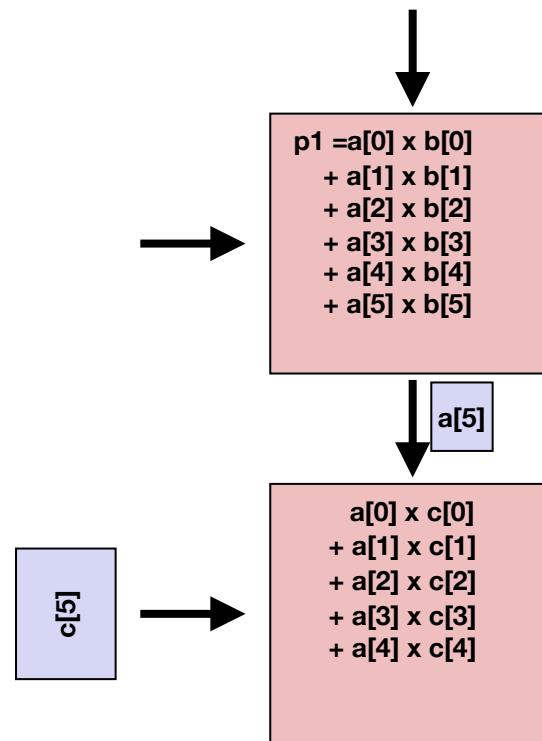
Matrix Vector Multiplication



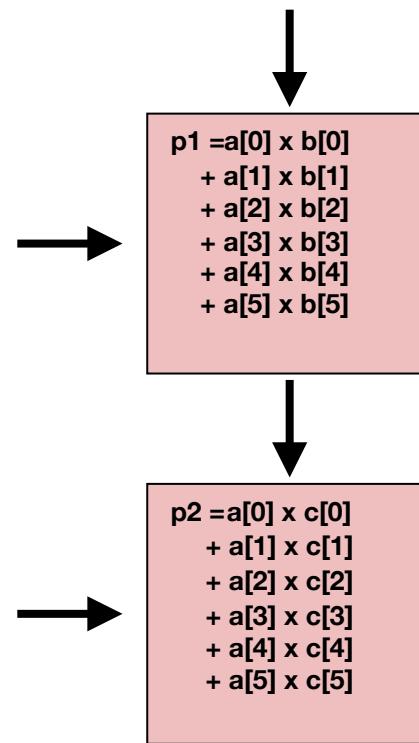
Matrix Vector Multiplication



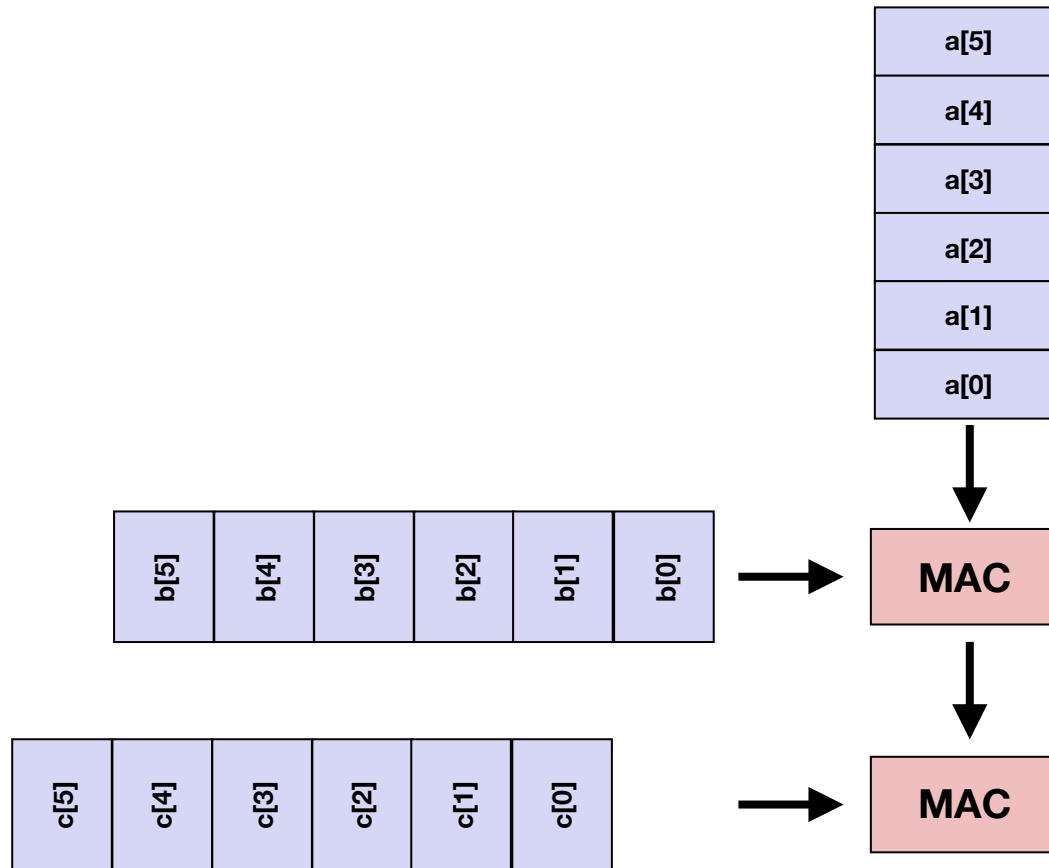
Matrix Vector Multiplication



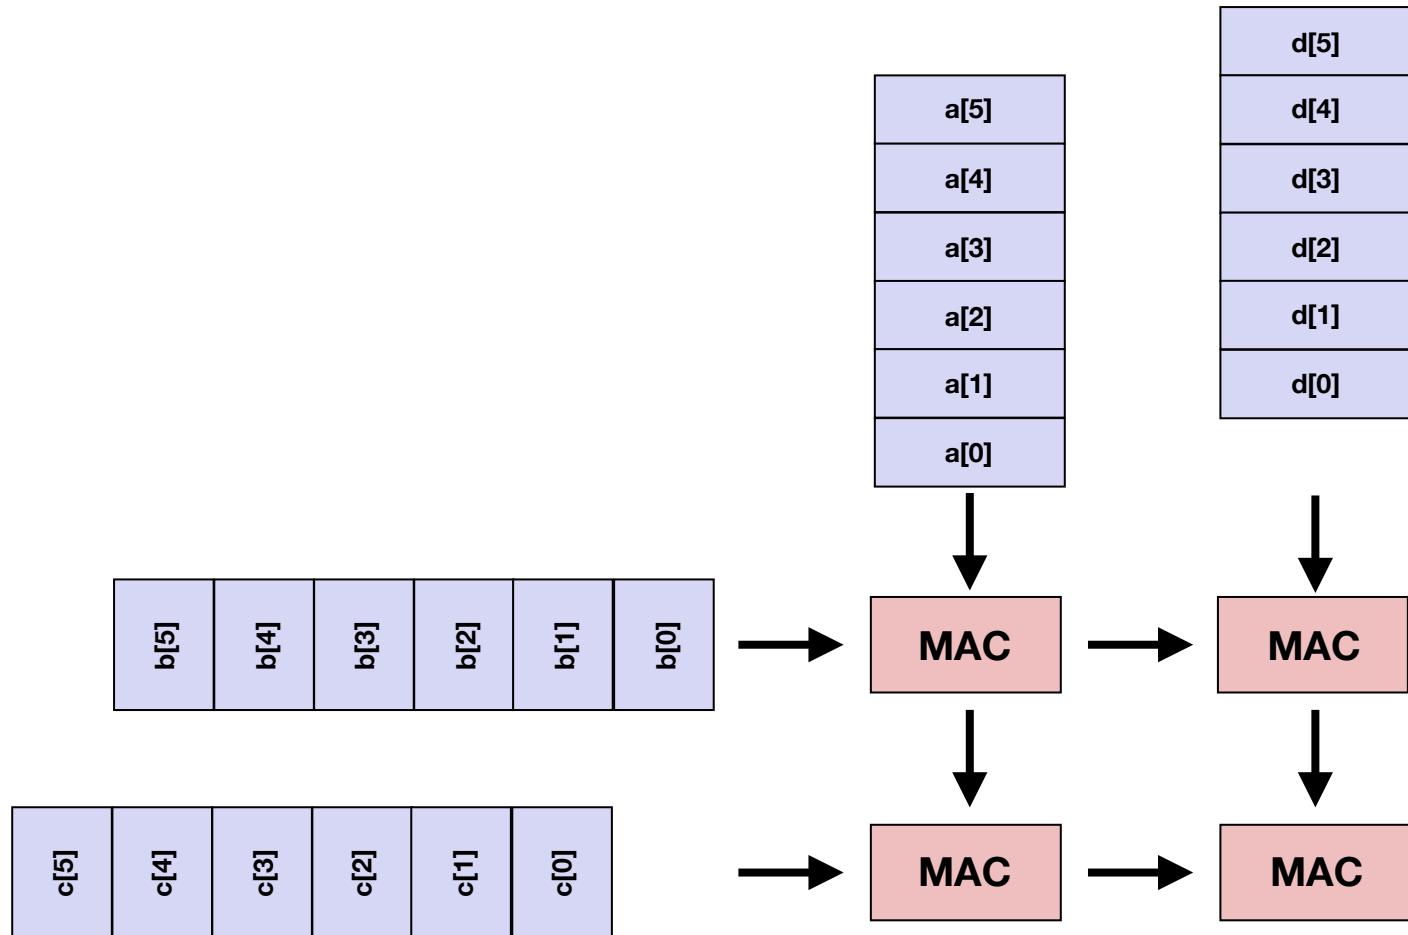
Matrix Vector Multiplication



Matrix Matrix Multiplication

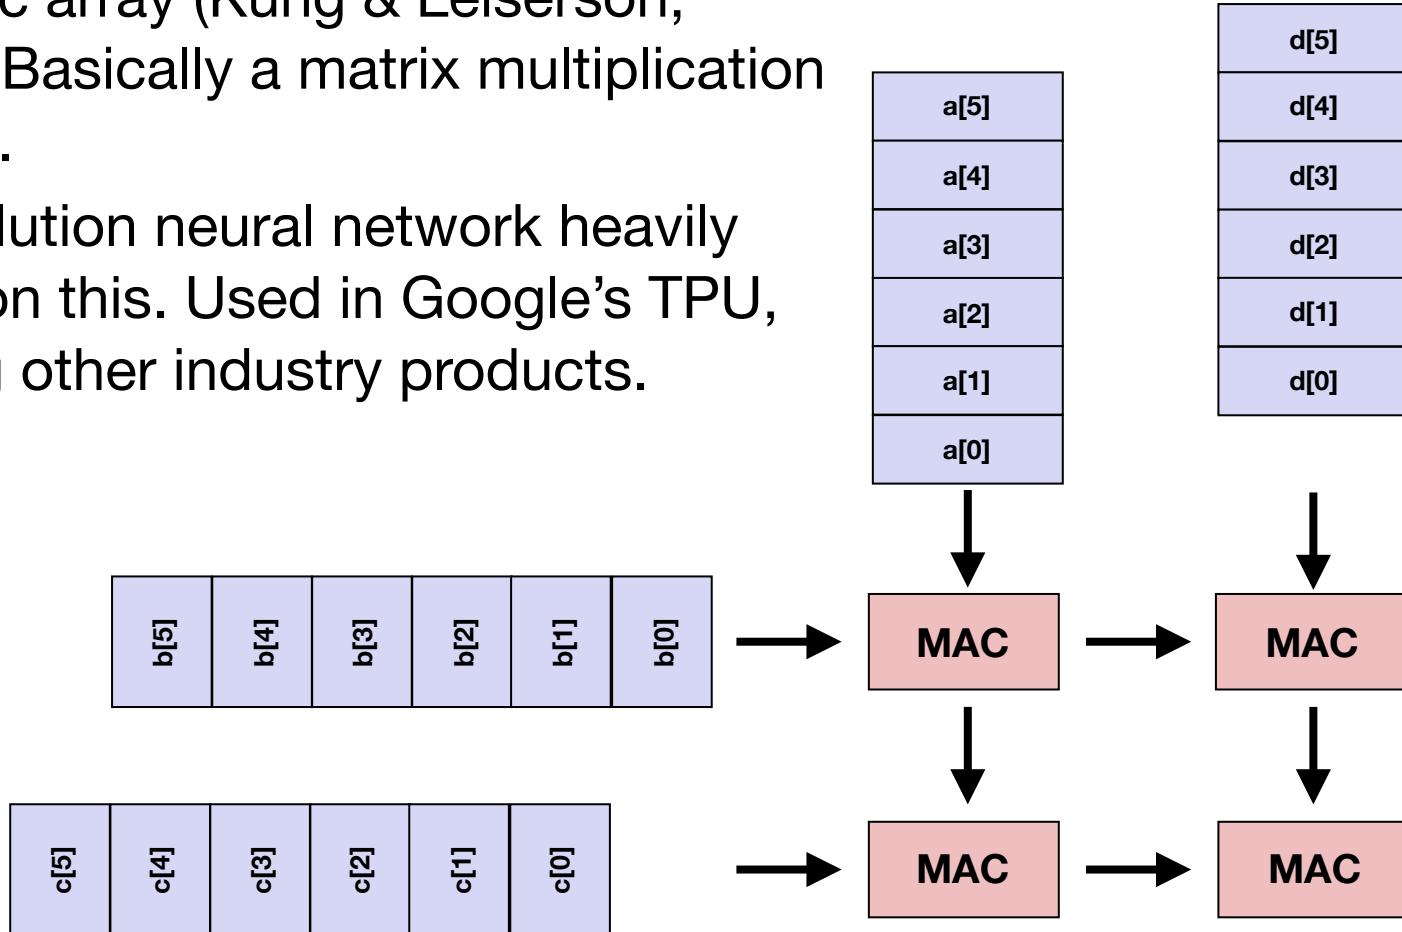


Matrix Matrix Multiplication



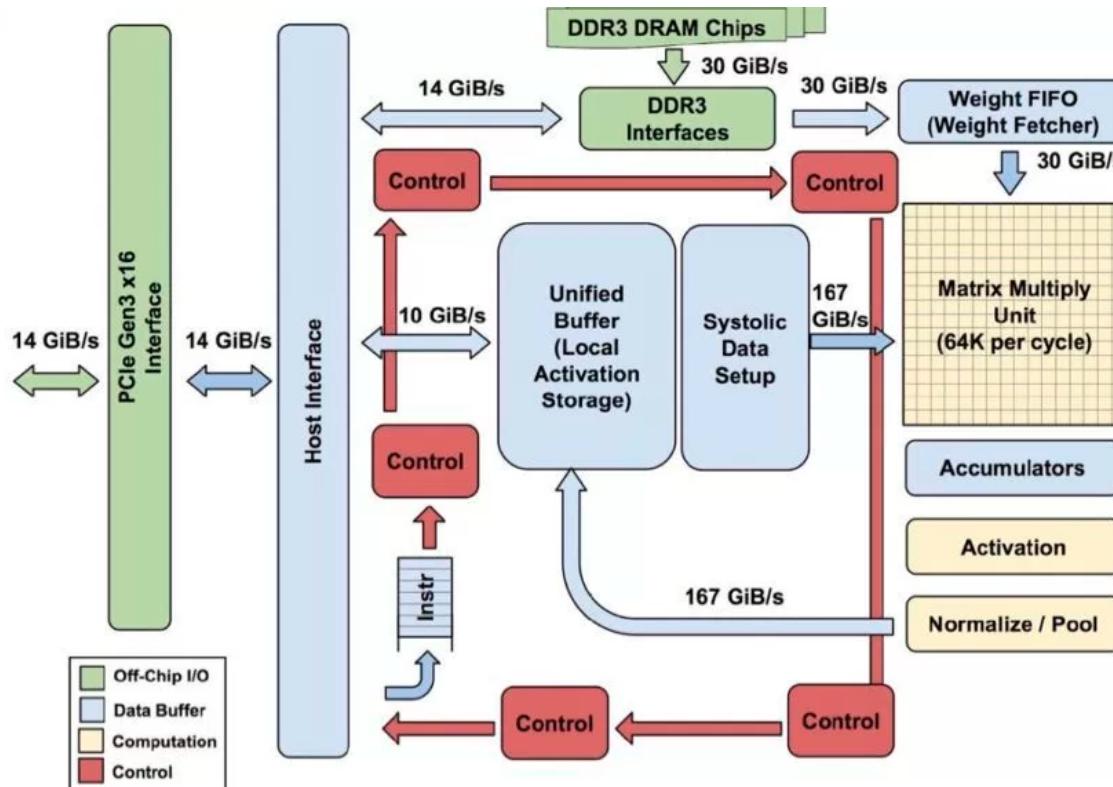
Matrix Matrix Multiplication

- Systolic array (Kung & Leiserson, 1978). Basically a matrix multiplication engine.
- Convolution neural network heavily relies on this. Used in Google's TPU, among other industry products.



Google Tensor Processing Unit

- Convolution in deep learning can be transformed to matrix multiplication.
- TPU: specialized processor (i.e., systolic array architecture) for tensor processing (matrix multiply)
 - 30x~80x more power-efficient than GPU



Another Domain: Video Compression



Another Domain: Video Compression

30-second video @ 1080p resolution (1920 x 1080 pixels per frame) @ 30 frames per second (FPS)
3 colors per pixel + 1 byte per color → 6.2 MB/frame → 6.2 MB x 30 s x 30 FPS = 5.2 GB total size
Actual H.264 video file size: 65.4 MB (**80-to-1 compression ratio**).

Compression/encoding done in real-time without you even realizing it!



Another Domain: Computational Photography

- Use computational algorithms to mimic a DSLR.
- Must be done in real-time. Executed on a dedicated Image Signal Processor (ISP).



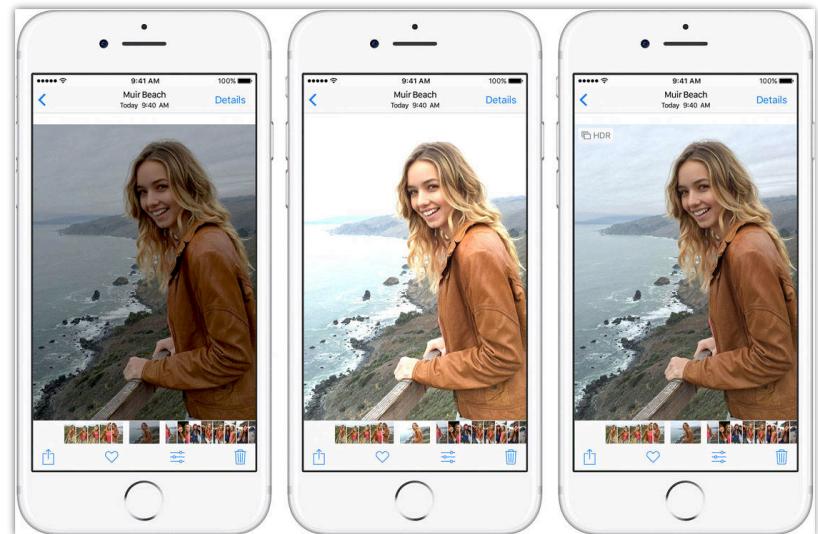
Conventional cameras



Today's “cameras”

Another Domain: Computational Photography

Portrait mode: simulate a large aperture



HDR mode: simulate a high dynamic range sensor

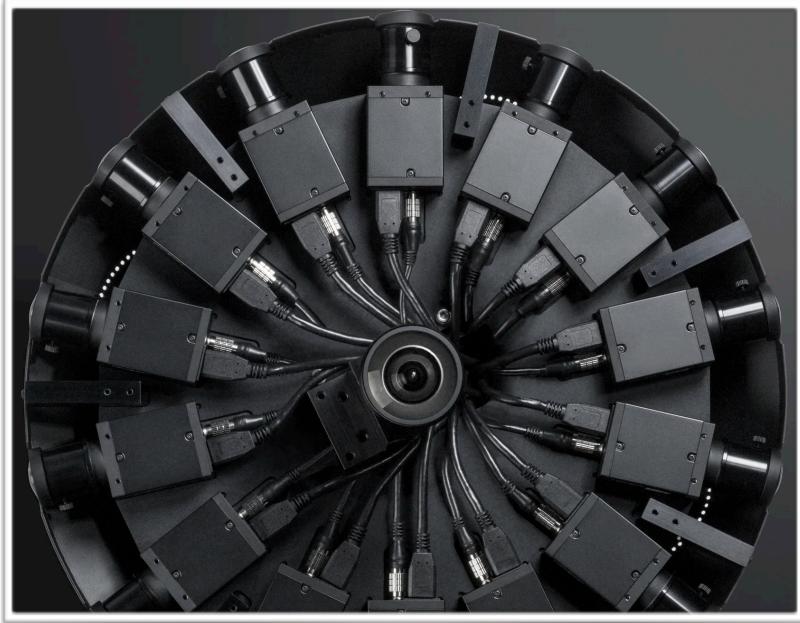
360° (VR/Panoramic) Videos and Photos



<https://www.cursosfotografiabarcelona.com/como-hacer-una-foto-esferica-360-grados-reflex/>

WWW.AIRPANO.COM

VR Video Capturing



Facebook Surround 360

Google Jump VR

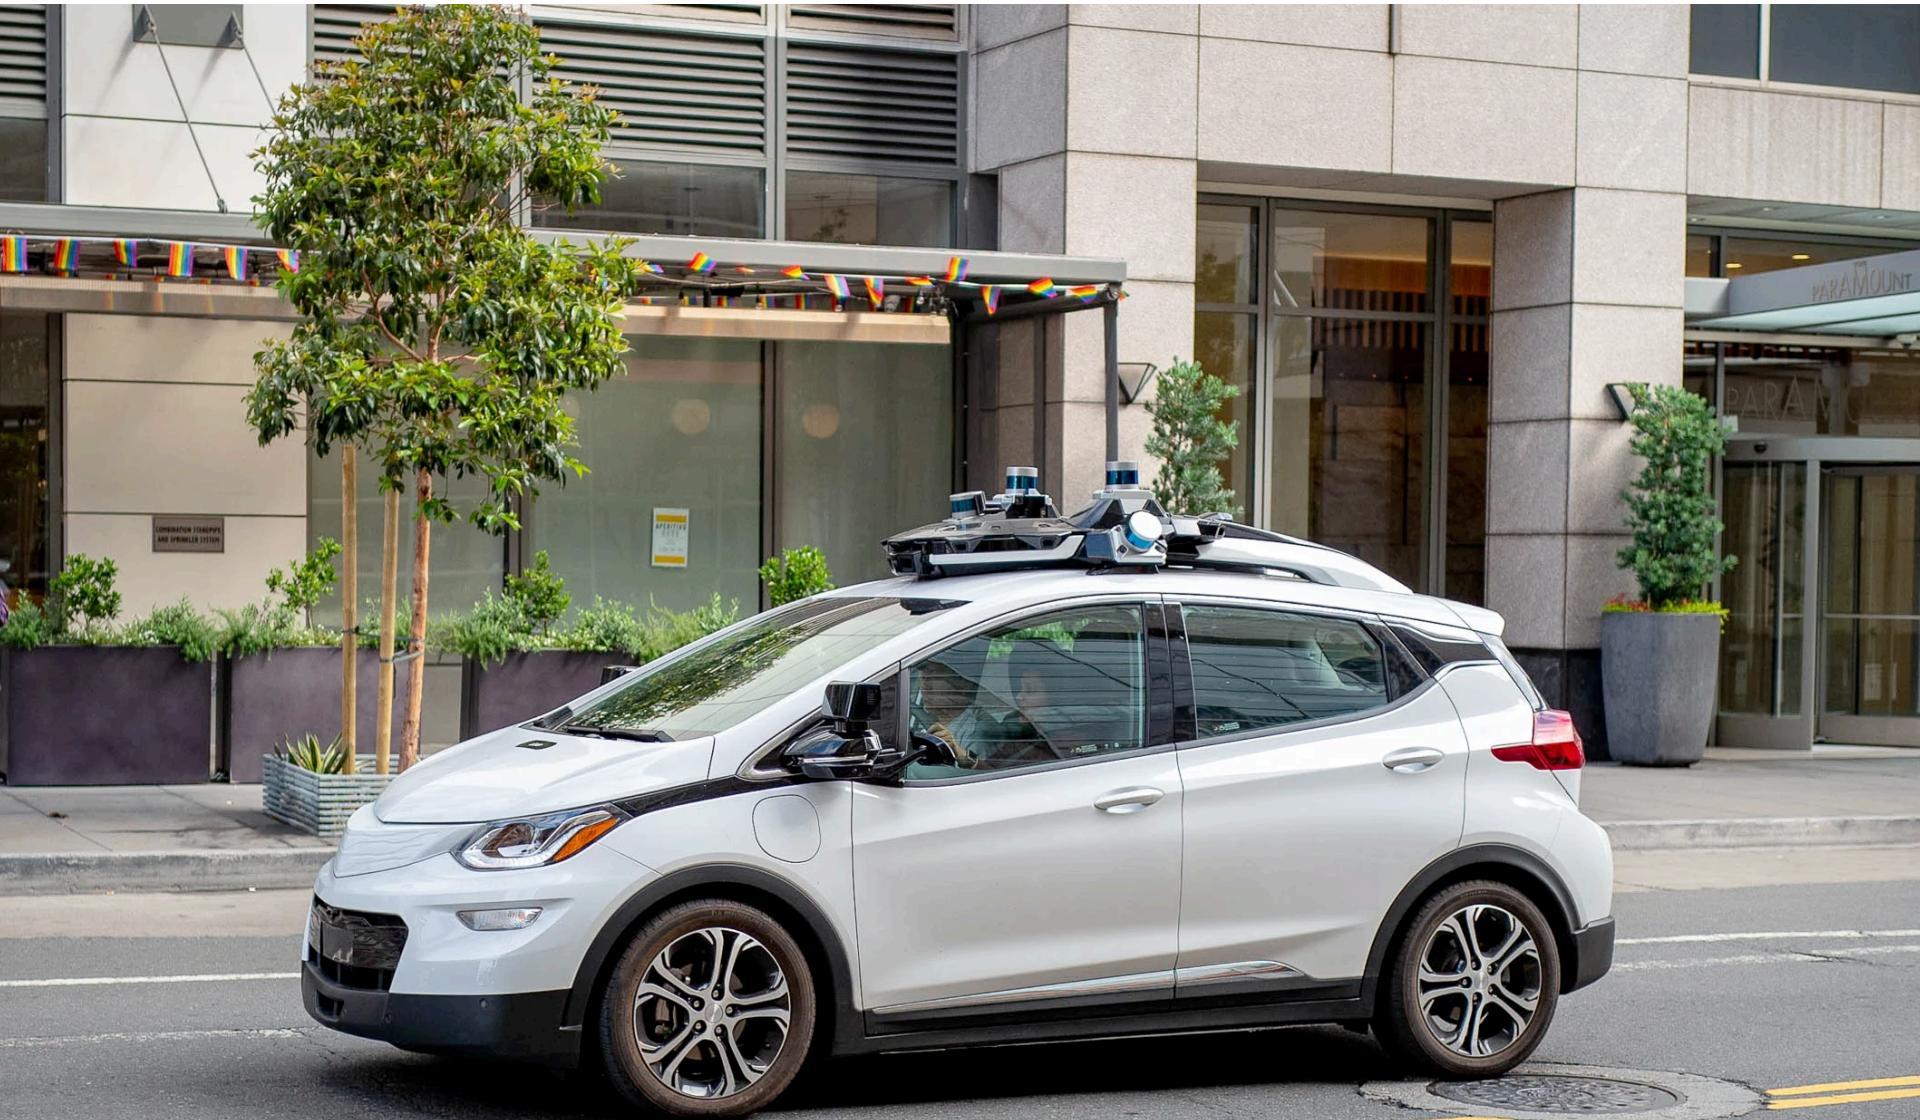


<https://static.googleusercontent.com/media/research.google.com/en//pubs/archive/45617.pdf>

<https://www.wired.com/2016/04/heres-360-degree-video-camera-facebook-giving-away/>

<https://twitter.com/yitechnology/status/918116570559336448>

Autonomous Machines



Autonomous Machines



Photorealistic Rendering



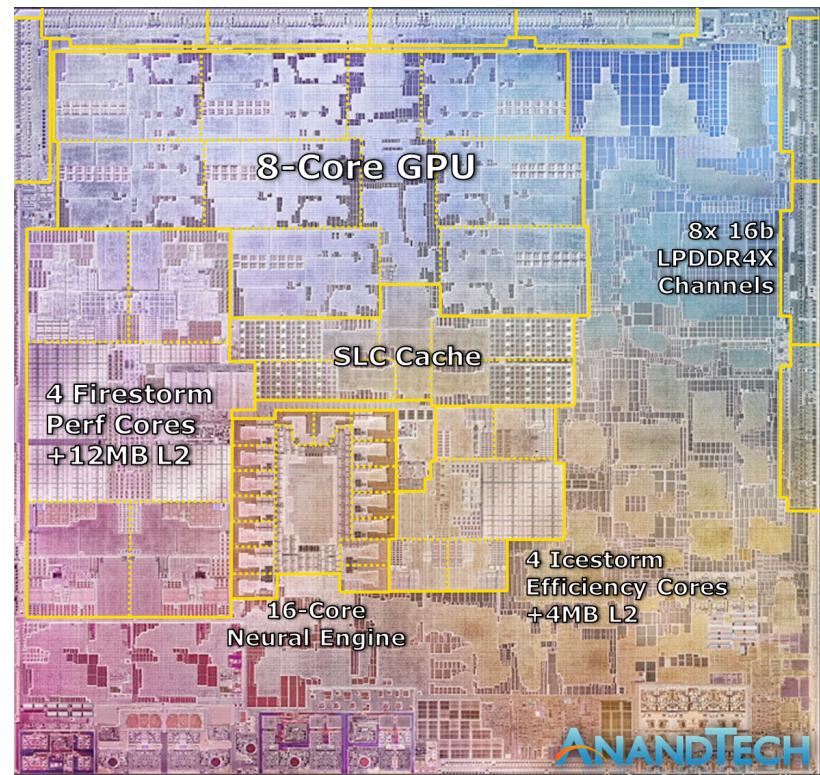
Photorealistic Rendering



Photorealistic Rendering

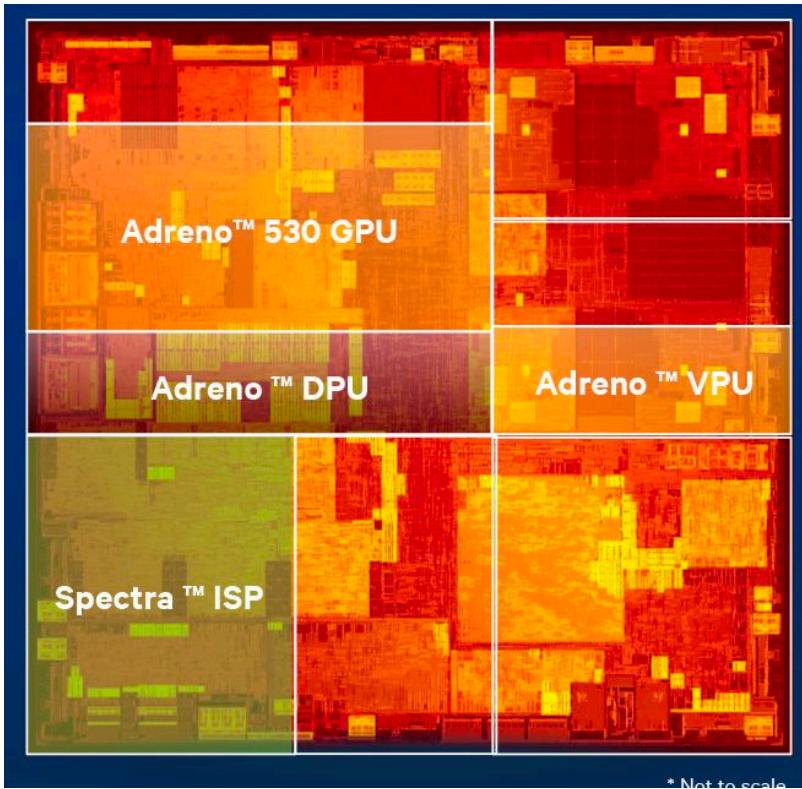


Today's Processor Chips are Full of Accelerators

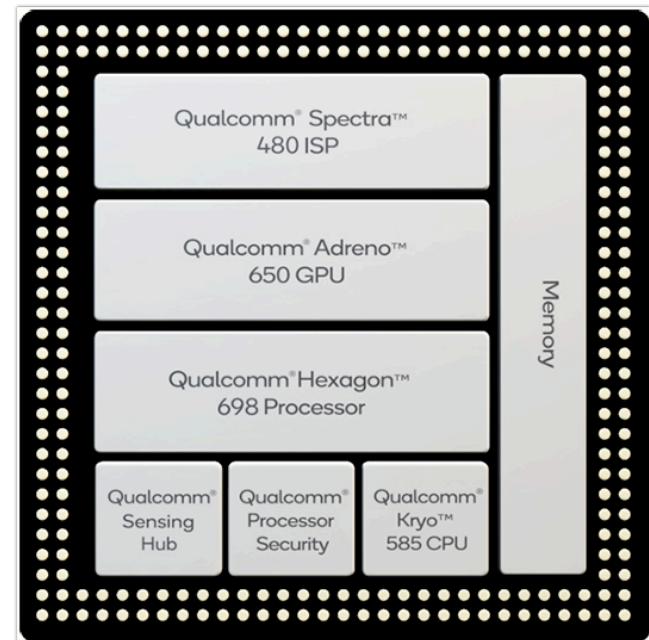


Today's Processor Chips are Full of Accelerators

Qualcomm
Snapdragon
820 SoC



Qualcomm
Snapdragon
835 SoC



Traditional Scope of Computer Systems

- Take a program and try to figure out how to best execute on the hardware

Problem

Algorithm

Program

Instruction Set
Architecture

Microarchitecture

Circuit

Real Scope of Computer Systems

- Understand the problem to be solved, design algorithms, understand algorithms characteristics to design the best computer systems.
- It's no longer enough to work with a given program without understanding it.

Problem

Algorithm

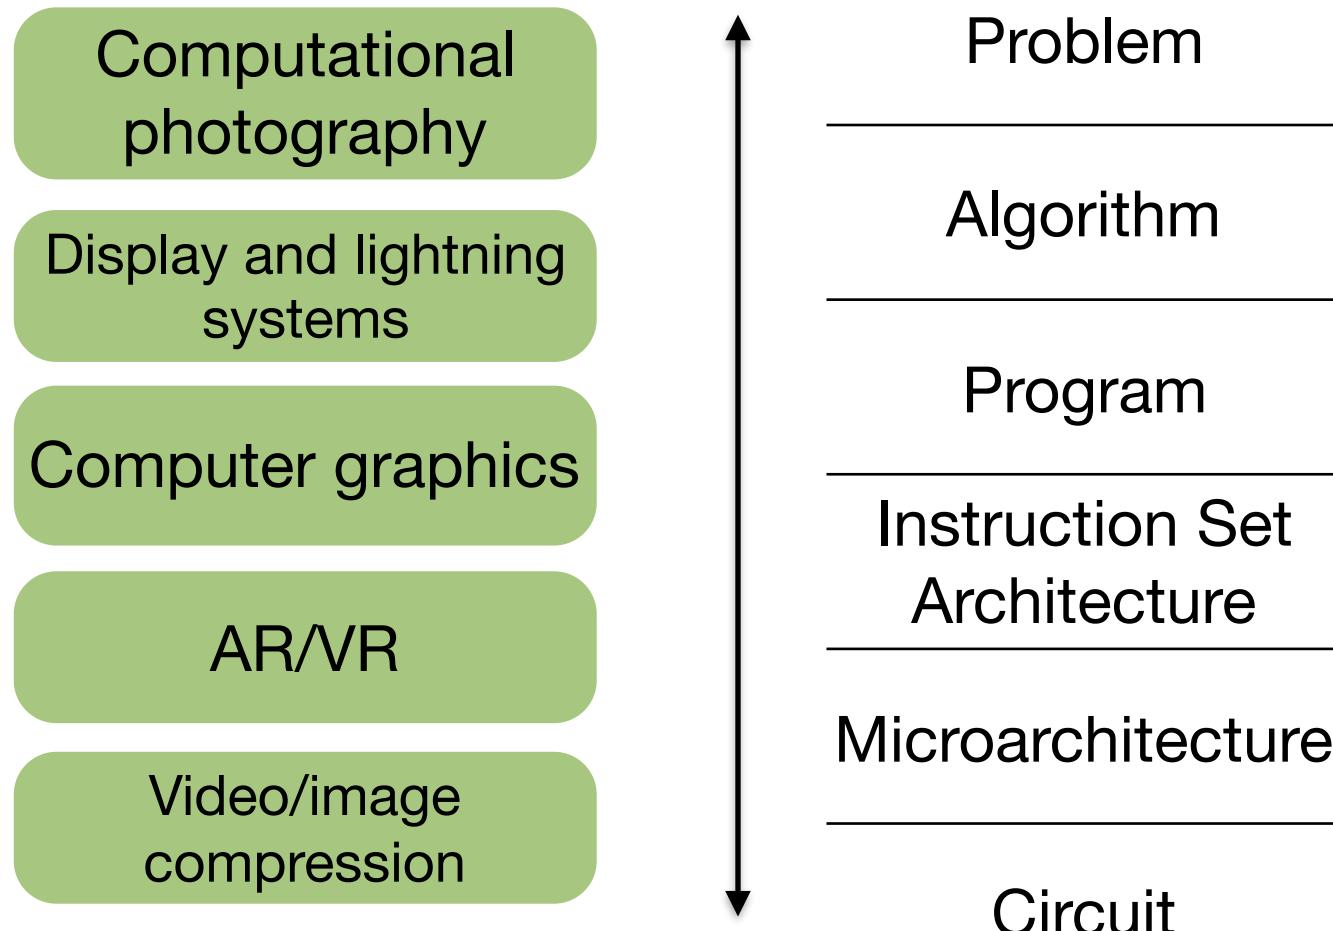
Program

Instruction Set
Architecture

Microarchitecture

Circuit

CSC 292/572: Mobile Visual Computing



The Most Important Take Away of 252

- “There is no magic.”
- Every thing can be derived from first principles. Trust your logical reasoning.
- Apply to virtually everything in science and engineering.

The Second Most Important Take Away of 252

- “Things don’t have to be this way.”
- As long as you don’t violate physics, you can design a computer however you want.
- But every design decision you make usually involves certain trade-offs. Be clear what your design goal is.

The Third Most Important Take Away of 252

- Virtually all computer system design practices follow a small set of basic principles.
- It is these basic principles that are important, not the practices.

